

Advanced Data Management for Data Analysis

Stefan Manegold

Data Management @ LIACS

Group leader Database Architectures
Centrum Wiskunde & Informatica (CWI)
Amsterdam

s.manegold@liacs.leidenuniv.nl

<http://www.cwi.nl/~manegold/>

ADM: Agenda

- 07.09.2022: Lecture 1: **Introduction**
- 14.09.2022: Lecture 2: **SQL Recap**
(plus Assignment 1 [in groups; 3 weeks]: TPC-H benchmark)
- 21.09.2022: Lecture 3: **Column-Oriented Database Systems (1/6) - Motivation & Basic Concepts**
- 28.09.2022: Lecture 4: **Column-Oriented Database Systems (2a/6) - Selected Execution Techniques (1/2)**
- 05.10.2022: Lecture 5: **Column-Oriented Database Systems (2b/6) - Selected Execution Techniques (2/2)**
(plus Assignment 3 [in groups; 3 weeks]: Compression techniques)
- 12.10.2022: Lecture 6: **Column-Oriented Database Systems (3/6) - Cache Conscious Joins**
- 19.10.2022: Lecture 7: **Column-Oriented Database Systems (4/6) - “Vectorized Execution”**
- 26.10.2022: **No lecture!**
- 02.11.2022: Lecture 8: **DuckDB: An embedded database for data science (1/2) (guest lecture & hands-on)**
(plus Assignment 2 [individual; 2 weeks]: Analysing NYC Cab dataset with DuckDB)
- 09.11.2022: Lecture 9: **DuckDB: An embedded database for data science (2/2) (guest lecture & hands-on)**
- 16.11.2022: Lecture 10: **Branch Misprediction & Predication**
(plus Assignment 4 [individual; 2 weeks]: Predication)
- 23.11.2022: Lecture 11: **Column-Oriented Database Systems (5/6) - Adaptive Indexing**
- 30.11.2022: Lecture 12: **Column-Oriented Database Systems (6/6) - Progressive Indexing**

ADM: Literature (1/2)

- **Column-Oriented Database Systems (1/6) - Motivation & Basic Concepts**

- “An overview of cantor: a new system for data analysis”. Ilkka Karasalo, Per Svensson. SSDBM 1983.
- “A decomposition storage model”. George P. Copeland, Setrag Khoshafian. SIGMOD Conference, 1985.
- “Cache Conscious Algorithms for Relational Query Processing”. Ambuj Shatdal, Chander Kant, Jeffrey F. Naughton. VLDB 1994.
- “MIL Primitives for Querying a Fragmented World”. Peter A. Boncz, Martin L. Kersten. VLDB J. 8(2): 101-119, 1999.
- “Database Architecture Optimized for the New Bottleneck: Memory Access”. Peter A. Boncz, Stefan Manegold, Martin L. Kersten. VLDB 1999.
- “DBMSs On A Modern Processor: Where Does Time Go?”. Anastassia Ailamaki, David J. DeWitt, Mark D. Hill, David A. Wood. VLDB 1999.
- “Weaving Relations for Cache Performance ("PAX")”. Anastassia Ailamaki, David J. DeWitt, Mark D. Hill, Marios Skounakis. VLDB 2001.
- “A Case for Fractured Mirrors”. Ravishankar Ramamurthy, David J. DeWitt, Qi Su. VLDB 2002.
- “Data Morphing: An Adaptive, Cache-Conscious Storage Technique”. Richard A. Hankins, Jignesh M. Patel. VLDB 2003.
- “Clotho: Decoupling Memory Page Layout from Storage Organization”. Minglong Shao, Jiri Schindler, Steven W. Schlosser, Anastassia Ailamaki, Gregory R. Ganger. VLDB 2004.
- “MonetDB-X100 - A DBMS In The CPU Cache”. Marcin Zukowski, Peter A. Boncz, Niels Nes, Sándor Héman. IEEE Data Eng. Bull. 28(2): 17-22, 2005.
- ““One size fits all”: an idea whose time has come and gone”. Michael Stonebraker, Ugur Çetintemel. ICDE 2005.
- “Performance Tradeoffs in Read-Optimized Databases”. Stavros Harizopoulos, Velen Liang, Daniel J. Abadi, Samuel Madden. VLDB 2006.
- ...

ADM: Literature (2/2)

- **Column-Oriented Database Systems (1/6) - Motivation & Basic Concepts** (*cont.*)
 - ...
 - “One Size Fits All? - Part 2: Benchmarking Results”. Michael Stonebraker, Chuck Bear, Ugur Çetintemel, Mitch Cherniack, Tingjian Ge, Nabil Hachem, Stavros Harizopoulos, John Lifter, Jennie Rogers, Stanley B. Zdonik. CIDR 2007.
 - “C-Store: A Column-oriented DBMS”. Michael Stonebraker, Daniel J. Abadi, Adam Batkin, Xuedong Chen, Mitch Cherniack, Miguel Ferreira, Edmond Lau, Amerson Lin, Samuel Madden, Elizabeth J. O’Neil, Patrick E. O’Neil, Alex Rasin, Nga Tran, Stanley B. Zdonik. VLDB 2005.
 - “Breaking the memory wall in MonetDB”. Peter A. Boncz, Martin L. Kersten, Stefan Manegold. Commun. ACM 51(12): 77-85, 2008.
 - “Column-Stores vs Row-Stores: How Different are They Really?”. Daniel J. Abadi, Samuel Madden, Nabil Hachem. SIGMOD Conference 2008.
 - “DSM vs. NSM: CPU performance tradeoffs in block-oriented query processing”. Marcin Zukowski, Niels Nes, Peter A. Boncz. DaMoN 2008.
 - “Fast Scans and Joins Using Flash Drives”. Mehul A. Shah, Stavros Harizopoulos, Janet L. Wiener, Goetz Graefe. DaMoN 2008.
 - “Read-Optimized Databases, In-Depth”. Allison L. Holloway, David J. DeWitt. Proc. VLDB Endow. 1(1): 502-513, 2008.
 - “Teaching an Old Elephant New Tricks”. Nicolas Bruno. CIDR 2009.
 - “Query Processing Techniques for Solid State Drives”. Dimitris Tsirogiannis, Stavros Harizopoulos, Mehul A. Shah, Janet L. Wiener, Goetz Graefe. SIGMOD Conference 2009.
 - “MonetDB: Two Decades of Research in Column-oriented Database Architectures”. Stratos Idreos, Fabian Groffen, Niels Nes, Stefan Manegold, K. Sjoerd Mullender, Martin L. Kersten. IEEE Data Eng. Bull. 35(1): 40-45, 2012.
 - “The Vertica Analytic Database: C-Store 7 Years Later”. Andrew Lamb, Matt Fuller, Ramakrishna Varadarajan, Nga Tran, Ben Vandiver, Lyric Doshi, Chuck Bear. Proc. VLDB Endow. 5(12): 1790-1801, 2012.

Why?

Motivation (early 1990s)

- Relational DBMSs dominate since the late 1970's / early 1980's
 - IBM DB2, MS SQL Server, Oracle, Ingres, ...
 - Transactional workloads (OLTP, row-wise access)
 - I/O based processing

Why?

Workload changes: Transactions (OLTP) vs ...

contract	client	date	name	price	city	product													
12302346	10042334		Eno		Redmond	Car													
37611373	10987097		Gotz		Berkeley	← Redmond	update query												
95371001	10032112		Chen		Seattle	House	lookup query												
51213123	10032423		Jones		Washington	Travel													
54535545	10087823		Smith		New York	House													
45447894	10013232		Doe		Boston	Car	insert query												

find client
10032112

OLTP queries:
access all
columns of
just one row.

Why?

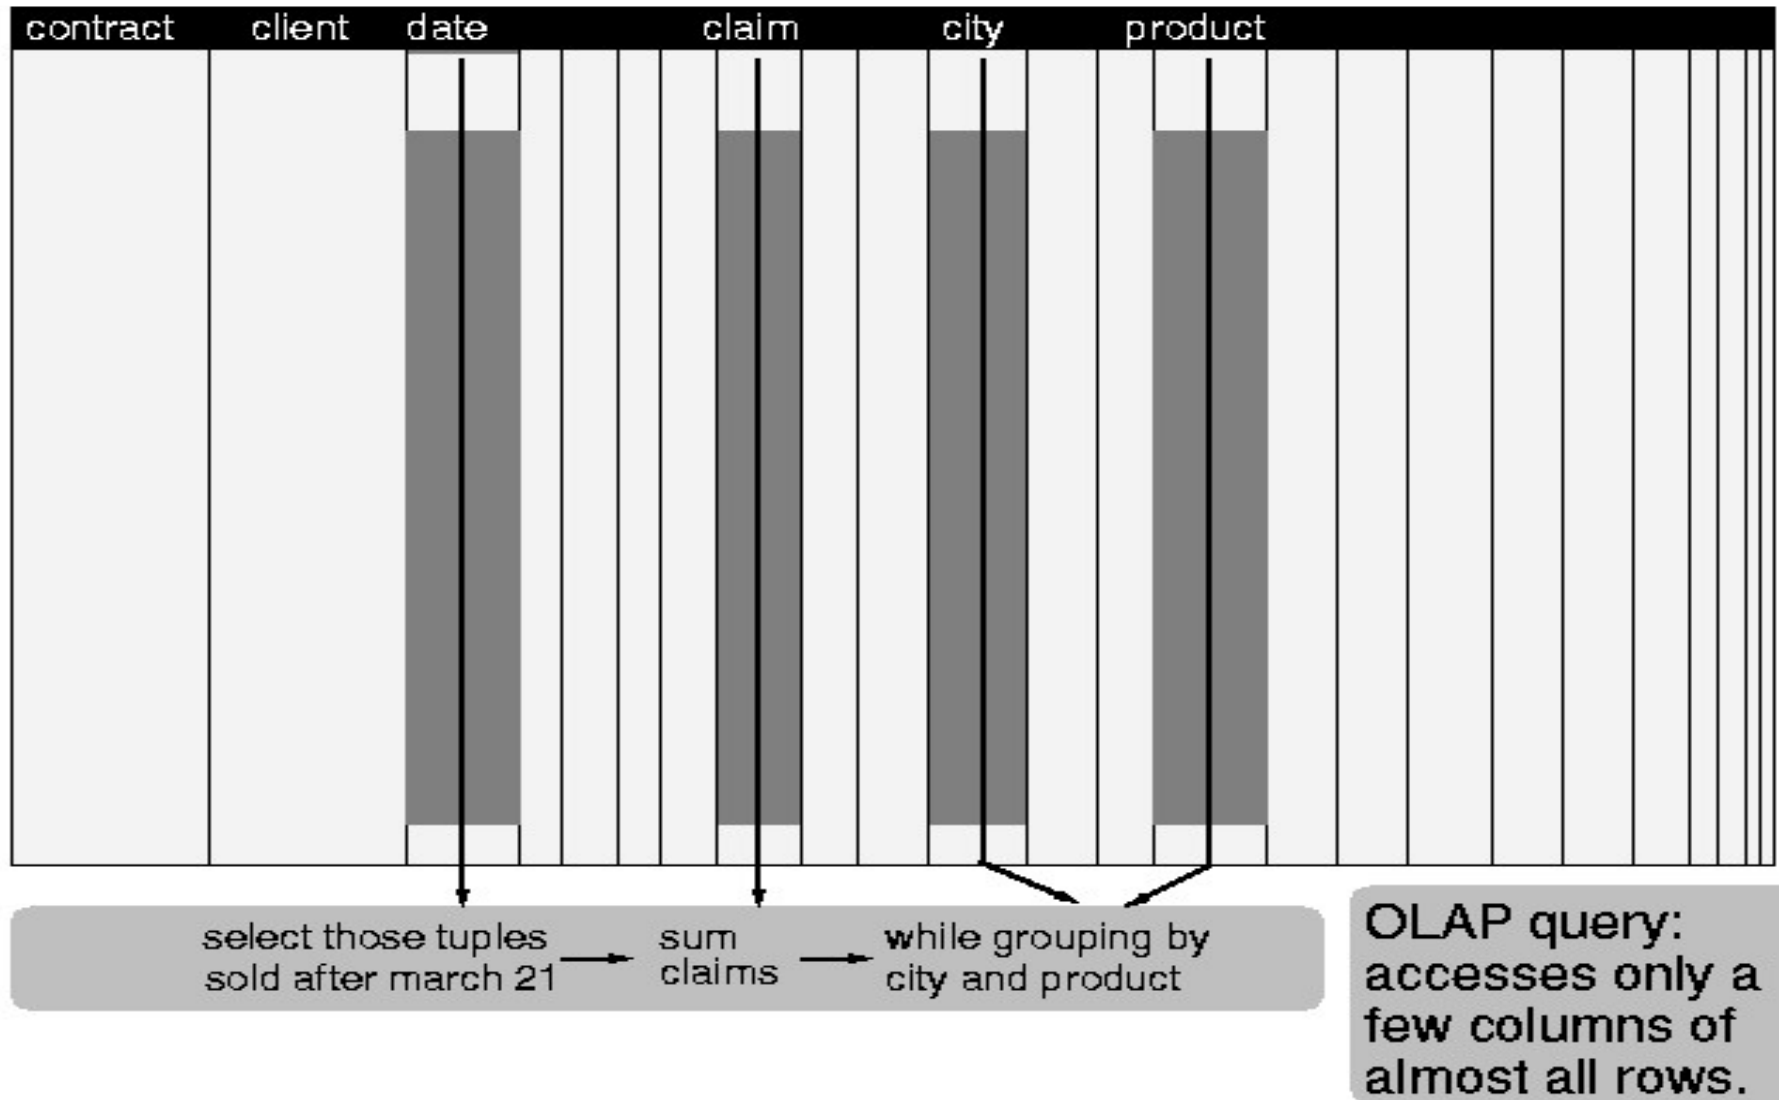


Motivation (early 1990s)

- Relational DBMSs dominate since the late 1970's / early 1980's
 - IBM DB2, MS SQL Server, Oracle, Ingres, ...
 - Transactional workloads (OLTP, row-wise access)
 - I/O based processing
- **But:**
 - Workloads change (early 1990s)
 - Hardware changes (late 1990s)
 - Data “explodes” (early 2000s)

Why?

Workload changes: ... vs OLAP, BI, Data Mining, ...




Why?



Databases hit The Memory Wall


- Detailed and exhaustive analysis for different workloads using 4 RDBMSs by Ailamaki, DeWitt, Hill, Wood in VLDB 1999: *“DBMSs On A Modern Processor: Where Does Time Go?”*
- CPU is 60%-90% idle, waiting for memory:
 - L1 data stalls
 - L1 instruction stalls
 - L2 data stalls
 - TLB stalls
 - Branch mispredictions
 - Resource stalls





Why?

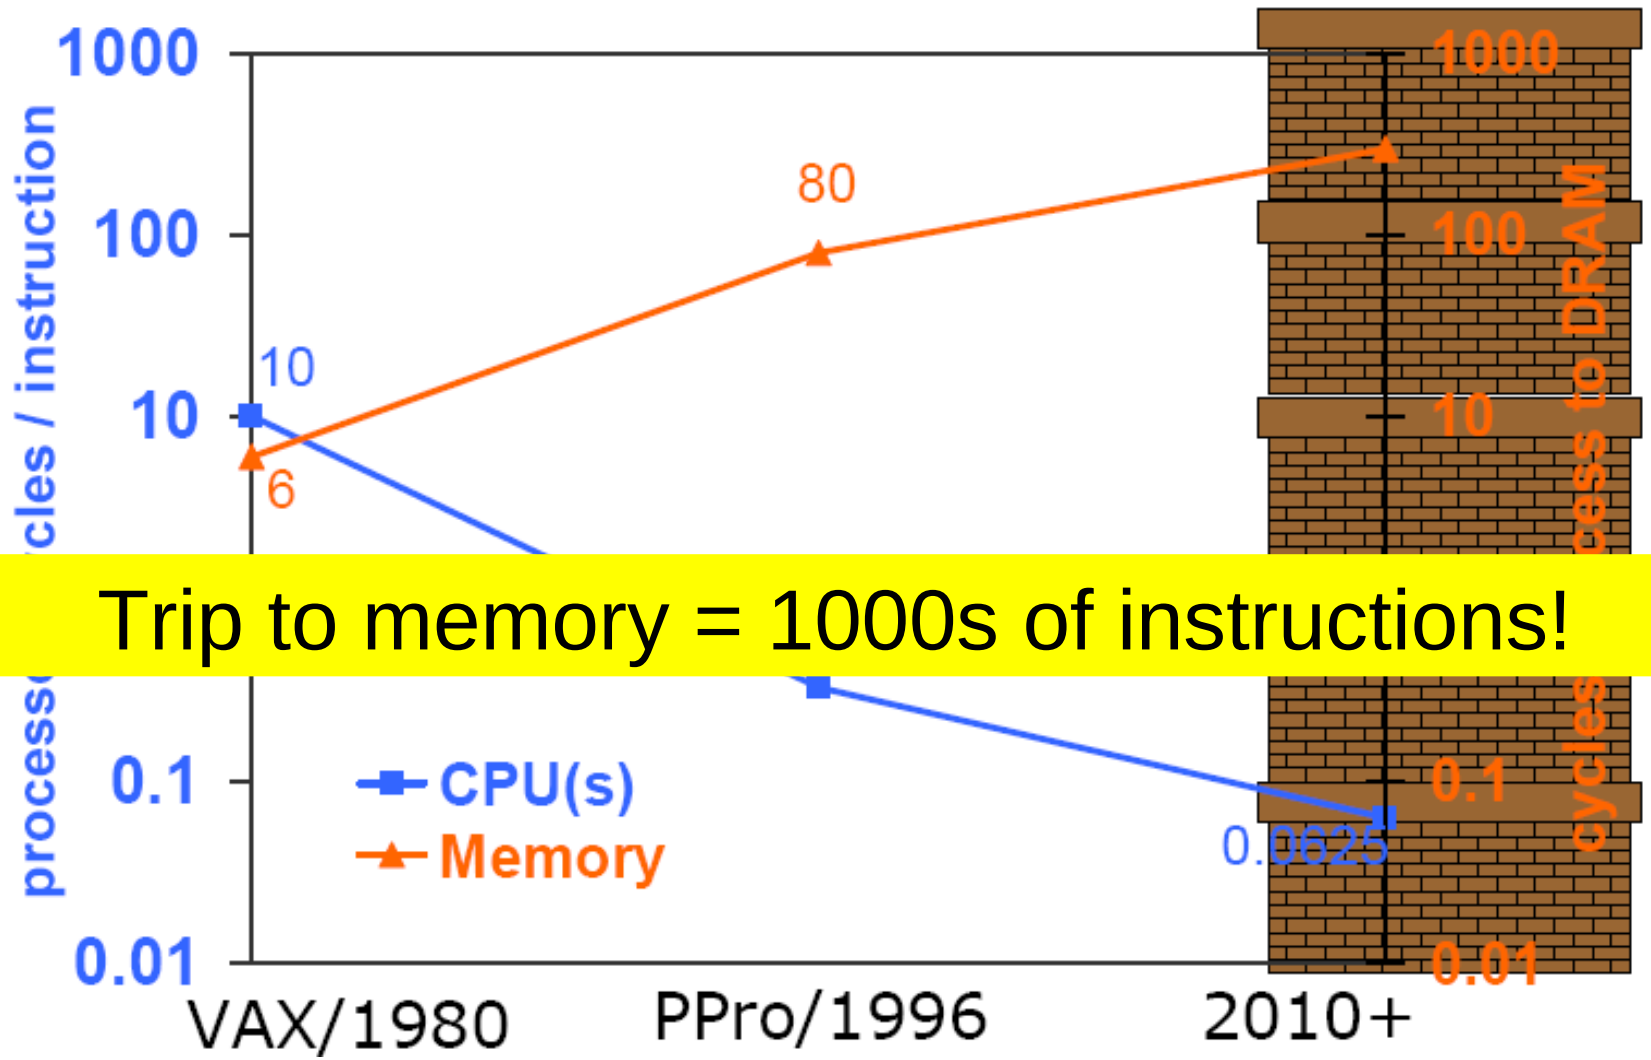
Hardware Changes: The Memory Wall





Why?

Hardware Changes: The Memory Wall




Trip to memory = 1000s of instructions!



CPU Architecture

Elements:

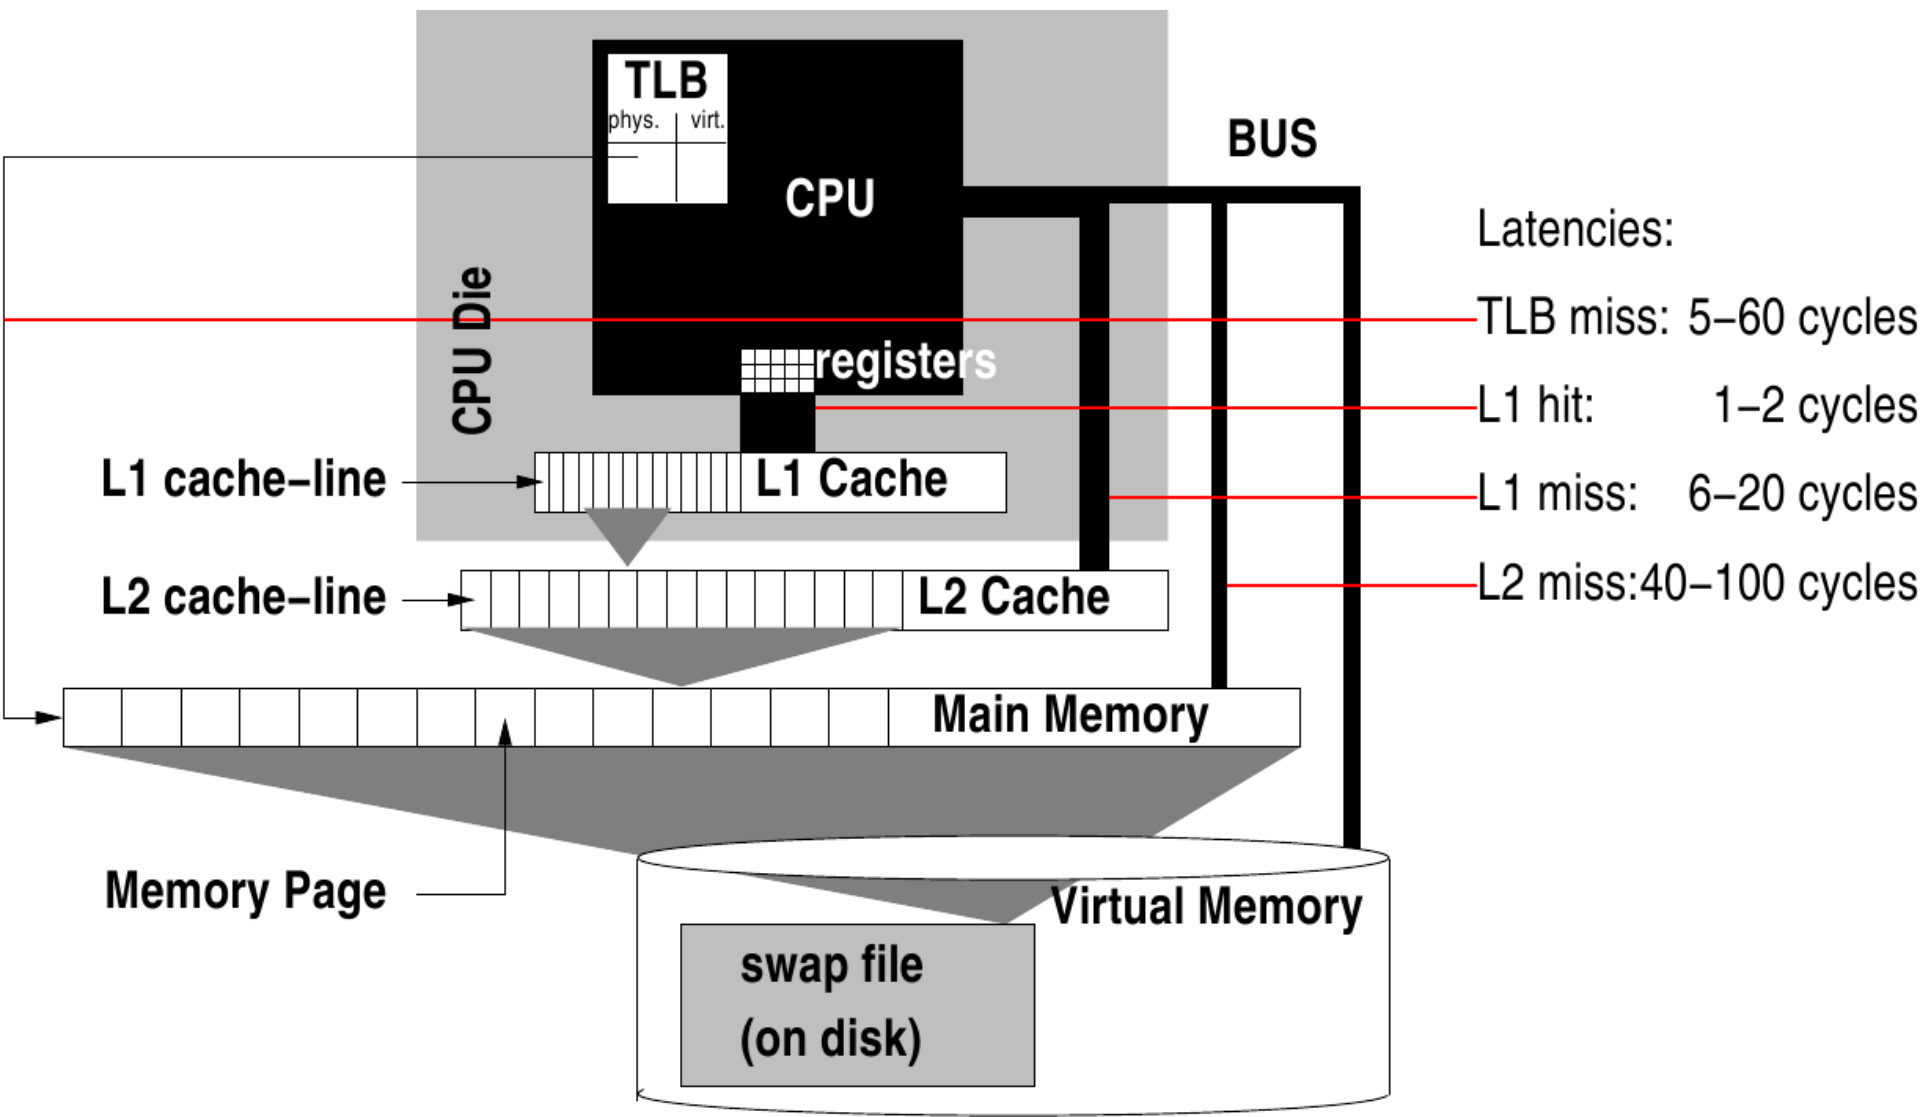
- Storage
 - CPU caches L1/L2/L3
- Registers
- Execution Unit(s)
 - Pipelined
 - SIMD





Why?

Hardware Changes: Memory Hierarchies






Why?

Hardware Changes: Memory Hierarchies

- Caches trade off capacity for speed
- Exploit instruction/data locality
- Demand fetch/wait for data

[ADH99]:

- Running top 4 database systems
- **At most 50% CPU utilization**





Why?


Hardware Changes: Memory Hierarchies

- Caches trade off capacity for speed
- Exploit instruction/data locality
- Demand fetch/wait for data

[ADH99]:

- Running top 4 database systems
- **At most 50% CPU utilization**

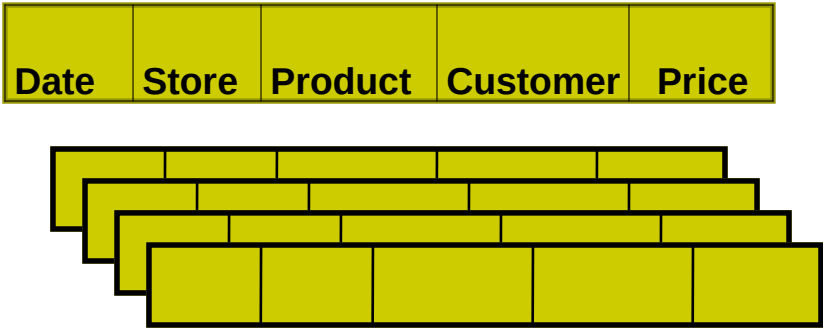
+Transition Lookaside Buffer (TLB)
Cache for VM address translation →
only 64 entries!





What is a column-store?

row-store

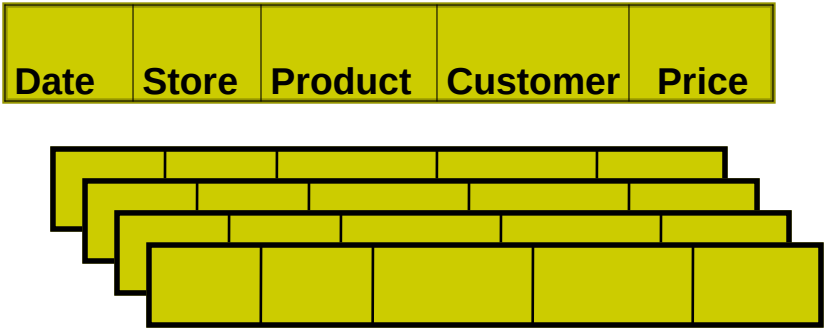


- + easy to add/modify a record
- might read in unnecessary data




What is a column-store?

row-store



- + easy to add/modify a record
- might read in unnecessary data

column-store

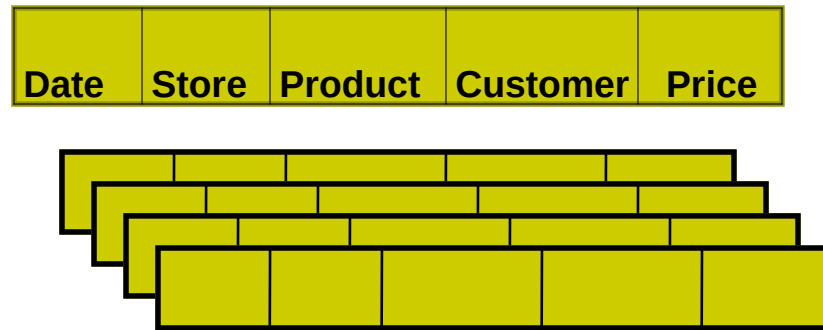


- + only need to read in relevant data
- tuple writes require multiple accesses




What is a column-store?

row-store



- + easy to add/modify a record
- might read in unnecessary data

column-store



- + only need to read in relevant data
- tuple writes require multiple accesses

=> suitable for read-mostly, read-intensive, large data repositories




Are these two fundamentally different?

- The only fundamental difference is the storage layout
- However: we need to look at the big picture



Are these two fundamentally different?

- The only fundamental difference is the storage layout
- However: we need to look at the big picture





Are these two fundamentally different?

- The only fundamental difference is the storage layout
- However: we need to look at the big picture

different storage layouts proposed

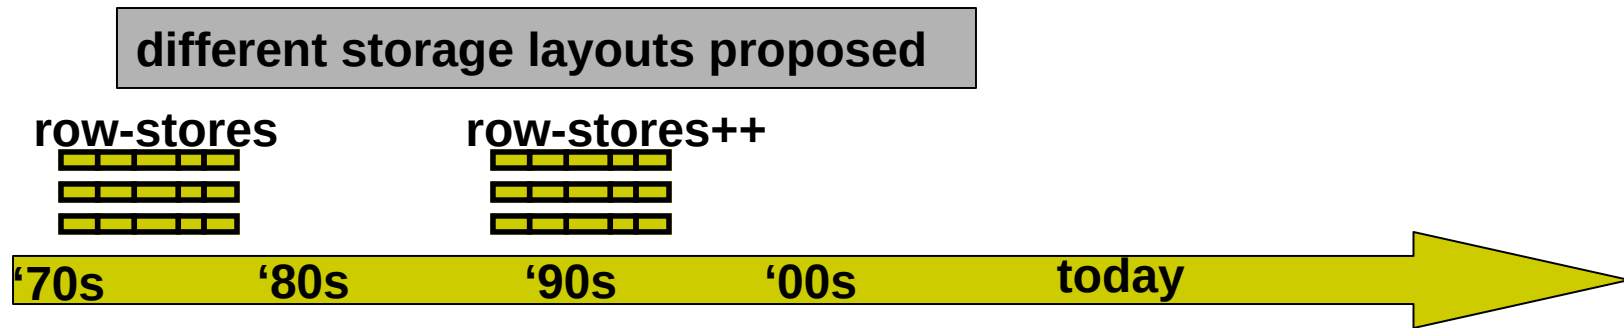
row-stores





Are these two fundamentally different?

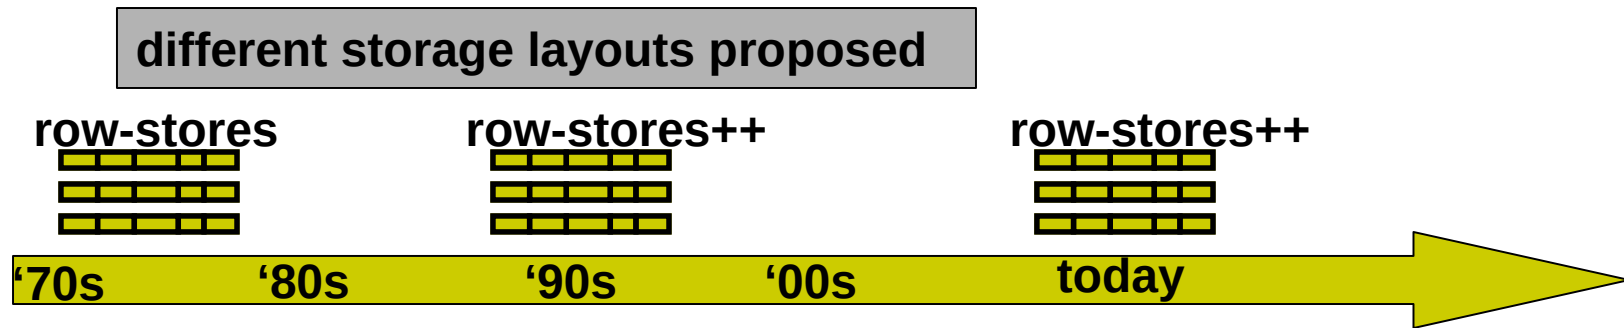
- The only fundamental difference is the storage layout
- However: we need to look at the big picture





Are these two fundamentally different?

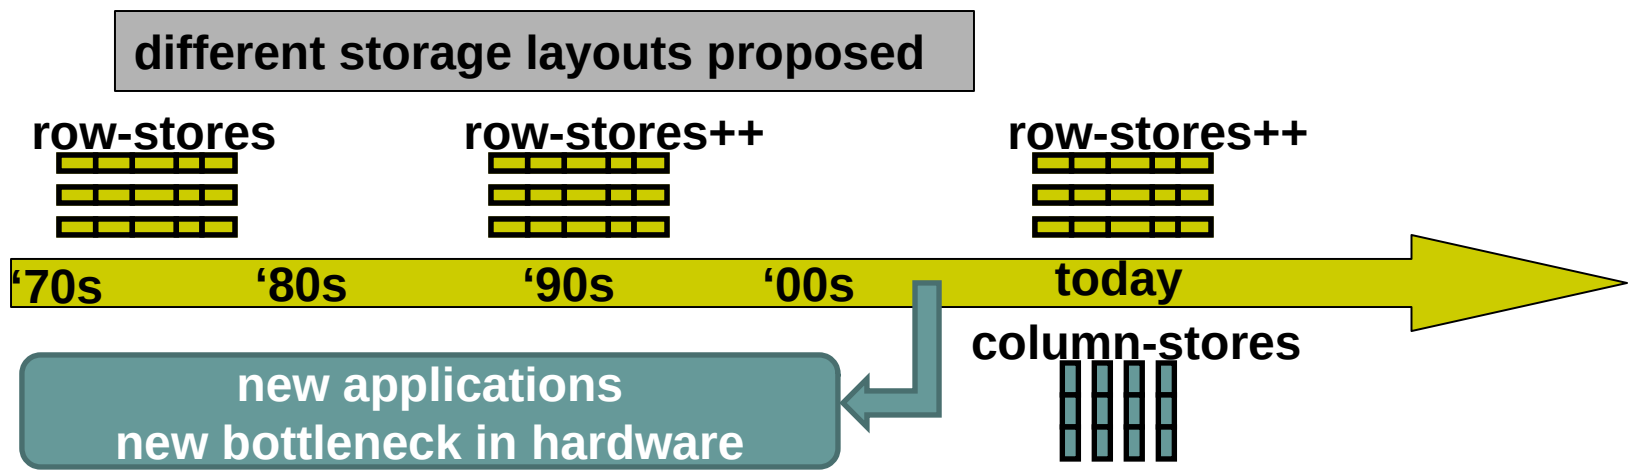
- The only fundamental difference is the storage layout
- However: we need to look at the big picture





Are these two fundamentally different?

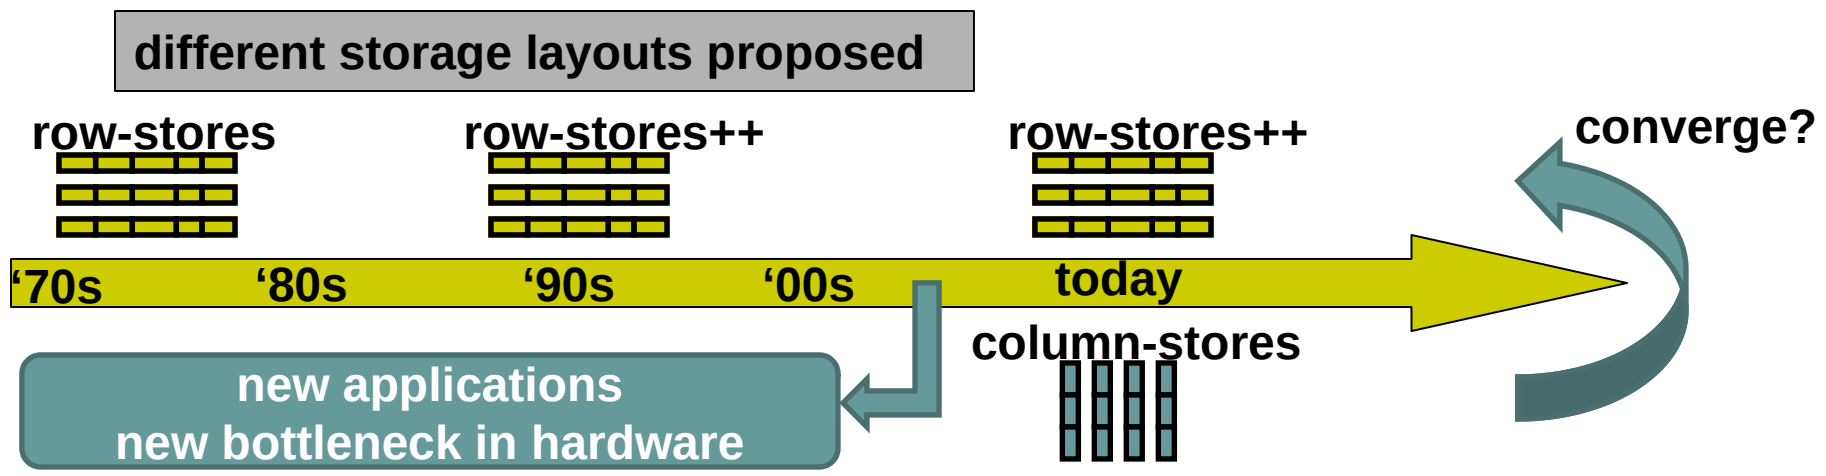
- The only fundamental difference is the storage layout
- However: we need to look at the big picture





Are these two fundamentally different?

- The only fundamental difference is the storage layout
- However: we need to look at the big picture




- How did we get here, and where we are heading?
- What are the column-specific optimizations?
- How do we improve CPU efficiency when operating on Cs?



Telco Data Warehousing example

- Typical DW installation
- Real-world example

“One Size Fits All? - Part 2: Benchmarking Results” Stonebraker et al. CIDR 2007






Telco Data Warehousing example

- Typical DW installation
- Real-world example

“One Size Fits All? - Part 2: Benchmarking Results” Stonebraker et al. CIDR 2007



QUERY 2

```
SELECT account.account_number,  
sum (usage.toll_airtime),  
sum (usage.toll_price)  
FROM usage, toll, source, account  
WHERE usage.toll_id = toll.toll_id  
AND usage.source_id = source.source_id  
AND usage.account_id = account.account_id  
AND toll.type_ind in ('AE', 'AA')  
AND usage.toll_price > 0  
AND source.type != 'CIBER'  
AND toll.rating_method = 'IS'  
AND usage.invoice_date = 20051013  
GROUP BY account.account_number
```


	Column-store	Row-store
Query 1	2.06	300
Query 2	2.20	300
Query 3	0.09	300
Query 4	5.24	300
Query 5	2.88	300

Why? Three main factors (next slides)

Telco example explained (1/3): *read efficiency*



row store




read pages containing entire rows

one row = 212 columns!

is this typical? (it depends)

column store



read only columns needed

in this example: 7 columns


caveats:

- “select * ” not any faster
- clever disk prefetching
- clever tuple reconstruction



Telco example explained (1/3): *read efficiency*

row store




read pages containing entire rows

one row = 212 columns!

is this typical? (it depends)

**What about vertical partitioning?
(it does not work with ad-hoc queries)**

column store



read only columns needed

in this example: 7 columns

caveats:

- “select * ” not any faster
- clever disk prefetching
- clever tuple reconstruction



Telco example explained (2/3): *compression efficiency*

- Columns compress better than rows
 - Typical row-store compression ratio 1 : 3
 - Column-store 1 : 10
- Why?
 - Rows contain values from different domains
=> more entropy, difficult to dense-pack
 - Columns exhibit significantly less entropy
 - Examples:

Male, Female, Female, Female, Male
1998, 1998, 1999, 1999, 1999, 2000
 - Caveat: CPU cost (use lightweight compression)



Telco example explained (3/3): *sorting & indexing efficiency*

- Compression and dense-packing free up space
 - Use multiple overlapping column collections
 - Sorted columns compress better
 - Range queries are faster
 - Use sparse clustered indexes

**What about heavily-indexed row-stores?
(works well for single column access,
cross-column joins become increasingly expensive)**




Additional opportunities for column-stores

- Block-tuple / vectorized processing
 - Easier to build block-tuple operators
 - Amortizes function-call cost, improves CPU cache performance
 - Easier to apply vectorized primitives
 - Software-based: bitwise operations
 - Hardware-based: SIMD
- Opportunities with compressed columns
 - *Avoid* decompression: operate directly on compressed
 - *Delay* decompression (and tuple reconstruction)
 - Also known as: *late materialization*
- Exploit columnar storage in other DBMS components
 - Physical design (both static and dynamic)

See: *Database Cracking*, from CWI

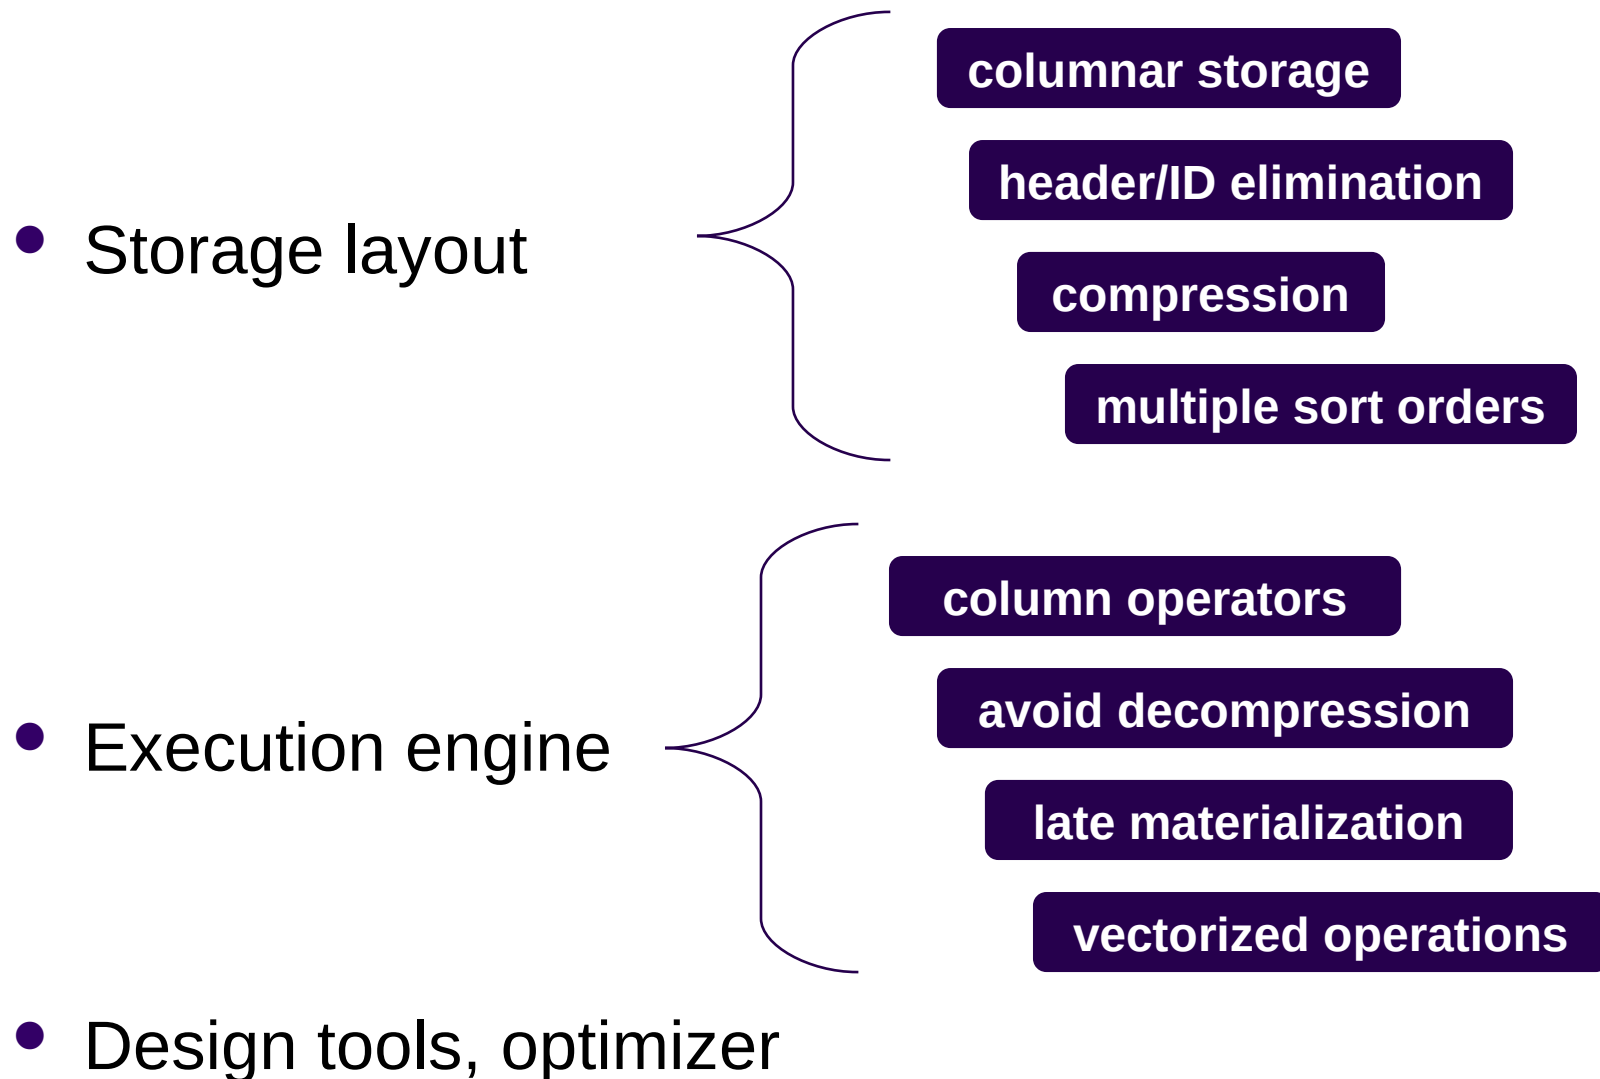
“Column-Stores vs Row-Stores: How Different are They Really?” Abadi, Hachem, and Madden. SIGMOD 2008.

Effect on C-Store performance





Summary of column-store key features





From DSM to Column-stores

70s -1985:

TOD: Time Oriented Database – Wiederhold et al.
"A Modular, Self-Describing Clinical Databank System,"
Computers and Biomedical Research, 1975
More 1970s: Transposed files, Lorie, Batory, Svensson.
"An overview of cantor: a new system for data analysis"
Karasalo, Svensson, SSDBM 1983

1985: DSM paper

"A decomposition storage model"
Copeland and Khoshafian. SIGMOD 1985.

1990s: Commercialization through SybaseIQ

Late 90s – 2000s: Focus on main-memory performance

- DSM "on steroids" [1997 – now] CWI: MonetDB
- Hybrid DSM/NSM [2001 – 2004] Wisconsin: PAX, Fractured Mirrors
Michigan: Data Morphing CMU: Clotho

2005 – : Re-birth of read-optimized DSM as "column-store"

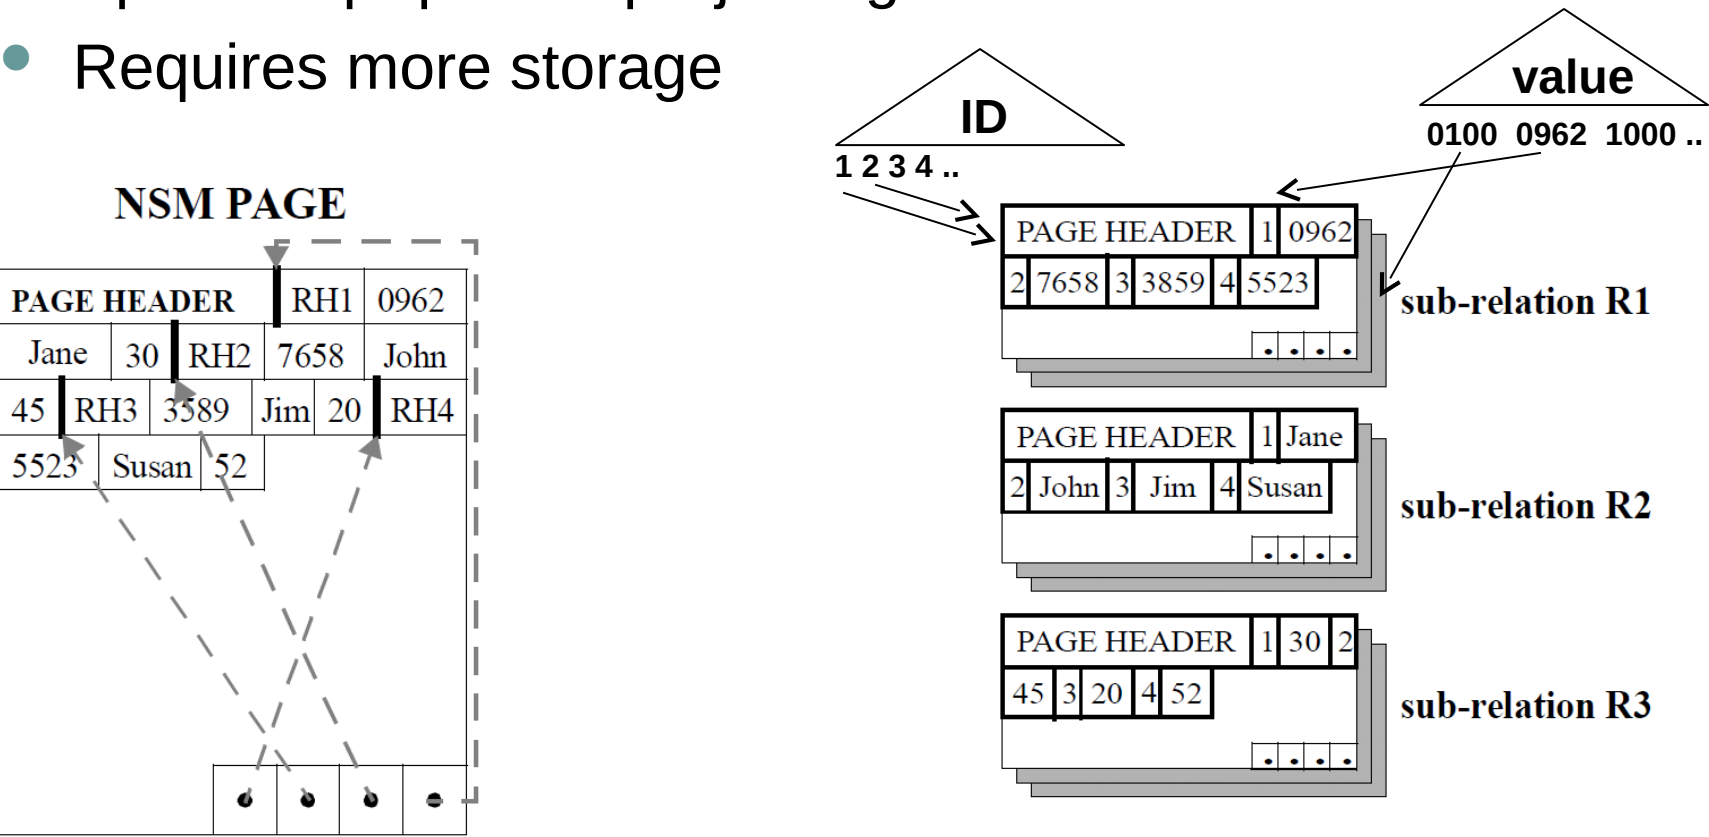
MIT: C-Store CWI: MonetDB/X100 10+ startups



The original DSM paper

“A decomposition storage model” Copeland and Khoshafian. SIGMOD 1985.

- Proposed as an alternative to NSM
- 2 indexes: clustered on ID, non-clustered on value
- Speeds up queries projecting few columns
- Requires more storage





Memory wall and PAX

- 90s: Cache-conscious research

from: “Cache Conscious Algorithms for Relational Query Processing.”
Shatdal, Kant, Naughton. VLDB 1994.

to: “Database Architecture Optimized for the New Bottleneck: Memory Access.”
Boncz, Manegold, Kersten. VLDB 1999.

and: “DBMSs on a modern processor: Where does time go?” Ailamaki, DeWitt, Hill, Wood. VLDB 1999.

- PAX: Partition Attributes Across

- Retains NSM I/O pattern
- Optimizes cache-to-RAM communication

“Weaving Relations for Cache Performance.”
Ailamaki, DeWitt, Hill, Skounakis, VLDB 2001.

PAX PAGE

PAGE HEADER				0962	7658
3859	5523				
Jane	John	Jim	Susan		
30	52	45	20		



More hybrid NSM/DSM schemes

- Dynamic PAX: Data Morphing


“Data morphing: an adaptive, cache-conscious storage technique.” Hankins, Patel, VLDB 2003.

- Clotho: custom layout using scatter-gather I/O

“Clotho: Decoupling Memory Page Layout from Storage Organization.” Shao, Schindler, Schlosser, Ailamaki, and Ganger. VLDB 2004.

- Fractured mirrors

- Smart mirroring with both NSM/DSM copies



“A Case For Fractured Mirrors.” Ramamurthy, DeWitt, Su, VLDB 2002.



MonetDB

- Late 1990s, CWI: Boncz, Manegold, and Kersten
- Motivation:
 - Main-memory
 - Improve computational efficiency by avoiding expression interpreter
 - DSM with virtual IDs natural choice
 - Developed new query execution algebra
- Initial contributions:
 - Pointed out memory-wall in DBMSs
 - Cache-conscious projections and joins
 - ...

Storing Relations in MonetDB

DSM => Column-store

contract	client	date	name	price	city	product						
12302346	10042334		Eno		Redmond	Car						
37611373	10987097		Gotz		Redmond	House						
51213123	10032423		Jones		Washington	Travel						
54535545	10087823		Smith		New York	House						
45447894	10013232		Doe		Boston	Car						
95371001	10032112		Chen		Seattle	House						

Front-End

logical
data model

mapping rules

physical
data model
(BATs)

Monet

oid	contract
1000	12302346
1001	37611373
1002	51213123
1003	54535545
1004	45447894
1005	95371001

oid	client
1000	10042334
1001	10987097
1002	10032423
1003	10087823
1004	10013232
1005	10032112

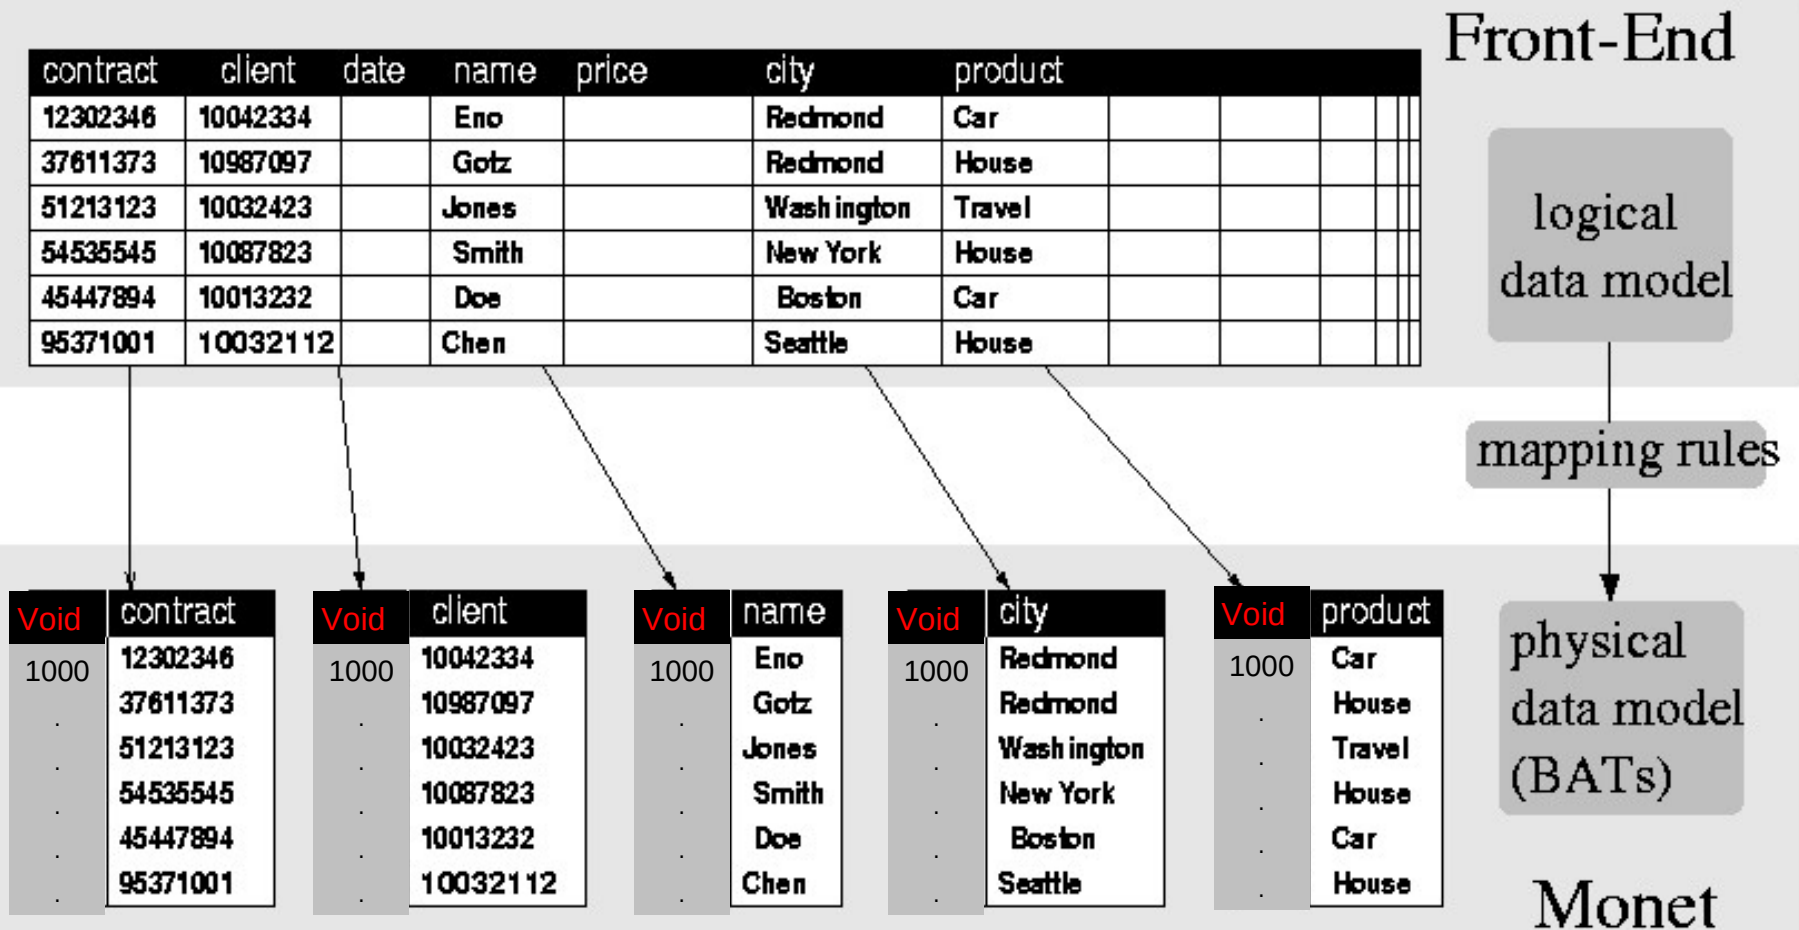
oid	name
1000	Eno
1001	Gotz
1002	Jones
1003	Smith
1004	Doe
1005	Chen

oid	city
1000	Redmond
1001	Redmond
1002	Washington
1003	New York
1004	Boston
1005	Seattle

oid	product
1000	Car
1001	House
1002	Travel
1003	House
1004	Car
1005	House

Storing Relations in MonetDB

DSM => Column-store



Virtual **OID**: seqbase=**1000** (increment=1)



2005: the (re)birth of column-stores

- New hardware and application realities
 - Faster CPUs, larger memories, disk bandwidth limit
 - Multi-terabyte Data Warehouses
- New approach: combine several techniques
 - Read-optimized, fast multi-column access, disk/CPU efficiency, light-weight compression
- C-store paper:
 - First comprehensive design description of a column-store
- MonetDB/X100
 - “proper” disk-based column store
- Explosion of new products

Column-Store History / Time-line

- 1985: **DSM** (Copeland et al.; SIGMOD 1985)
- 1992: First ideas and kernel for MonetDB (Kersten)
- 1993: MonetDB is born
- 1993: KDB (first commercial DSM system (?))
- 1995: Sybase IQ
- 2002: MonetDB goes open-source
- 2004?: Stonebraker et al. start “C-Store” project and coin DSM as “**Column-Store**”
- 2006?: Stonebraker et al. found “Vertica”; end of “C-Store” as research project
- 2008: Zukowski, Boncz, et al. (CWI) found VectorWise (based on *MonetDB/X100*)
- 2010: INGRES (now called Actian) acquires VectorWise
- 2011: HP acquires Vertica
- 201?: SAP HANA, IBM BLINK -> ISAO -> BLU, Oracle Database In-Memory
Microsoft SQL Server Column-store indexes (“Apollo”), ...
- 2018: Raasveldt & Mühleisen (CWI) start developing DuckDB as embeddable analytical DBMS (based on columnar storage & vectorized execution):
“*SQLite for analytics*”



Performance tradeoffs: columns vs. rows

DSM traditionally was not favored by technology trends
How has this changed?

- Optimized DSM in “Fractured Mirrors,” 2002
- “Apples-to-apples” comparison “Performance Tradeoffs in Read-Optimized Databases” Harizopoulos, Liang, Abadi, Madden, VLDB’06
- Follow-up study “Read-Optimized Databases, In-Depth” Holloway, DeWitt, VLDB’08
- Main-memory DSM vs. NSM “DSM vs. NSM: CPU performance tradeoffs in block-oriented query processing” Boncz, Zukowski, Nes, DaMoN’08
- Flash-disks: a come-back for PAX?
 - “Fast Scans and Joins Using Flash Drives” Shah, Harizopoulos, Wiener, Graefe. DaMoN’08
 - “Query Processing Techniques for Solid State Drives” Tsirogiannis, Harizopoulos, Shah, Wiener, Graefe, SIGMOD’09



Applications for column-stores

- Data Warehousing
 - High end (clustering)
 - Mid end/Mass Market
 - Personal Analytics
- Data Mining
 - E.g. Proximity
- Google BigTable
- RDF
 - Semantic web data management
- Information retrieval
 - Terabyte TREC
- Scientific datasets
 - SciDB initiative
 - SLOAN Digital Sky Survey on MonetDB



List of column-store systems

- Cantor (history)
- Sybase IQ
- SenSage (former Addamark Technologies)
- Kdb
- 1010data
- MonetDB
- C-Store/Vertica
- X100/VectorWise
- KickFire
- SAP Business Accelerator, SAP HANA
- Infobright
- ParAccel
- Exasol




Simulate a Column-Store inside a Row-Store

Date	Store	Product	Customer	Price
01/01	BOS	Table	Mesa	\$20
01/01	NYC	Chair	Lutz	\$13
01/01	BOS	Bed	Mudd	\$79

Option A: Vertical Partitioning

Date		Store		Product		Customer		Price	
TID	Value	TID	Value	TID	Value	TID	Value	TID	Value
1	01/01	1	BOS	1	Table	1	Mesa	1	\$20
2	01/01	2	NYC	2	Chair	2	Lutz	2	\$13
3	01/01	3	BOS	3	Bed	3	Mudd	3	\$79

Option B: Index Every Column





Simulate a Column-Store inside a Row-Store

Date	Store	Product	Customer	Price
01/01	BOS	Table	Mesa	\$20
01/01	NYC	Chair	Lutz	\$13
01/01	BOS	Bed	Mudd	\$79

Option A: Vertical Partitioning


Date		
Value	StartPos	Length
01/01	1	3

Can explicitly run-length encode date

“Teaching an Old Elephant New Tricks.”
Bruno, CIDR 2009.

Store		Product		Customer		Price	
TID	Value	TID	Value	TID	Value	TID	Value
1	BOS	1	Table	1	Mesa	1	\$20
2	NYC	2	Chair	2	Lutz	2	\$13
3	BOS	3	Bed	3	Mudd	3	\$79

Option B: Index Every Column

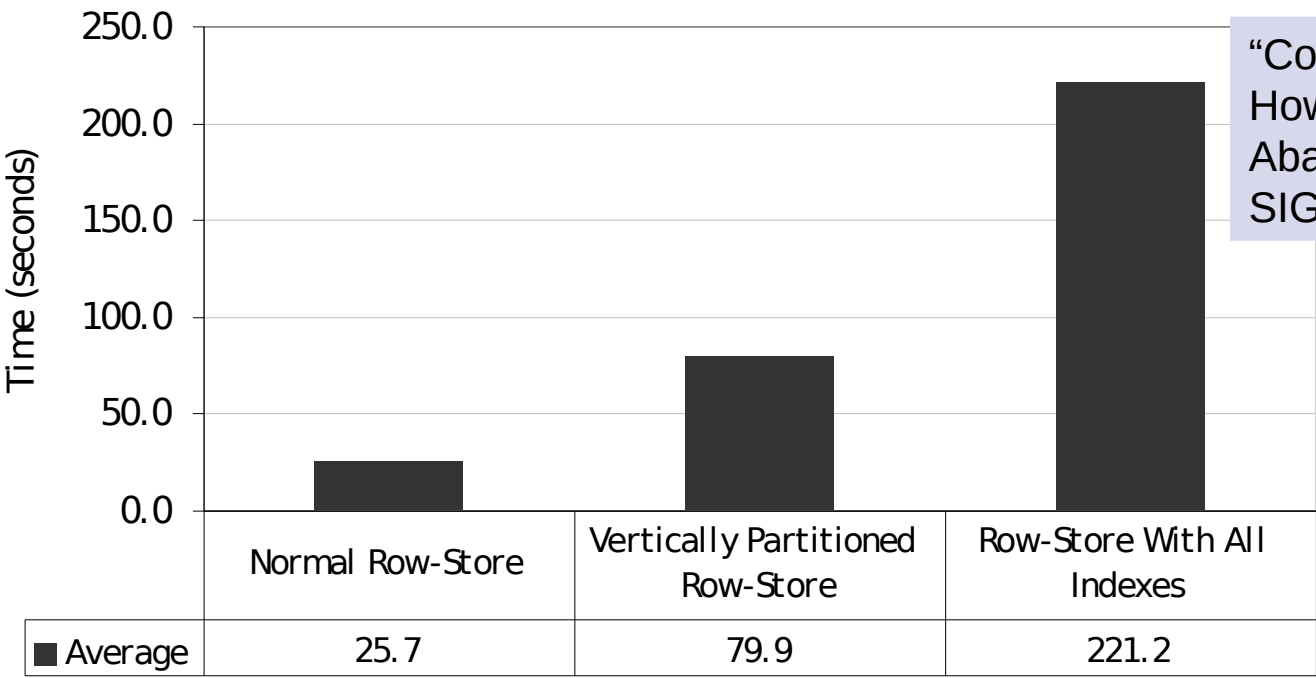




Experiments

- Star Schema Benchmark (SSBM)
- Implemented by professional DBA
- Original row-store plus 2 column-store simulations on same row-store product

Adjoined Dimension Column Index (ADC Index) to Improve Star Schema Query Performance”. O’Neil et. al. ICDE 2008.



“Column-Stores vs Row-Stores: How Different are They Really?” Abadi, Hachem, and Madden. SIGMOD 2008.



What's Going On? Vertical Partitions

- Vertical partitions in row-stores:
 - Work well when workload is known
 - ..and queries access disjoint sets of columns
 - See automated physical design
- Do not work well as full-columns
 - TupleID overhead significant
 - Excessive joins

Tuple Header	TID	Column Data
	1	
	2	
	3	

“Column-Stores vs. Row-Stores:
How Different Are They Really?”
Abadi, Madden, and Hachem.
SIGMOD 2008.

Queries touch 3-4 foreign keys in fact table, 1-2 numeric columns

Complete fact table takes up ~4 GB (compressed)

Vertically partitioned tables take up 0.7-1.1 GB (compressed)



What's Going On? All Indexes Case

- Tuple construction

- Common type of query:

```
SELECT store_name, SUM(revenue)
FROM Facts, Stores
WHERE fact.store_id = stores.store_id
      AND stores.country = "Canada"
GROUP BY store_name
```

- Result of lower part of query plan is a set of TIDs that passed all predicates
 - Need to extract SELECT attributes at these TIDs
 - BUT: index maps value to TID
 - You really want to map TID to value (i.e., a vertical partition)
- Tuple construction is SLOW



So....

- All indexes approach is a poor way to simulate a column-store
- Problems with vertical partitioning are NOT fundamental
 - Store tuple header in a separate partition
 - Allow virtual TIDs
 - Combine clustered indexes, vertical partitioning
- So can row-stores simulate column-stores?



So....

- All indexes approach is a poor way to simulate a column-store
- Problems with vertical partitioning are NOT fundamental
 - Store tuple header in a separate partition
 - Allow virtual TIDs
 - Combine clustered indexes, vertical partitioning
- So can row-stores simulate column-stores?
 - Might be possible, BUT:
 - Need better support for vertical partitioning at the storage layer
 - Need support for column-specific optimizations at the executor level
 - Full integration: buffer pool, transaction manager, ...

ADM: Literature (1/2)

- **Column-Oriented Database Systems (1/6) - Motivation & Basic Concepts**

- “An overview of cantor: a new system for data analysis”. Ilkka Karasalo, Per Svensson. SSDBM 1983.
- “A decomposition storage model”. George P. Copeland, Setrag Khoshafian. SIGMOD Conference, 1985.
- “Cache Conscious Algorithms for Relational Query Processing”. Ambuj Shatdal, Chander Kant, Jeffrey F. Naughton. VLDB 1994.
- “MIL Primitives for Querying a Fragmented World”. Peter A. Boncz, Martin L. Kersten. VLDB J. 8(2): 101-119, 1999.
- “Database Architecture Optimized for the New Bottleneck: Memory Access”. Peter A. Boncz, Stefan Manegold, Martin L. Kersten. VLDB 1999.
- “DBMSs On A Modern Processor: Where Does Time Go?”. Anastassia Ailamaki, David J. DeWitt, Mark D. Hill, David A. Wood. VLDB 1999.
- “Weaving Relations for Cache Performance ("PAX")”. Anastassia Ailamaki, David J. DeWitt, Mark D. Hill, Marios Skounakis. VLDB 2001.
- “A Case for Fractured Mirrors”. Ravishankar Ramamurthy, David J. DeWitt, Qi Su. VLDB 2002.
- “Data Morphing: An Adaptive, Cache-Conscious Storage Technique”. Richard A. Hankins, Jignesh M. Patel. VLDB 2003.
- “Clotho: Decoupling Memory Page Layout from Storage Organization”. Minglong Shao, Jiri Schindler, Steven W. Schlosser, Anastassia Ailamaki, Gregory R. Ganger. VLDB 2004.
- “MonetDB-X100 - A DBMS In The CPU Cache”. Marcin Zukowski, Peter A. Boncz, Niels Nes, Sándor Héman. IEEE Data Eng. Bull. 28(2): 17-22, 2005.
- ““One size fits all”: an idea whose time has come and gone”. Michael Stonebraker, Ugur Çetintemel. ICDE 2005.
- “Performance Tradeoffs in Read-Optimized Databases”. Stavros Harizopoulos, Velen Liang, Daniel J. Abadi, Samuel Madden. VLDB 2006.
- ...

ADM: Literature (2/2)

- **Column-Oriented Database Systems (1/6) - Motivation & Basic Concepts** (*cont.*)
 - ...
 - “One Size Fits All? - Part 2: Benchmarking Results”. Michael Stonebraker, Chuck Bear, Ugur Çetintemel, Mitch Cherniack, Tingjian Ge, Nabil Hachem, Stavros Harizopoulos, John Lifter, Jennie Rogers, Stanley B. Zdonik. CIDR 2007.
 - “C-Store: A Column-oriented DBMS”. Michael Stonebraker, Daniel J. Abadi, Adam Batkin, Xuedong Chen, Mitch Cherniack, Miguel Ferreira, Edmond Lau, Amerson Lin, Samuel Madden, Elizabeth J. O’Neil, Patrick E. O’Neil, Alex Rasin, Nga Tran, Stanley B. Zdonik. VLDB 2005.
 - “Breaking the memory wall in MonetDB”. Peter A. Boncz, Martin L. Kersten, Stefan Manegold. Commun. ACM 51(12): 77-85, 2008.
 - “Column-Stores vs Row-Stores: How Different are They Really?”. Daniel J. Abadi, Samuel Madden, Nabil Hachem. SIGMOD Conference 2008.
 - “DSM vs. NSM: CPU performance tradeoffs in block-oriented query processing”. Marcin Zukowski, Niels Nes, Peter A. Boncz. DaMoN 2008.
 - “Fast Scans and Joins Using Flash Drives”. Mehul A. Shah, Stavros Harizopoulos, Janet L. Wiener, Goetz Graefe. DaMoN 2008.
 - “Read-Optimized Databases, In-Depth”. Allison L. Holloway, David J. DeWitt. Proc. VLDB Endow. 1(1): 502-513, 2008.
 - “Teaching an Old Elephant New Tricks”. Nicolas Bruno. CIDR 2009.
 - “Query Processing Techniques for Solid State Drives”. Dimitris Tsirogiannis, Stavros Harizopoulos, Mehul A. Shah, Janet L. Wiener, Goetz Graefe. SIGMOD Conference 2009.
 - “MonetDB: Two Decades of Research in Column-oriented Database Architectures”. Stratos Idreos, Fabian Groffen, Niels Nes, Stefan Manegold, K. Sjoerd Mullender, Martin L. Kersten. IEEE Data Eng. Bull. 35(1): 40-45, 2012.
 - “The Vertica Analytic Database: C-Store 7 Years Later”. Andrew Lamb, Matt Fuller, Ramakrishna Varadarajan, Nga Tran, Ben Vandiver, Lyric Doshi, Chuck Bear. Proc. VLDB Endow. 5(12): 1790-1801, 2012.

ADM: Agenda

- 07.09.2022: Lecture 1: **Introduction**
- 14.09.2022: Lecture 2: **SQL Recap**
(plus Assignment 1 [in groups; 3 weeks]: TPC-H benchmark)
- 21.09.2022: Lecture 3: **Column-Oriented Database Systems (1/6) - Motivation & Basic Concepts**
- 28.09.2022: Lecture 4: **Column-Oriented Database Systems (2a/6) - Selected Execution Techniques (1/2)**
- 05.10.2022: Lecture 5: **Column-Oriented Database Systems (2b/6) - Selected Execution Techniques (2/2)**
(plus Assignment 3 [in groups; 3 weeks]: Compression techniques)
- 12.10.2022: Lecture 6: **Column-Oriented Database Systems (3/6) - Cache Conscious Joins**
- 19.10.2022: Lecture 7: **Column-Oriented Database Systems (4/6) - “Vectorized Execution”**
- 26.10.2022: **No lecture!**
- 02.11.2022: Lecture 8: **DuckDB: An embedded database for data science (1/2) (guest lecture & hands-on)**
(plus Assignment 2 [individual; 2 weeks]: Analysing NYC Cab dataset with DuckDB)
- 09.11.2022: Lecture 9: **DuckDB: An embedded database for data science (2/2) (guest lecture & hands-on)**
- 16.11.2022: Lecture 10: **Branch Misprediction & Predication**
(plus Assignment 4 [individual; 2 weeks]: Predication)
- 23.11.2022: Lecture 11: **Column-Oriented Database Systems (5/6) - Adaptive Indexing**
- 30.11.2022: Lecture 12: **Column-Oriented Database Systems (6/6) - Progressive Indexing**

ADM: Literature

- **Column-Oriented Database Systems (2/6) - Selected Execution Techniques**
 - Compression
 - “Compressing Relations and Indexes”. Goldstein, Ramakrishnan, Shaft. ICDE’98.
 - “Query optimization in compressed database systems”. Chen, Gehrke, Korn. SIGMOD’01.
 - “Super-Scalar RAM-CPU Cache Compression”. Zukowski, Heman, Nes, Boncz. ICDE’06.
 - “Integrating Compression and Execution in Column-Oriented Database Systems”. Abadi, Madden, Ferreira. SIGMOD’06.
 - “Improved Word-Aligned Binary Compression for Text Indexing”. Ahn, Moffat. TKDE’06.
 - Tuple Materialization
 - “Materialization Strategies in a Column-Oriented DBMS”. Abadi, Myers, DeWitt, Madden. ICDE’07.
 - “Column-Stores vs Row-Stores: How Different are They Really?”. Abadi, Madden, Hachem. SIGMOD’08.
 - “Query Processing Techniques for Solid State Drives”. Tsirogiannis, Harizopoulos Shah, Wiener, Graefe. SIGMOD’09.
 - “Self-organizing tuple reconstruction in column-stores”. Idreos, Manegold, Kersten. SIGMOD’09.
 - Join
 - “Fast Joins using Join Indices”. Li and Ross. VLDBJ 8:1-24, 1999.