

# **Advanced Data Management for Data Analysis**

*Stefan Manegold*

Data Management @ LIACS

Group leader Database Architectures  
Centrum Wiskunde & Informatica (CWI)  
Amsterdam

[s.manegold@liacs.leidenuniv.nl](mailto:s.manegold@liacs.leidenuniv.nl)

<http://www.cwi.nl/~manegold/>

# ADM: Agenda

- 07.09.2022: Lecture 1: **Introduction**
- 14.09.2022: Lecture 2: **SQL Recap**  
(plus Assignment 1 [in groups; 3 weeks]: TPC-H benchmark)
- 21.09.2022: Lecture 3: **Column-Oriented Database Systems (1/6) - Motivation & Basic Concepts**
- 28.09.2022: Lecture 4: **Column-Oriented Database Systems (2a/6) - Selected Execution Techniques (1/2)**
- 05.10.2022: Lecture 5: **Column-Oriented Database Systems (2b/6) - Selected Execution Techniques (2/2)**  
(plus Assignment 2 [in groups; 3 weeks]: Compression techniques)
- 12.10.2022: Lecture 6: **Column-Oriented Database Systems (3/6) - Cache Conscious Joins**
- 19.10.2022: Lecture 7: **Column-Oriented Database Systems (4/6) - “Vectorized Execution”**
- 26.10.2022: **No lecture!**
- 02.11.2022: Lecture 8: **DuckDB: An embedded database for data science (1/2)** (guest lecture & **hands-on**)  
(plus Assignment 3 [individual; 2 weeks]: Analysing NYC Cab dataset with DuckDB)
- 09.11.2022: Lecture 9: **DuckDB: An embedded database for data science (2/2)** (guest lecture & **hands-on**)
- 16.11.2022: Lecture 10: **Branch Misprediction & Predication**  
(plus Assignment 4 [individual; 2 weeks]: Predication)
- 23.11.2022: Lecture 11: **Column-Oriented Database Systems (5/6) - Adaptive Indexing**
- 30.11.2022: Lecture 12: **Column-Oriented Database Systems (6/6) - Progressive Indexing**

# ADM: Literature

- **Column-Oriented Database Systems (2/6) - Selected Execution Techniques**
  - Compression
    - “Compressing Relations and Indexes”. Goldstein, Ramakrishnan, Shaft. ICDE’98.
    - “Query optimization in compressed database systems”. Chen, Gehrke, Korn. SIGMOD’01.
    - “Super-Scalar RAM-CPU Cache Compression”. Zukowski, Heman, Nes, Boncz. ICDE’06.
    - “Integrating Compression and Execution in Column-Oriented Database Systems”. Abadi, Madden, Ferreira. SIGMOD’06.
    - “Improved Word-Aligned Binary Compression for Text Indexing”. Ahn, Moffat. TKDE’06.
  - Tuple Materialization
    - “Materialization Strategies in a Column-Oriented DBMS”. Abadi, Myers, DeWitt, Madden. ICDE’07.
    - “Column-Stores vs Row-Stores: How Different are They Really?”. Abadi, Madden, Hachem. SIGMOD’08.
    - “Query Processing Techniques for Solid State Drives”. Tsirogiannis, Harizopoulos Shah, Wiener, Graefe. SIGMOD’09.
    - “Self-organizing tuple reconstruction in column-stores”. Idreos, Manegold, Kersten. SIGMOD’09.
  - Join
    - “Fast Joins using Join Indices”. Li and Ross. VLDBJ 8:1-24, 1999.

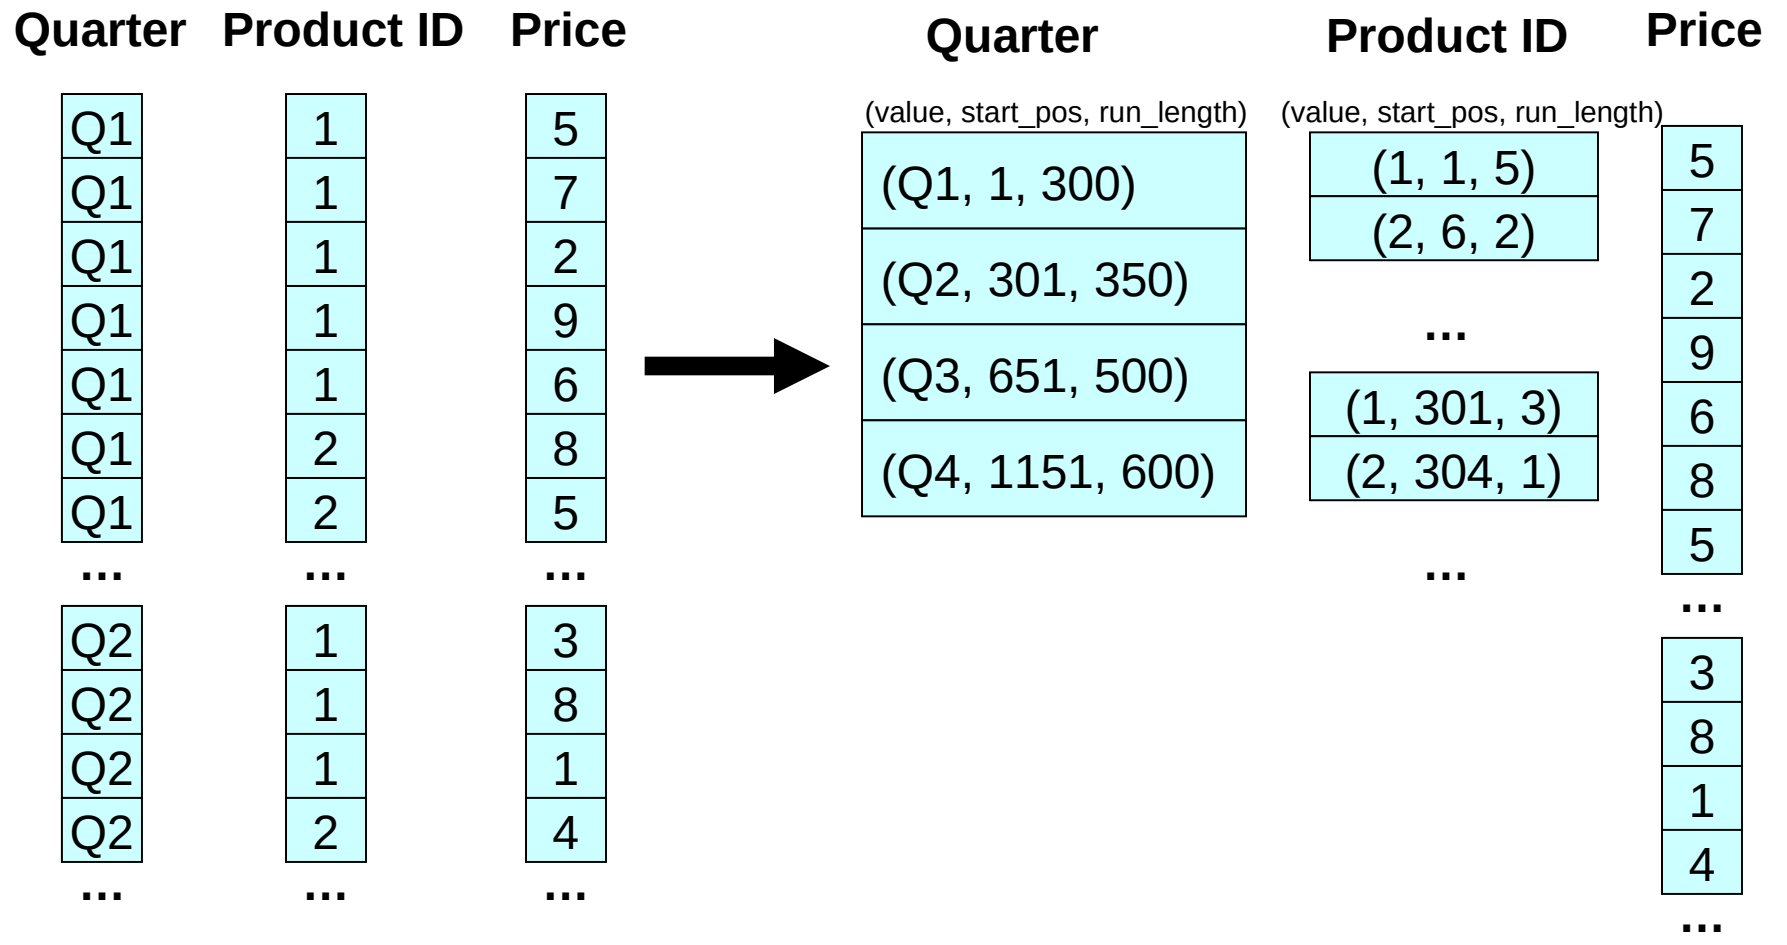


# Compression

- **Trades I/O for CPU**
  - Smaller data size on storage  
=> less data to read (or write)
  - Extra computation (CPU work) required to (de-)compress
- **Increased column-store opportunities:**
  - Higher data value locality in column stores
  - Techniques such as run length encoding far more useful
  - Can use extra space to store multiple copies of data in different sort orders
- **Sample “light-weight” encoding techniques:**
  - Run-length, bit-vector, dictionary, frame-of-reference, differential



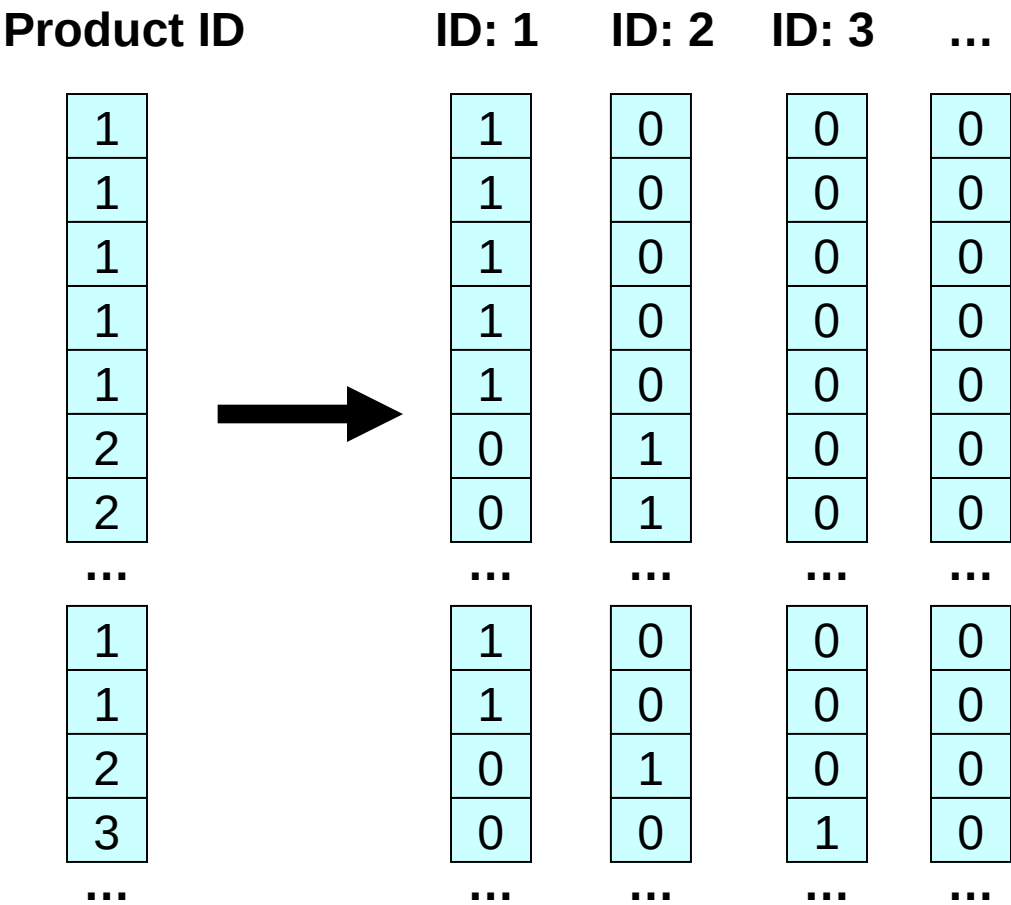
# Run-length Encoding (RLE)





# Bit-vector Encoding

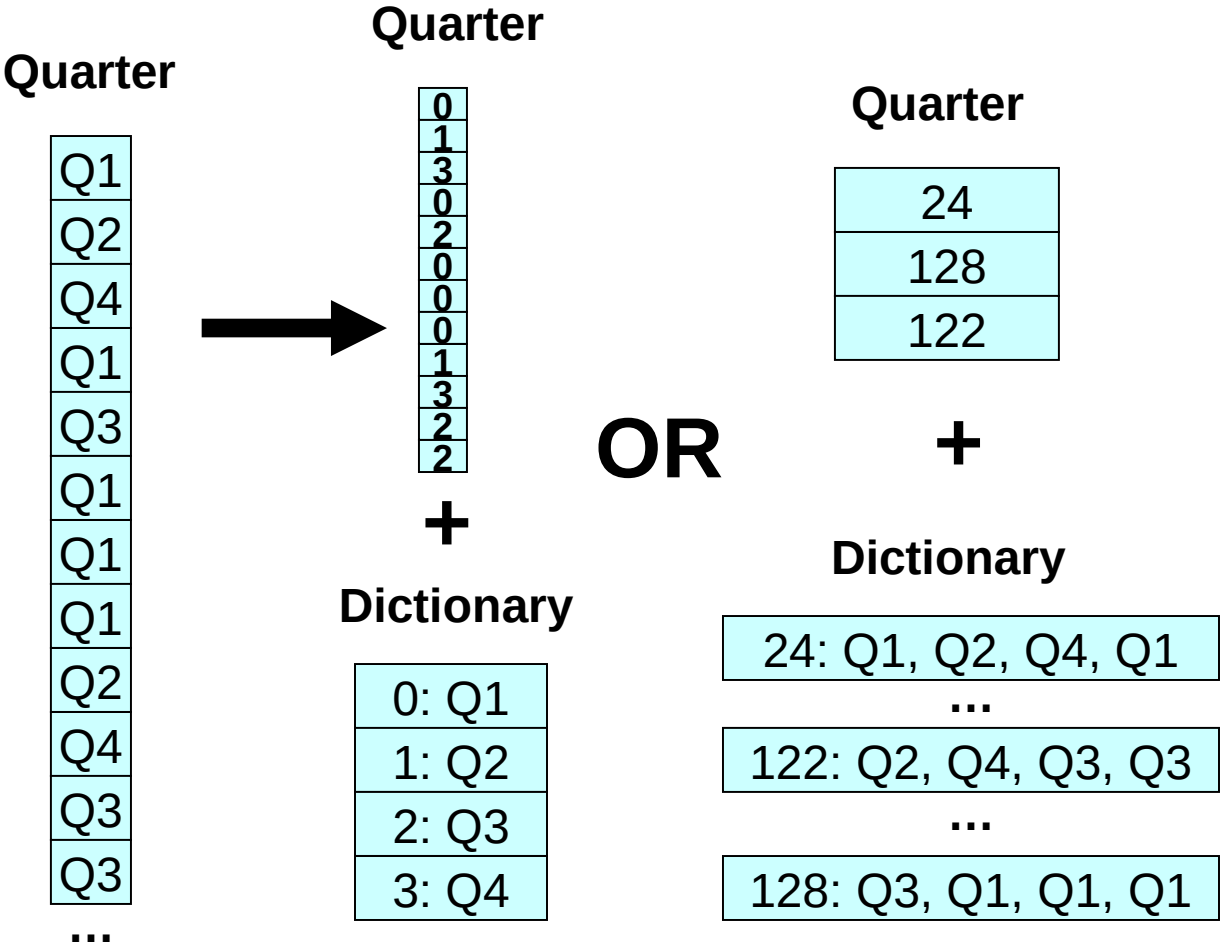
- For each unique value,  $v$ , in column  $c$ , create bit-vector  $b_v$ 
  - $b_v[i] = 1$  if  $c[i] = v$
- Good for columns with few unique values
- Each bit-vector can be further compressed if sparse





# Dictionary Encoding (DIC)

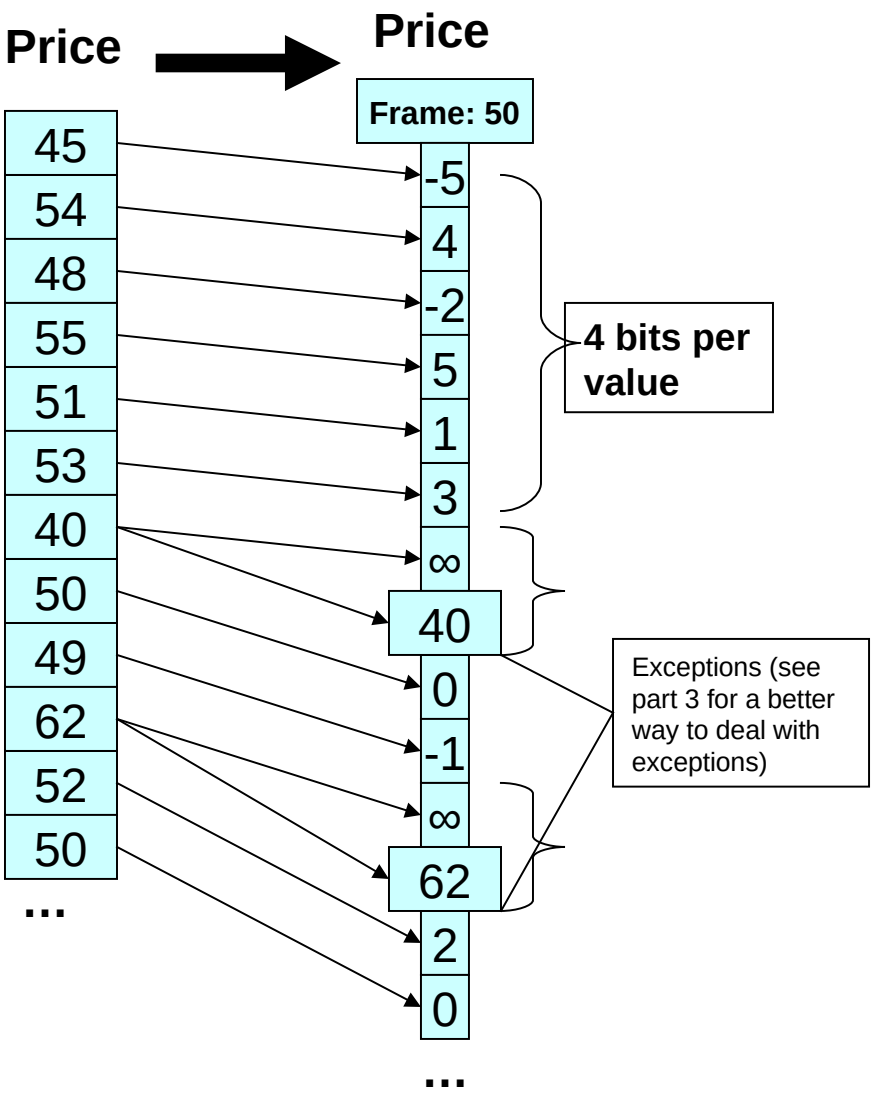
- For each unique value create dictionary entry
- Dictionary can be per-block or per-column
- Column-stores have the advantage that dictionary entries may encode multiple values at once





# Frame Of Reference Encoding (FOR)

- Encodes values as b bit offset from chosen frame of reference
- Special escape code (e.g. all bits set to 1) indicates a difference larger than can be stored in b bits
  - After escape code, original (uncompressed) value is written



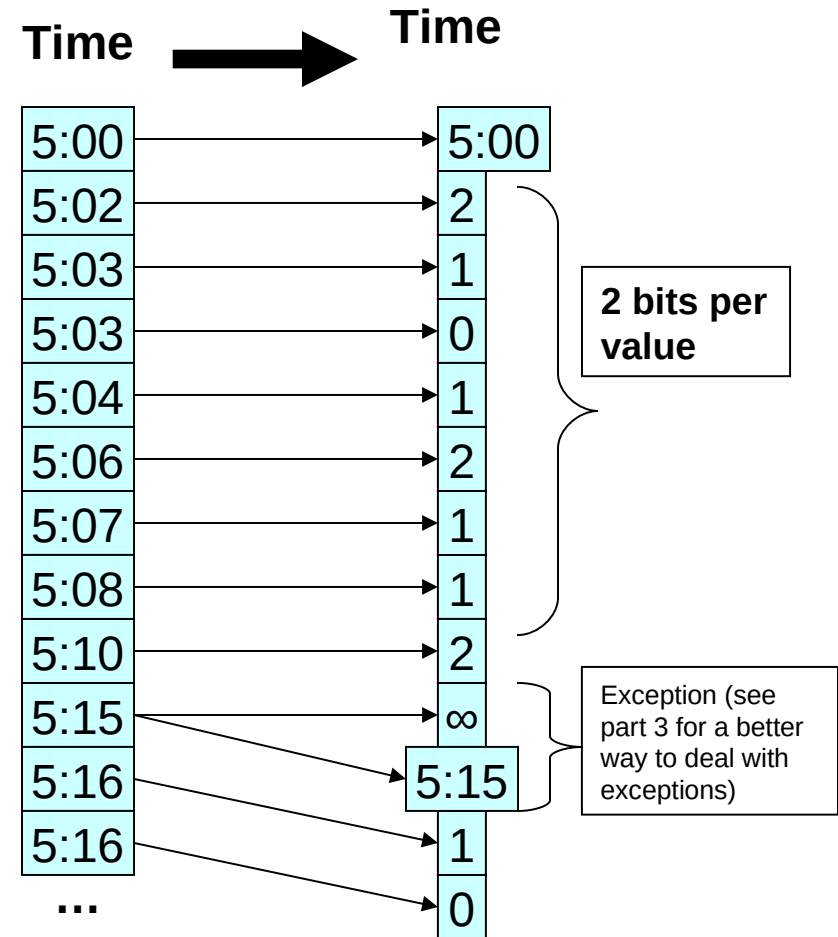
“Compressing Relations and Indexes”  
Goldstein, Ramakrishnan, Shaft,  
ICDE’98





# Differential Encoding (DIF)

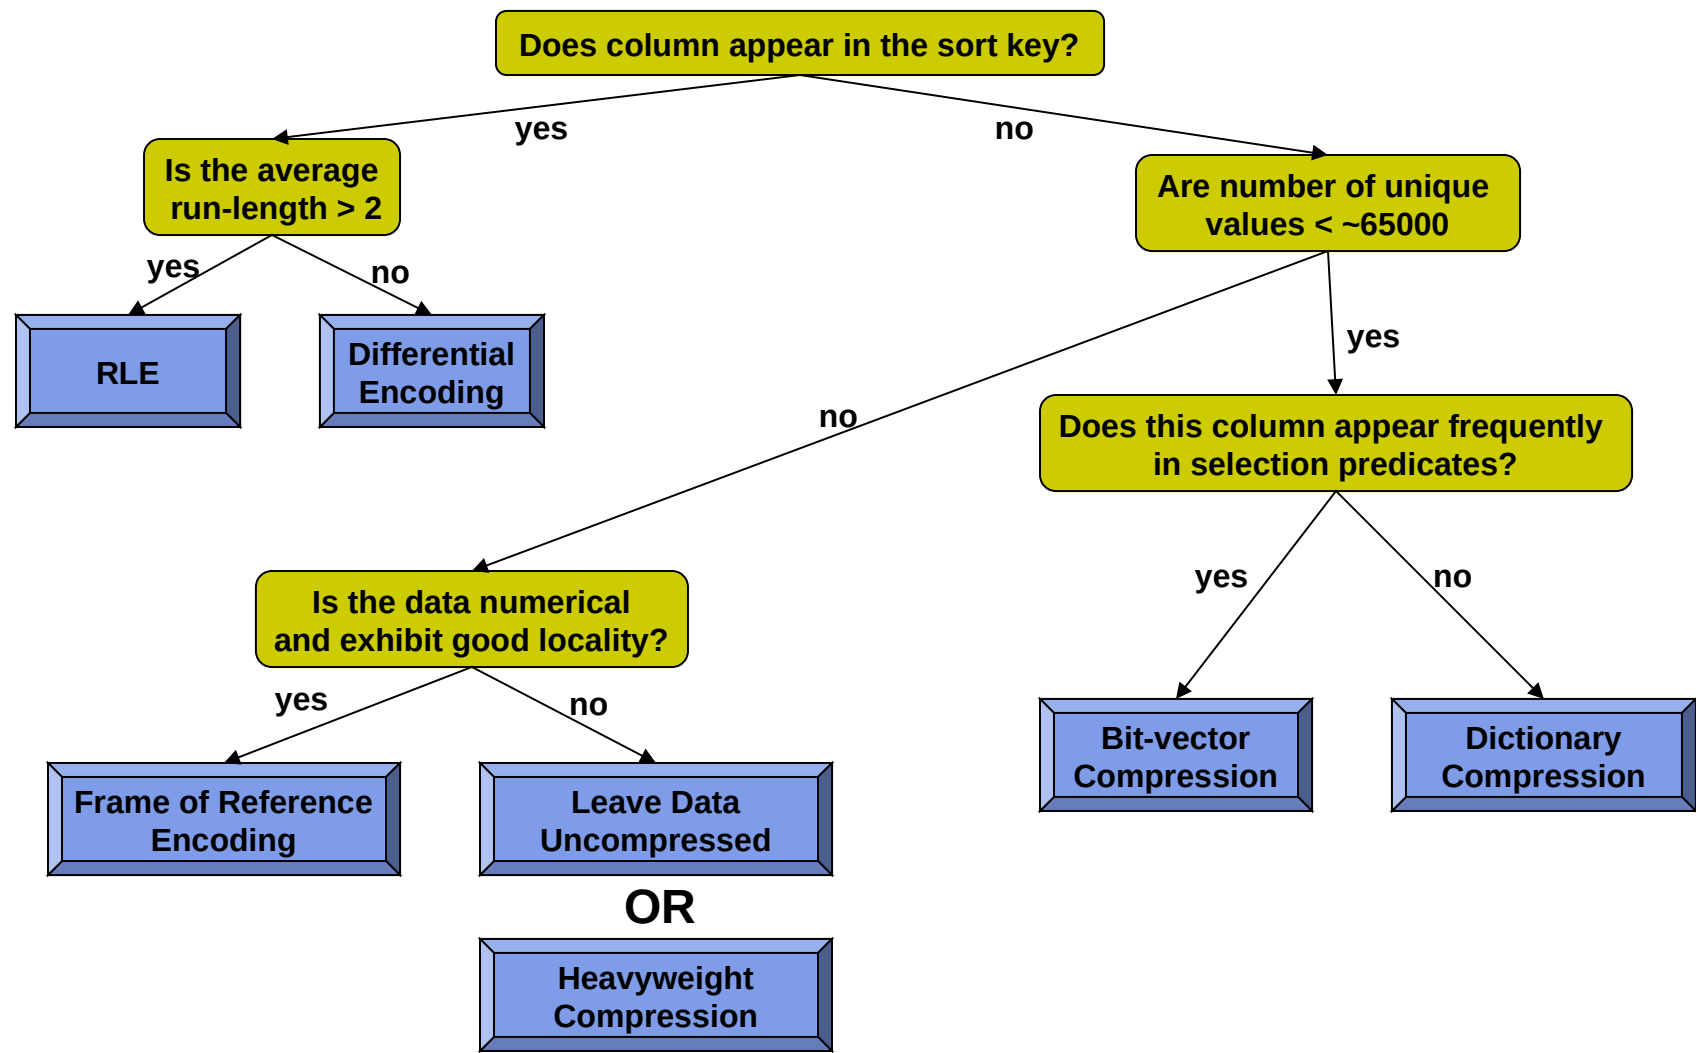
- Encodes values as  $b$  bit offset from previous value
- Special escape code (just like frame of reference encoding) indicates a difference larger than can be stored in  $b$  bits
  - After escape code, original (uncompressed) value is written
- Performs well on columns containing increasing/decreasing sequences
  - inverted lists
  - timestamps
  - object IDs
  - sorted / clustered columns



“Improved Word-Aligned Binary Compression for Text Indexing”  
Ahn, Moffat, TKDE’06



# What Compression Scheme To Use?





# Heavy-Weight Compression Schemes

Algorithm	Decompression Bandwidth
BZIP	10 MB/s
ZLIB	80 MB/s
LZO	300 MB/s

- Modern disk arrays can achieve > 1GB/s
- 1/3 CPU for decompression ➔ 3GB/s needed



# Heavy-Weight Compression Schemes

Algorithm	Decompression Bandwidth
BZIP	10 MB/s
ZLIB	80 MB/s
LZO	300 MB/s

- Modern disk arrays can achieve > 1GB/s
- 1/3 CPU for decompression → 3GB/s needed
- Lightweight compression schemes are better
- Even better: operate directly on compressed data

“Integrating Compression and Execution in Column-Oriented Database Systems” Abadi et. al, SIGMOD '06



# Operating Directly on Compressed Data

- I/O - CPU tradeoff is no longer a tradeoff
- Reduces memory–CPU bandwidth requirements
- Opens up possibility of operating on multiple records at once

## Examples

- $\text{SUM}_i(\text{rle-compressed column}[i]) \rightarrow \text{SUM}_g(\text{count}[g] * \text{value}[g])$
- $(\text{country} == \text{"Asia"}) \rightarrow \text{countryCode} == 6$   

strcmpSIMD



# Operating Directly on Compressed Data

Quarter	Product ID		
	1	2	3
	...	...	...
(Q1, 1, 300)	1	0	0
	0	0	1
(Q2, 301, 6)	0	1	0
	1	0	0
(Q3, 307, 500)	1	0	0
	0	0	1
(Q4, 807, 600)	0	1	0
	1	0	0
	0	0	1
	0	1	0
	0	0	1
	...	...	...

```
SELECT ProductID, Count(*)
FROM table

WHERE Quarter = Q2

GROUP BY ProductID
```



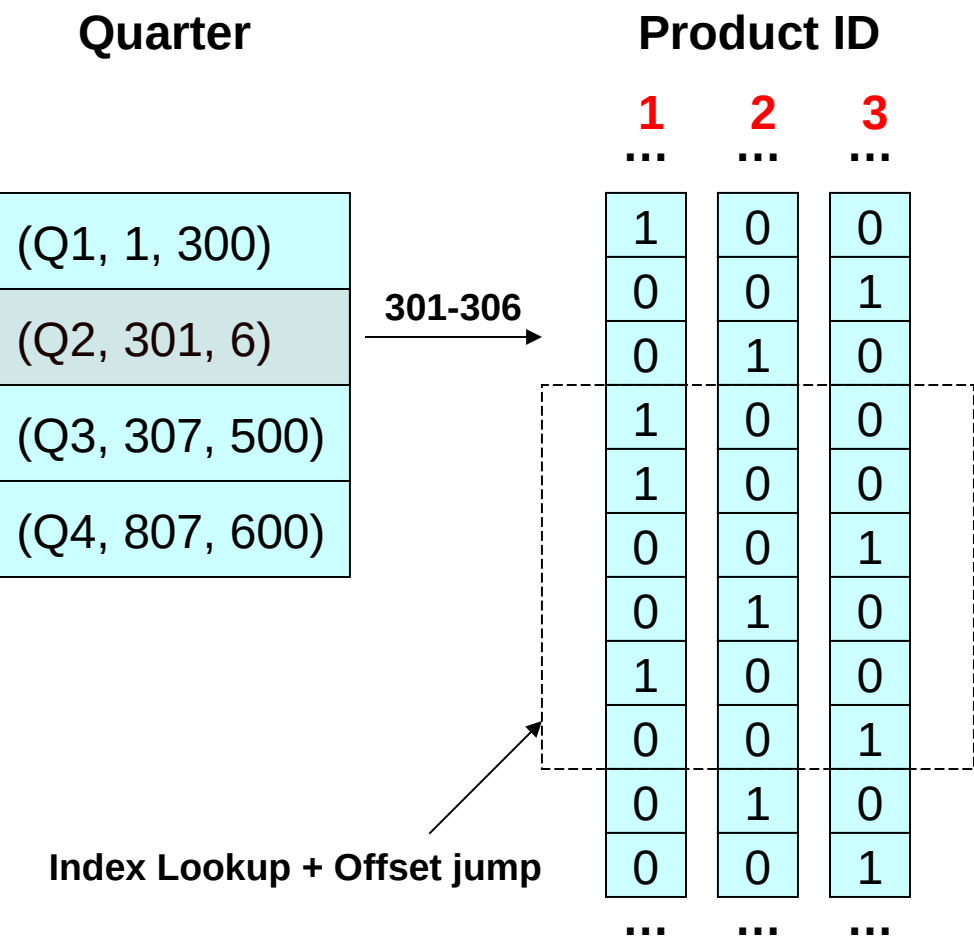
# Operating Directly on Compressed Data

Quarter	Product ID		
	1 ...	2 ...	3 ...
(Q1, 1, 300)	1	0	0
(Q2, 301, 6)	0	0	1
	0	1	0
(Q3, 307, 500)	1	0	0
	1	0	0
(Q4, 807, 600)	0	0	1
	0	1	0
	1	0	0
	0	0	1
	0	1	0
	0	0	1
	...	...	...

```
SELECT ProductID, Count(*)  
FROM table  
  
WHERE Quarter = Q2  
  
GROUP BY ProductID
```



# Operating Directly on Compressed Data

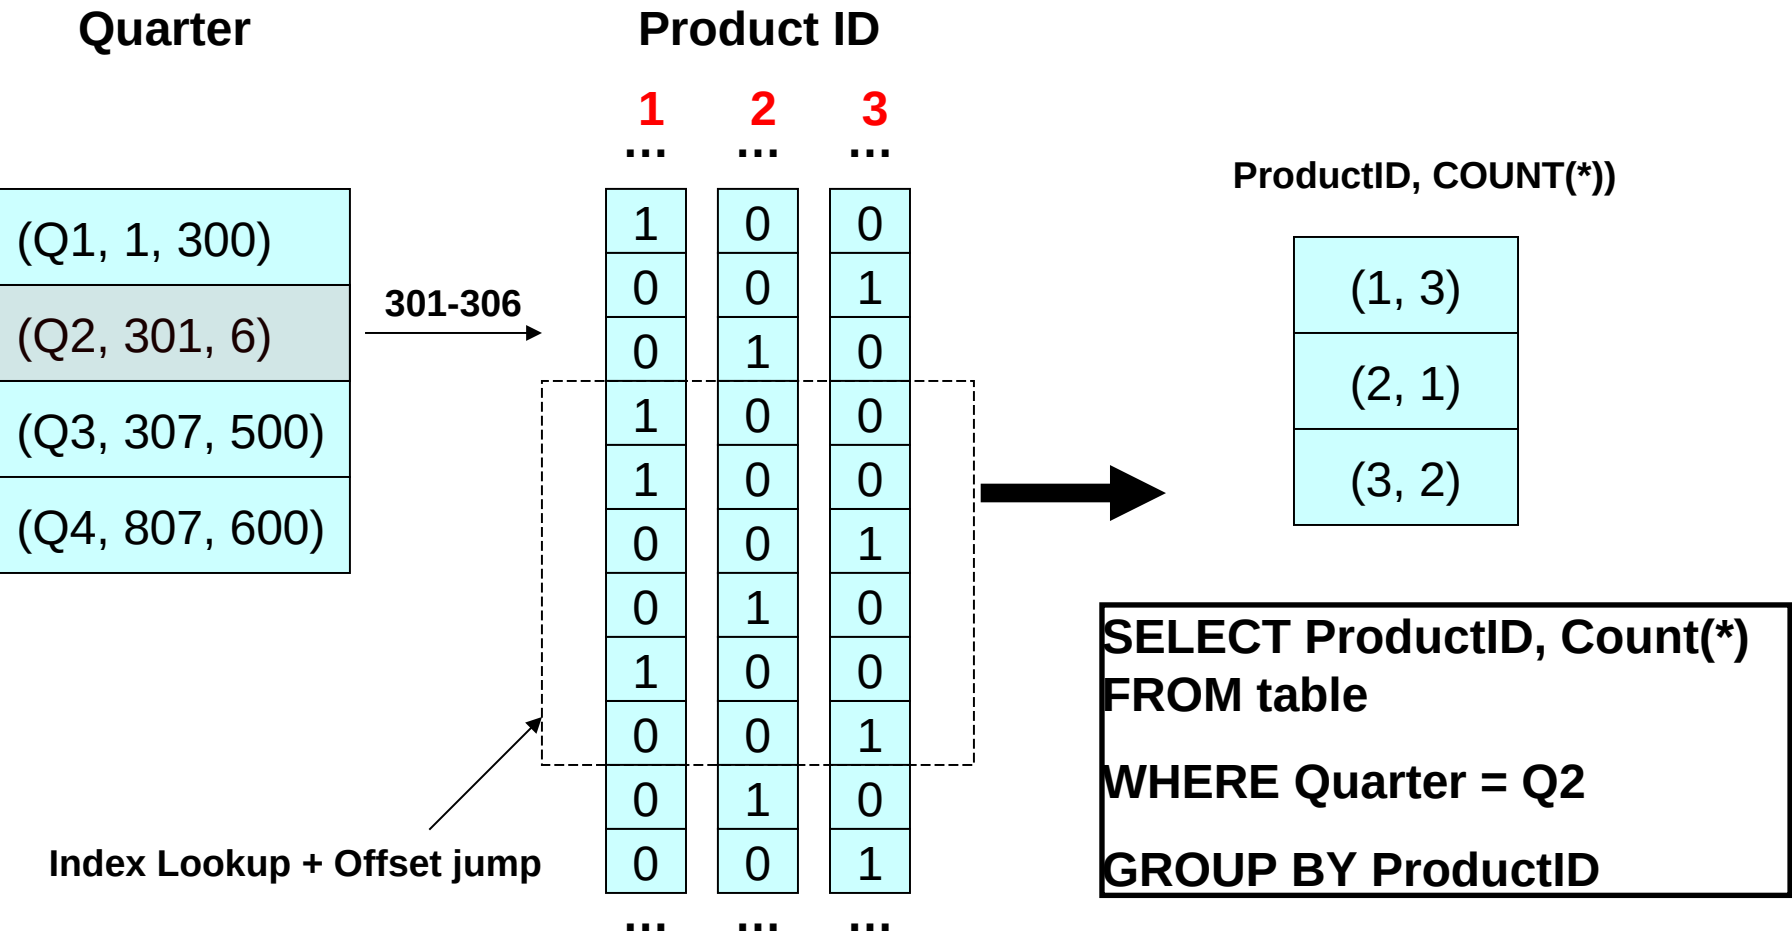


```
SELECT ProductID, Count(*)
FROM table
WHERE Quarter = Q2
GROUP BY ProductID
```





# Operating Directly on Compressed Data



# ADM: Literature

- **Column-Oriented Database Systems (2/6) - Selected Execution Techniques**
  - Compression
    - “Compressing Relations and Indexes”. Goldstein, Ramakrishnan, Shaft. ICDE’98.
    - “Query optimization in compressed database systems”. Chen, Gehrke, Korn. SIGMOD’01.
    - “Super-Scalar RAM-CPU Cache Compression”. Zukowski, Heman, Nes, Boncz. ICDE’06.
    - “Integrating Compression and Execution in Column-Oriented Database Systems”. Abadi, Madden, Ferreira. SIGMOD’06.
    - “Improved Word-Aligned Binary Compression for Text Indexing”. Ahn, Moffat. TKDE’06.
  - Tuple Materialization
    - “Materialization Strategies in a Column-Oriented DBMS”. Abadi, Myers, DeWitt, Madden. ICDE’07.
    - “Column-Stores vs Row-Stores: How Different are They Really?”. Abadi, Madden, Hachem. SIGMOD’08.
    - “Query Processing Techniques for Solid State Drives”. Tsirogiannis, Harizopoulos Shah, Wiener, Graefe. SIGMOD’09.
    - “Self-organizing tuple reconstruction in column-stores”. Idreos, Manegold, Kersten. SIGMOD’09.
  - Join
    - “Fast Joins using Join Indices”. Li and Ross. VLDBJ 8:1-24, 1999.

# ADM: Agenda

- 07.09.2022: Lecture 1: **Introduction**
- 14.09.2022: Lecture 2: **SQL Recap**  
(plus Assignment 1 [in groups; 3 weeks]: TPC-H benchmark)
- 21.09.2022: Lecture 3: **Column-Oriented Database Systems (1/6) - Motivation & Basic Concepts**
- 28.09.2022: Lecture 4: **Column-Oriented Database Systems (2a/6) - Selected Execution Techniques (1/2)**
- 05.10.2022: Lecture 5: **Column-Oriented Database Systems (2b/6) - Selected Execution Techniques (2/2)**  
(plus Assignment 2 [in groups; 3 weeks]: Compression techniques)
- 12.10.2022: Lecture 6: **Column-Oriented Database Systems (3/6) - Cache Conscious Joins**
- 19.10.2022: Lecture 7: **Column-Oriented Database Systems (4/6) - “Vectorized Execution”**
- 26.10.2022: **No lecture!**
- 02.11.2022: Lecture 8: **DuckDB: An embedded database for data science (1/2) (guest lecture & hands-on)**  
(plus Assignment 3 [individual; 2 weeks]: Analysing NYC Cab dataset with DuckDB)
- 09.11.2022: Lecture 9: **DuckDB: An embedded database for data science (2/2) (guest lecture & hands-on)**
- 16.11.2022: Lecture 10: **Branch Misprediction & Predication**  
(plus Assignment 4 [individual; 2 weeks]: Predication)
- 23.11.2022: Lecture 11: **Column-Oriented Database Systems (5/6) - Adaptive Indexing**
- 30.11.2022: Lecture 12: **Column-Oriented Database Systems (6/6) - Progressive Indexing**

# ADM: Literature

- **Column-Oriented Database Systems (2/6) - Selected Execution Techniques**

- Compression

- “Compressing Relations and Indexes”. Goldstein, Ramakrishnan, Shaft. ICDE’98.
- “Query optimization in compressed database systems”. Chen, Gehrke, Korn. SIGMOD’01.
- “Super-Scalar RAM-CPU Cache Compression”. Zukowski, Heman, Nes, Boncz. ICDE’06.
- “Integrating Compression and Execution in Column-Oriented Database Systems”. Abadi, Madden, Ferreira. SIGMOD’06.
- “Improved Word-Aligned Binary Compression for Text Indexing”. Ahn, Moffat. TKDE’06.

- Tuple Materialization

- “Materialization Strategies in a Column-Oriented DBMS”. Abadi, Myers, DeWitt, Madden. ICDE’07.
- “Column-Stores vs Row-Stores: How Different are They Really?”. Abadi, Madden, Hachem. SIGMOD’08.
- “Query Processing Techniques for Solid State Drives”. Tsirogiannis, Harizopoulos Shah, Wiener, Graefe. SIGMOD’09.
- “Self-organizing tuple reconstruction in column-stores”. Idreos, Manegold, Kersten. SIGMOD’09.

- Join

- “Fast Joins using Join Indices”. Li and Ross. VLDBJ 8:1-24, 1999.