

Advanced Data Management for Data Analysis

Stefan Manegold

Data Management @ LIACS

Group leader Database Architectures
Centrum Wiskunde & Informatica (CWI)
Amsterdam

s.manegold@liacs.leidenuniv.nl

<http://www.cwi.nl/~manegold/>

ADM: Agenda

- 07.09.2022: Lecture 1: **Introduction**
- 14.09.2022: Lecture 2: **SQL Recap**
(plus Assignment 1 [in groups; 3 weeks]: TPC-H benchmark)
- 21.09.2022: Lecture 3: **Column-Oriented Database Systems (1/6) - Motivation & Basic Concepts**
- 28.09.2022: Lecture 4: **Column-Oriented Database Systems (2a/6) - Selected Execution Techniques (1/2)**
- 05.10.2022: Lecture 5: **Column-Oriented Database Systems (2b/6) - Selected Execution Techniques (2/2)**
(plus Assignment 2 [in groups; 3 weeks]: Compression techniques)
- 12.10.2022: Lecture 6: **Column-Oriented Database Systems (3/6) - Cache Conscious Joins**
- 19.10.2022: Lecture 7: **Column-Oriented Database Systems (4/6) - “Vectorized Execution”**
- 26.10.2022: **No lecture!**
- 02.11.2022: Lecture 8: **DuckDB: An embedded database for data science (1/2) (guest lecture & hands-on)**
(plus Assignment 3 [individual; 2 weeks]: Analysing NYC Cab dataset with DuckDB)
- 09.11.2022: Lecture 9: **DuckDB: An embedded database for data science (2/2) (guest lecture & hands-on)**
- 16.11.2022: Lecture 10: **Branch Misprediction & Predication**
(plus Assignment 4 [individual; 2 weeks]: Predication)
- 23.11.2022: Lecture 11: **Column-Oriented Database Systems (5/6) - Adaptive Indexing**
- 30.11.2022: Lecture 12: **Column-Oriented Database Systems (6/6) - Progressive Indexing**

ADM: Literature

- **Column-Oriented Database Systems (2/6) - Selected Execution Techniques**
 - Compression
 - “Compressing Relations and Indexes”. Goldstein, Ramakrishnan, Shaft. ICDE’98.
 - “Query optimization in compressed database systems”. Chen, Gehrke, Korn. SIGMOD’01.
 - “Super-Scalar RAM-CPU Cache Compression”. Zukowski, Heman, Nes, Boncz. ICDE’06.
 - “Integrating Compression and Execution in Column-Oriented Database Systems”. Abadi, Madden, Ferreira. SIGMOD’06.
 - “Improved Word-Aligned Binary Compression for Text Indexing”. Ahn, Moffat. TKDE’06.
 - Tuple Materialization
 - “Materialization Strategies in a Column-Oriented DBMS”. Abadi, Myers, DeWitt, Madden. ICDE’07.
 - “Column-Stores vs Row-Stores: How Different are They Really?”. Abadi, Madden, Hachem. SIGMOD’08.
 - “Query Processing Techniques for Solid State Drives”. Tsirogiannis, Harizopoulos Shah, Wiener, Graefe. SIGMOD’09.
 - “Self-organizing tuple reconstruction in column-stores”. Idreos, Manegold, Kersten. SIGMOD’09.
 - Join
 - “Fast Joins using Join Indices”. Li and Ross. VLDBJ 8:1-24, 1999.




When should columns be projected?

- Where should column projection operators be placed in a query plan?
 - Row-store:
 - Column projection involves removing unneeded columns from tuples
 - Generally done as early as possible
 - Column-store:
 - Operation is almost completely opposite from a row-store
 - Column projection involves reading needed columns from storage and extracting values for a listed set of tuples
 - This process is called “materialization”
 - Early materialization (EM): project columns at beginning of query plan
 - Straightforward since there is a one-to-one mapping across columns
 - Late materialization (LM): wait as long as possible for projecting columns
 - More complicated since selection and join operators on one column obfuscates mapping to other columns from same table
 - Most column-stores construct tuples at column projection time
 - Many database interfaces expect output in regular tuples (rows)
 - Rest of discussion will focus on this case



When should tuples be constructed?

Solution 1: Create Rows first (EM)




```
QUERY:
SELECT custID,SUM(price)
FROM table
WHERE (prodID = 4) AND
      (storeID = 1)
GROUP BY custID
```

- But:
 - Need to construct ALL tuples
 - Need to decompress data
 - Poor memory bandwidth utilization



When should tuples be constructed?

Solution 2: Operate on columns



QUERY:

SELECT custID,SUM(price)

FROM table

WHERE (prodID = 4) AND

(storeID = 1)


GROUP BY custID

4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price



When should tuples be constructed?

Solution 2: Operate on columns




```
QUERY:
SELECT custID,SUM(price)
FROM table
WHERE (prodID = 4) AND
      (storeID = 1)
GROUP BY custID
```

4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price



When should tuples be constructed?

Solution 2: Operate on columns



```
QUERY:
SELECT custID,SUM(price)
FROM table
WHERE (prodID = 4) AND
      (storeID = 1)
GROUP BY custID
```

4	2
4	1
4	3
4	1

prodID storeID



4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80

prodID storeID custID price



When should tuples be constructed?

Solution 2: Operate on columns



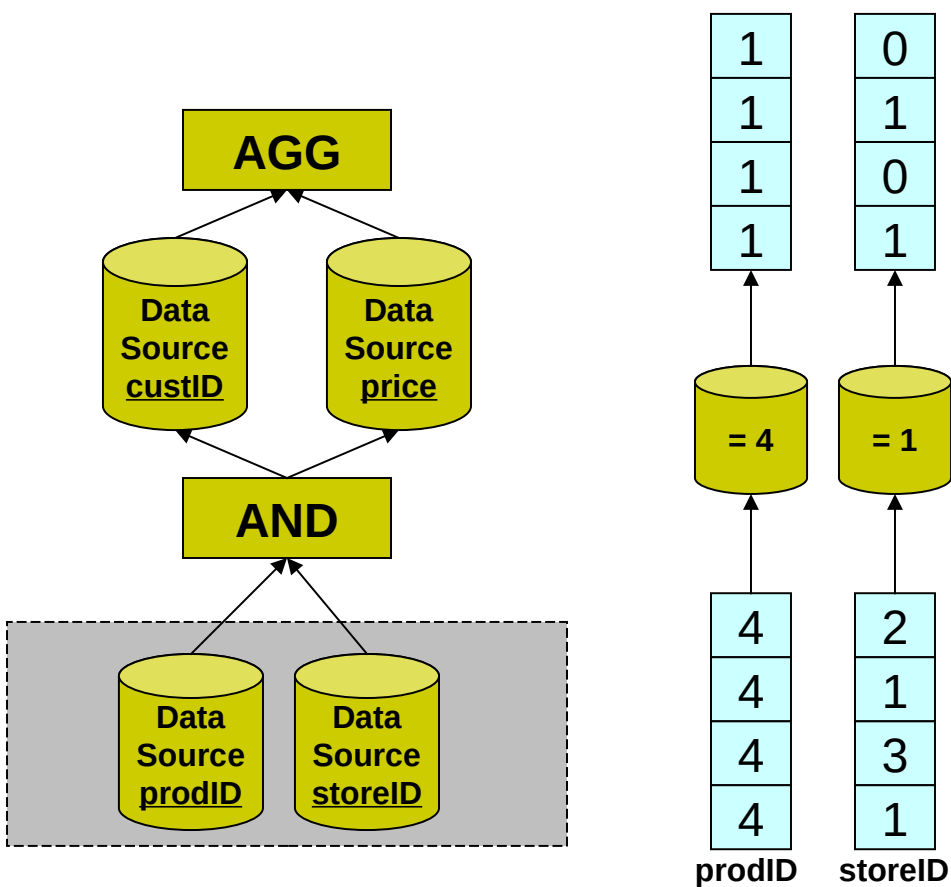
QUERY:
SELECT custID,SUM(price)
FROM table
WHERE (prodID = 4) AND
 (storeID = 1)
GROUP BY custID

4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price



When should tuples be constructed?

Solution 2: Operate on columns



QUERY:

```
SELECT custID, SUM(price)
FROM table
WHERE (prodID = 4) AND
      (storeID = 1)
GROUP BY custID
```


4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80

prodID storeID custID price



When should tuples be constructed?

Solution 2: Operate on columns



1	0
1	1
1	0
1	1

QUERY:

SELECT custID,SUM(price)

FROM table

WHERE (prodID = 4) AND

(storeID = 1)


GROUP BY custID

4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price



When should tuples be constructed?

Solution 2: Operate on columns



QUERY:

SELECT custID,SUM(price)

FROM table

WHERE (prodID = 4) AND

(storeID = 1)

GROUP BY custID



1	0
1	1
1	0
1	1

4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price



When should tuples be constructed?

Solution 2: Operate on columns





QUERY:
SELECT custID,SUM(price)
FROM table
WHERE (prodID = 4) AND
 (storeID = 1)
GROUP BY custID

4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price



When should tuples be constructed?

Solution 2: Operate on columns



QUERY:


```
SELECT custID,SUM(price)
FROM table
WHERE (prodID = 4) AND
      (storeID = 1)
GROUP BY custID
```

4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price



When should tuples be constructed?

Solution 2: Operate on columns



0
1
0
1

QUERY:

SELECT custID,SUM(price)

FROM table

WHERE (prodID = 4) AND

(storeID = 1)


GROUP BY custID

4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price



When should tuples be constructed?

Solution 2: Operate on columns



QUERY:
SELECT custID,SUM(price)
FROM table
WHERE (prodID = 4) **AND**
 (storeID = 1)
GROUP BY custID

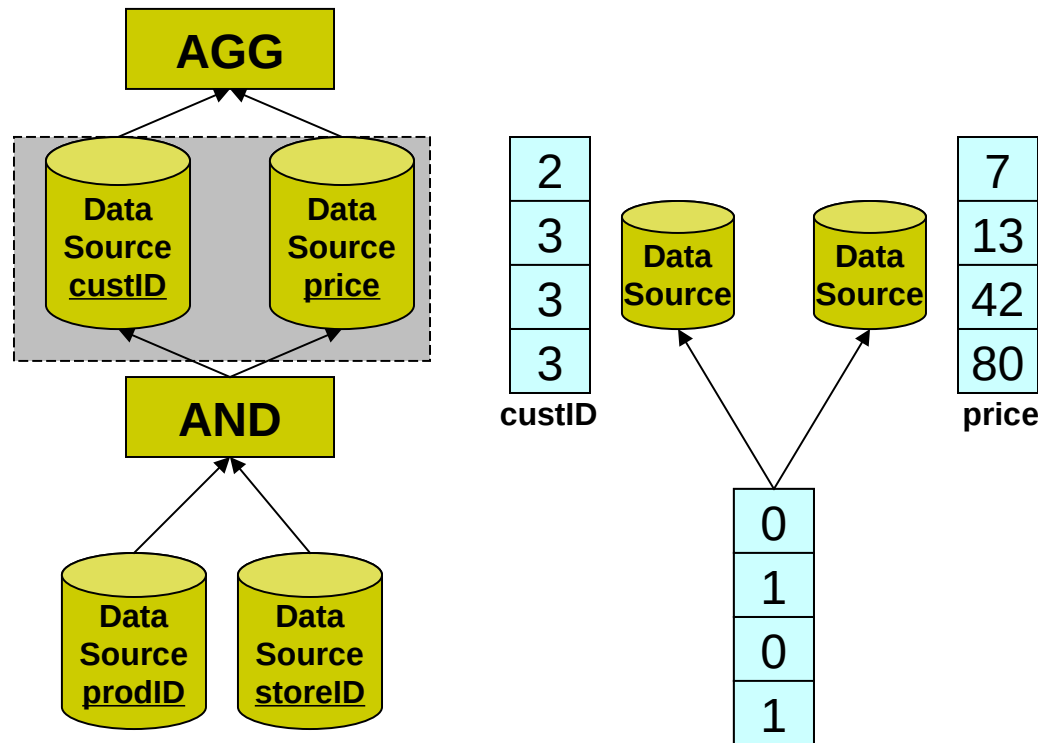
0
1
0
1

4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price



When should tuples be constructed?

Solution 2: Operate on columns



```
QUERY:
SELECT custID,SUM(price)
FROM table
WHERE (prodID = 4) AND
      (storeID = 1)
GROUP BY custID
```

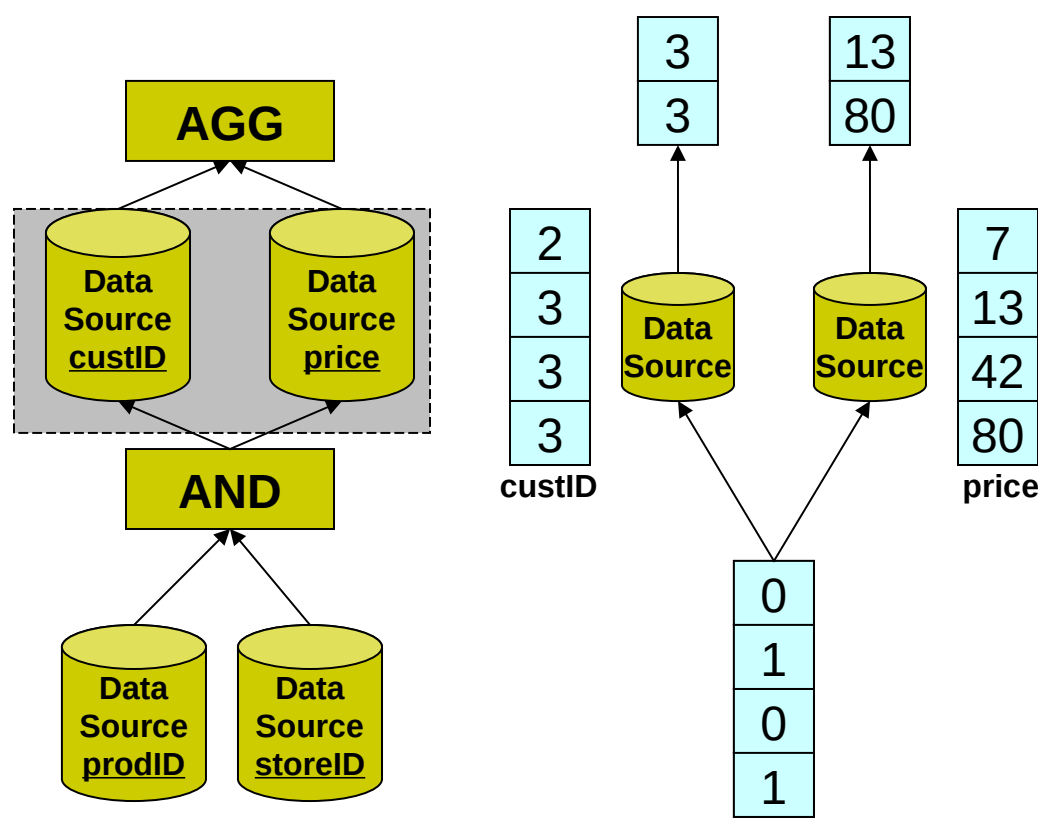
4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80

prodID storeID custID price



When should tuples be constructed?

Solution 2: Operate on columns



```
QUERY:
SELECT custID,SUM(price)
FROM table
WHERE (prodID = 4) AND
      (storeID = 1)
GROUP BY custID
```


4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80

prodID storeID custID price



When should tuples be constructed?

Solution 2: Operate on columns



3	13
3	80

QUERY:

SELECT custID,SUM(price)

FROM table

WHERE (prodID = 4) AND

(storeID = 1)


GROUP BY custID

4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price



When should tuples be constructed?

Solution 2: Operate on columns



```
QUERY:
SELECT custID,SUM(price)
FROM table
WHERE (prodID = 4) AND
      (storeID = 1)
GROUP BY custID
```

3
3



13
80

4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price



When should tuples be constructed?

Solution 2: Operate on columns





```
QUERY:
SELECT custID,SUM(price)
FROM table
WHERE (prodID = 4) AND
      (storeID = 1)
GROUP BY custID
```

4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price



When should tuples be constructed?

Solution 2: Operate on columns




QUERY:
SELECT custID,SUM(price)
FROM table
WHERE (prodID = 4) **AND**
 (storeID = 1)
GROUP BY custID

4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price



When should tuples be constructed?

Solution 2: Operate on columns



3	93
---	----

QUERY:

SELECT custID,SUM(price)

FROM table

WHERE (prodID = 4) AND

(storeID = 1)

GROUP BY custID

4	2	2	7
4	1	3	13
4	3	3	42
4	1	3	80
prodID	storeID	custID	price



RISC Relational Algebra (MonetDB)

people_id		people_name		people_age	
(void)	(int)	(void)	(str)	(void)	(int)
0	101	0	Alice	0	22
1	102	1	Ivan	1	37
2	104	2	Peggy	2	45
3	105	3	Victor	3	25
4	108	4	Eve	4	19
5	109	5	Walter	5	31
6	112	6	Trudy	6	27
7	113	7	Bob	7	29
8	114	8	Zoe	8	42
9	115	9	Charlie	9	35



RISC Relational Algebra (MonetDB)


```
SELECT  id, name, (age-30)*50 as bonus
FROM    people
WHERE   age > 30
```

people_id		people_name		people_age	
(void)	(int)	(void)	(str)	(void)	(int)
0	101	0	Alice	0	22
1	102	1	Ivan	1	37
2	104	2	Peggy	2	45
3	105	3	Victor	3	25
4	108	4	Eve	4	19
5	109	5	Walter	5	31
6	112	6	Trudy	6	27
7	113	7	Bob	7	29
8	114	8	Zoe	8	42
9	115	9	Charlie	9	35



RISC Relational Algebra (MonetDB)


```
SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30
```





RISC Relational Algebra (MonetDB)

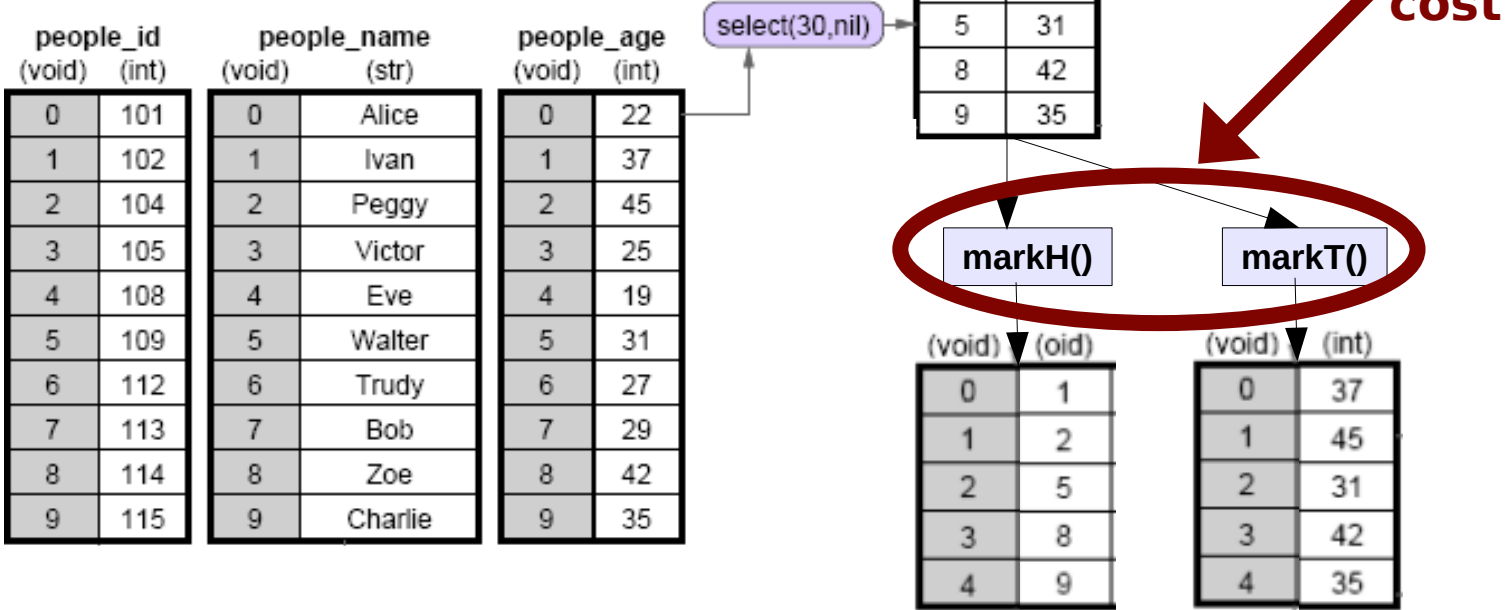
```
SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30
```





RISC Relational Algebra (MonetDB)

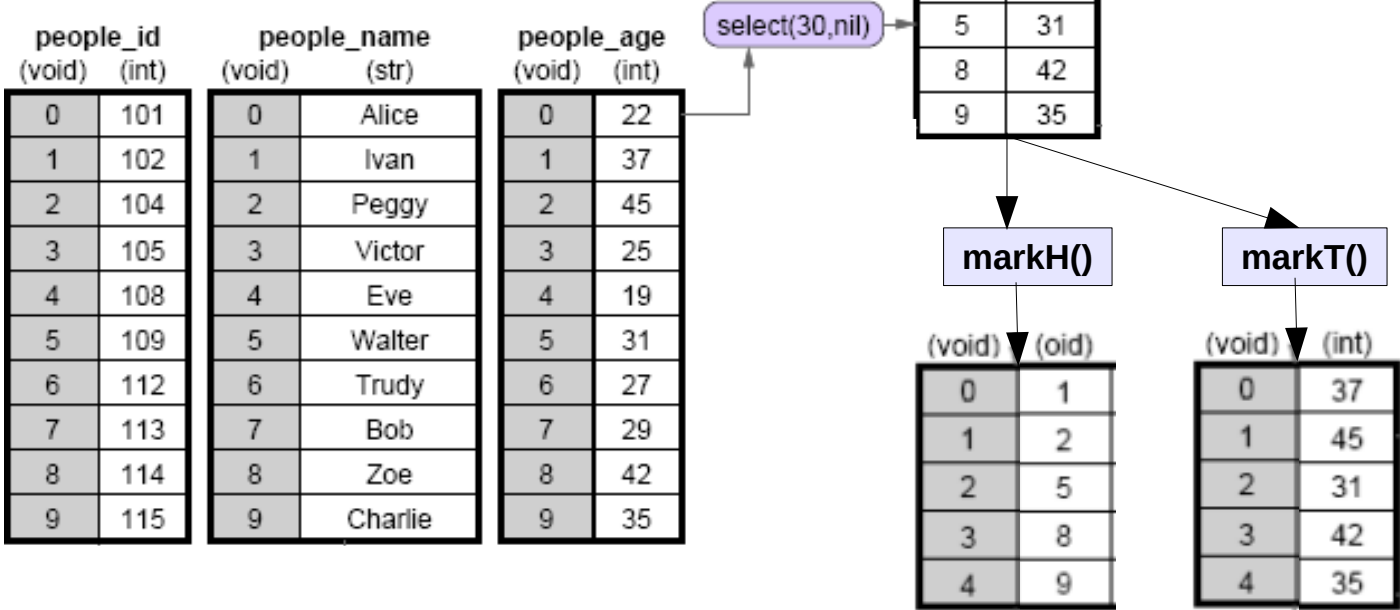
```
SELECT id, name, (age-30)*50 as bonus
FROM   people
WHERE  age > 30
```





RISC Relational Algebra (MonetDB)

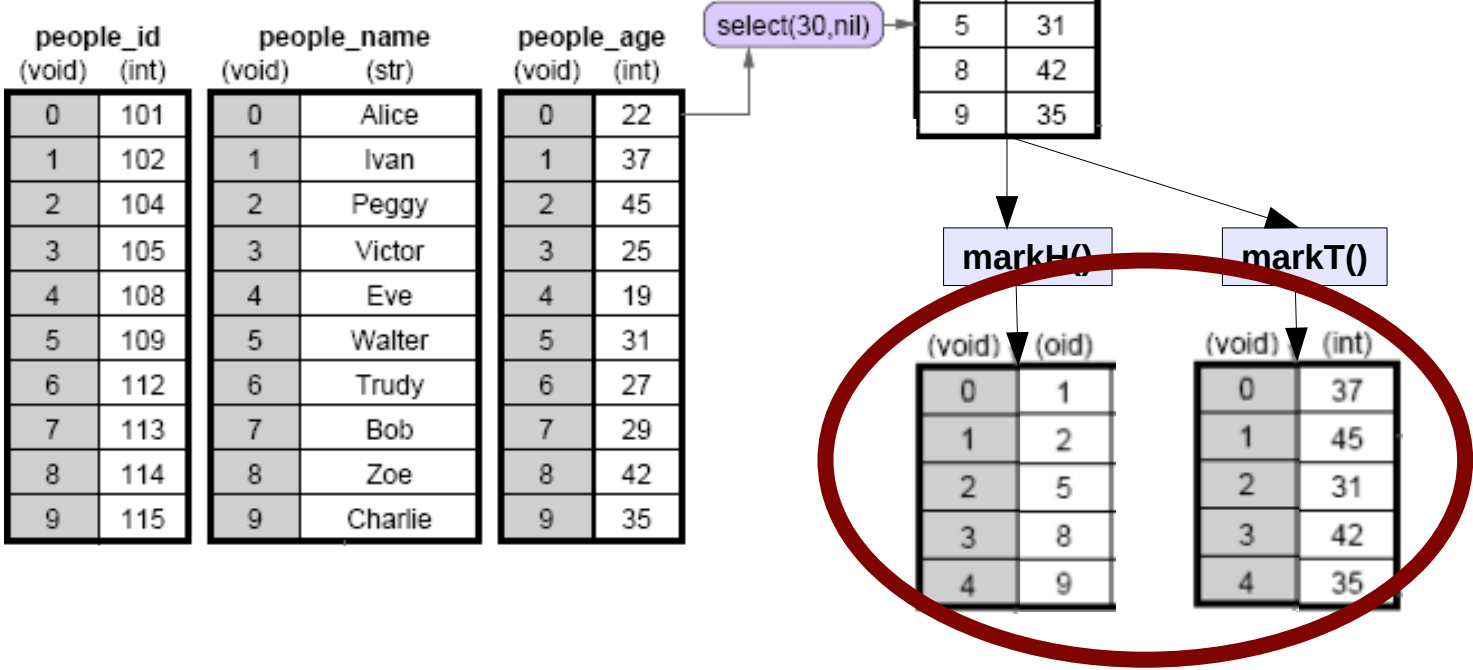
```
SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30
```





RISC Relational Algebra (MonetDB)

```
SELECT id, name, (age-30)*50 as bonus
FROM   people
WHERE  age > 30
```




VIEWS
(not materialized)



RISC Relational Algebra (MonetDB)

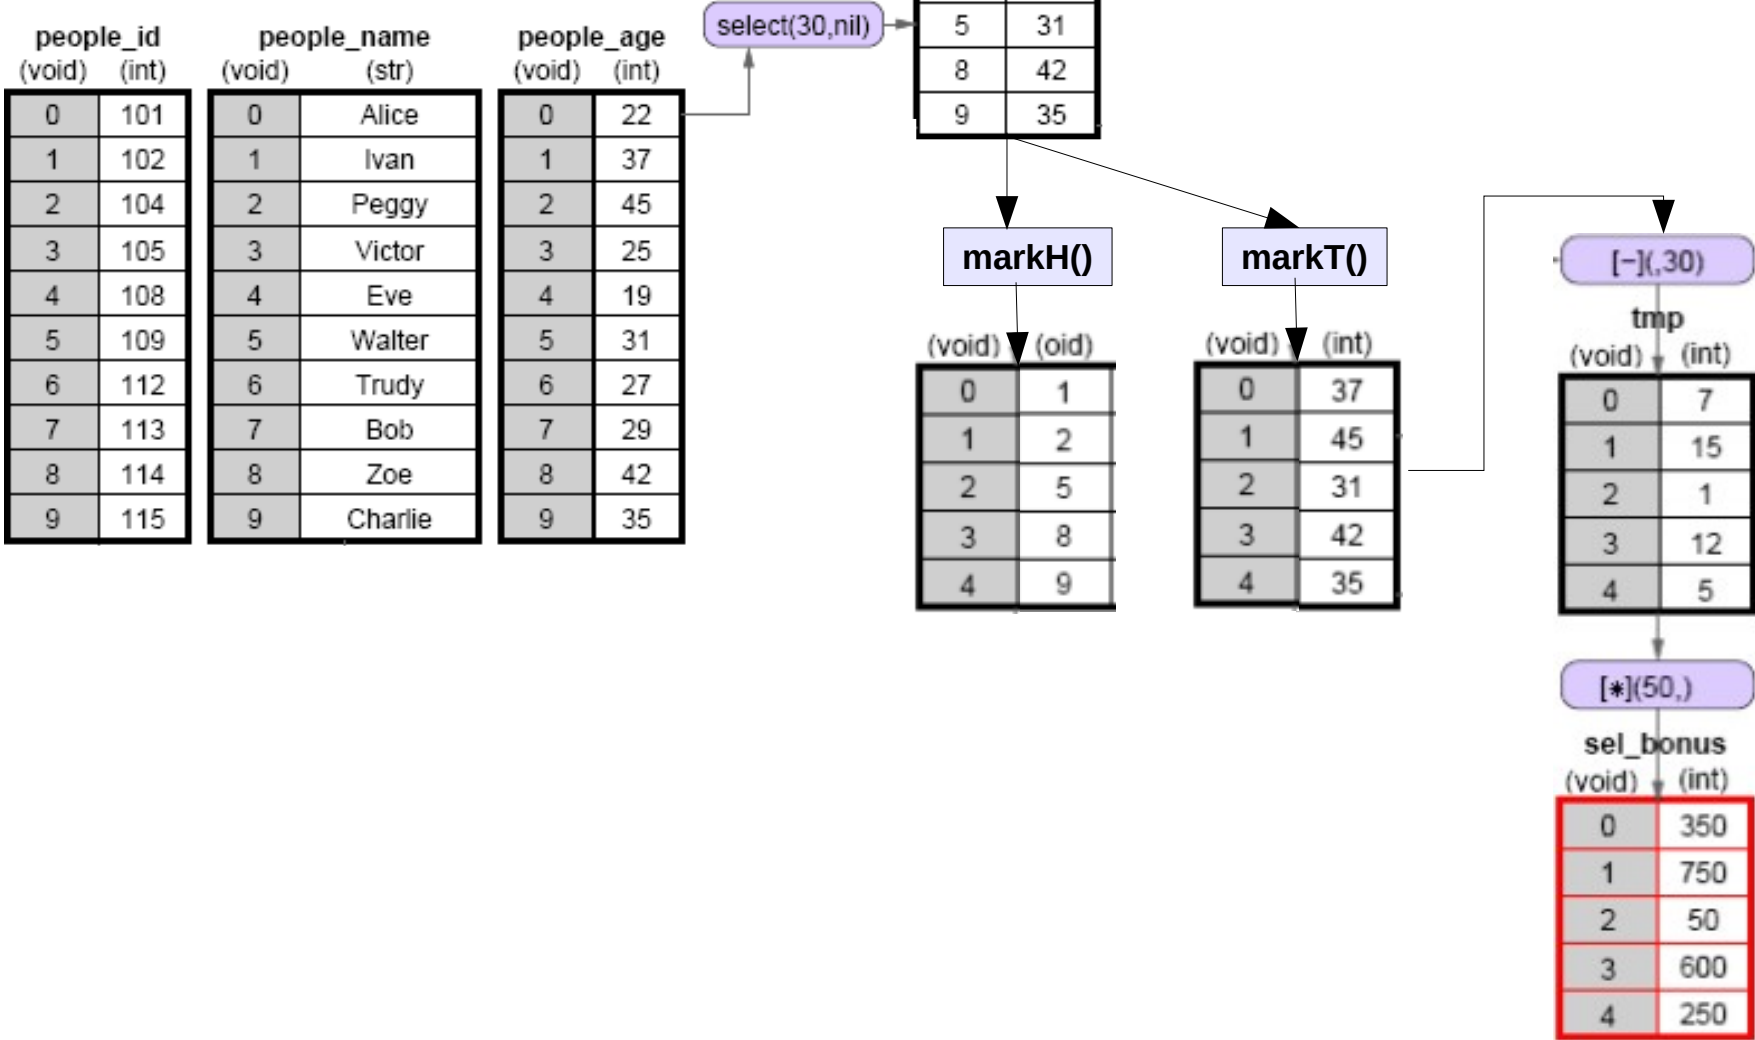
```
SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30
```





RISC Relational Algebra (MonetDB)

```
SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30
```

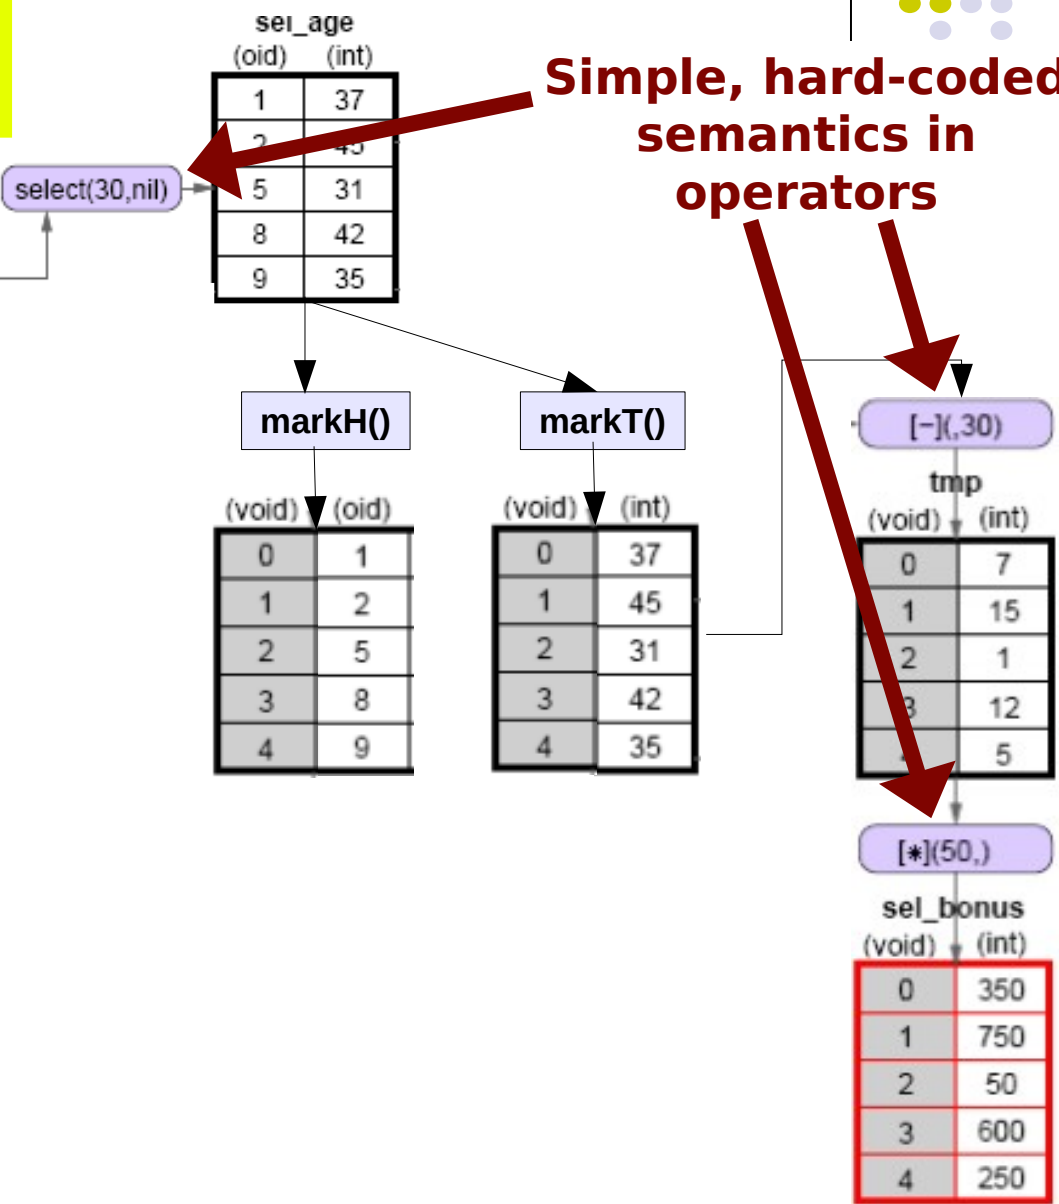




RISC Relational Algebra (MonetDB)

```
SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30
```

people_id		people_name		people_age	
(void)	(int)	(void)	(str)	(void)	(int)
0	101	0	Alice	0	22
1	102	1	Ivan	1	37
2	104	2	Peggy	2	45
3	105	3	Victor	3	25
4	108	4	Eve	4	19
5	109	5	Walter	5	31
6	112	6	Trudy	6	27
7	113	7	Bob	7	29
8	114	8	Zoe	8	42
9	115	9	Charlie	9	35



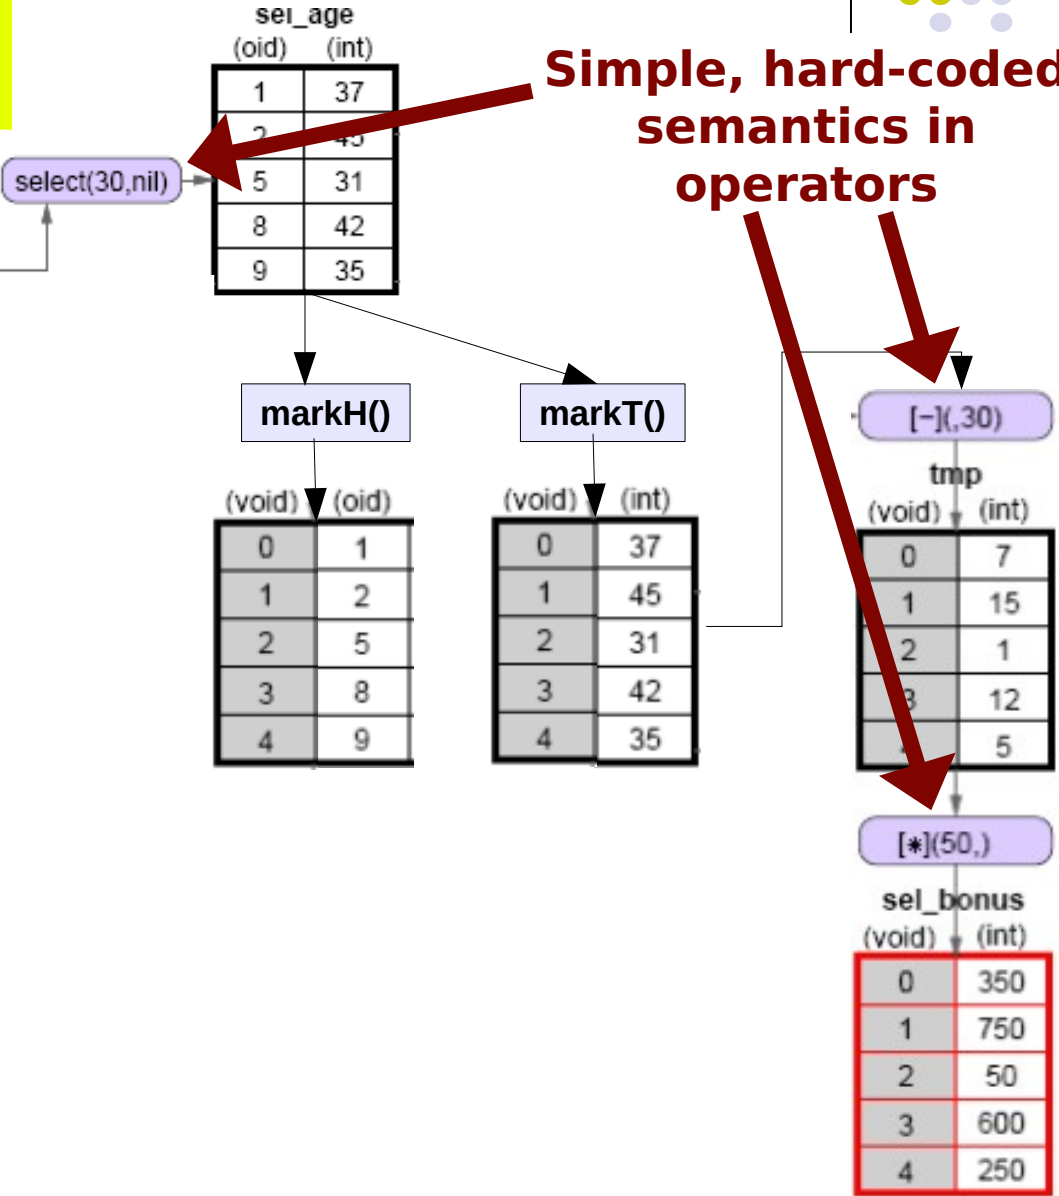


RISC Relational Algebra (MonetDB)

```
SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30
```

```
batcalc_minus_int(int* res,
                  int* col,
                  int val,
                  int n)
{
    for(i=0; i<n; i++)
        res[i] = col[i] - val;
}
```

7	113	7	Bob	7	29
8	114	8	Zoe	8	42
9	115	9	Charlie	9	35



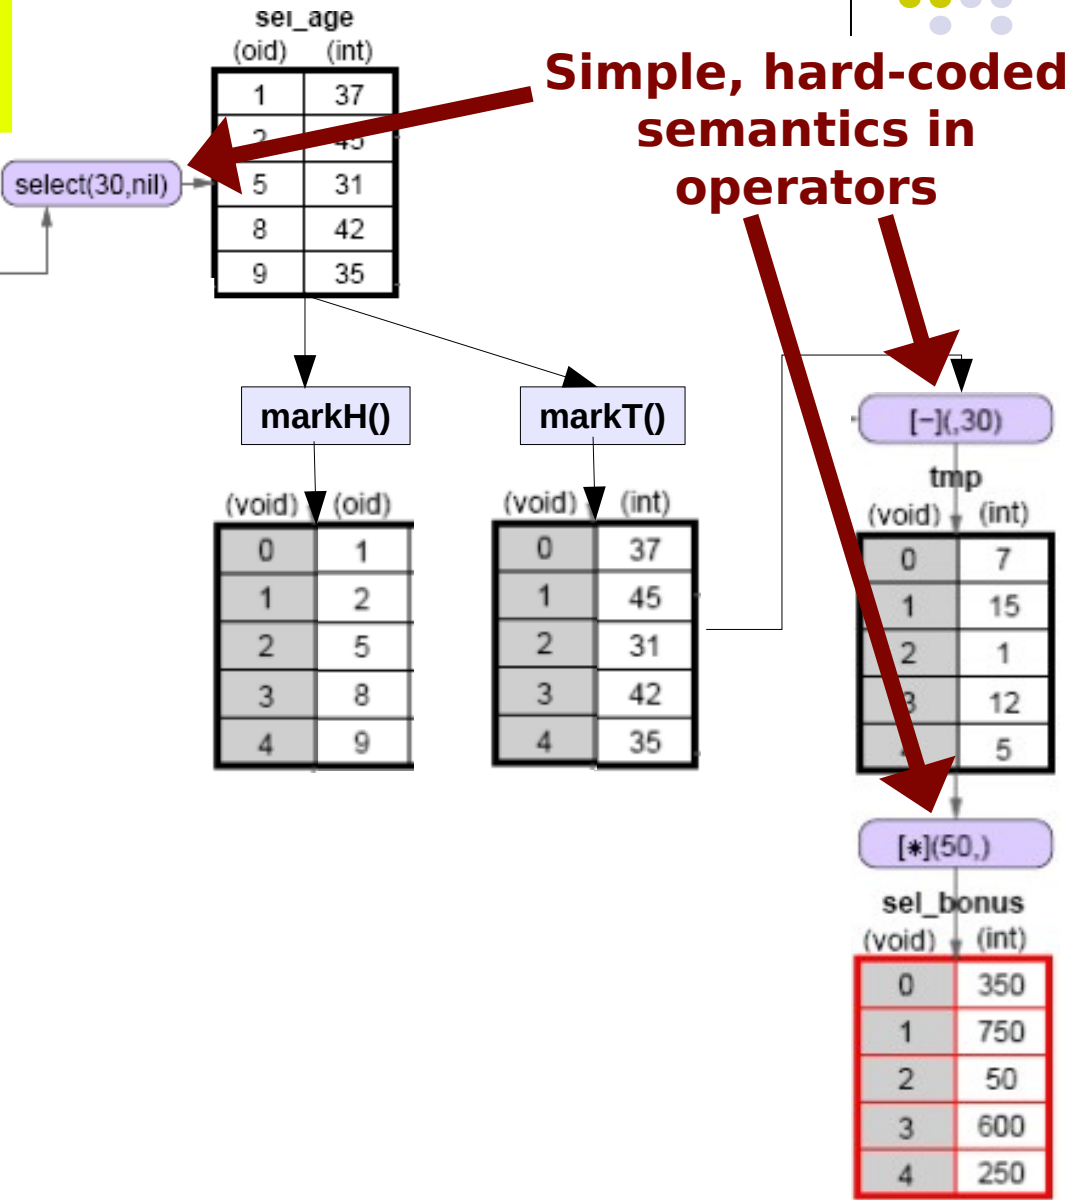


RISC Relational Algebra (MonetDB)

```
SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30
```

```
batcalc_minus_int(int* res,
                  int* col,
                  int val,
                  int n)
{
    for(i=0; i<n; i++)
        res[i] = col[i] - val;
}
```

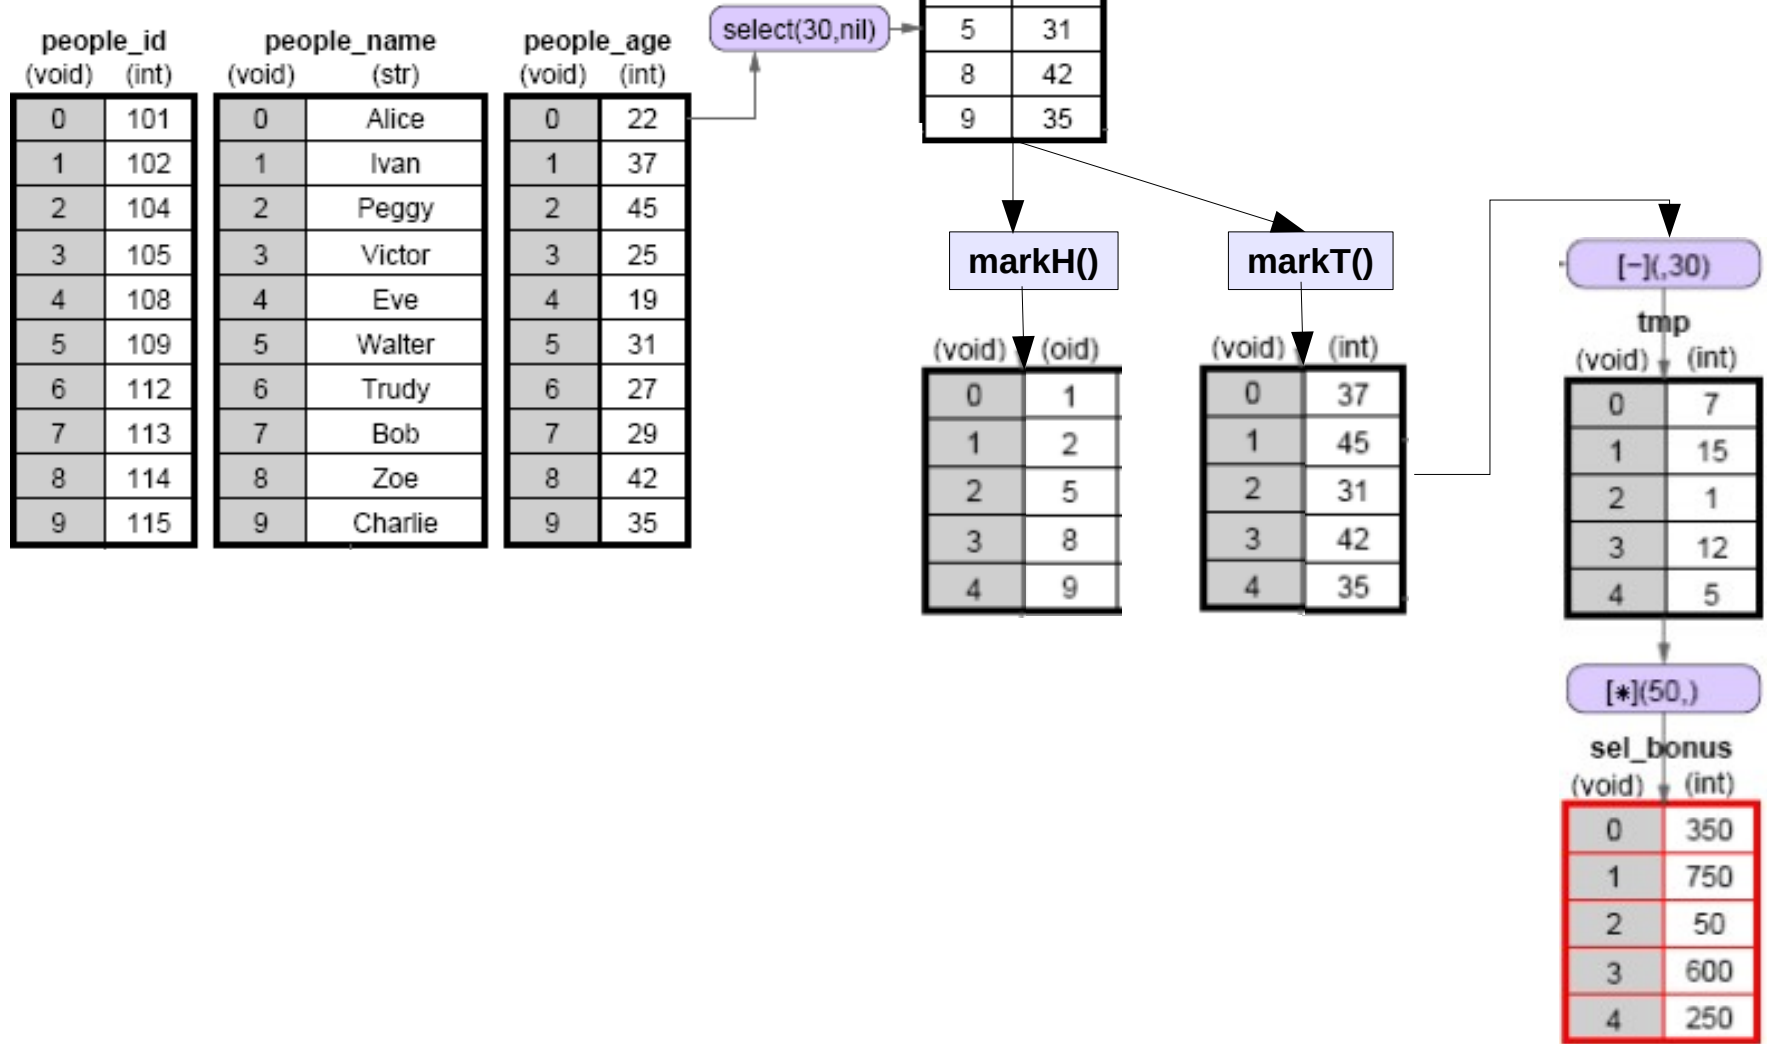
- CPU ☺? Give it “nice” code !**
- few dependencies (control,data)
 - CPU gets out-of-order execution
 - compiler can e.g. generate SIMD
- One loop for an entire column**
- no per-tuple interpretation
 - arrays: no record navigation
 - better instruction cache locality





RISC Relational Algebra (MonetDB)

```
SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30
```





RISC Relational Algebra (MonetDB)

```
SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30
```

people_id	people_name	people_age
(void) (int)	(void) (str)	(void) (int)
0 101	0 Alice	0 22
1 102	1 Ivan	1 37
2 104	2 Peggy	2 45
3 105	3 Victor	3 25
4 108	4 Eve	4 19
5 109	5 Walter	5 31
6 112	6 Trudy	6 27
7 113	7 Bob	7 29
8 114	8 Zoe	8 42
9 115	9 Charlie	9 35

select(30,nil)

sel_age
(oid) (int)
1 37
2 45
5 31
8 42
9 35

MATERIALIZED
intermediate
results

markH()

(void)	(oid)
0	1
1	2
2	5
3	8
4	9

markT()

(void)	(int)
0	37
1	45
2	31
3	42
4	35

[-](,30)

tmp
(void) (int)
0 7
1 15
2 1
3 12
4 5


[*](50,)

sel_bonus
(void) (int)
0 350
1 750
2 50
3 600
4 250



RISC Relational Algebra (MonetDB)

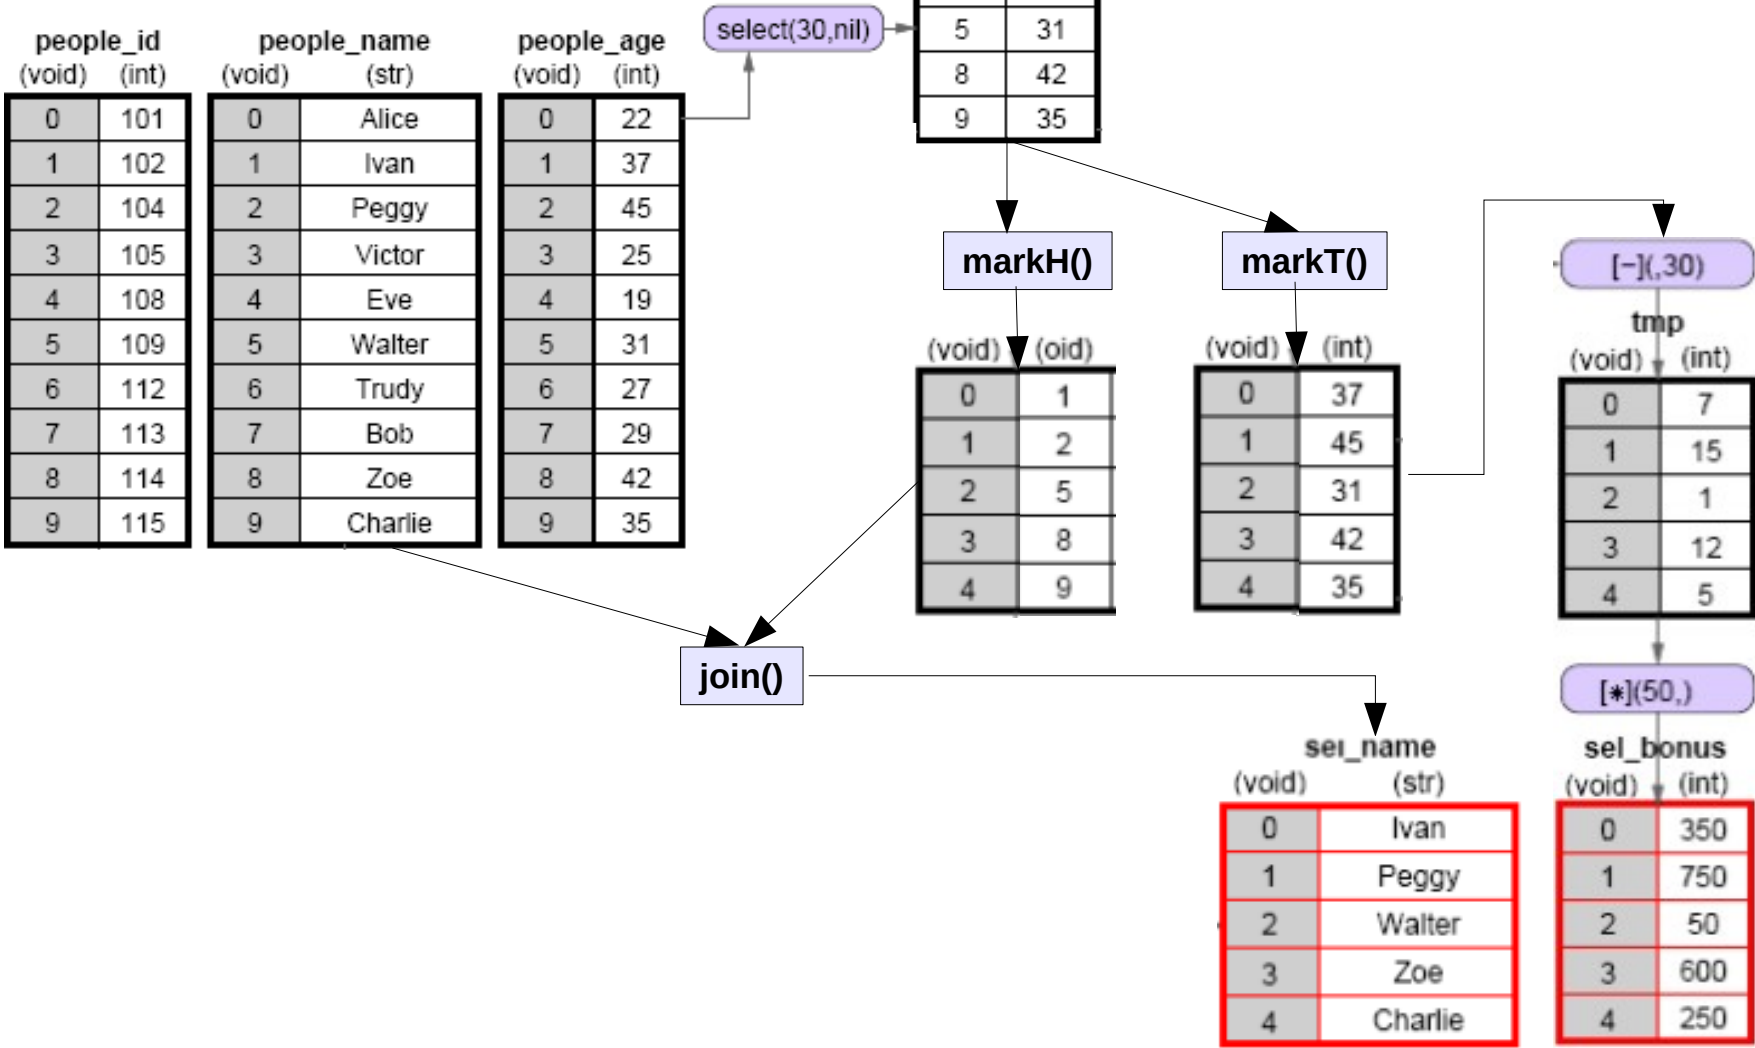
```
SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30
```





RISC Relational Algebra (MonetDB)

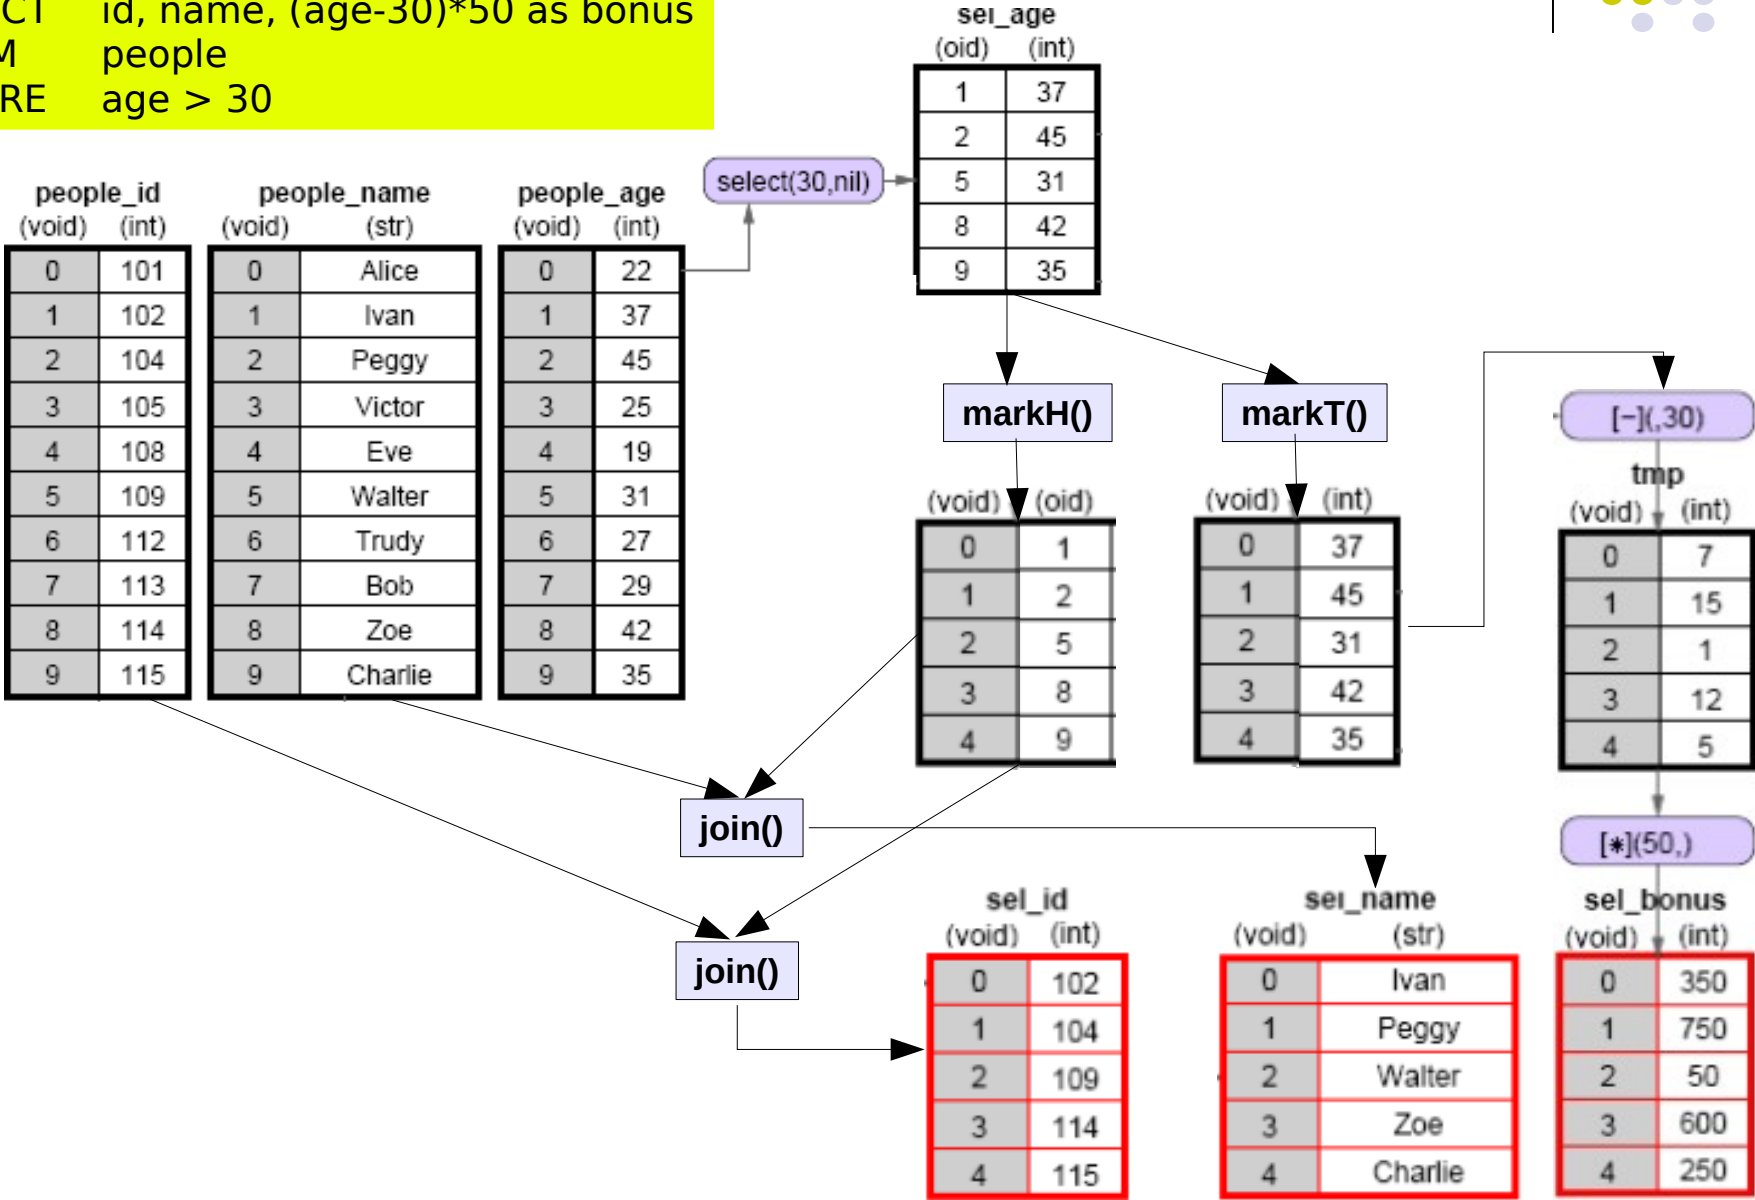
```
SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30
```





RISC Relational Algebra (MonetDB)

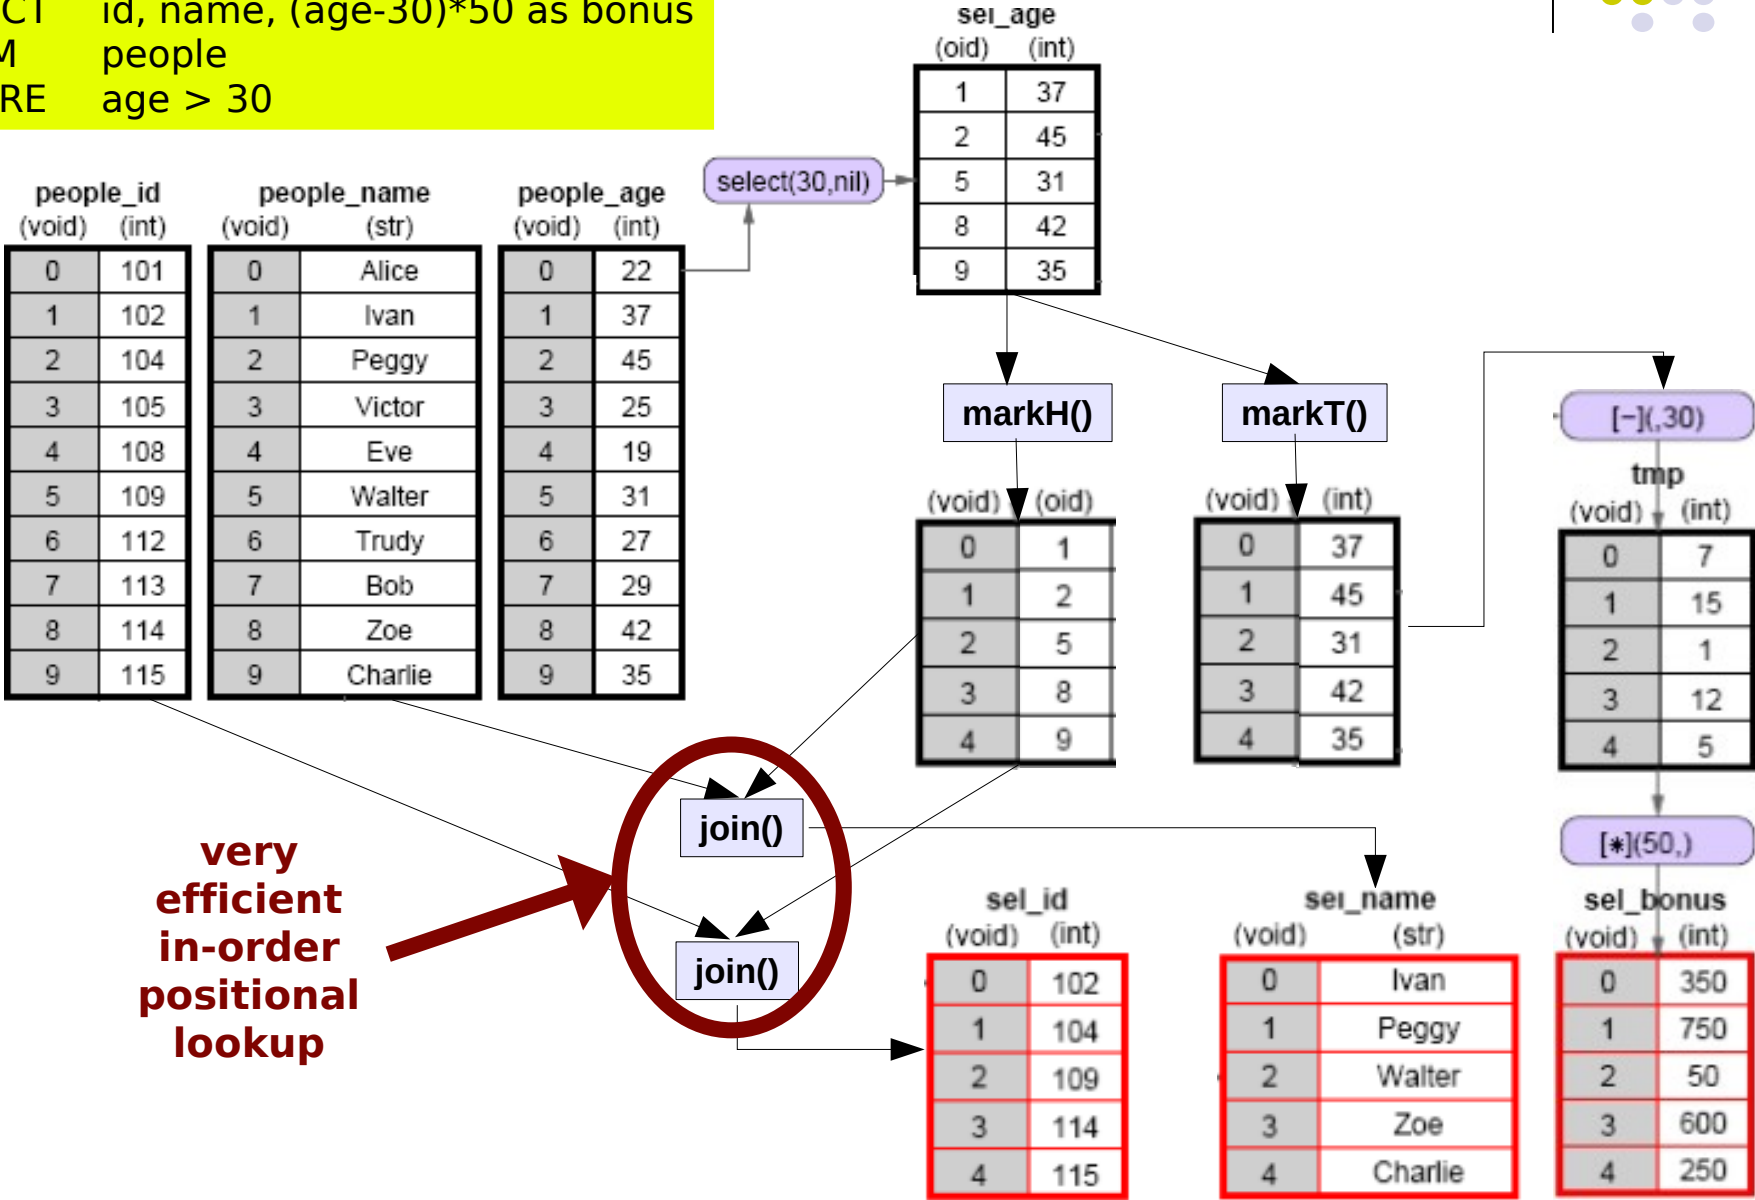
```
SELECT id, name, (age-30)*50 as bonus
FROM   people
WHERE  age > 30
```





RISC Relational Algebra (MonetDB)


```
SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30
```





RISC Relational Algebra (MonetDB)

```
SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30
```

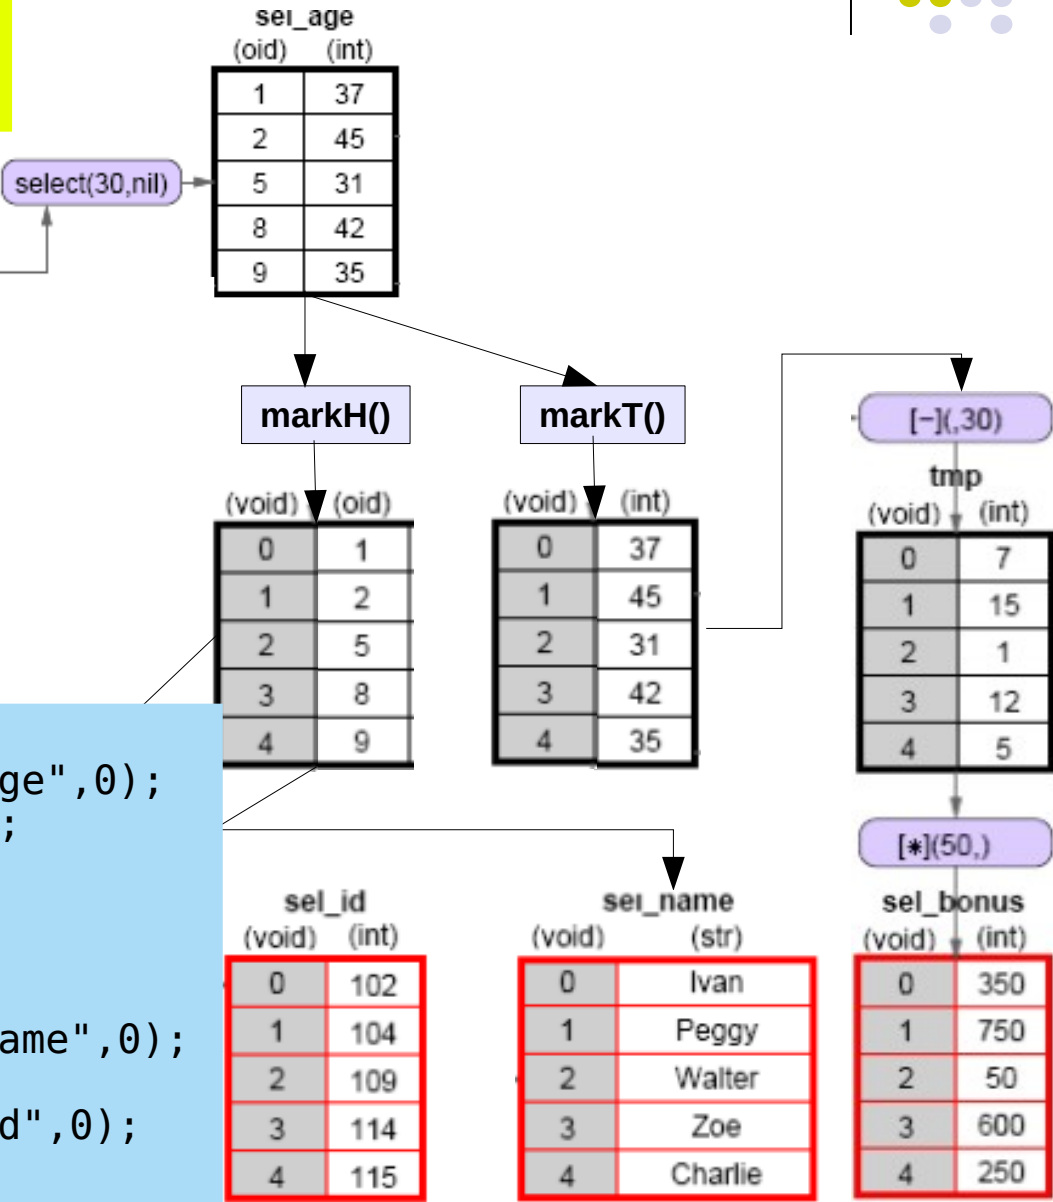




RISC Relational Algebra (MonetDB)

SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30

people_id	people_name	people_age
(void) (int)	(void) (str)	(void) (int)
0 101	0 Alice	0 22
1 102	1 Ivan	1 37
2 104	2 Peggy	2 45
3 105	3 Victor	3 25
4 108	4 Eve	4 19
5 109	5 Walter	5 31
6 112	6 Trudy	6 27
7 113	7 Bob	7 29
8 114	8 Zoe	8 42
9 115	9 Charlie	9 35



```
_01 := sql.bind("sys", "people", "age", 0);  
_02 := algebra.select(_01, 30, nil);  
_03 := algebra.markH(_02, 0);  
_04 := algebra.markT(_02, 0);  
_05 := batcalc.-(_04, 30);  
_06 := batcalc.*(_05, 50);  
_07 := sql.bind("sys", "people", "name", 0);  
_08 := algebra.join(_03, _07);  
_09 := sql.bind("sys", "people", "id", 0);  
_10 := algebra.join(_03, _09);
```



RISC Relational Algebra (MonetDB)

SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30

people_id	people_name	people_age
(void) (int)	(void) (str)	(void) (int)
0 101	0 Alice	0 22
1 102	1 Ivan	1 37
2 104	2 Peggy	2 45
3 105	3 Victor	3 25
4 108	4 Eve	4 19
5 109	5 Walter	5 31
6 112	6 Trudy	6 27
7 113	7 Bob	7 29
8 114	8 Zoe	8 42
9 115	9 Charlie	9 35

select(30,nil)

sel_age
(oid) (int)
1 37
2 45
5 31
8 42
9 35

markH()

(void) (oid)
0 1
1 2
2 5
3 8
4 9

markT()

(void) (int)
0 37
1 45
2 31
3 42
4 35

[-](,30)

tmp
(void) (int)
0 7
1 15
2 1
3 12
4 5

[(*)(50,)

sel_bonus
(void) (int)
0 350
1 750
2 50
3 600
4 250

sel_id

(void) (int)
0 102
1 104
2 109
3 114
4 115

sel_name

(void) (str)
0 Ivan
1 Peggy
2 Walter
3 Zoe
4 Charlie

```
_01 := sql.bind("sys", "people", "age", 0);  
_02 := algebra.select(_01, 30, nil);  
_03 := algebra.markH(_02, 0);  
_04 := algebra.markT(_02, 0);  
_05 := batcalc.-(_04, 30);  
_06 := batcalc.*(_05, 50);  
_07 := sql.bind("sys", "people", "name", 0);  
_08 := algebra.project(_03, _07);  
_09 := sql.bind("sys", "people", "id", 0);  
_10 := algebra.project(_03, _09);
```

MonetDB Front-end: SQL



```
PLAN SELECT a FROM t WHERE c < 10;
```

```
project (  
  select (  
    table(sys.t) [ t.a, t.c, t.%TID% NOT NULL ]  
  ) [ t.c < convert(10) ]  
) [ t.a ]
```

MonetDB Front-end: SQL



EXPLAIN SELECT a FROM t WHERE c < 10;

```
function user.s1_1():void;
barrier _55 := language.dataflow();
  _02:bat[:void,:int] := sql.bind("sys","t","c",0);
  _07:bat[:oid, :int] := algebra.thetaselect(_02,10,"<");
  _11:bat[:void,:oid] := algebra.markH(_07,0@0);
  _12:bat[:oid, :int] := sql.bind("sys","t","a",0);
  _14:bat[:void,:int] := algebra.project(_11,_12);
exit _55;
  _15 := sql.resultSet(1,1,_14);
  sql.rsColumn(_15,"sys.t","a","int",32,0,_14);
  _21 := io.stdout();
  sql.exportResult(_21,_15);
end s1_1;
```

MonetDB Front-end: SQL



PLAN SELECT a, z FROM t, s WHERE t.c = s.x;

PLAN SELECT a, z FROM t join s on t.c = s.x;

```
project (  
  join (  
    table(sys.t) [ t.a, t.c, t.%TID% NOT NULL ],  
    table(sys.s) [ s.x, s.z, s.%TID% NOT NULL ]  
  ) [ t.c = s.x ]  
) [ t.a, s.z ]
```

MonetDB Front-end: SQL




EXPLAIN SELECT **a**, **z** FROM t, s WHERE **t.c** = **s.x**;

```
function user.s2_1():void;
  barrier _73 := language.dataflow();
  _02:bat[:void,:int] := sql.bind("sys","t","c",0);
  _07:bat[:void,:int] := sql.bind("sys","s","x",0);
  _10:bat[:int,:void] := bat.reverse(_07);
  _11:bat[:oid, :oid] := algebra.join(_02,_10);
  _14:bat[:void,:oid] := algebra.markH(_11,0@0);
  _15:bat[:void,:int] := sql.bind("sys","t","a",0);
  _17:bat[:void,:int] := algebra.project(_14,_15);
  _20:bat[:oid,:void] := algebra.markT(_18,0@0);
  _21:bat[:void,:int] := sql.bind("sys","s","z",0);
  _23:bat[:void,:int] := algebra.project(_20,_21);
  exit _73;
  _24 := sql.resultSet(2,1,_17);
  sql.rsColumn(_24,"sys.t","a","int",32,0,_17);
  sql.rsColumn(_24,"sys.s","z","int",32,0,_23);
  _33 := io.stdout();
  sql.exportResult(_33,_24);
end s2_1;
```


“Materialization Strategies in a Column-Oriented DBMS”
Abadi, Myers, DeWitt, and Madden. ICDE 2007.



For plans without joins, late materialization is a win



QUERY:

```
SELECT  $C_1$ , SUM( $C_2$ )
```

```
FROM table
```

```
WHERE ( $C_1$  < CONST) AND  
        ( $C_2$  < CONST)
```


```
GROUP BY  $C_1$ 
```

- Ran on 2 compressed columns from TPC-H scale 10 data

“Materialization Strategies in a Column-Oriented DBMS”
Abadi, Myers, DeWitt, and Madden. ICDE 2007.



Even on uncompressed data, late materialization is still a win



QUERY:

SELECT C_1 , SUM(C_2)

FROM table

**WHERE ($C_1 < \text{CONST}$) AND
($C_2 < \text{CONST}$)**



GROUP BY C_1

- **Materializing late still works best**



What about for plans with joins?



Select R1.B, R1.C, R2.E, R2.H, R3.F
From R1, R2, R3
Where R1.A = R2.D AND R2.G = R3.K





What about for plans with joins?



Select R1.B, R1.C, R2.E, R2.H, R3.F
From R1, R2, R3
Where R1.A = R2.D AND R2.G = R3.K





What about for plans with joins?



Select R1.B, R1.C, R2.E, R2.H, R3.F
From R1, R2, R3
Where R1.A = R2.D AND R2.G = R3.K





What about for plans with joins?



Select R1.B, R1.C, R2.E, R2.H, R3.F
From R1, R2, R3
Where R1.A = R2.D AND R2.G = R3.K





What about for plans with joins?



Select R1.B, R1.C, R2.E, R2.H, R3.F
From R1, R2, R3
Where R1.A = R2.D AND R2.G = R3.K





What about for plans with joins?



Select R1.B, R1.C, R2.E, R2.H, R3.F
From R1, R2, R3
Where R1.A = R2.D AND R2.G = R3.K





What about for plans with joins?

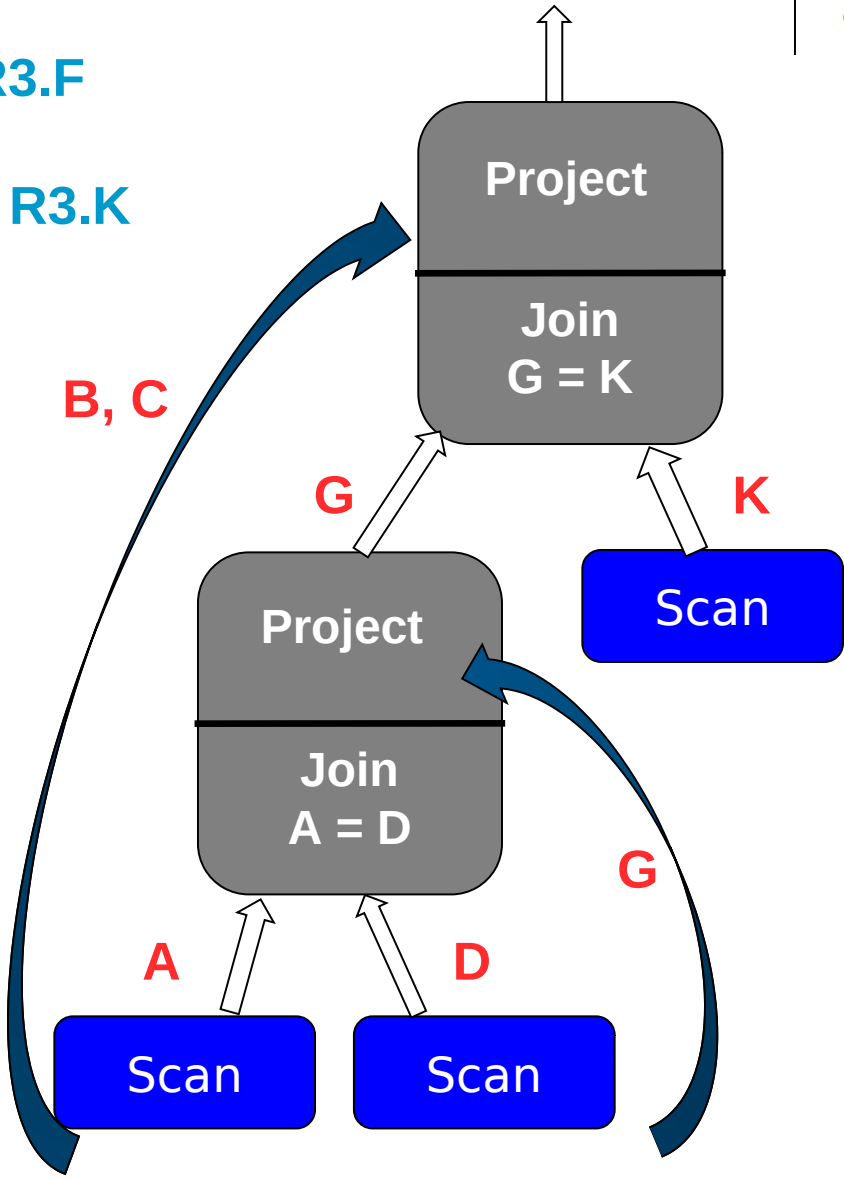

Select R1.B, R1.C, R2.E, R2.H, R3.F
From R1, R2, R3
Where R1.A = R2.D AND R2.G = R3.K





What about for plans with joins?

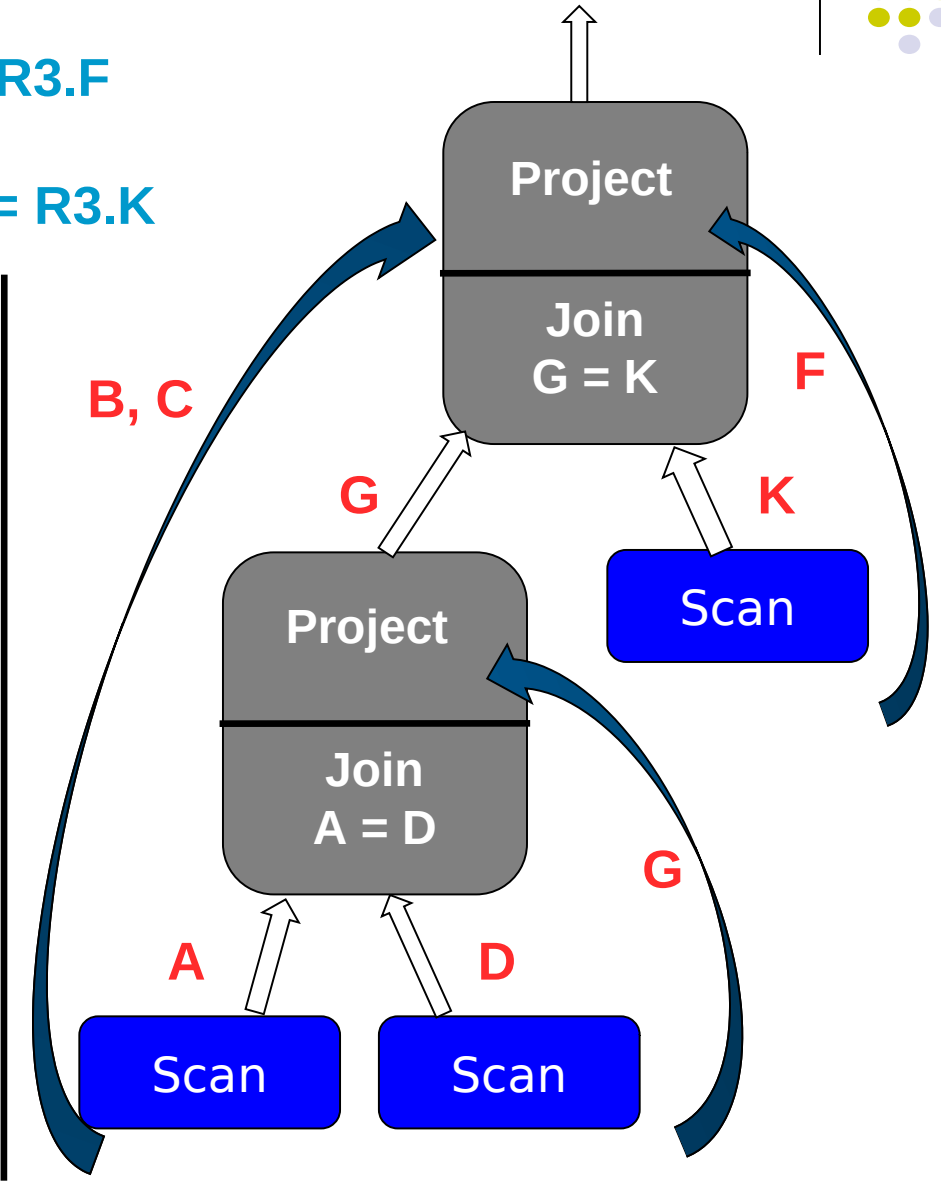

Select R1.B, R1.C, R2.E, R2.H, R3.F
From R1, R2, R3
Where R1.A = R2.D AND R2.G = R3.K





What about for plans with joins?

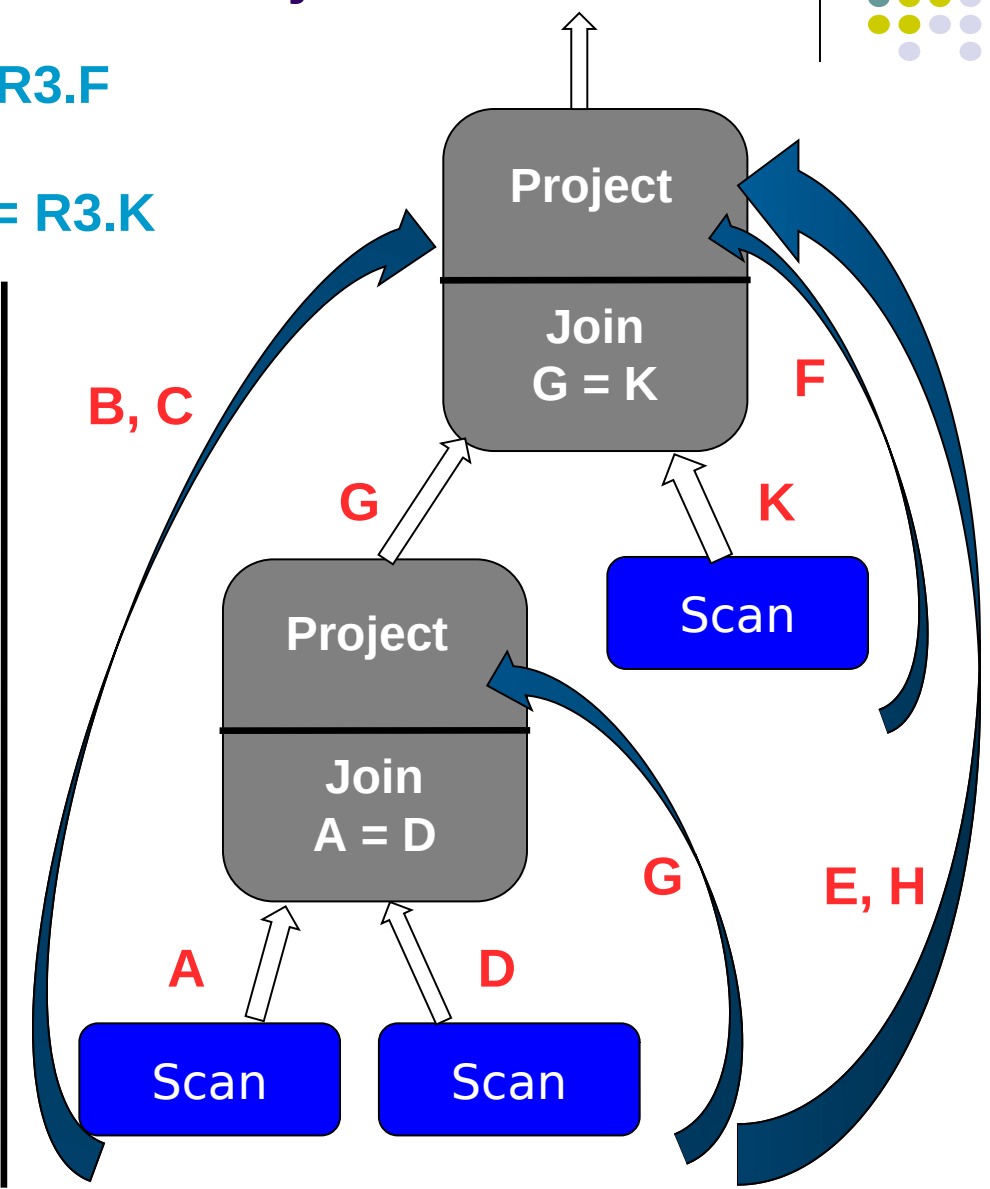

Select R1.B, R1.C, R2.E, R2.H, R3.F
From R1, R2, R3
Where R1.A = R2.D AND R2.G = R3.K





What about for plans with joins?



Select R1.B, R1.C, R2.E, R2.H, R3.F
From R1, R2, R3
Where R1.A = R2.D AND R2.G = R3.K





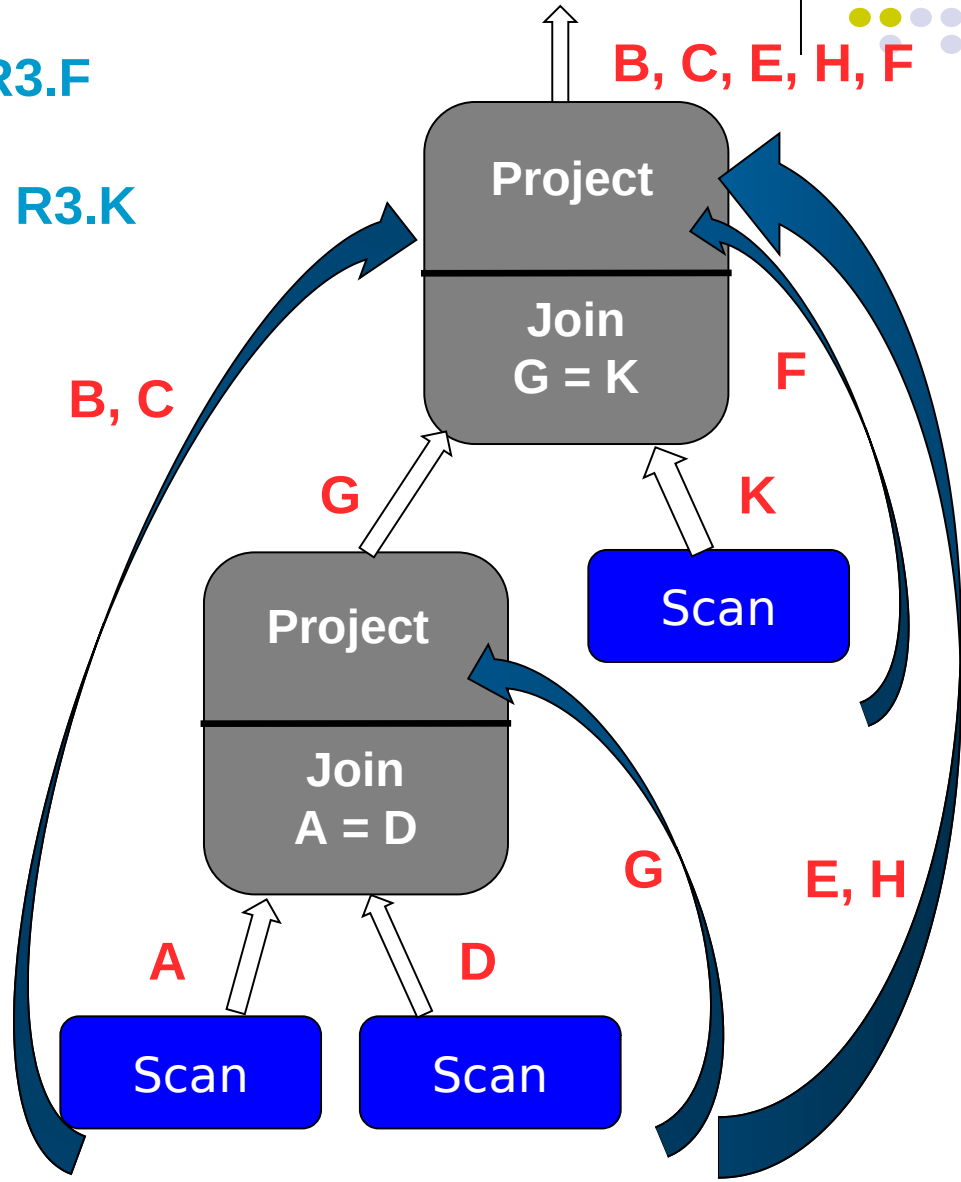

What about for plans with joins?

Select R1.B, R1.C, R2.E, R2.H, R3.F
From R1, R2, R3
Where R1.A = R2.D AND R2.G = R3.K



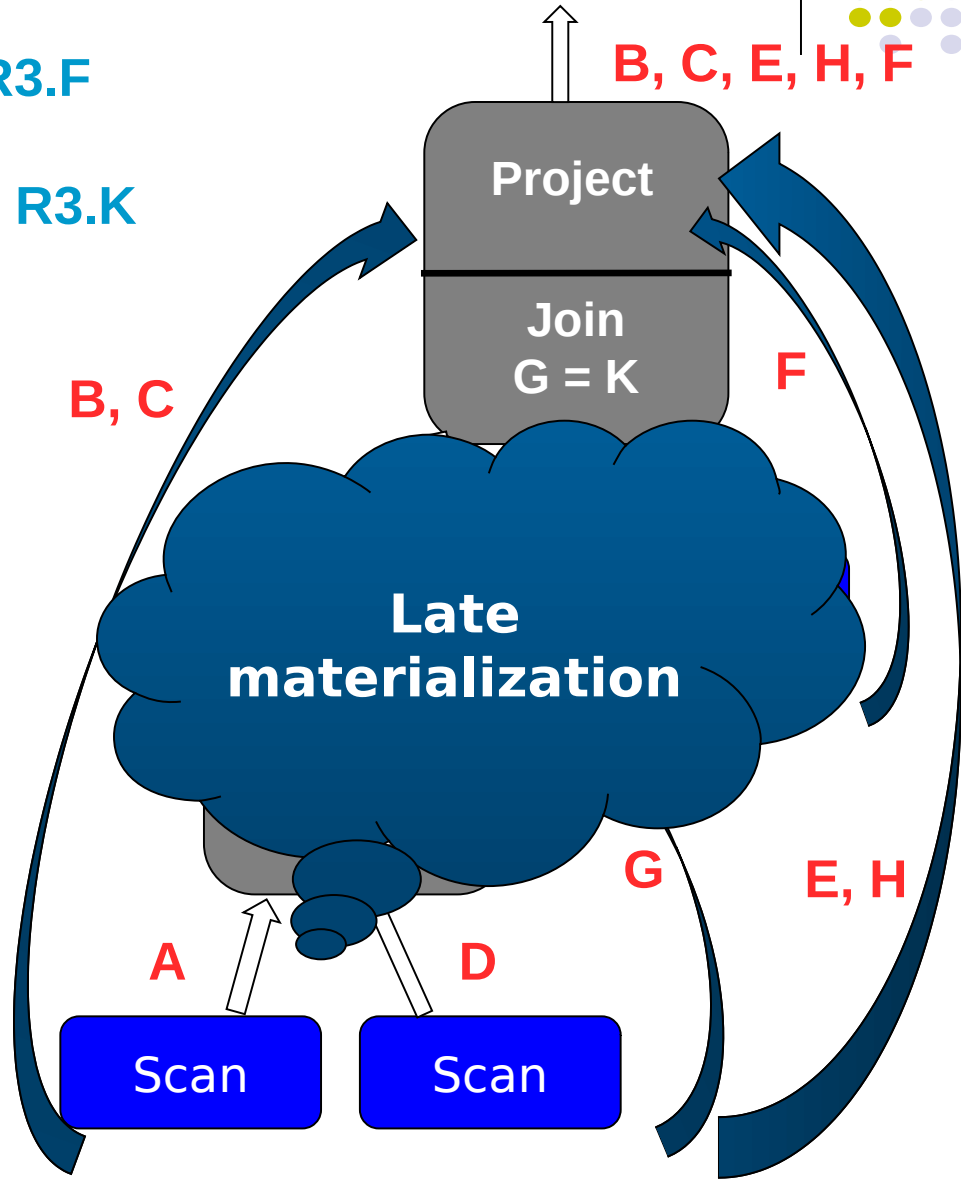

What about for plans with joins?

Select R1.B, R1.C, R2.E, R2.H, R3.F
From R1, R2, R3
Where R1.A = R2.D AND R2.G = R3.K



What about for plans with joins?

Select R1.B, R1.C, R2.E, R2.H, R3.F
From R1, R2, R3
Where R1.A = R2.D AND R2.G = R3.K





Early Materialization Example

QUERY:

SELECT C.lastName,SUM(F.price)

FROM facts AS F, customers AS C

WHERE F.custID = C.custID

GROUP BY C.lastName

(4,1,4)	12	6	2	7
	1	1	3	13
	11	2	3	42
	1	1	3	80

prodID storeID quantity custID price

Facts



1	Green
2	White
3	Brown

custID lastName

Customers



Early Materialization Example



QUERY:

```
SELECT C.lastName,SUM(F.price)
FROM facts AS F, customers AS C
WHERE F.custID = C.custID
GROUP BY C.lastName
```

prodID storeID quantity custID price

custID lastName

Facts

Customers



Early Materialization Example

2	7
3	13
3	42
3	80

1	Green
2	White
3	Brown

QUERY:
SELECT C.lastName,SUM(F.price)
FROM facts AS F, customers AS C
WHERE F.custID = C.custID
GROUP BY C.lastName



Early Materialization Example

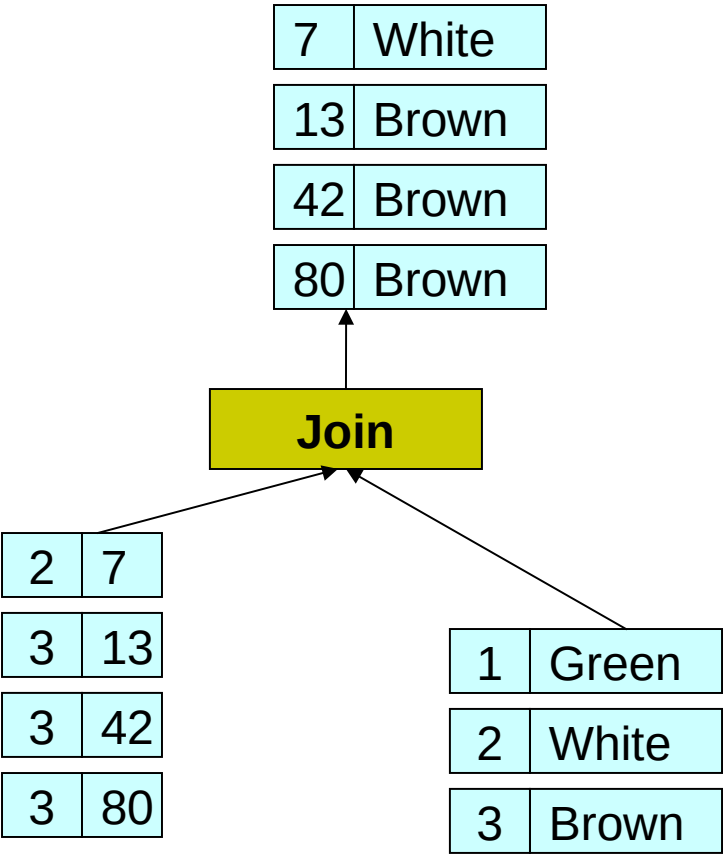
QUERY:
SELECT C.lastName,SUM(F.price)
FROM facts AS F, customers AS C
WHERE F.custID = C.custID
GROUP BY C.lastName

2	7
3	13
3	42
3	80

1	Green
2	White
3	Brown



Early Materialization Example



QUERY:

```
SELECT C.lastName,SUM(F.price)
FROM facts AS F, customers AS C
WHERE F.custID = C.custID
GROUP BY C.lastName
```



Late Materialization Example

QUERY:
SELECT C.lastName,SUM(F.price)
FROM facts AS F, customers AS C
WHERE F.custID = C.custID
GROUP BY C.lastName

(4,1,4)

12	6	2	7
1	1	3	13
11	2	1	42
1	1	3	80

prodID storeID quantity custID price

Facts

1	Green
2	White
3	Brown

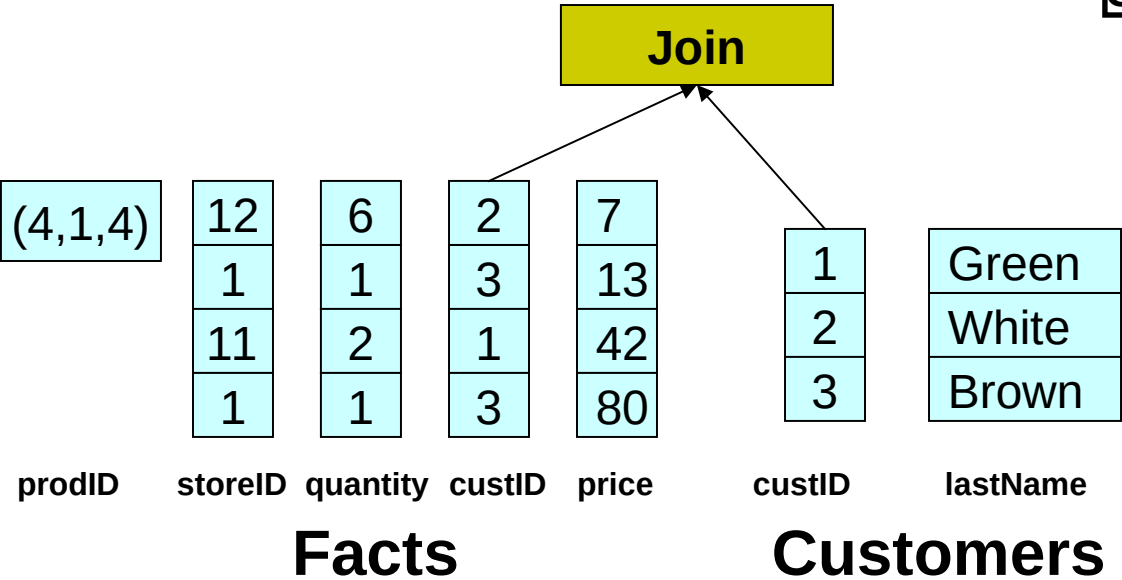
custID lastName

Customers



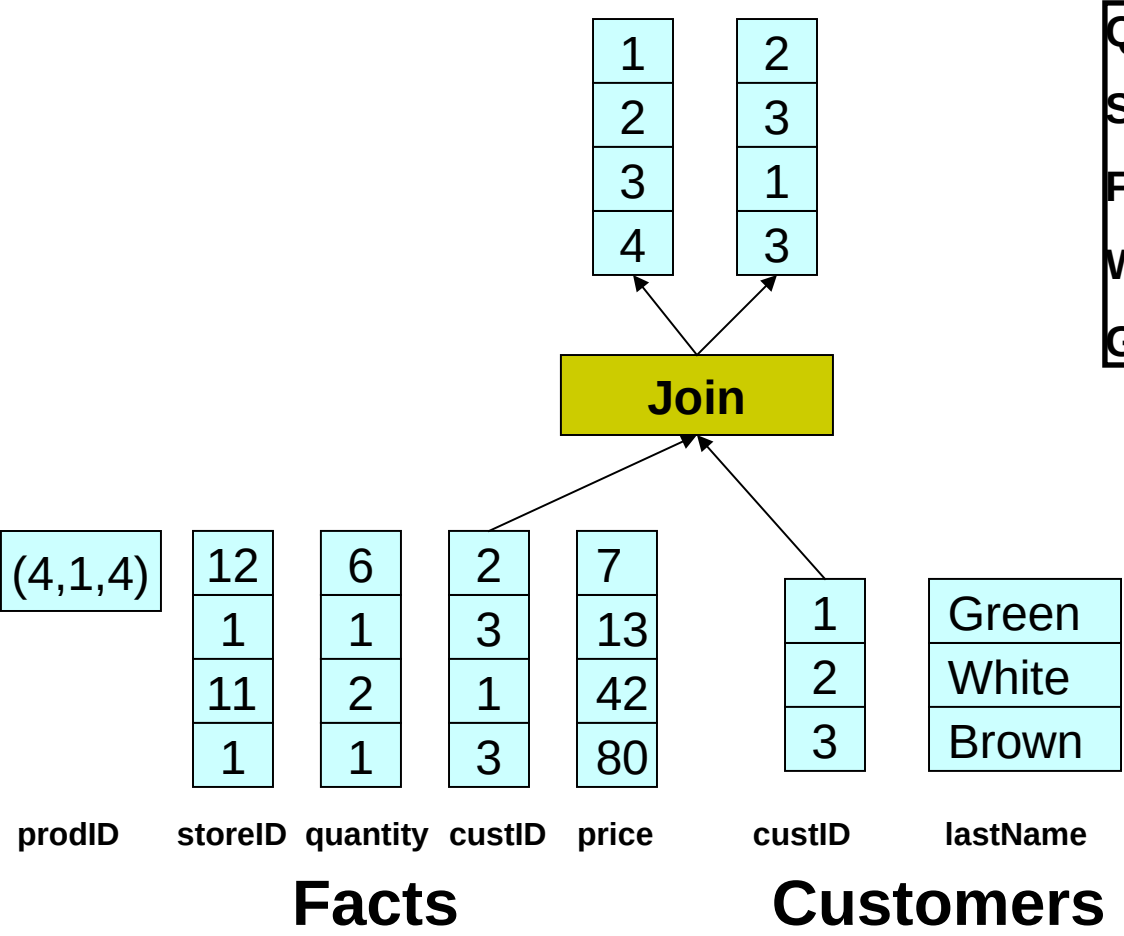
Late Materialization Example

QUERY:
SELECT C.lastName,SUM(F.price)
FROM facts AS F, customers AS C
WHERE F.custID = C.custID
GROUP BY C.lastName





Late Materialization Example

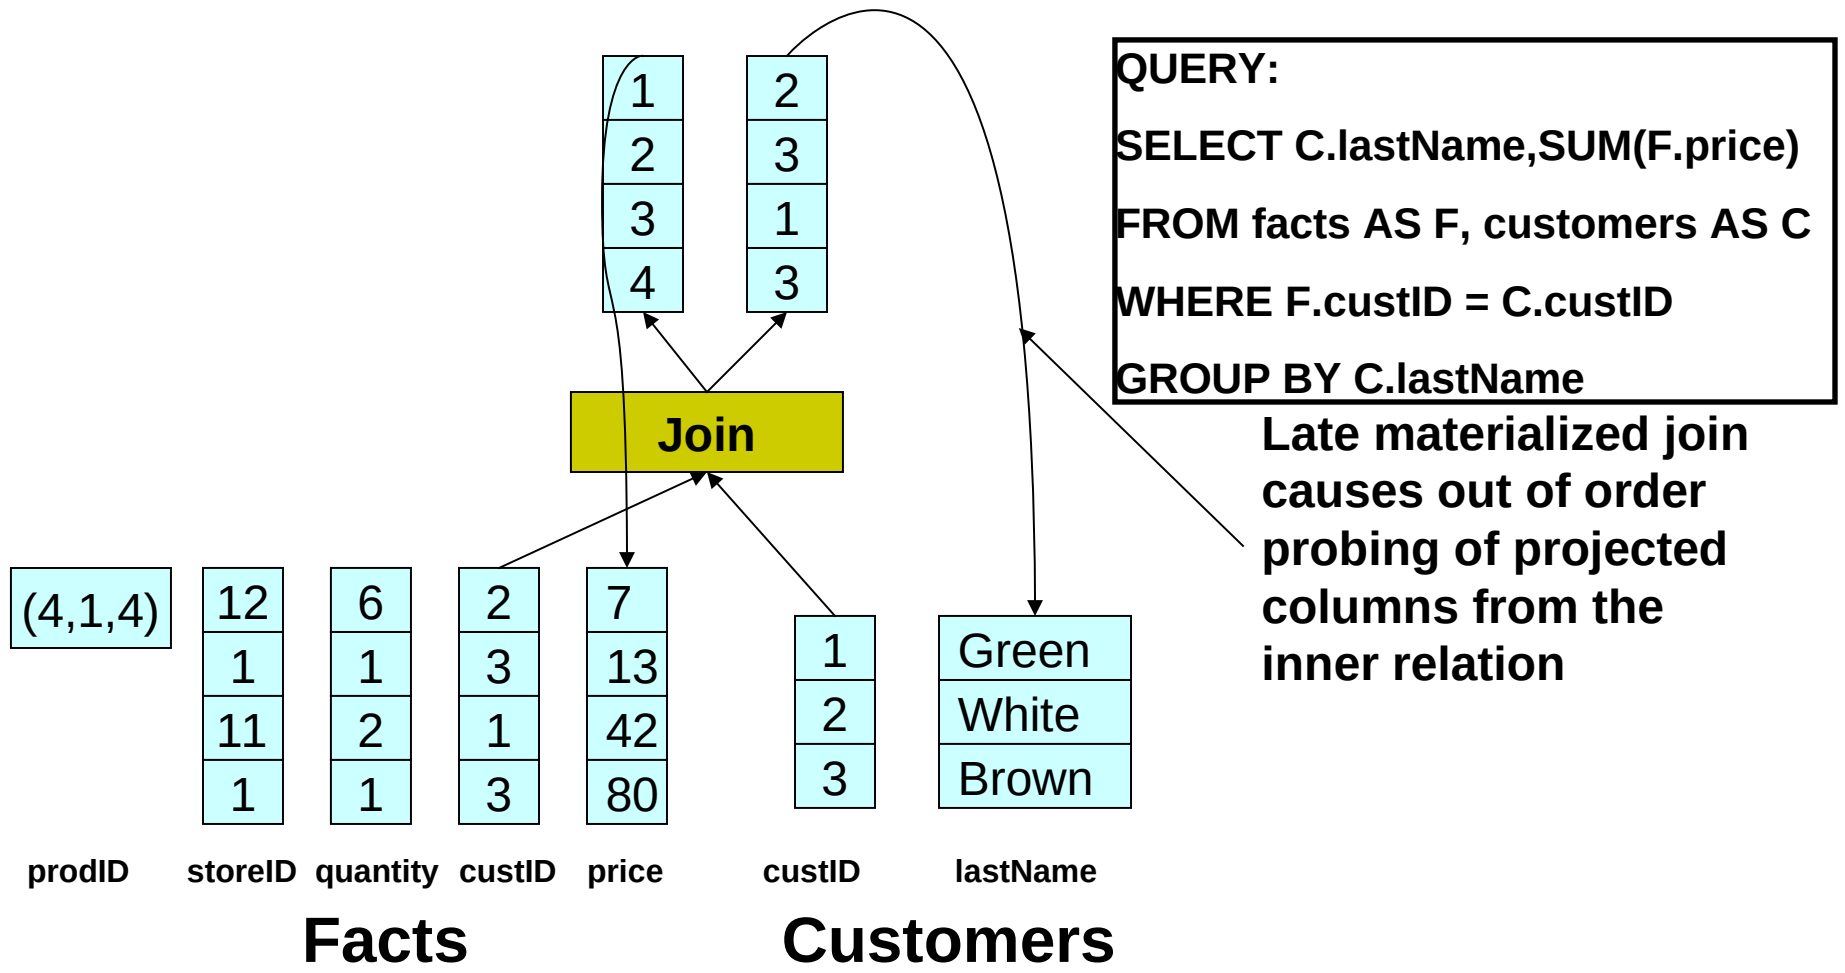


QUERY:

```
SELECT C.lastName,SUM(F.price)
FROM facts AS F, customers AS C
WHERE F.custID = C.custID
GROUP BY C.lastName
```



Late Materialization Example





Late Materialized Join Performance

- Naïve LM join (can be) about 2X slower than EM join on typical queries (due to random I/O)
 - This number is very dependent on
 - Amount of memory available
 - Number of projected attributes
 - Join cardinality
- But we can do better
 - Invisible Join
 - Jive/Flash Join
 - Radix cluster/decluster join



Invisible Join

“Column-Stores vs Row-Stores: How Different are They Really?” Abadi, Madden, and Hachem. SIGMOD 2008.

- Designed for typical joins when data is modeled using a star schema
 - One (big) “fact” table is joined with multiple (small) “dimension” tables
- Typical query (Star Schema Benchmark (SSBM)):

```
select c_nation, s_nation, d_year,  
       sum(lo_revenue) as revenue  
from lineorder, customer, supplier, date  
where lo_custkey = c_custkey  
      and lo_suppkey = s_suppkey  
      and lo_orderdate = d_datekey  
      and c_region = 'ASIA'  
      and s_region = 'ASIA'  
      and d_year >= 1992 and d_year <= 1997  
group by c_nation, s_nation, d_year  
order by d_year asc, revenue desc;
```



“Column-Stores vs Row-Stores: How Different are They Really?” Abadi, Madden, and Hachem. SIGMOD 2008.

Invisible Join

Apply “region = ‘Asia’” On Customer Table

custkey	region	nation	...
1	ASIA	CHINA	...
2	ASIA	INDIA	...
3	ASIA	INDIA	...
4	EUROPE	FRANCE	...



Hash Table (or bit-map)
Containing Keys 1, 2 and 3

Apply “region = ‘Asia’” On Supplier Table

suppkey	region	nation	...
1	ASIA	RUSSIA	...
2	EUROPE	SPAIN	...
3	ASIA	JAPAN	...



Hash Table (or bit-map)
Containing Keys 1, 3

Apply “year in [1992,1997]” On Date Table

dateid	year	...
01011997	1997	...
01021997	1997	...
01031997	1997	...




Hash Table Containing
Keys 01011997, 01021997,
and 01031997




“Column-Stores vs Row-Stores:
How Different are They Really?”
Abadi et. al. SIGMOD 2008.

Diagram illustrating the bitwise AND operation:

A		B		C
1		1		1
1		0		0
0		1		0
1		1		1
0		0		0
1		0		0
1		0		0

**Hash Table Containing
Keys 01011997,
01021997, and 01031997**







“Column-Stores vs Row-Stores: How Different are They Really?” Abadi, Madden, and Hachem. SIGMOD 2008.

Invisible Join

Apply “region = ‘Asia’” On Customer Table

custkey	region	nation	...
1	ASIA	CHINA	...
2	ASIA	INDIA	...
3	ASIA	INDIA	...
4	EUROPE	FRANCE	...



Hash Table (or bit-map)
Containing Keys 1, 2 and 3

Apply “region = ‘Asia’” On Supplier Table

suppkey	region	nation	...
1	ASIA	RUSSIA	...
2	EUROPE	SPAIN	...
3	ASIA	JAPAN	...



Hash Table (or bit-map)
Containing Keys 1, 3

Apply “year in [1992,1997]” On Date Table

dateid	year	...
01011997	1997	...
01021997	1997	...
01031997	1997	...



Hash Table Containing
Keys 01011997, 01021997,
and 01031997



“Column-Stores vs Row-Stores: How Different are They Really?” Abadi, Madden, and Hachem. SIGMOD 2008.

Invisible Join

Apply “region = ‘Asia’” On Customer Table

custkey	region	nation	...
1	ASIA	CHINA	...
2	ASIA	INDIA	...
3	ASIA	INDIA	...
4	EUROPE	FRANCE	...



~~Hash Table (or bit-map)
Containing Keys 1, 2 and 3~~

Range [1-3]
(between-predicate rewriting)

Apply “region = ‘Asia’” On Supplier Table

suppkey	region	nation	...
1	ASIA	RUSSIA	...
2	EUROPE	SPAIN	...
3	ASIA	JAPAN	...



Hash Table (or bit-map)
Containing Keys 1, 3

Apply “year in [1992,1997]” On Date Table

dateid	year	...
01011997	1997	...
01021997	1997	...
01031997	1997	...




Hash Table Containing
Keys 01011997, 01021997,
and 01031997


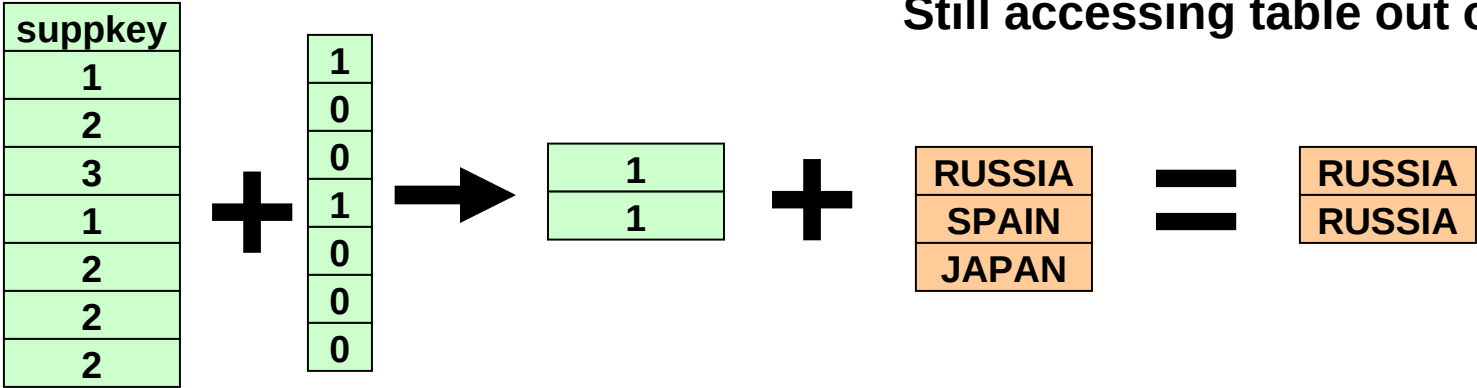


Invisible Join

- **Bottom Line**
 - Many data warehouses model data using star/snowflake schemes
 - Joins of one (fact) table with many dimension tables is common
 - Invisible join takes advantage of this by making sure that the table that can be accessed in position order is the fact table for each join
 - Position lists from the fact table are then intersected (in position order)
 - This reduces the amount of data that must be accessed out of order from the dimension tables
 - “Between-predicate rewriting” trick not relevant for this discussion




Still accessing table out of order





Jive/Flash Join




Still accessing table out of order

“Fast Joins using Join Indices”. Li and Ross, VLDBJ 8:1-24, 1999.

“Query Processing Techniques for Solid State Drives”. Tsirogiannis, Harizopoulos et. al. SIGMOD 2009.


Jive/Flash Join





Jive/Flash Join

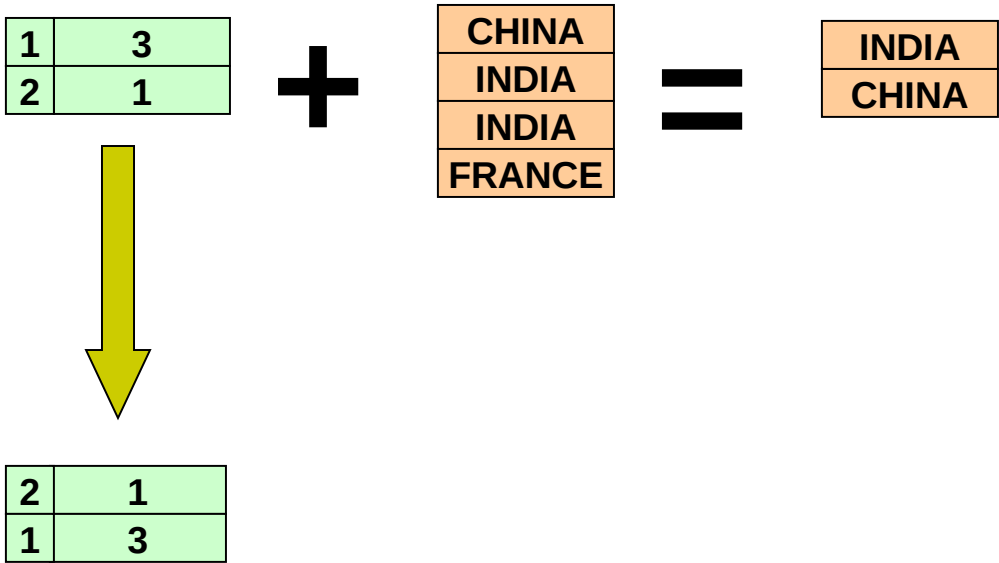
- 1. Add column with dense ascending integers from 1





Jive/Flash Join

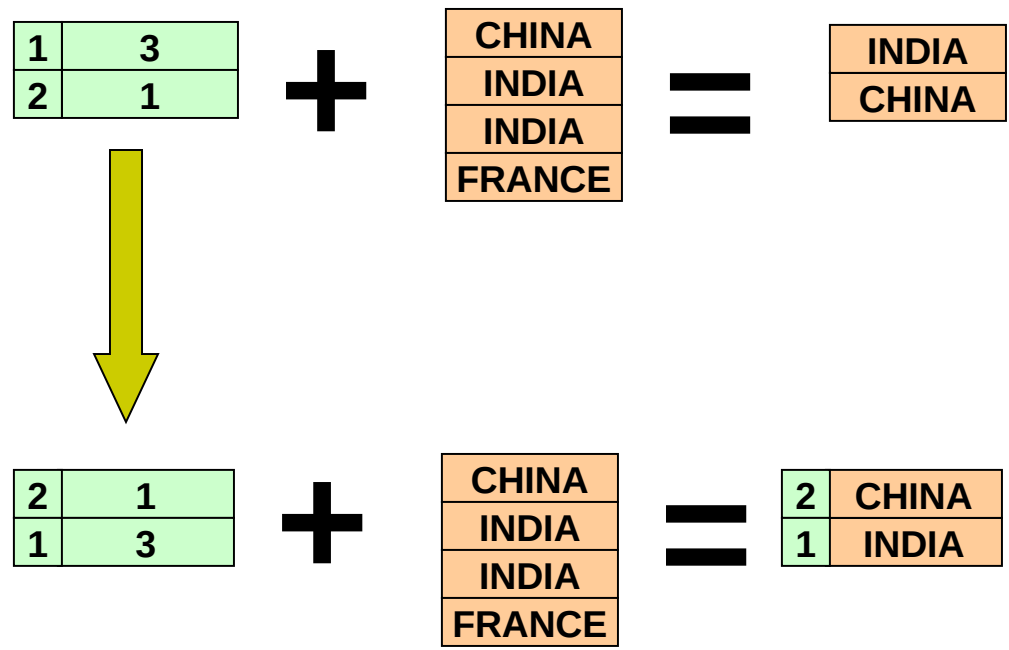
- 1. Add column with dense ascending integers from 1
- 2. Sort new position list by second column





Jive/Flash Join

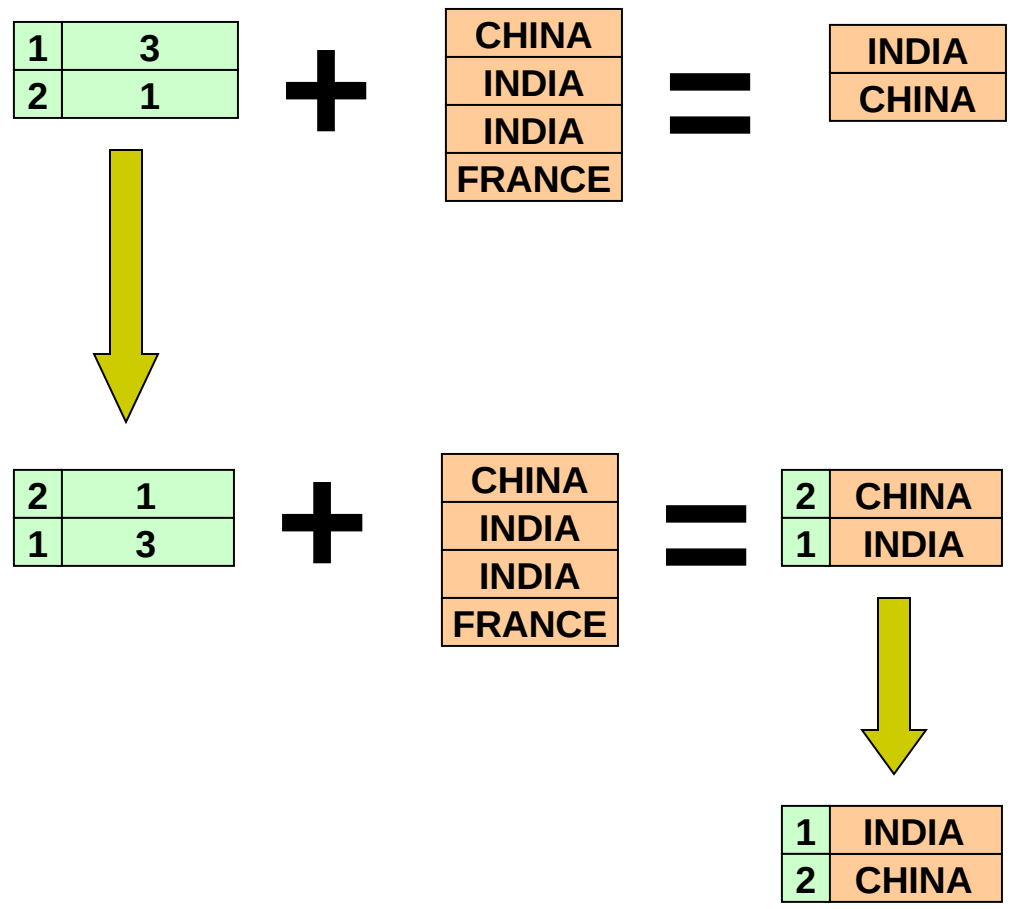
- 1. Add column with dense ascending integers from 1
- 2. Sort new position list by second column
- 3. Probe projected column in order using new sorted position list, keeping first column from position list around





Jive/Flash Join

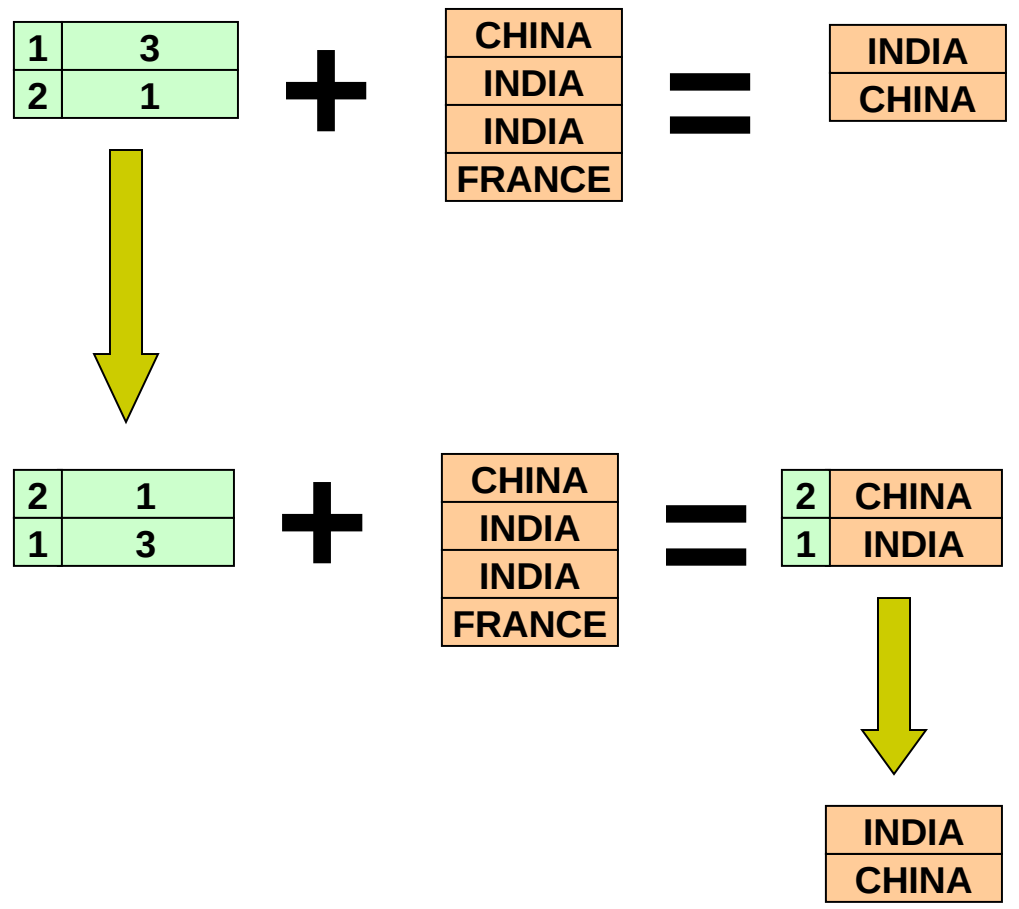
- 1. Add column with dense ascending integers from 1
- 2. Sort new position list by second column
- 3. Probe projected column in order using new sorted position list, keeping first column from position list around
- 4. Sort new result by first column





Jive/Flash Join

- 1. Add column with dense ascending integers from 1
- 2. Sort new position list by second column
- 3. Probe projected column in order using new sorted position list, keeping first column from position list around
- 4. Sort new result by first column
- 5. Remove first column





Jive/Flash Join

- **Bottom Line**
 - Instead of probing projected columns from inner table out of order:
 - Sort join index
 - Probe projected columns in order
 - Sort result using an added column
 - LM vs EM tradeoffs:
 - LM has the extra sorts (EM accesses all columns in order)
 - LM only has to fit join columns into memory (EM needs join columns and all projected columns)
 - Results in big memory and CPU savings (see part 3 for why there is CPU savings)
 - LM only has to materialize relevant columns
 - In many cases LM advantages outweigh disadvantages
 - LM would be a clear winner if not for those pesky sorts ... can we do better?

ADM: Literature

- **Column-Oriented Database Systems (2/6) - Selected Execution Techniques**

- Compression

- “Compressing Relations and Indexes”. Goldstein, Ramakrishnan, Shaft. ICDE’98.
- “Query optimization in compressed database systems”. Chen, Gehrke, Korn. SIGMOD’01.
- “Super-Scalar RAM-CPU Cache Compression”. Zukowski, Heman, Nes, Boncz. ICDE’06.
- “Integrating Compression and Execution in Column-Oriented Database Systems”. Abadi, Madden, Ferreira. SIGMOD’06.
- “Improved Word-Aligned Binary Compression for Text Indexing”. Ahn, Moffat. TKDE’06.

- Tuple Materialization

- “Materialization Strategies in a Column-Oriented DBMS”. Abadi, Myers, DeWitt, Madden. ICDE’07.
- “Column-Stores vs Row-Stores: How Different are They Really?”. Abadi, Madden, Hachem. SIGMOD’08.
- “Query Processing Techniques for Solid State Drives”. Tsirogiannis, Harizopoulos Shah, Wiener, Graefe. SIGMOD’09.
- “Self-organizing tuple reconstruction in column-stores”. Idreos, Manegold, Kersten. SIGMOD’09.

- Join

- “Fast Joins using Join Indices”. Li and Ross. VLDBJ 8:1-24, 1999.

ADM: Agenda

- 07.09.2022: Lecture 1: **Introduction**
- 14.09.2022: Lecture 2: **SQL Recap**
(plus Assignment 1 [in groups; 3 weeks]: TPC-H benchmark)
- 21.09.2022: Lecture 3: **Column-Oriented Database Systems (1/6) - Motivation & Basic Concepts**
- 28.09.2022: Lecture 4: **Column-Oriented Database Systems (2a/6) - Selected Execution Techniques (1/2)**
- 05.10.2022: Lecture 5: **Column-Oriented Database Systems (2b/6) - Selected Execution Techniques (2/2)**
(plus Assignment 2 [in groups; 3 weeks]: Compression techniques)
- 12.10.2022: Lecture 6: **Column-Oriented Database Systems (3/6) - Cache Conscious Joins**
- 19.10.2022: Lecture 7: **Column-Oriented Database Systems (4/6) - “Vectorized Execution”**
- 26.10.2022: **No lecture!**
- 02.11.2022: Lecture 8: **DuckDB: An embedded database for data science (1/2) (guest lecture & hands-on)**
(plus Assignment 3 [individual; 2 weeks]: Analysing NYC Cab dataset with DuckDB)
- 09.11.2022: Lecture 9: **DuckDB: An embedded database for data science (2/2) (guest lecture & hands-on)**
- 16.11.2022: Lecture 10: **Branch Misprediction & Predication**
(plus Assignment 4 [individual; 2 weeks]: Predication)
- 23.11.2022: Lecture 11: **Column-Oriented Database Systems (5/6) - Adaptive Indexing**
- 30.11.2022: Lecture 12: **Column-Oriented Database Systems (6/6) - Progressive Indexing**

ADM: Literature

- **Column-Oriented Database Systems (3/6) - Cache Conscious Joins**
 - “Cache Conscious Algorithms for Relational Query Processing”. Shatdal, Kant, Naughton. VLDB’94.
 - “Fast Joins using Join Indices”. Li and Ross. VLDBJ 8:1-24, 1999.
 - “Optimizing main-memory join on modern hardware”. Boncz, Manegold, Kersten, TKDE 14(4): 709-730, 2002.
 - “Database Architecture Optimized for the New Bottleneck: Memory Access”. Boncz, Manegold, Kersten. VLDB’99.
 - “What Happens During a Join? Dissecting CPU and Memory Optimization Effects”. Manegold, Boncz, Kersten. VLDB’00.
 - “Optimizing database architecture for the new bottleneck: memory access”. Manegold, Boncz, Kersten. VLDB J. 9(3): 231-246, 2000.
 - “Generic Database Cost Models for Hierarchical Memory Systems”. Manegold, Boncz, Kersten. VLDB’02.
 - “Cache-Conscious Radix-Declasser Projections”. Manegold, Boncz, Nes. VLDB’04.