

# **Advanced Data Management for Data Analysis**

*Stefan Manegold*

Data Management @ LIACS

Group leader Database Architectures  
Centrum Wiskunde & Informatica (CWI)  
Amsterdam

[s.manegold@liacs.leidenuniv.nl](mailto:s.manegold@liacs.leidenuniv.nl)

<http://www.cwi.nl/~manegold/>

# ADM: Agenda

- 07.09.2022: Lecture 1: **Introduction**
- 14.09.2022: Lecture 2: **SQL Recap**  
(plus Assignment 1 [in groups; 3 weeks]: TPC-H benchmark)
- 21.09.2022: Lecture 3: **Column-Oriented Database Systems (1/6) - Motivation & Basic Concepts**
- 28.09.2022: Lecture 4: **Column-Oriented Database Systems (2a/6) - Selected Execution Techniques (1/2)**
- 05.10.2022: Lecture 5: **Column-Oriented Database Systems (2b/6) - Selected Execution Techniques (2/2)**  
(plus Assignment 2 [in groups; 4 weeks]: Compression techniques)
- 12.10.2022: Lecture 6: **Column-Oriented Database Systems (3/6) - Cache Conscious Joins**
- 19.10.2022: Lecture 7: **Column-Oriented Database Systems (4/6) - “Vectorized Execution”**
- 26.10.2022: **No lecture!**
- 02.11.2022: Lecture 8: **DuckDB: An embedded database for data science (1/2) (guest lecture & hands-on)**  
(plus Assignment 3 [individual; 2 weeks]: Analysing NYC Cab dataset with DuckDB)
- 09.11.2022: Lecture 9: **DuckDB: An embedded database for data science (2/2) (guest lecture & hands-on)**
- 16.11.2022: Lecture 10: **Branch Misprediction & Predication**  
(plus Assignment 4 [individual; 2 weeks]: Predication)
- 23.11.2022: Lecture 11: **Column-Oriented Database Systems (5/6) - Adaptive Indexing**
- 30.11.2022: Lecture 12: **Column-Oriented Database Systems (6/6) - Progressive Indexing**

# ADM: Literature

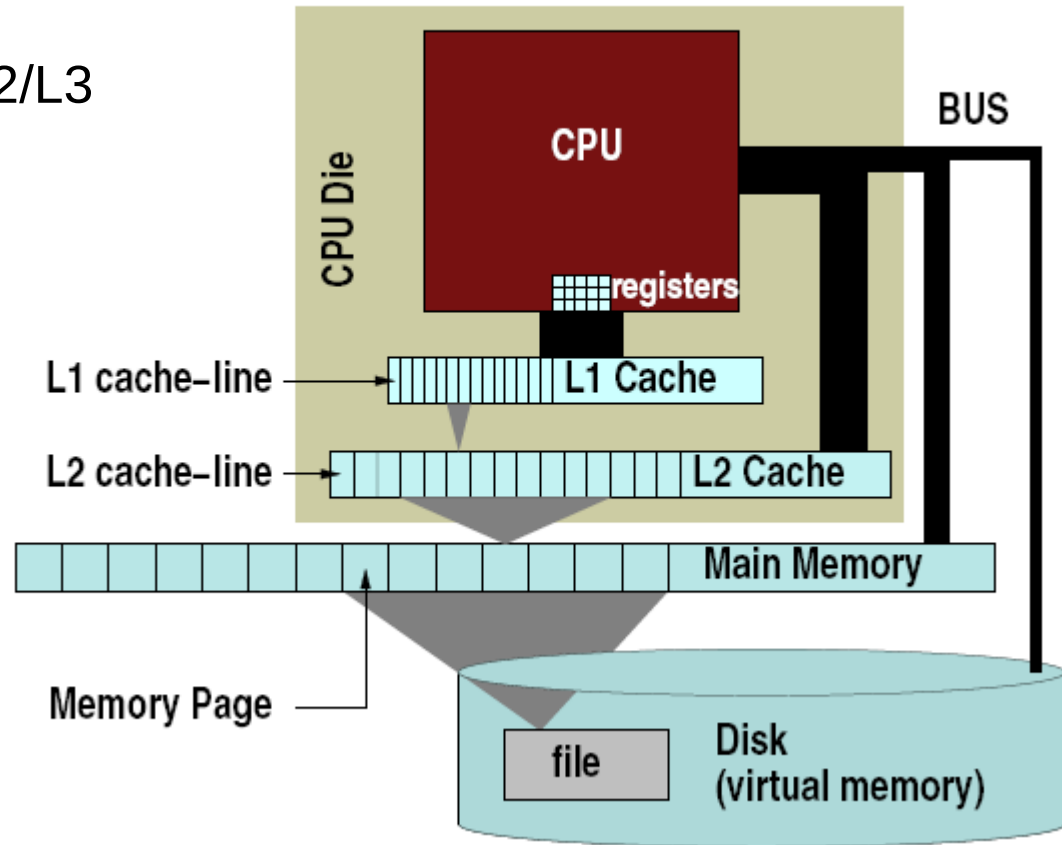
- **Column-Oriented Database Systems (4/6) - “Vectorized Execution”**
  - “MonetDB/X100: Hyper-Pipelining Query Execution”. Boncz, Zukowski, Nes. CIDR’05.
  - “Buffering Database Operations for Enhanced Instruction Cache Performance”. Zhou and Ross. SIGMOD’04.
  - “Block oriented processing of relational database operations in modern computer architectures”. Padmanabhan, Malkemus, Agarwal. ICDE’01.
  - “Balancing Vectorized Query Execution with Bandwidth Optimized Storage”. Zukowski. PhD Thesis. CWI 2008.

# CPU Architecture



## Elements:

- Storage
  - CPU caches L1/L2/L3
- Registers
- Execution Unit(s)
  - Pipelined
  - SIMD



# CPU Metrics

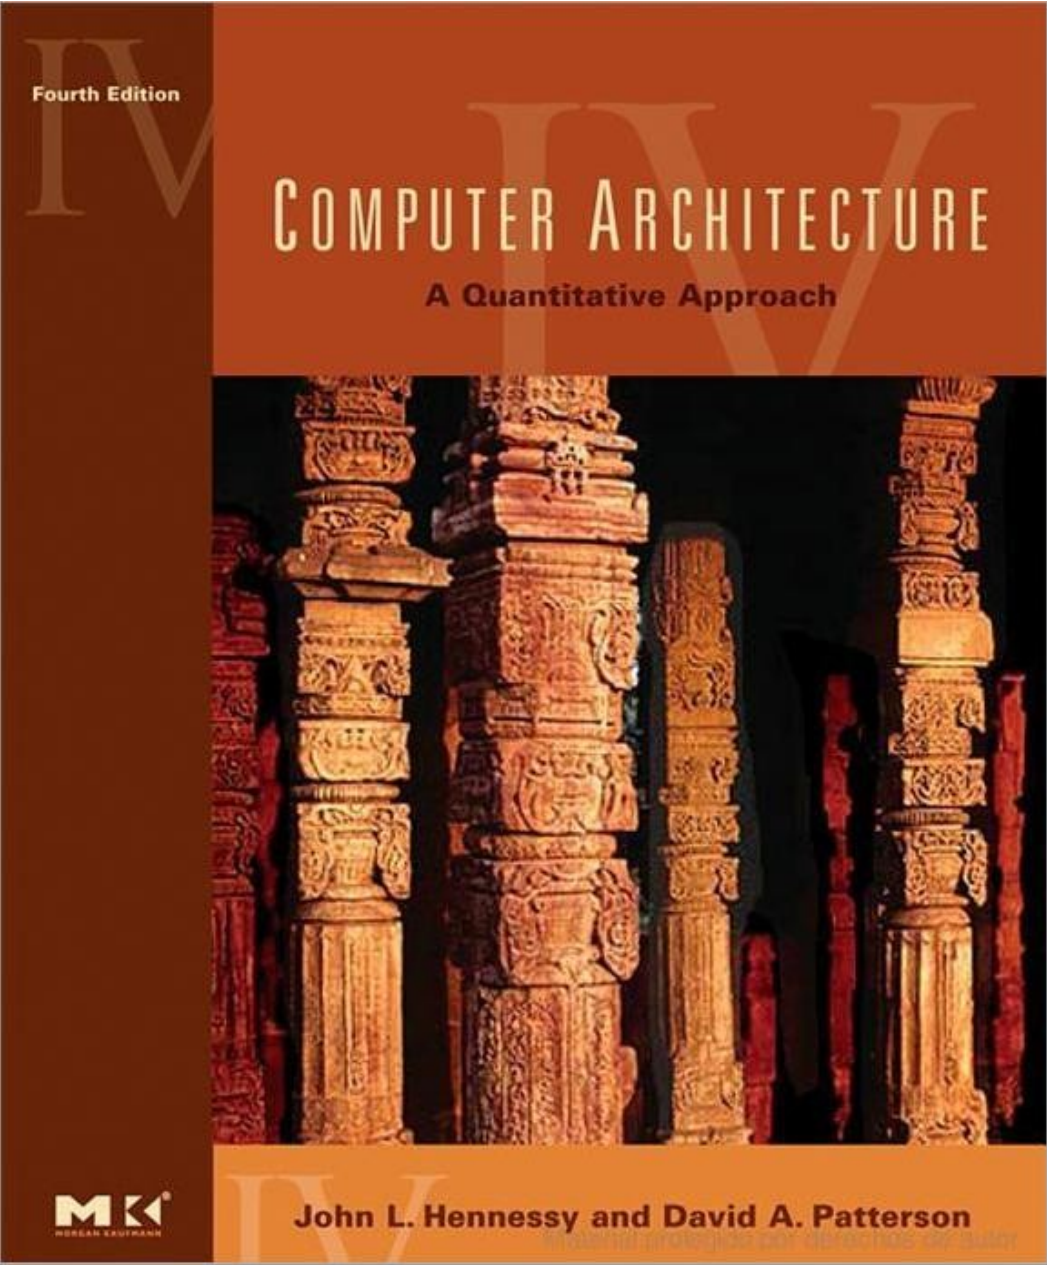


Processor	16-bit address/ bus, micro-coded	32-bit address/ bus, micro-coded	5-stage pipeline, on-chip I&D caches FPU	2-way super-scalar, 64-bit bus	Out-of-order, 3-way super-scalar	Out-of-order, super-pipelined, on-chip L2 cache	Multi-core
Product	80286	80386	80486	Pentium	PentiumPro	Pentium4	CoreDuo
Year	1982	1985	1989	1993	1997	2001	2006
Transistors (thousands)	134	275	1,200	3,100	5,500	42,000	151,600
Latency (clocks)	6	5	5	5	10	22	12
Bus width (bits)	16	32	32	64	64	64	64
Clock rate (MHz)	12.5	16	25	66	200	1500	2333
Bandwidth (MIPS)	2	6	25	132	600	4500	21000
Latency (ns)	320	313	200	76	50	15	5

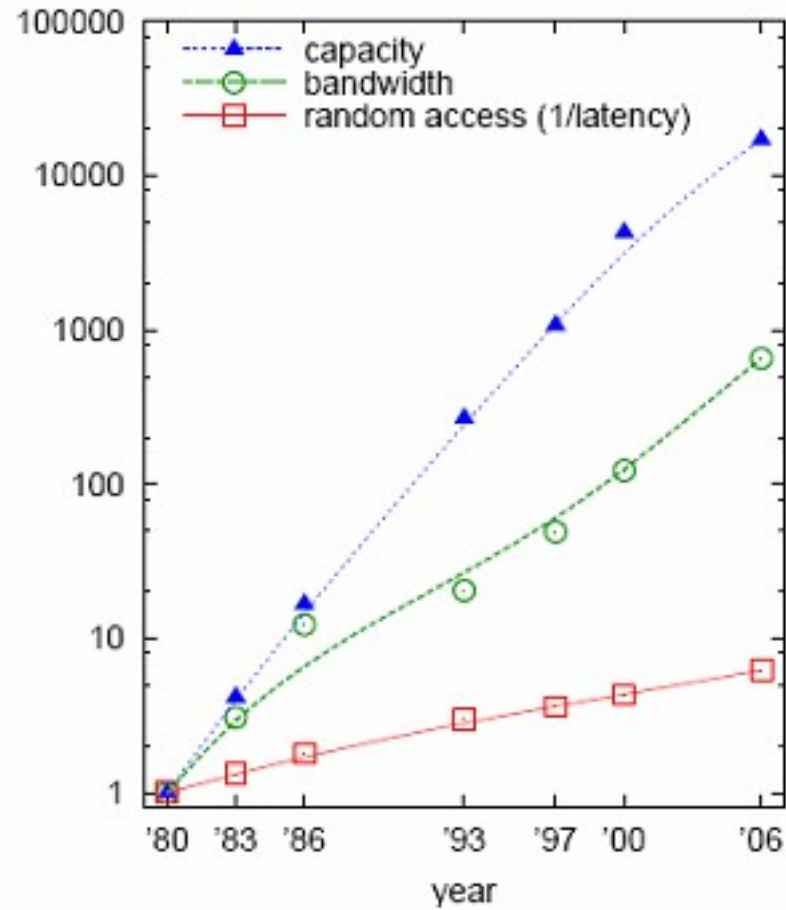
# CPU Metrics



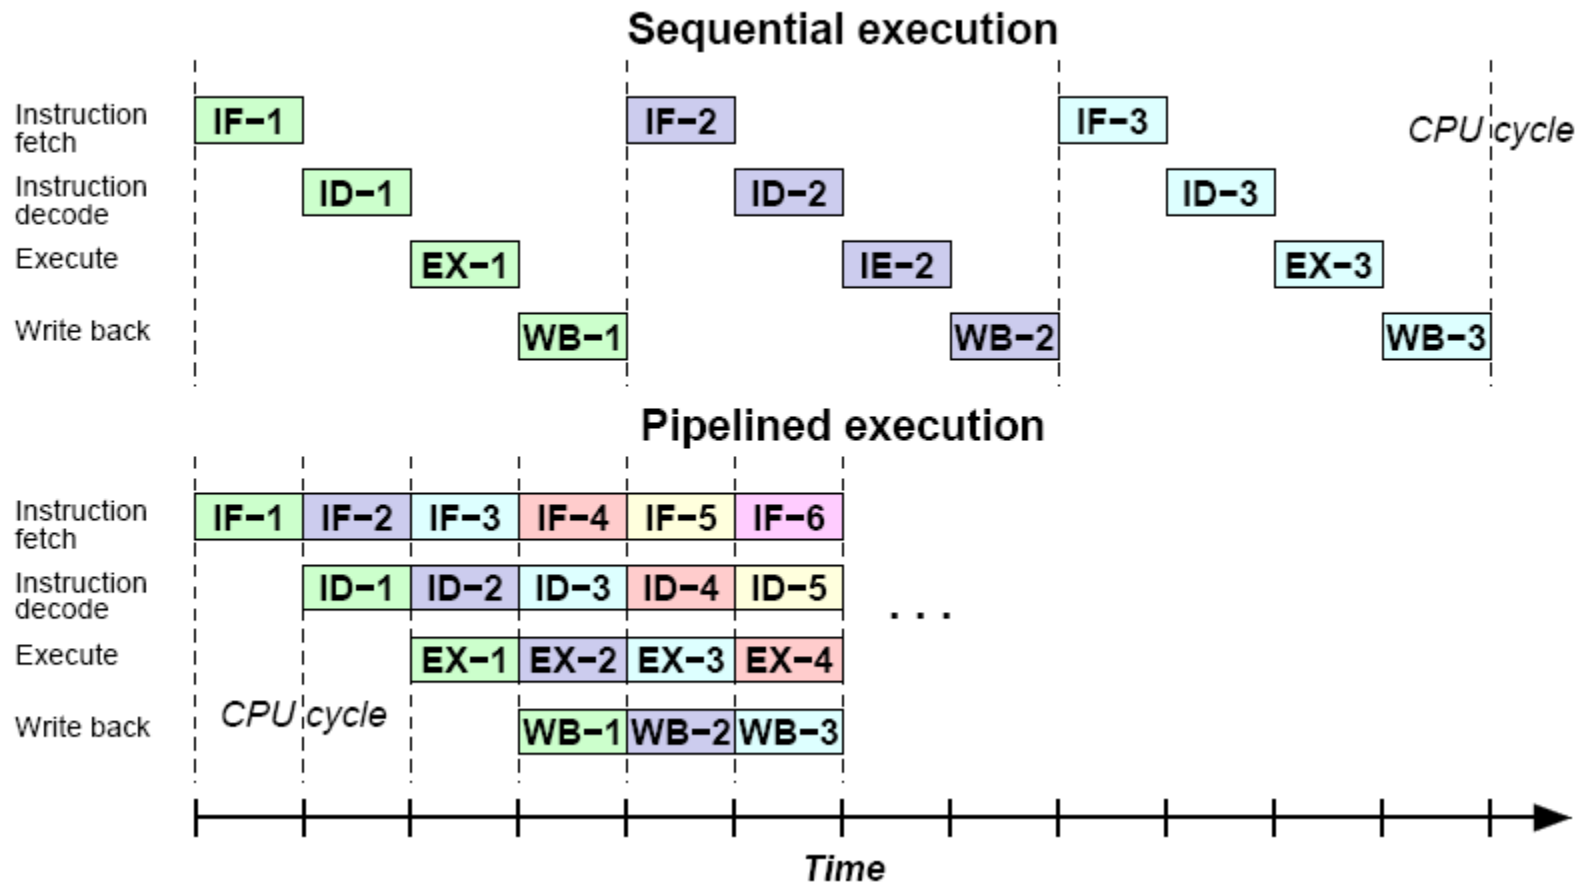
Processor	16-bit address/, bus, micro-coded	32 b mi co	core
Product	80286	80	Duo
Year	1982	19	6
Transistors (thousands)	134	2	300
Latency (clocks)	6		
Bus width (bits)	16	3	
Clock rate (MHz)	12.5	1	3
Bandwidth (MIPS)	2		00
Latency (ns)	320	3	



# DRAM Metrics



# Super-Scalar Execution (pipelining)



(See also [https://en.wikipedia.org/wiki/Instruction\\_pipelining](https://en.wikipedia.org/wiki/Instruction_pipelining) )

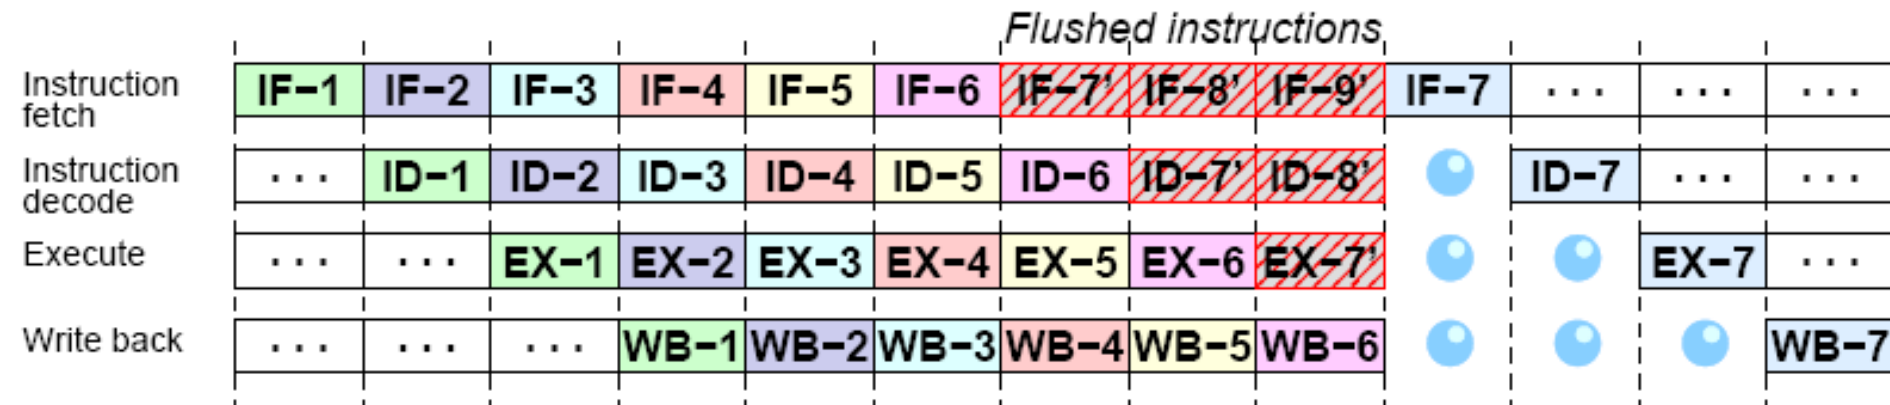




# Hazards

- Data hazards
  - Dependencies between instructions
  - L1 data cache misses
- Control Hazards
  - Branch mispredictions
  - Computed branches (late binding)
  - L1 instruction cache misses

Result: bubbles in the pipeline



Out-of-order execution addresses data hazards

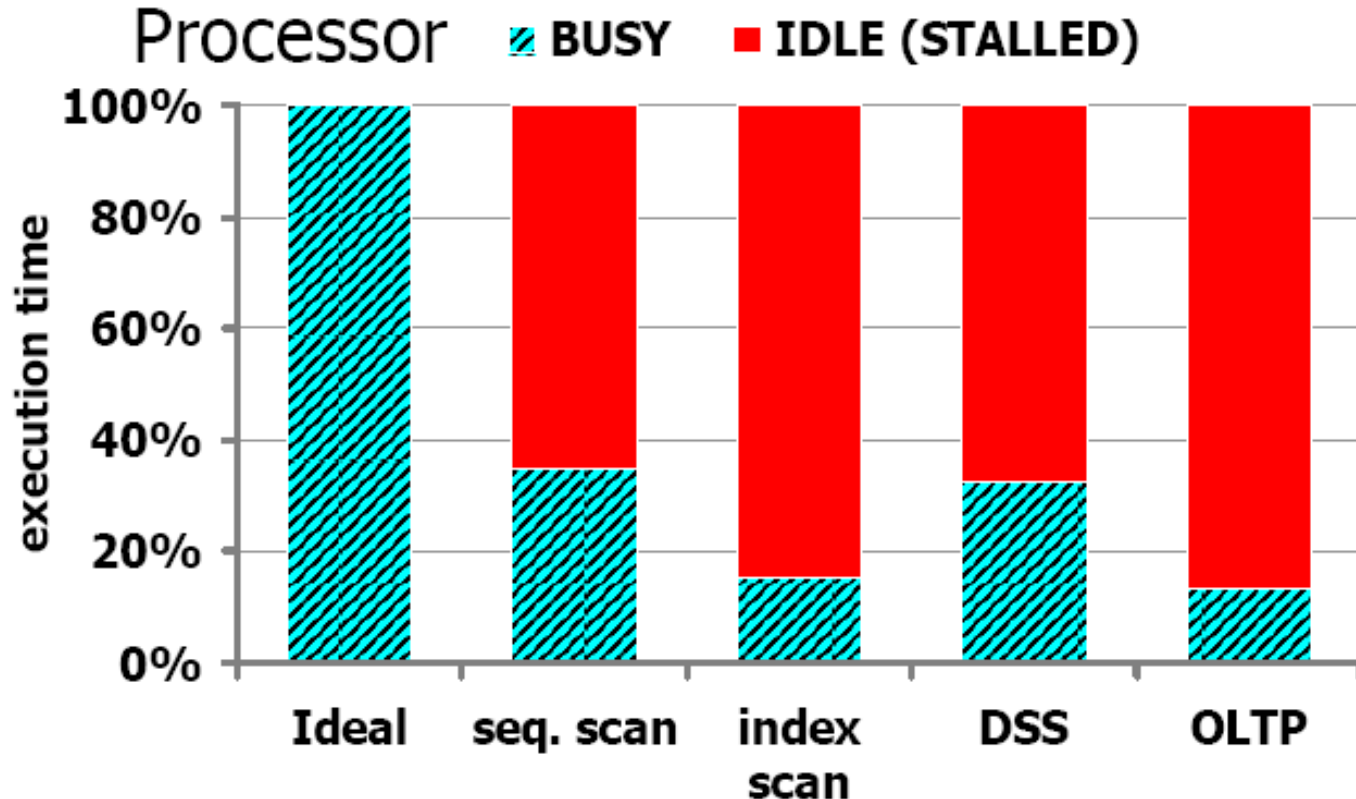
- control hazards typically more expensive

(See also [https://en.wikipedia.org/wiki/Instruction\\_pipelining](https://en.wikipedia.org/wiki/Instruction_pipelining) )



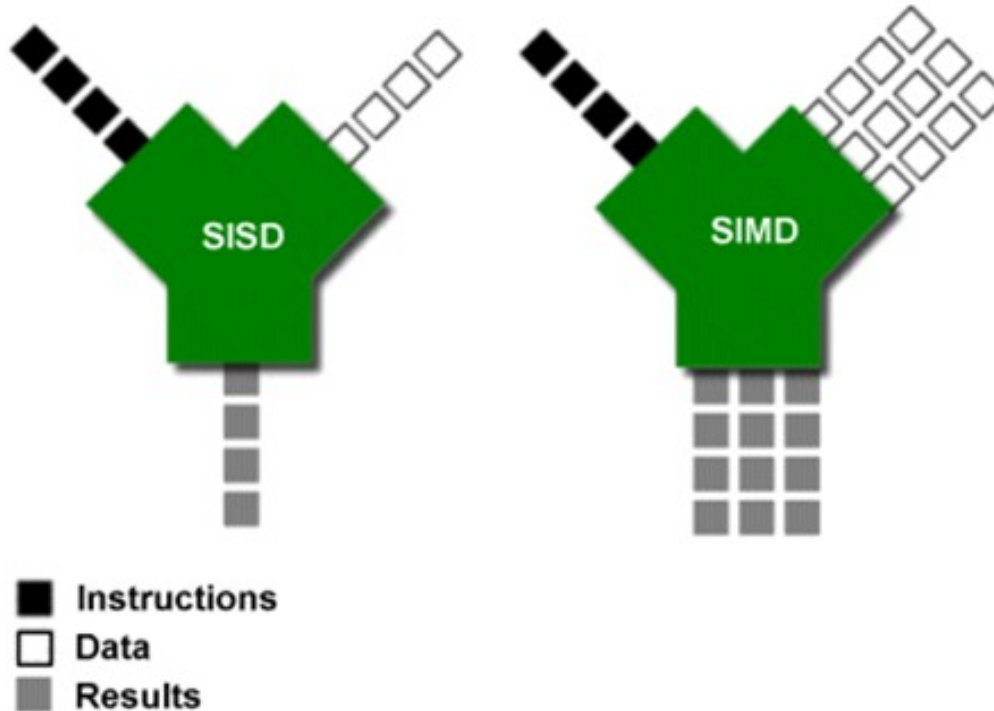
# Database Architecture causes Hazards

- DB workload execution on a modern computer



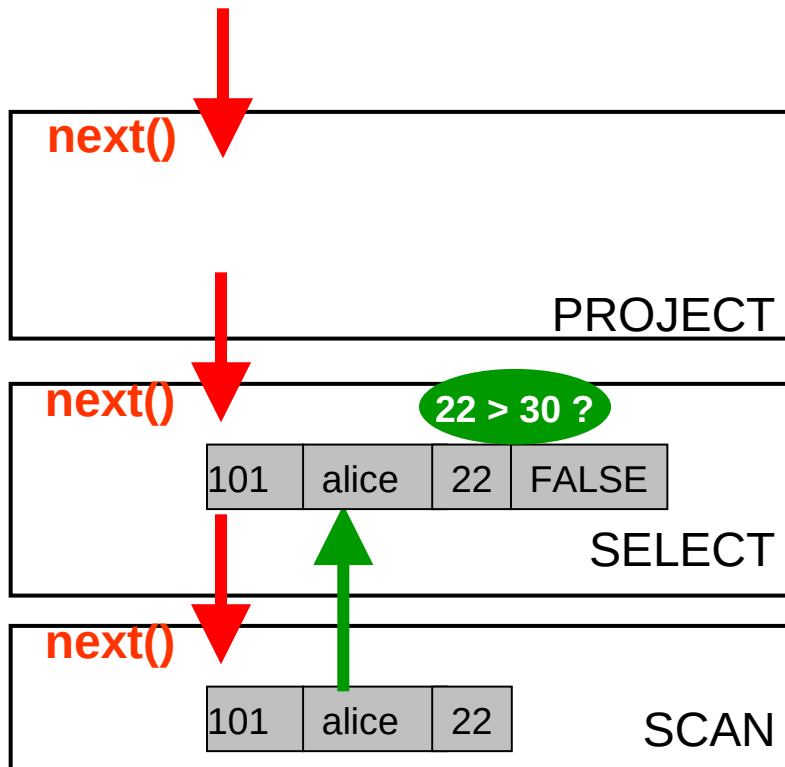
“DBMSs On A Modern Processor: Where Does Time Go? ”  
Ailamaki, DeWitt, Hill, Wood, VLDB’99

# SIMD



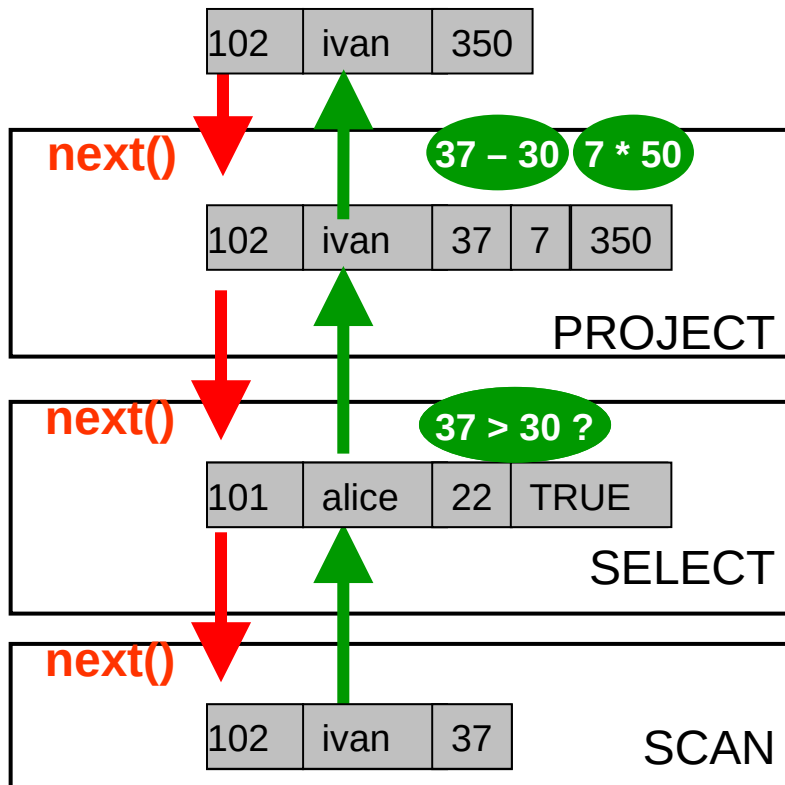
- Single Instruction Multiple Data
  - Same operation applied on a vector of values
  - MMX: 64 bits, SSE: 128bits, AVX: 256bits
  - SSE, e.g. multiply 8 short integers

# A Look at the Query Pipeline



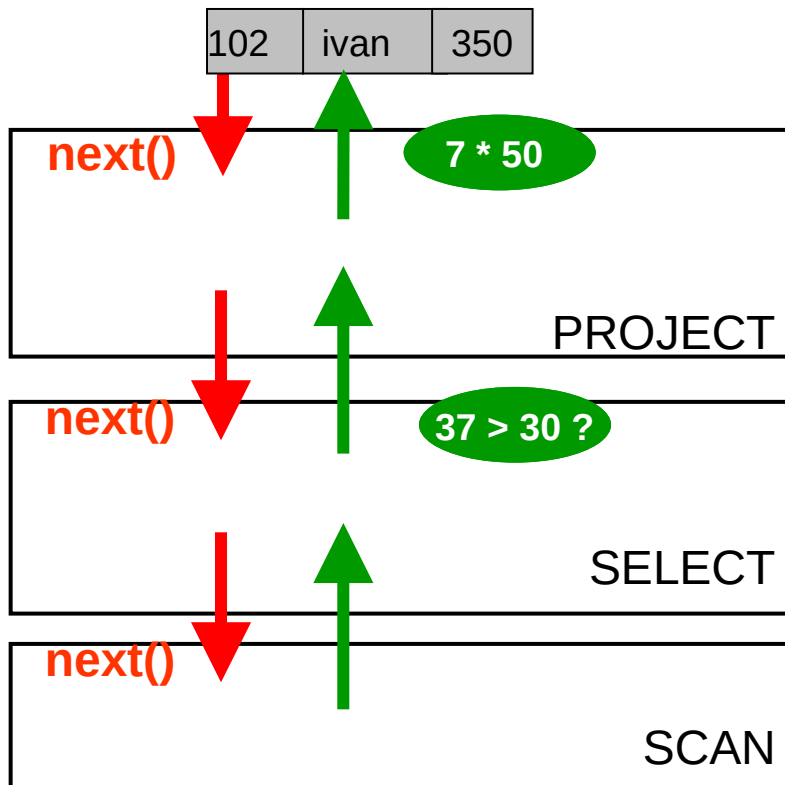
```
SELECT id, name
      (age-30)*50 AS bonus
FROM   employee
WHERE  age > 30
```

# A Look at the Query Pipeline



```
SELECT id, name
      (age-30)*50 AS bonus
FROM   employee
WHERE  age > 30
```

# A Look at the Query Pipeline



## Operators

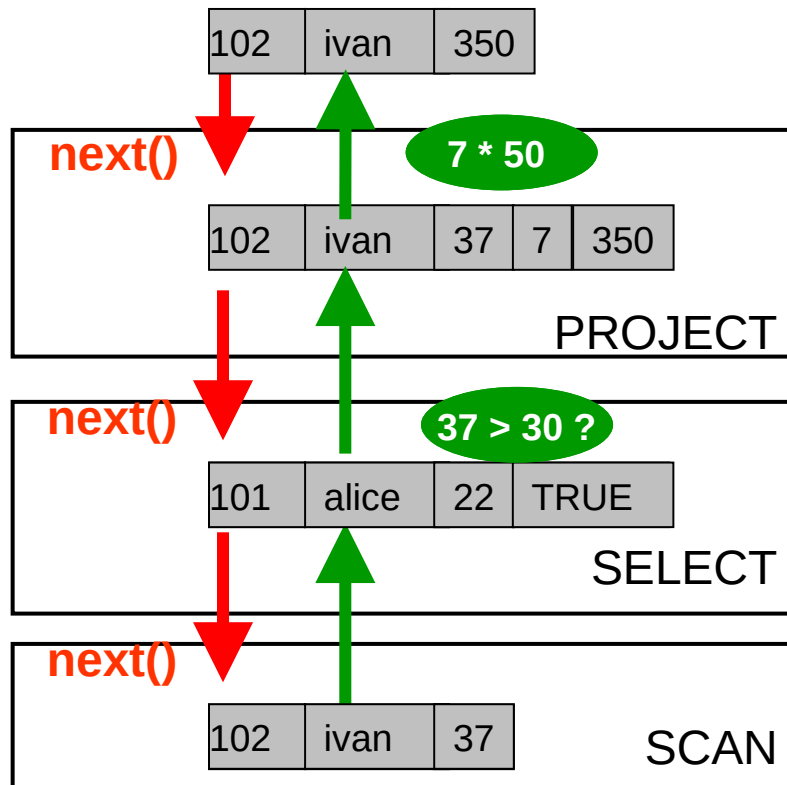
Iterator interface

-open()

-**next()**: tuple

-close()

# A Look at the Query Pipeline



## Primitives

Provide computational functionality

All arithmetic allowed in expressions,  
e.g. Multiplication

$7 * 50$

`mult(int,int) → int`



# Database Architecture causes Hazards

- Data hazards

- Dependencies between instructions
- L1 data cache misses

SIMD

Out-of-order Execution

work on one tuple at a time

Large Tree/Hash Structures

Code footprint of all operators in query plan exceeds L1 cache

Data-dependent conditions

Next() late binding method calls

Tree, List, Hash traversal

Complex NSM record navigation

- Control Hazards

- Branch mispredictions
- Computed branches (late binding)
- L1 instruction cache misses

PROJECT

SELECT

SCAN

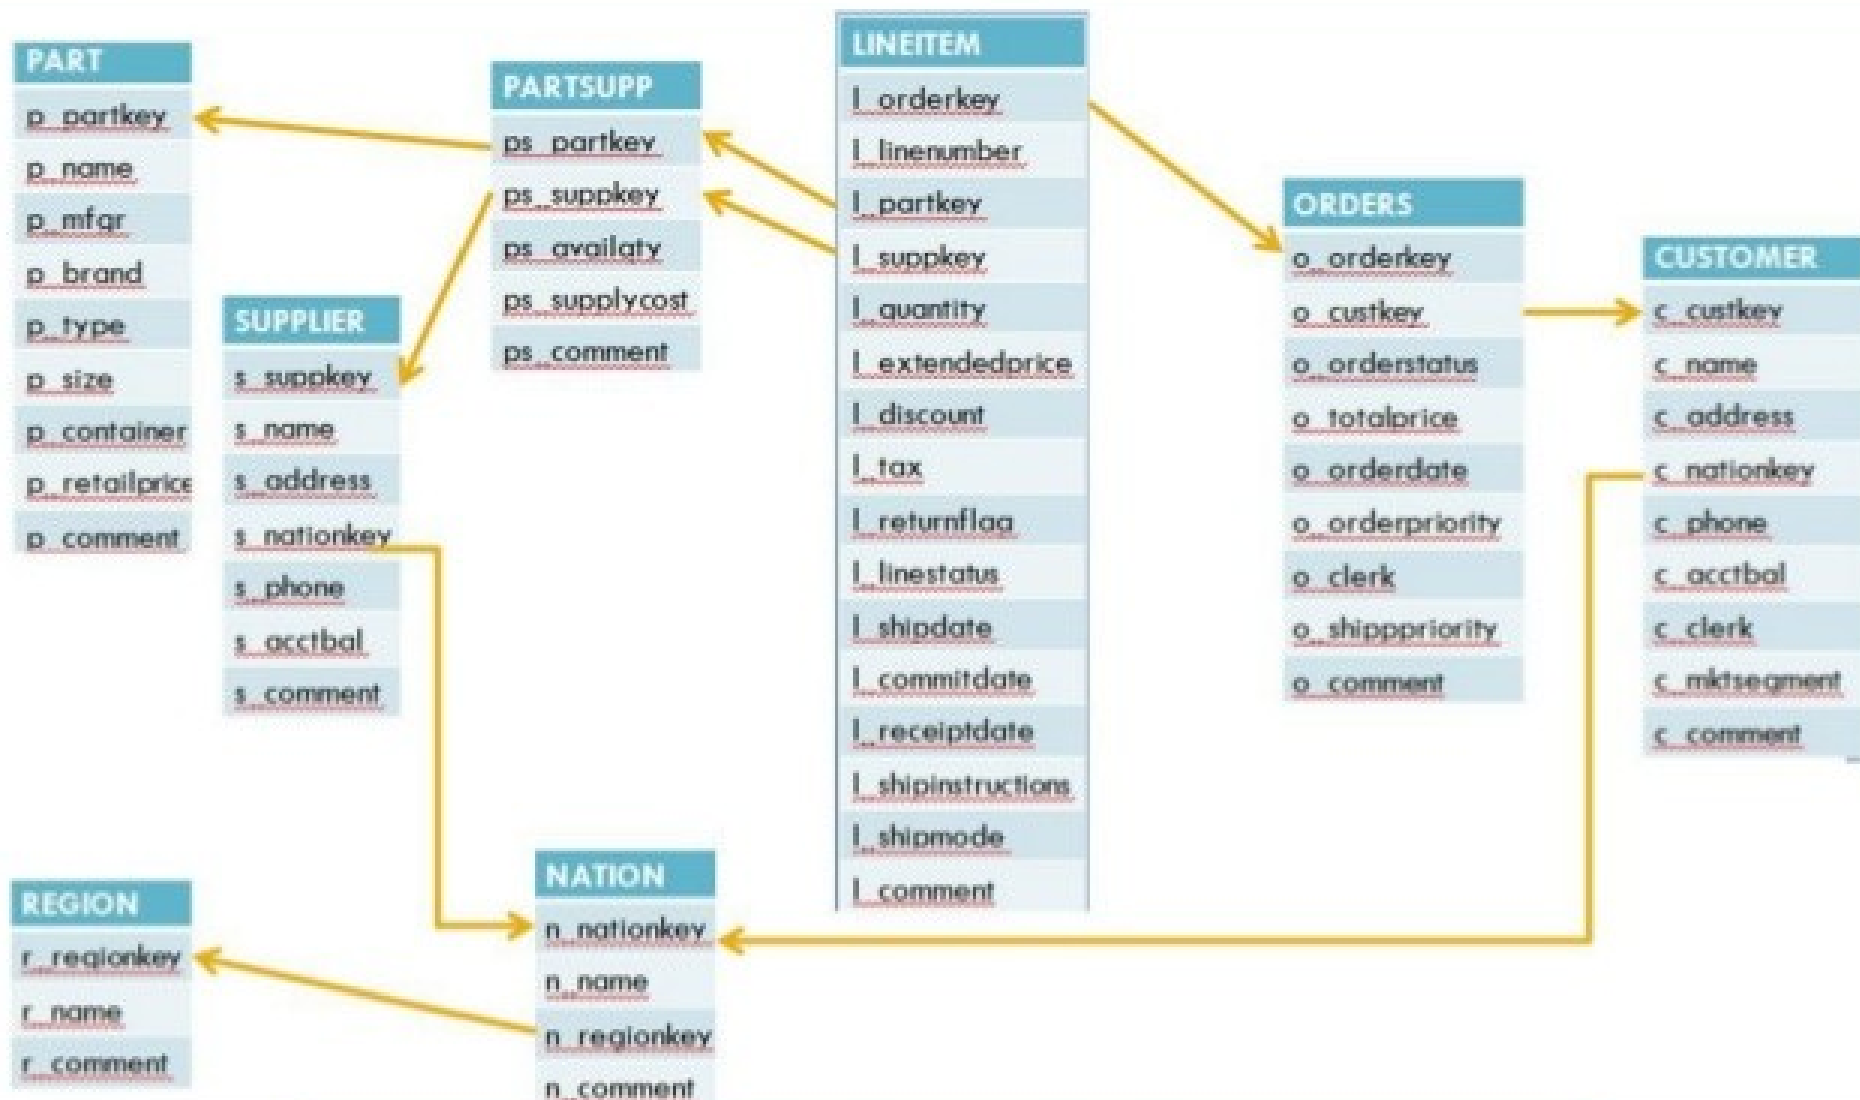
## Operators

Iterator interface

-open()  
-next(): tuple  
-close()



# TPC-H Database Schema



# TPC-H Query 1



```
select
    l_returnflag,
    l_linestatus,
    sum(l_quantity) as sum_qty,
    sum(l_extendedprice) as sum_base_price,
    sum(l_extendedprice * (1 - l_discount)) as sum_disc_price,
    sum(l_extendedprice * (1 - l_discount) * (1 + l_tax)) as sum_charge,
    avg(l_quantity) as avg_qty,
    avg(l_extendedprice) as avg_price,
    avg(l_discount) as avg_disc,
    count(*) as count_order
from
    lineitem
where
    l_shipdate <= date '1998-12-01' - interval '90' day (3)
group by
    l_returnflag,
    l_linestatus
order by
    l_returnflag,
    l_linestatus;
```



# DBMS Computational Efficiency

TPC-H 1GB, query 1

- selects 98% of fact table, computes net prices and aggregates all
- Results:
  - C program: ?
  - MySQL: 26.2s
  - DBMS “X”: 28.1s

“MonetDB/X100: Hyper-Pipelining Query Execution ” Boncz, Zukowski, Nes, CIDR’05



# DBMS Computational Efficiency

TPC-H 1GB, query 1

- selects 98% of fact table, computes net prices and aggregates all
- Results:
  - C program: **0.2s**
  - MySQL: 26.2s
  - DBMS “X”: 28.1s

“MonetDB/X100: Hyper-Pipelining Query Execution ” Boncz, Zukowski, Nes, CIDR’05

# a column-store

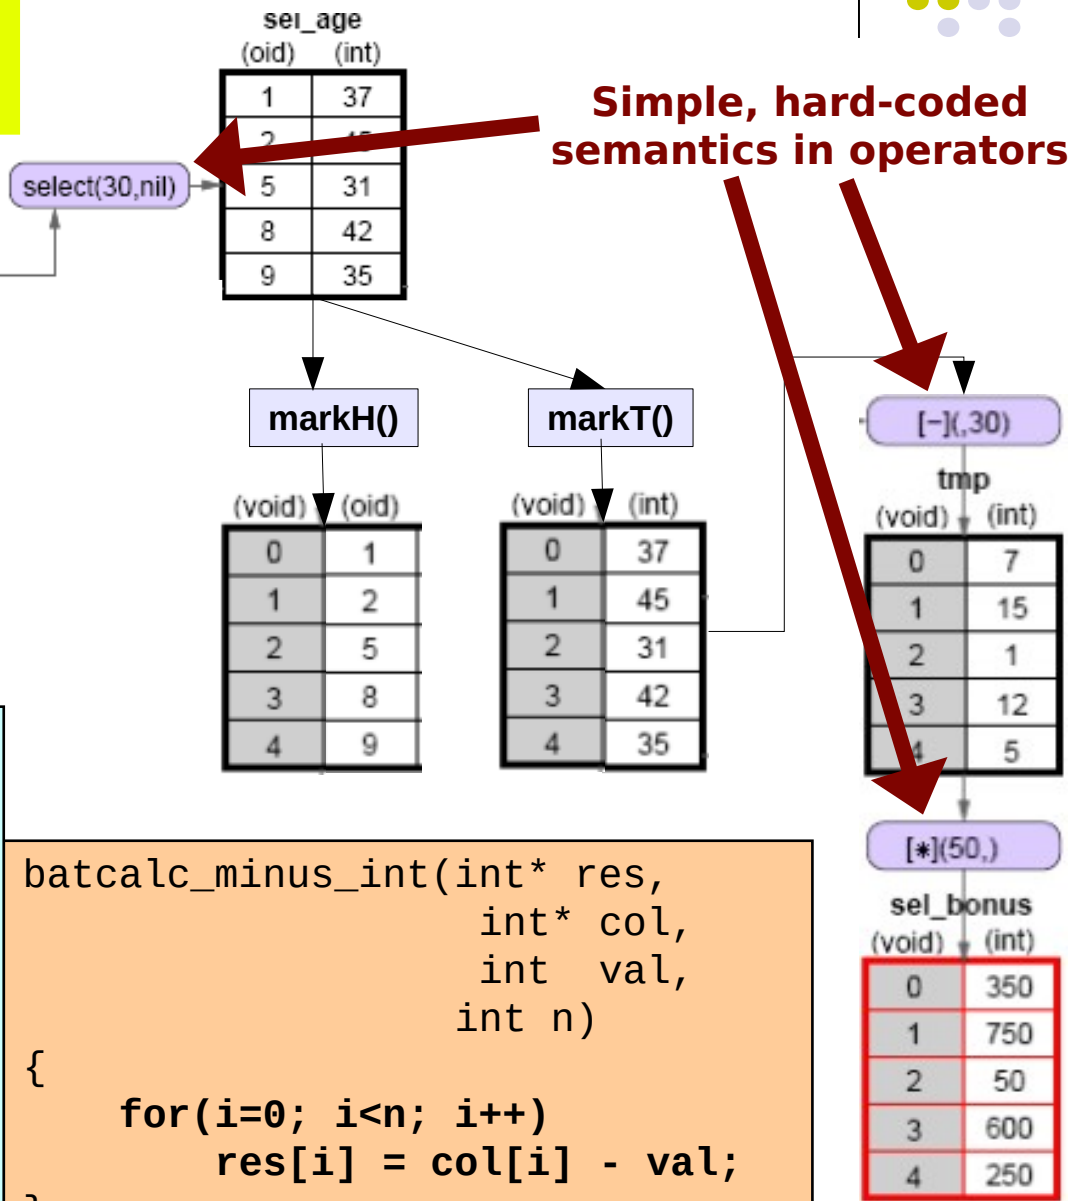


- ~~“save disk I/O when scan-intensive queries need a few columns”~~
- “avoid an expression interpreter to improve computational efficiency”



```
SELECT id, name, (age-30)*50 as bonus
FROM people
WHERE age > 30
```

people_id		people_name		people_age	
(void)	(int)	(void)	(str)	(void)	(int)
0	101	0	Alice	0	22
1	102	1	Ivan	1	37
2	104	2	Peggy	2	45
3	105	3	Victor	3	25
4	108	4	Eve	4	19
5	109	5	Walter	5	31
6	112	6	Trudy	6	27
7	113	7	Bob	7	29
8	114	8	Zoe	8	42
9	115	9	Charlie	9	35



**CPU ☺? Give it “nice” code !**

- few dependencies (control,data)
- CPU gets out-of-order execution
- compiler can e.g. generate SIMD

**One loop for an entire column**

- no per-tuple interpretation
- arrays: no record navigation
- better instruction cache locality

```
batcalc_minus_int(int* res,
                  int* col,
                  int val,
                  int n)
{
    for(i=0; i<n; i++)
        res[i] = col[i] - val;
}
```



- ~~“save disk I/O when scan-intensive queries need a few columns”~~
- “avoid an expression interpreter to improve computational efficiency”

How?

- RISC query algebra: hard-coded semantics
  - Decompose complex expressions in multiple operations
- Operators only handle **simple arrays**
  - No code that handles slotted buffered record layout
- Relational algebra becomes **array manipulation language**
  - Often SIMD for free
  - Plus: use of *cache-conscious* algorithms for Sort/Aggr/Join

# *monetdb* a Faustian pact



- You want efficiency
  - Simple hard-coded operators
- I take scalability
  - Result materialization

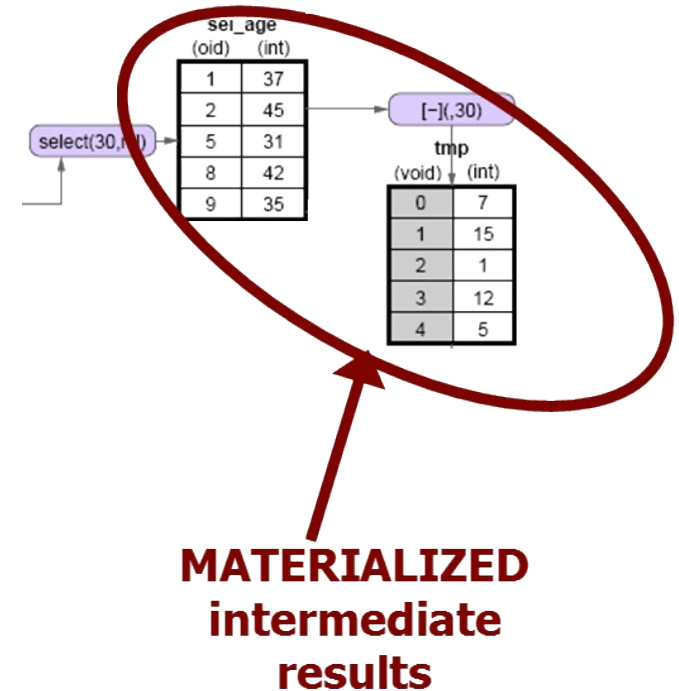
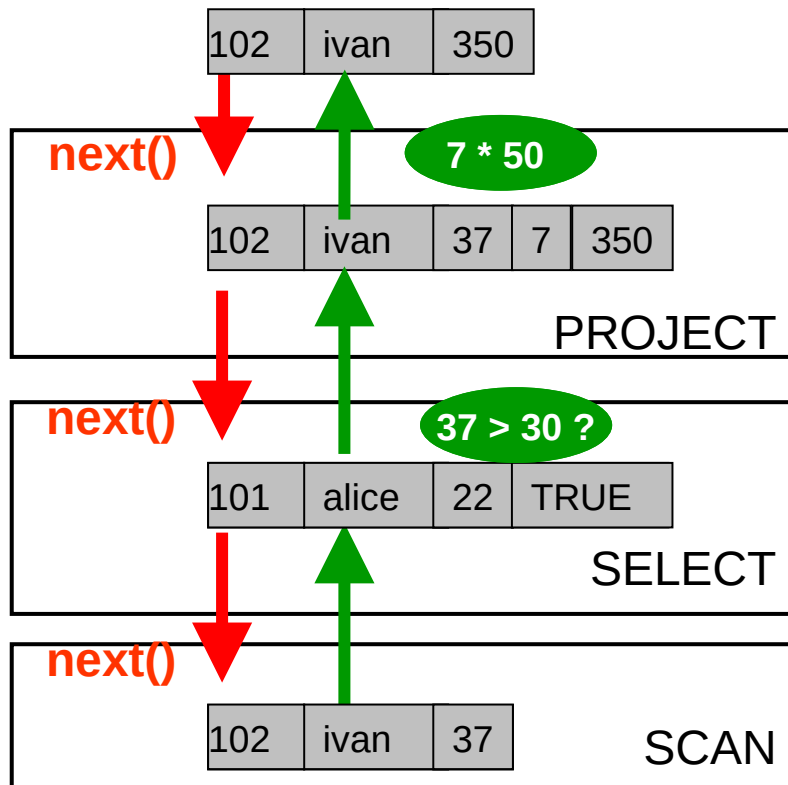
■ C program:	0.2s
■ MonetDB:	3.7s
■ MySQL:	26.2s
■ DBMS "X":	28.1s





# Pipelining vs

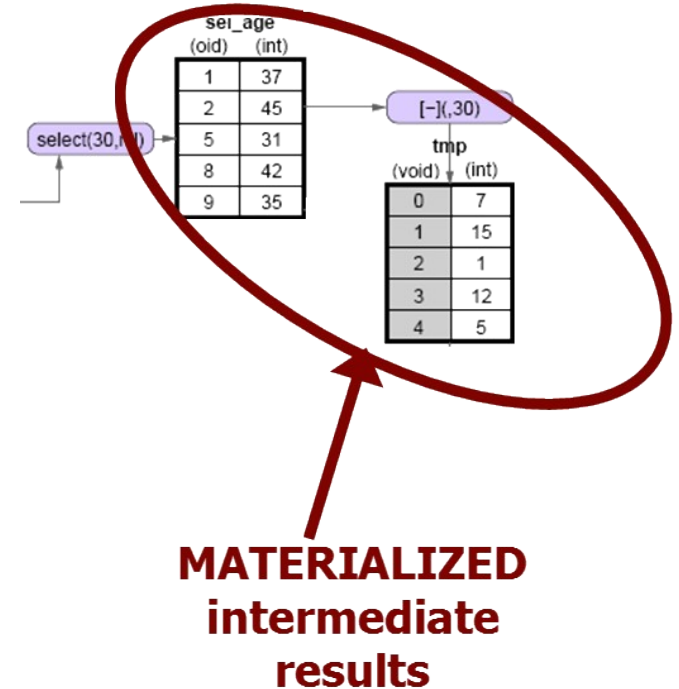
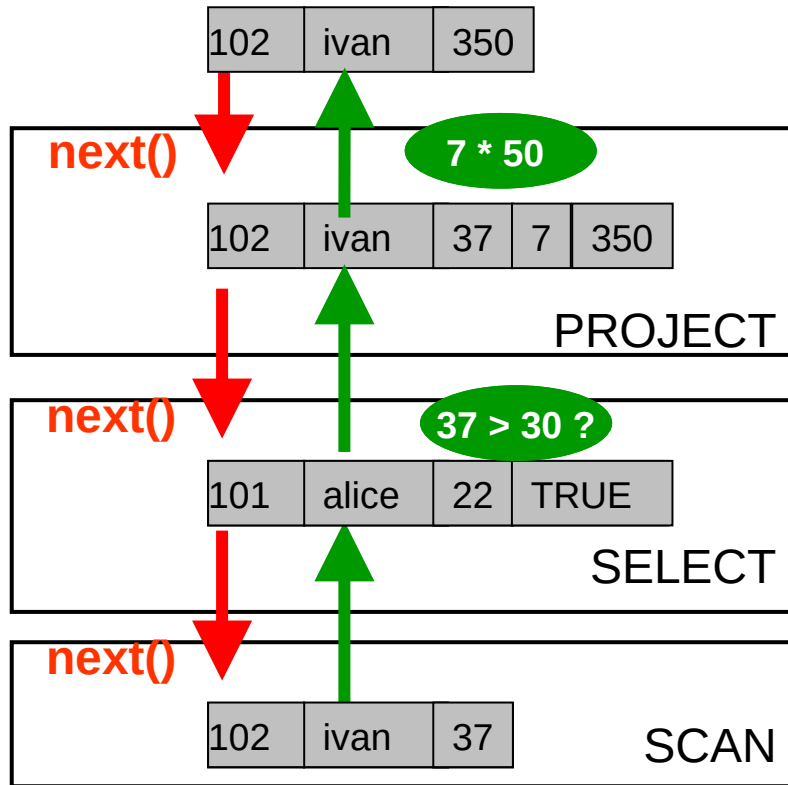
# Materialization



# Pipelining

vs

# Materialization



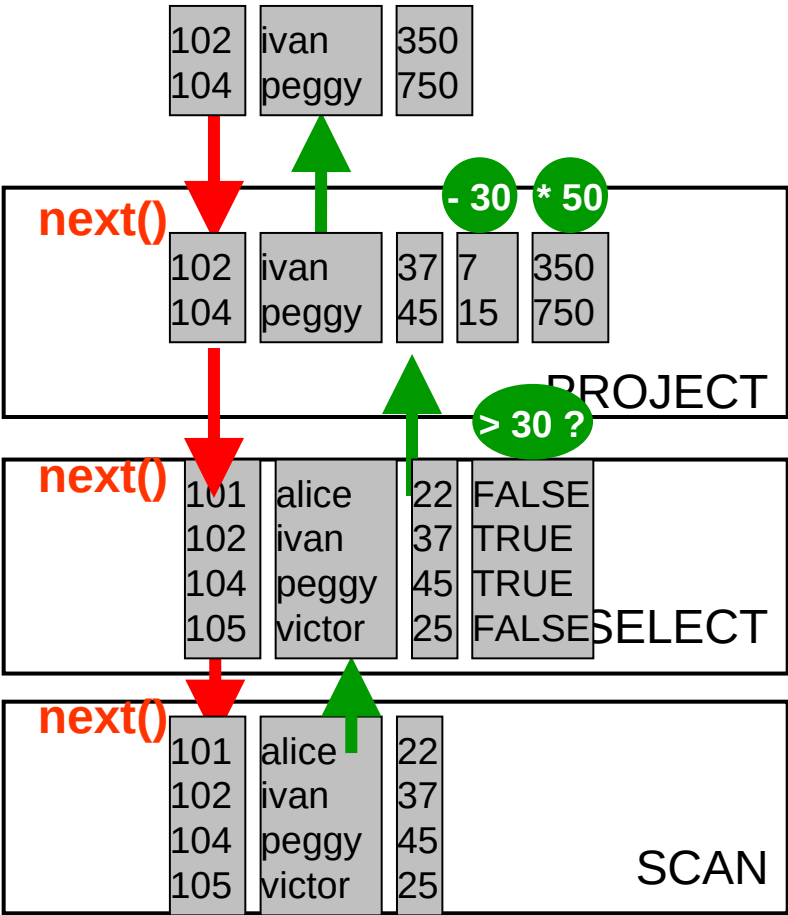
=> **vectorized query processing**



Observations:

next() called much less often → more time spent in primitives less in overhead

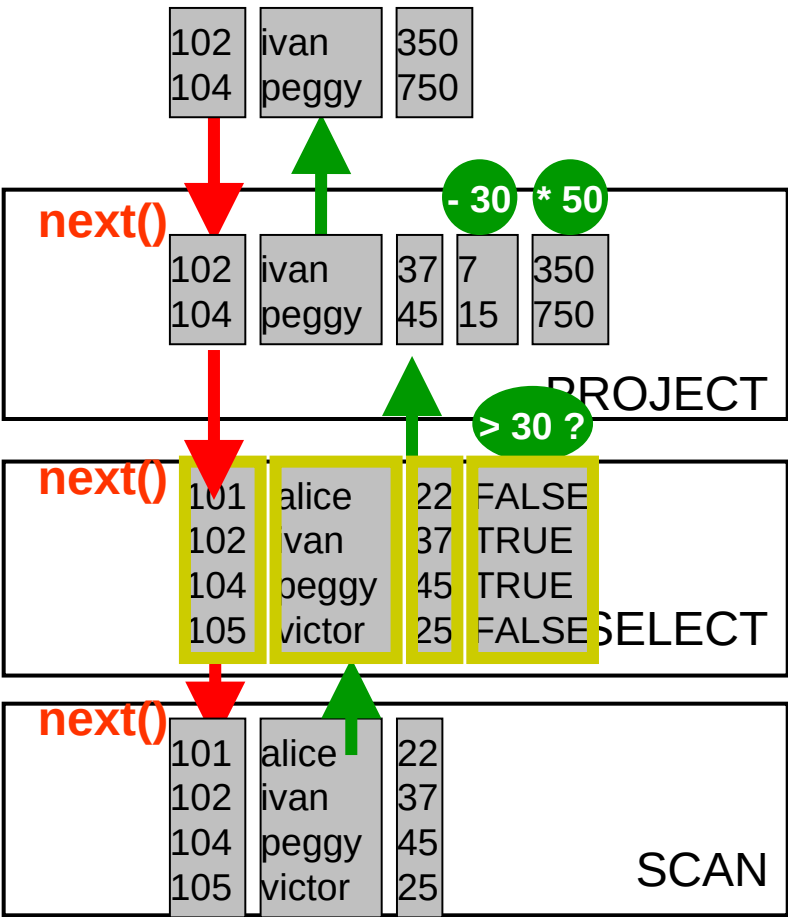
primitive calls process an array of values in a loop:





“Vectorized In Cache Processing”

vector = array of ~1000  
processed in a tight loop  
CPU cache Resident





## Observations:

`next()` called much less often →  
more time spent in **primitives** less  
in **overhead**

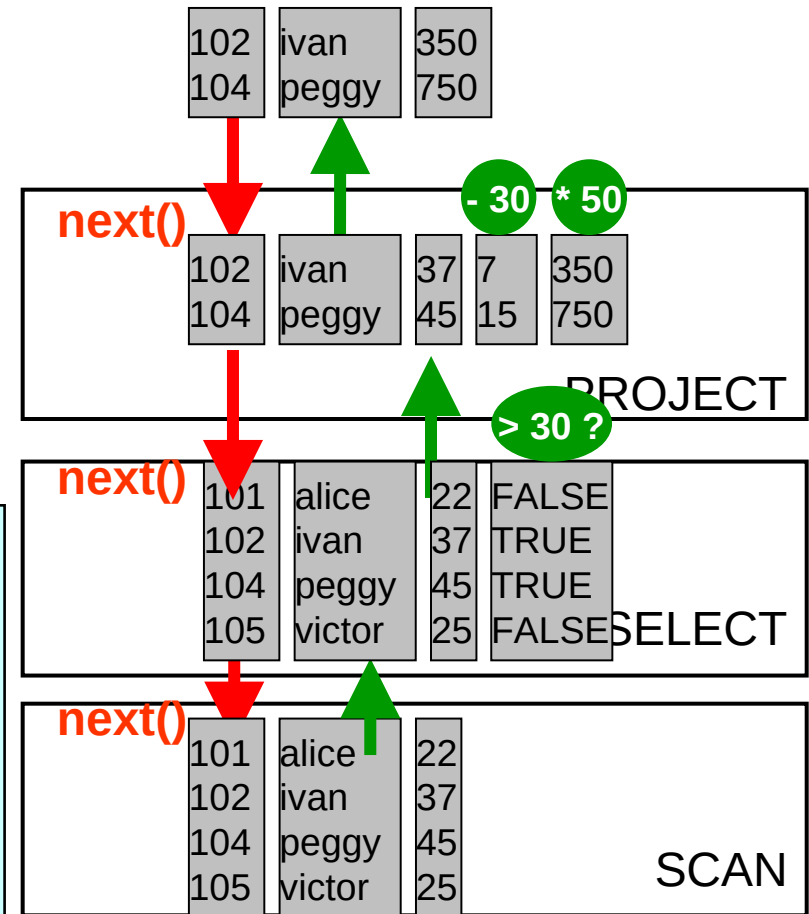
primitive calls process an array of  
values in a **loop**:

### CPU Efficiency depends on “nice” code

- out-of-order execution
- few dependencies (control,data)
- compiler support

### Compilers like simple loops over arrays

- loop-pipelining
- automatic SIMD





## Observations:

`next()` called much less often →  
more time spent in **primitives** less  
in **overhead**

primitive calls process an array of  
values in a **loop**:

### CPU Efficiency depends on “nice” code

- out-of-order execution
- few dependencies (control,data)
- compiler support

### Compilers like simple loops over arrays

- loop-pipelining
- automatic SIMD

> 30 ?

FALSE  
TRUE  
TRUE  
FALSE

```
for(i=0; i<n; i++)  
    res[i] = (col[i] > x)
```

- 30

7  
15

```
for(i=0; i<n; i++)  
    res[i] = (col[i] - x)
```

\* 50

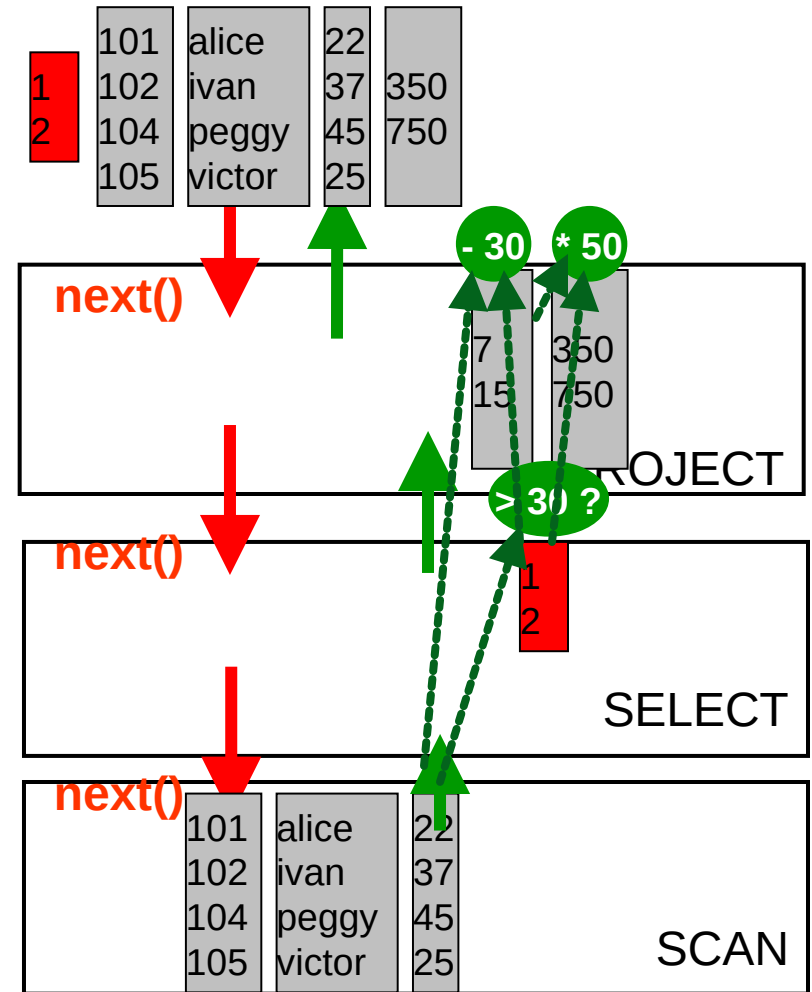
350  
750

```
for(i=0; i<n; i++)  
    res[i] = (col[i] * x)
```



## Tricks being played:

- Late materialization
- Materialization avoidance using **selection vectors**

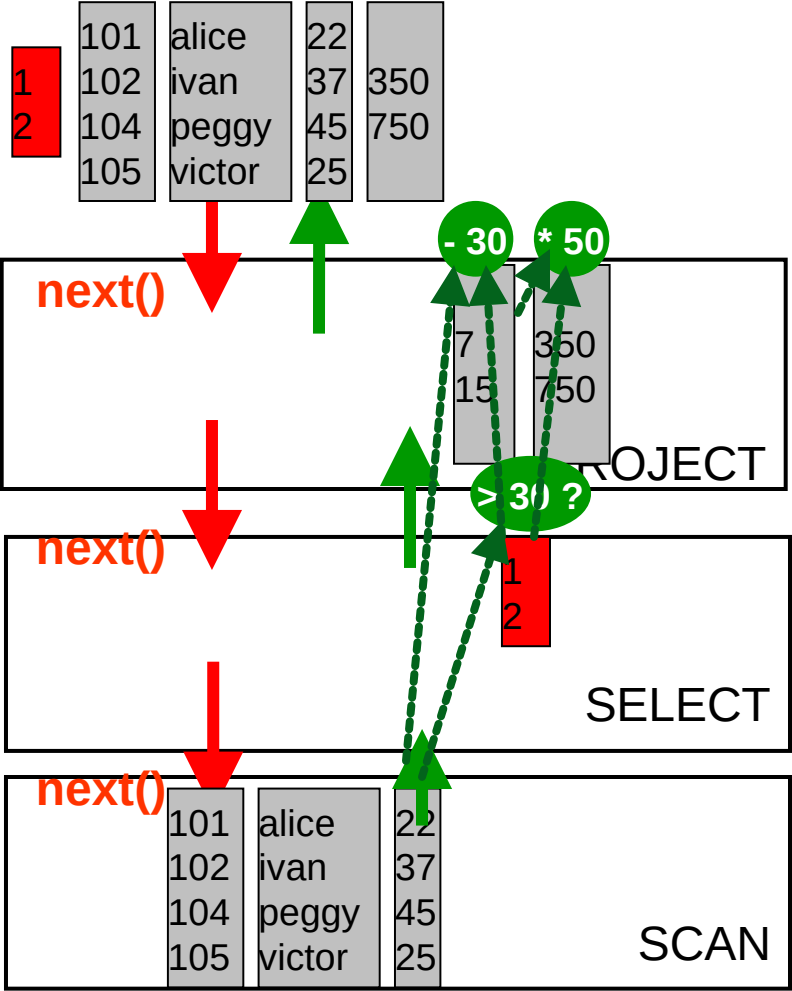




```
map_mul_flt_val_flt_col(  
    float *res,  
    int* sel,  
    float val,  
    float *col, int n)  
{  
    for(int i=0; i<n; i++)  
        res[i] = val * col[sel[i]];  
}
```

selection vectors used to reduce vector copying

contain selected positions







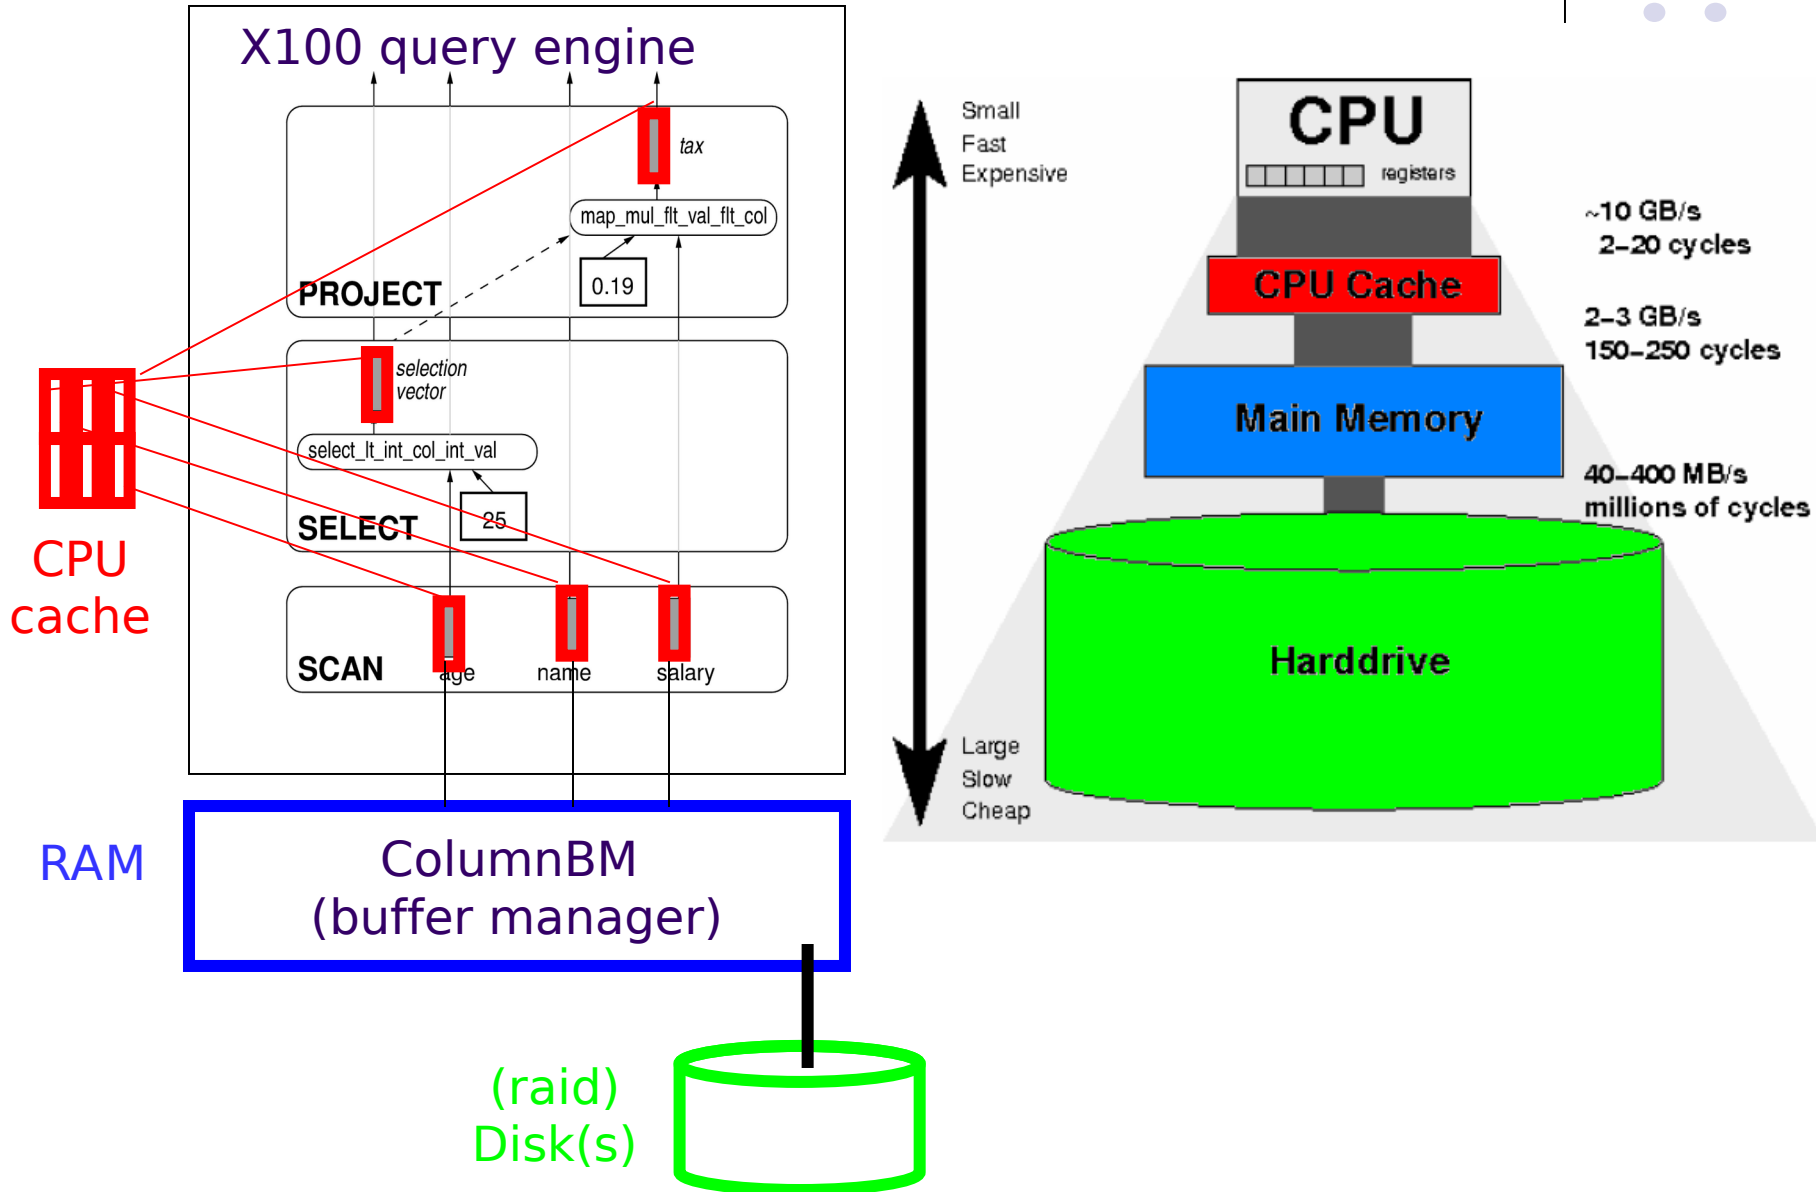
- Both efficiency
  - Vectorized primitives
- and scalability..
  - Pipelined query evaluation

■ C program:	0.2s
■ MonetDB/X100:	0.6s
■ MonetDB:	3.7s
■ MySQL:	26.2s
■ DBMS “X”:	28.1s



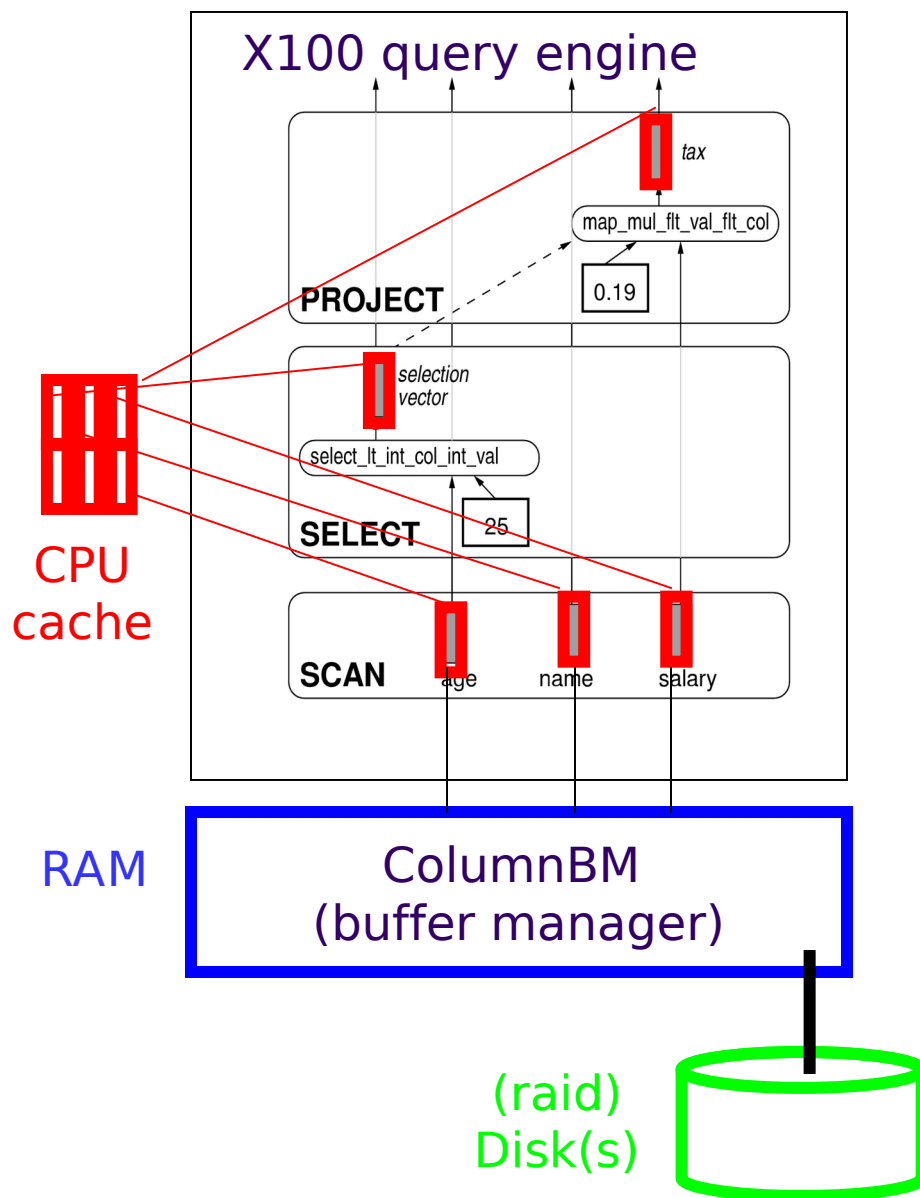
# Memory Hierarchy

“MonetDB/X100: Hyper-Pipelining Query Execution” Boncz, Zukowski, Nes, CIDR’05





# Memory Hierarchy



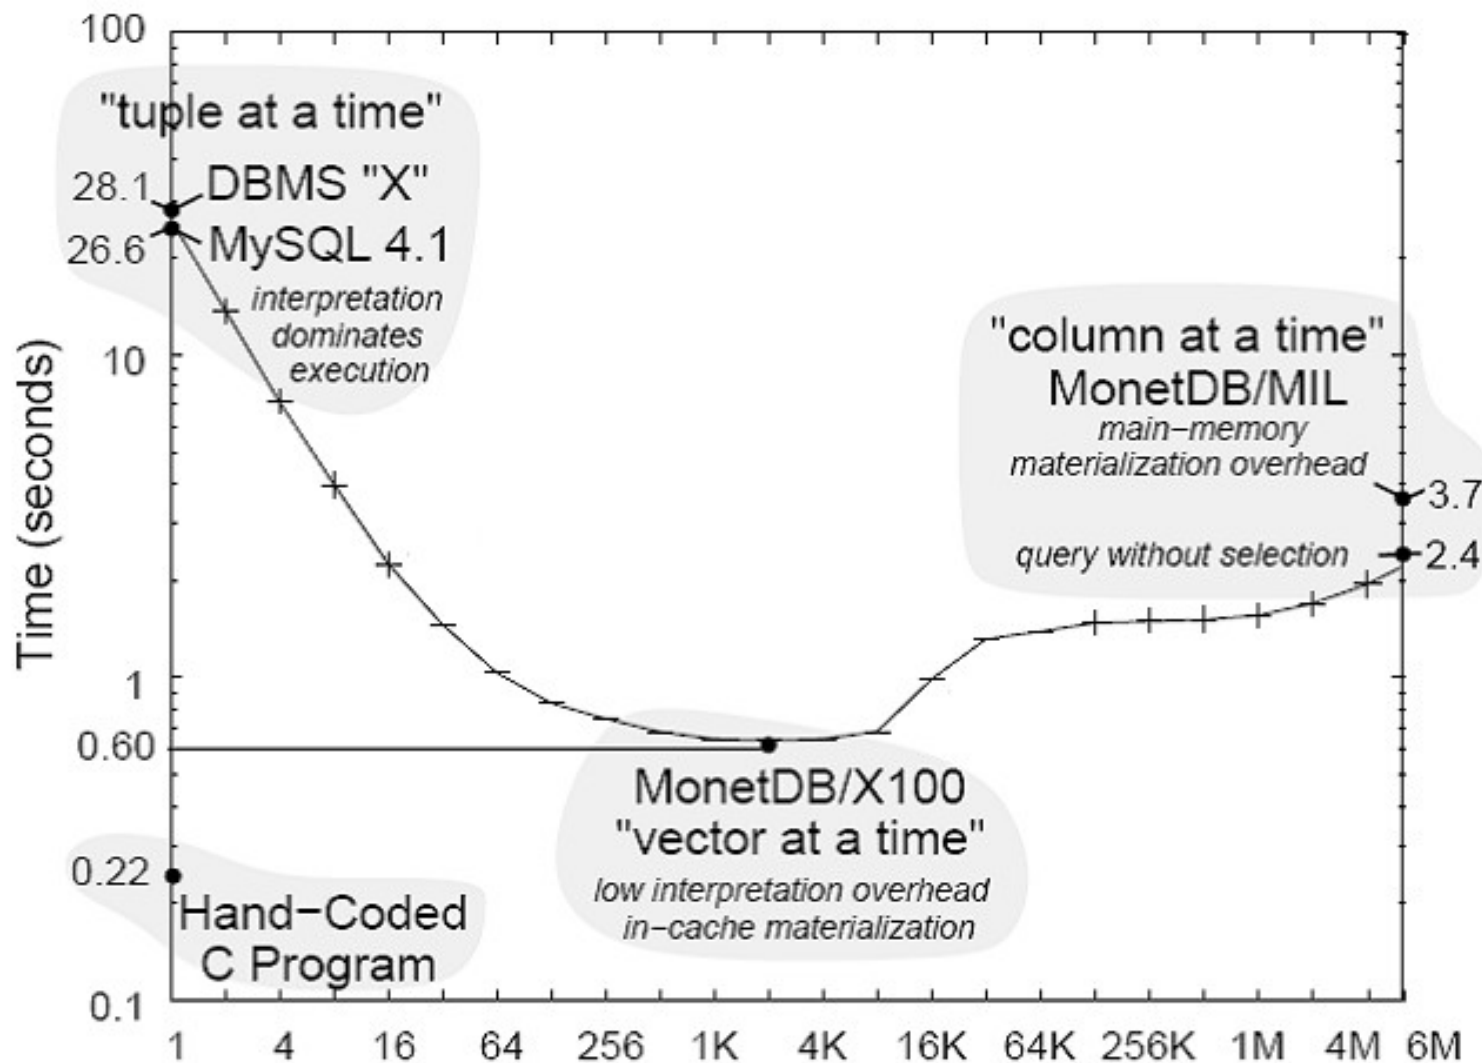
Vectors are only the in-cache representation

RAM & disk representation might actually be different

(vectorwise uses both PAX & DSM)

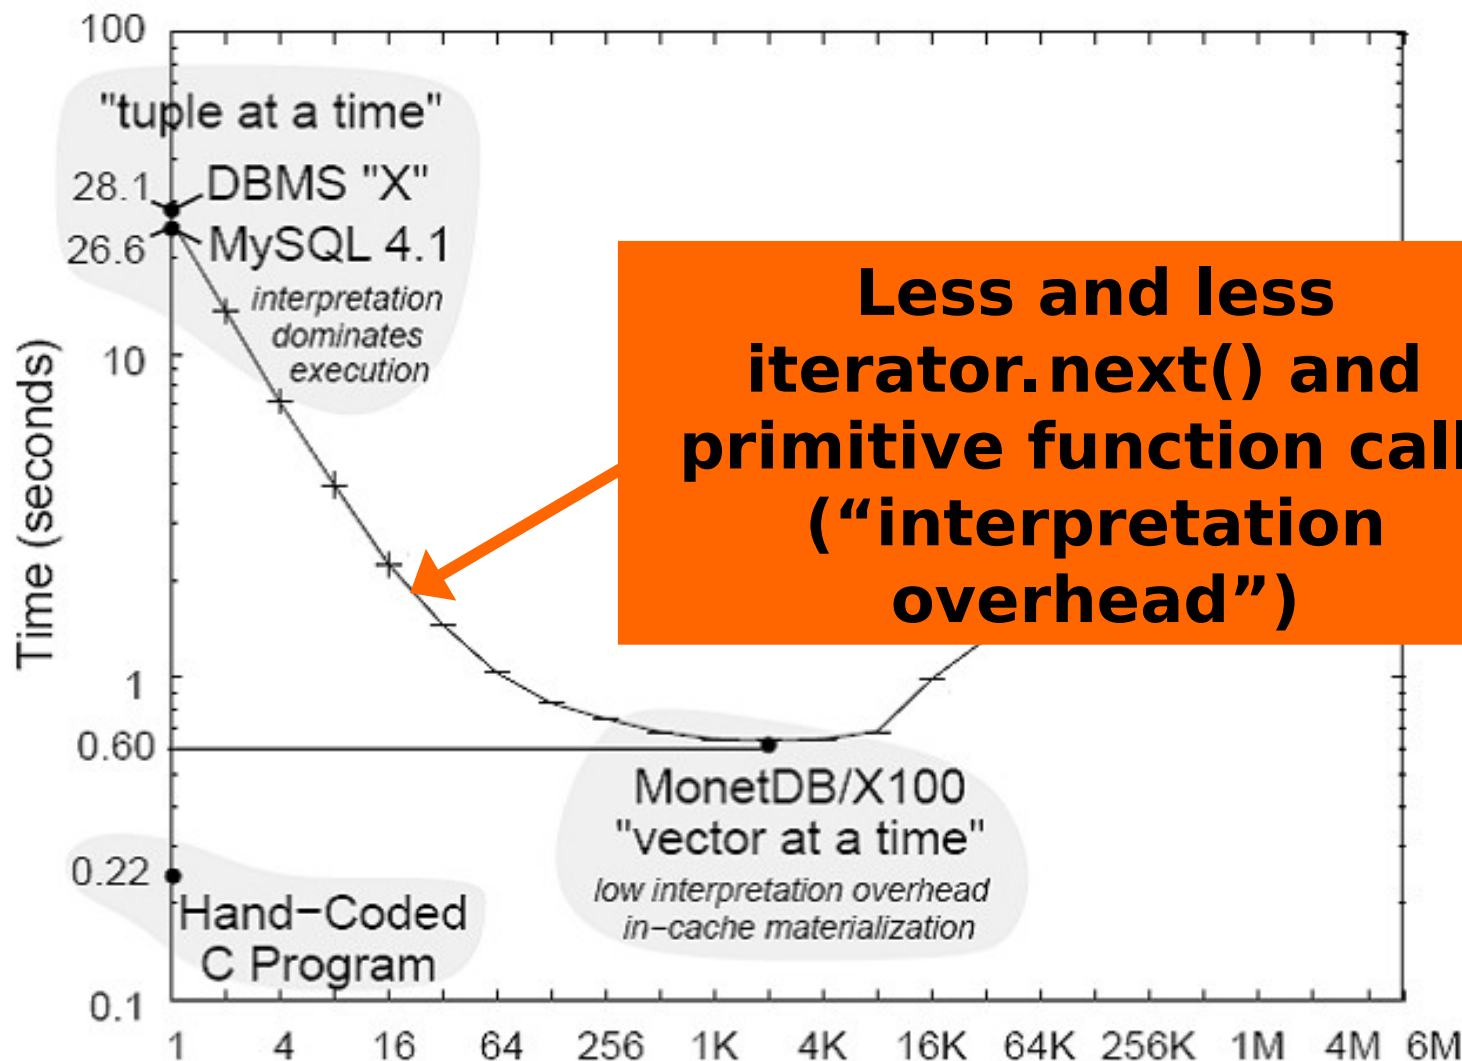


# Varying the Vector size





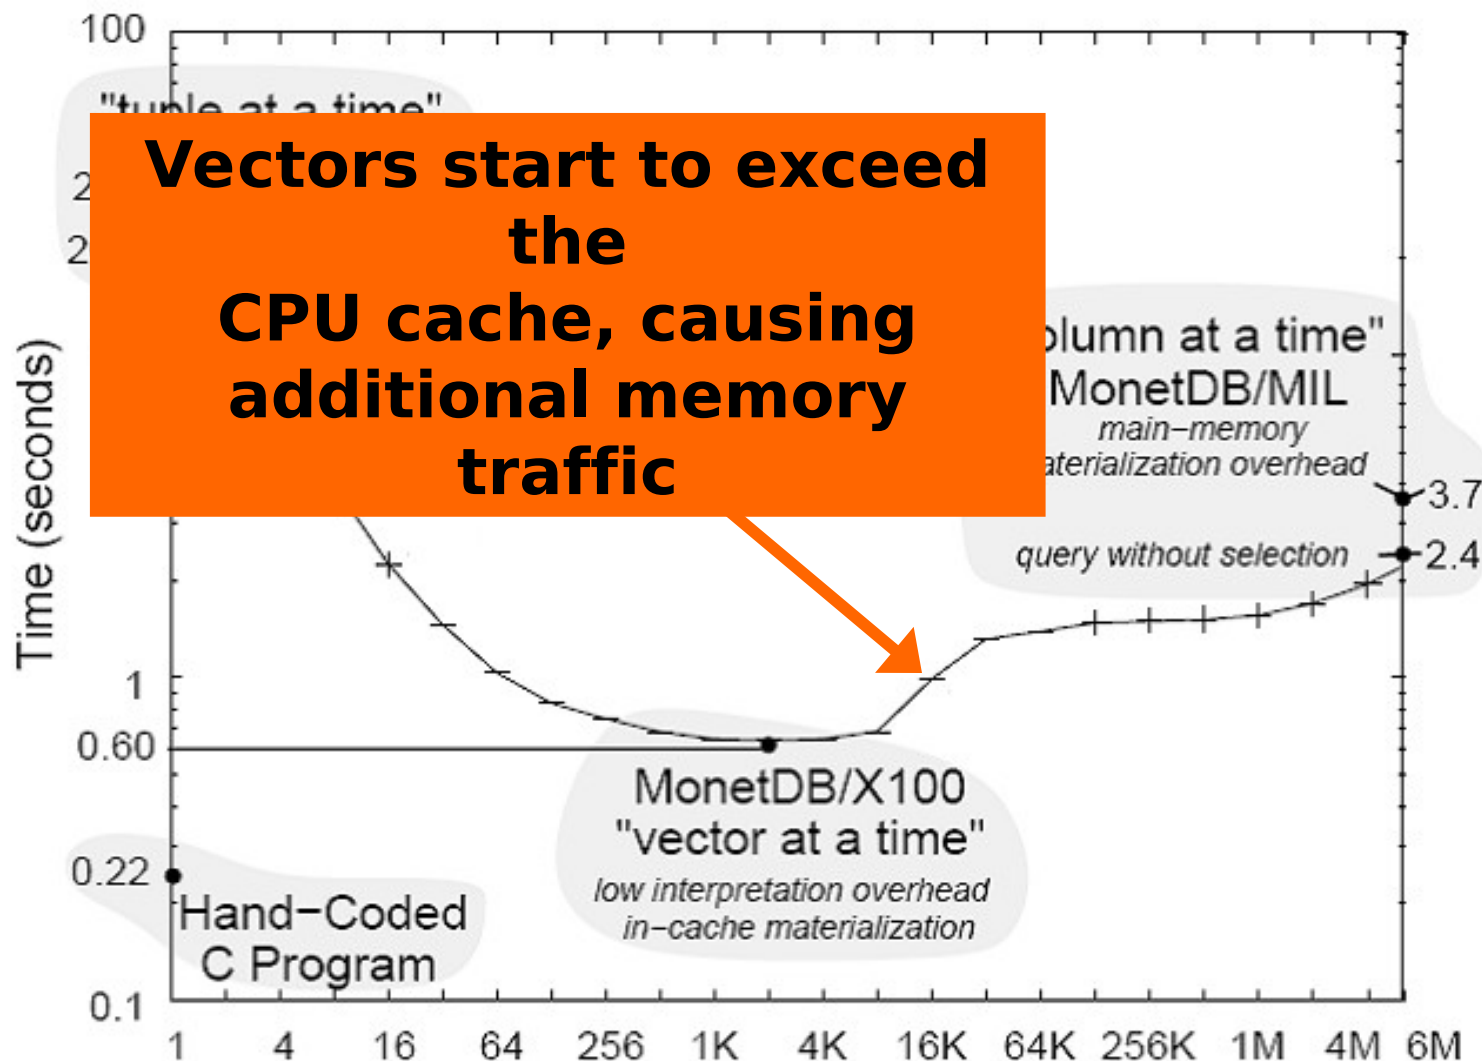
# Varying the Vector size



**Less and less  
iterator.next() and  
primitive function calls  
("interpretation  
overhead")**

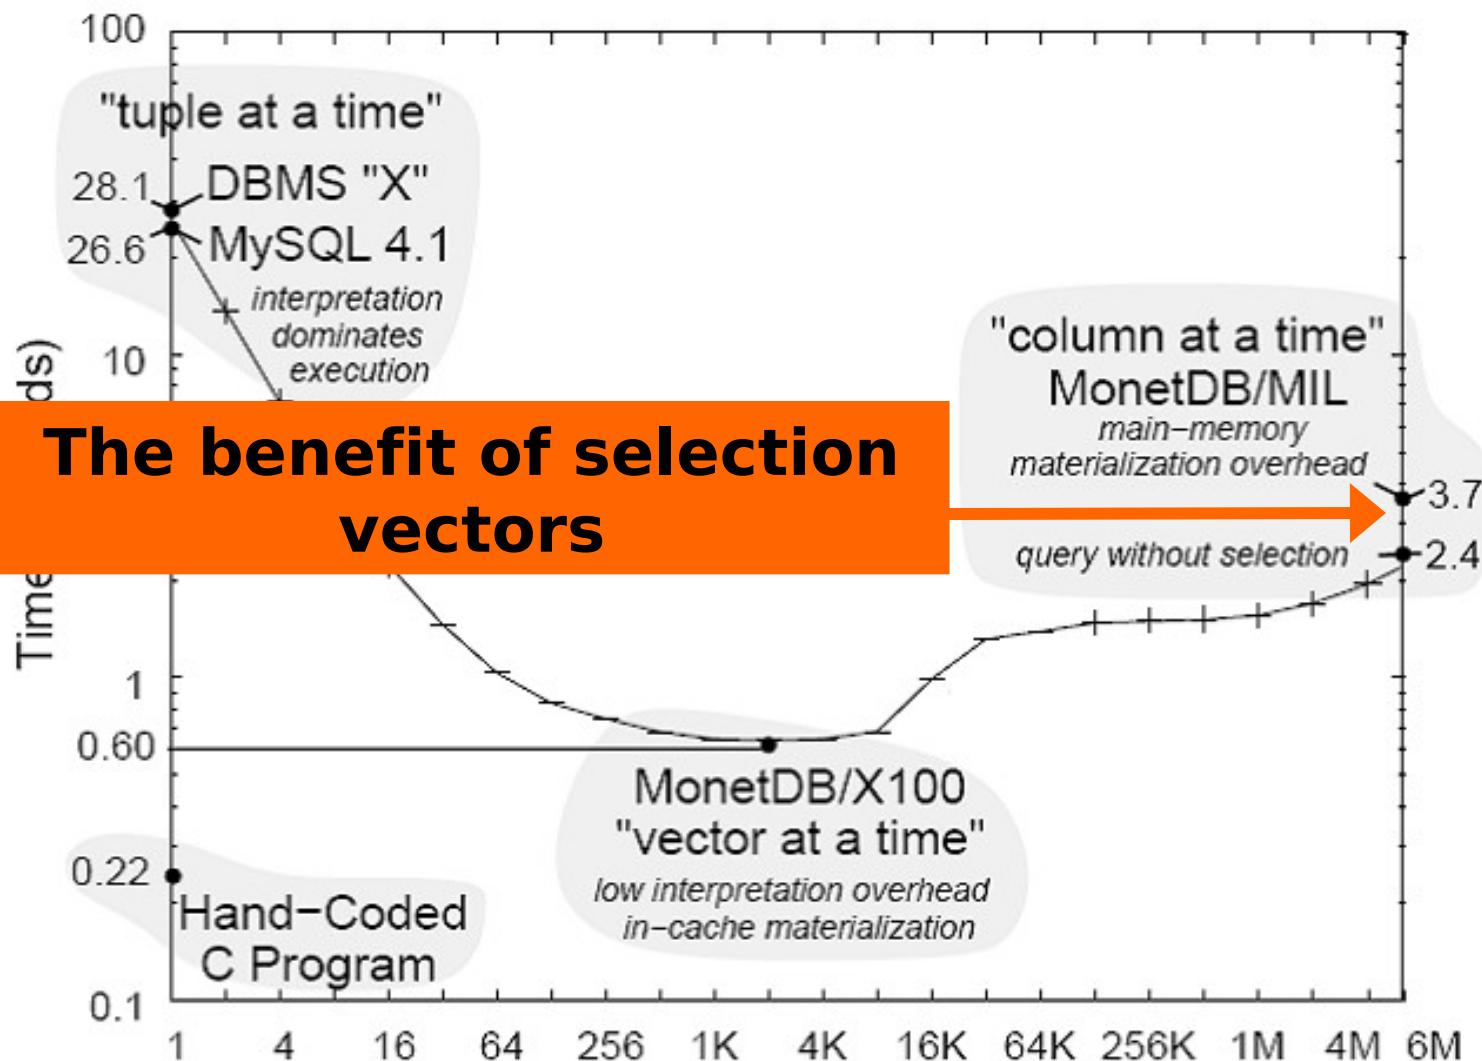


# Varying the Vector size





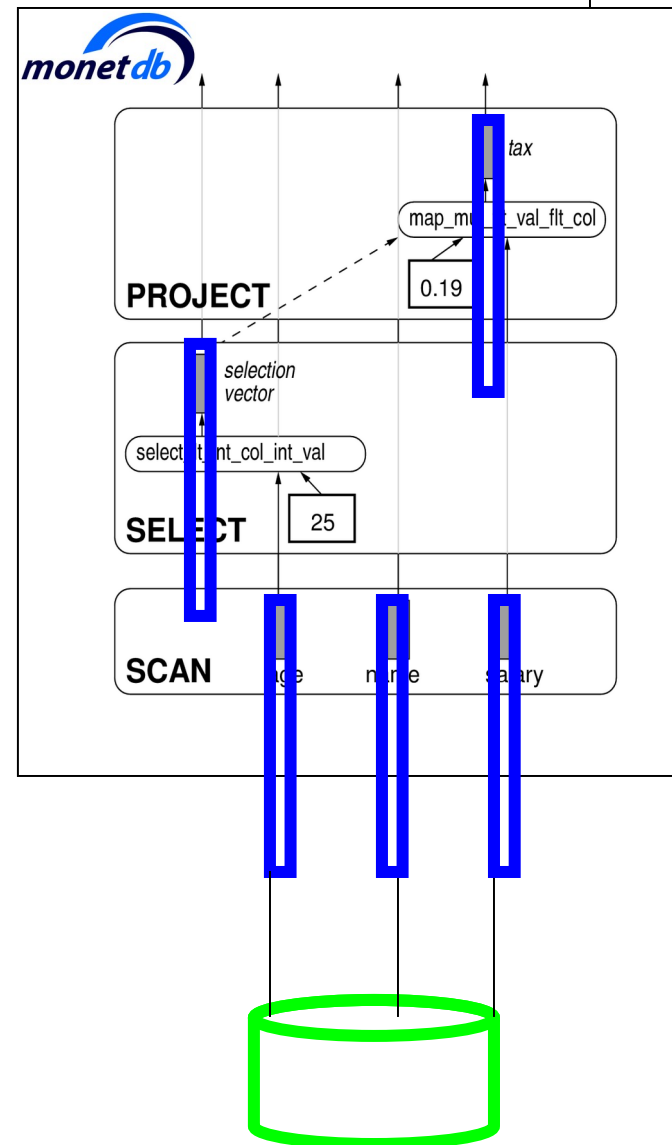
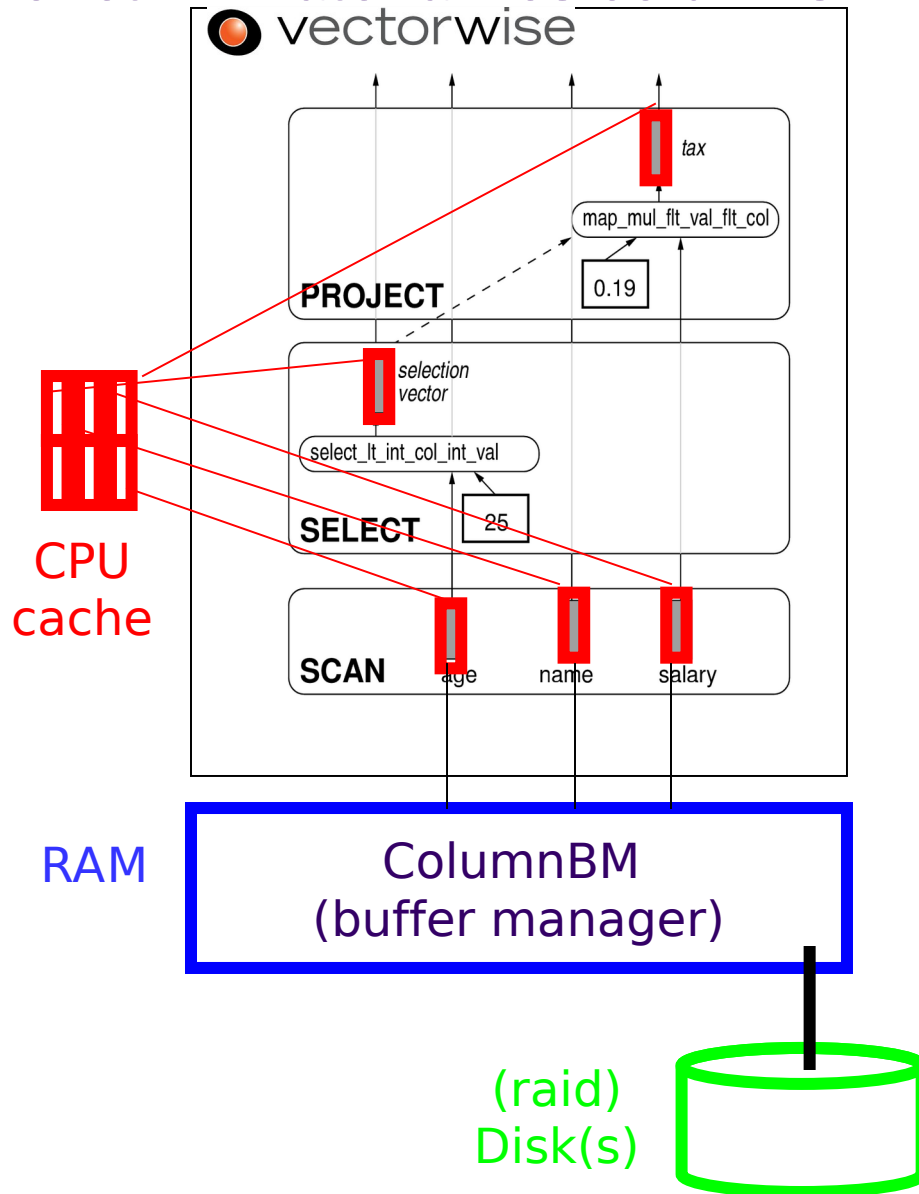
# Varying the Vector size







## MonetDB materializes columns







# Benefits of Vectorized Processing

- **100x less Function Calls**
  - `iterator.next()`, primitives
- **No Instruction Cache Misses**
  - High locality in the primitives
- **Less Data Cache Misses**
  - Cache-conscious data placement
- **No Tuple Navigation**
  - Primitives are record-oblivious, only see arrays
- **Vectorization allows algorithmic optimization**
  - Move activities out of the loop (“strength reduction”)
- **Compiler-friendly function bodies**
  - Loop-pipelining, automatic SIMD

“Buffering Database Operations for Enhanced Instruction Cache Performance”  
Zhou, Ross, SIGMOD’04

“Block oriented processing of relational database operations in modern computer architectures”  
Padmanabhan, Malkemus, Agarwal, ICDE’01



# Vectorizing Relational Operators

- Project
- Select
  - Exploit selectivities, test buffer overflow
- Aggregation
  - Ordered, Hashed
- Sort
  - Radix-sort nicely vectorizes
- Join
  - Merge-join + Hash-join

# ADM: Literature

- **Column-Oriented Database Systems (4/6) - “Vectorized Execution”**
  - “MonetDB/X100: Hyper-Pipelining Query Execution”. Boncz, Zukowski, Nes. CIDR’05.
  - “Buffering Database Operations for Enhanced Instruction Cache Performance”. Zhou and Ross. SIGMOD’04.
  - “Block oriented processing of relational database operations in modern computer architectures”. Padmanabhan, Malkemus, Agarwal. ICDE’01.
  - “Balancing Vectorized Query Execution with Bandwidth Optimized Storage”. Zukowski. PhD Thesis. CWI 2008.

# ADM: Agenda

- 07.09.2022: Lecture 1: **Introduction**
- 14.09.2022: Lecture 2: **SQL Recap**  
(plus Assignment 1 [in groups; 3 weeks]: TPC-H benchmark)
- 21.09.2022: Lecture 3: **Column-Oriented Database Systems (1/6) - Motivation & Basic Concepts**
- 28.09.2022: Lecture 4: **Column-Oriented Database Systems (2a/6) - Selected Execution Techniques (1/2)**
- 05.10.2022: Lecture 5: **Column-Oriented Database Systems (2b/6) - Selected Execution Techniques (2/2)**  
(plus Assignment 2 [in groups; 4 weeks]: Compression techniques)
- 12.10.2022: Lecture 6: **Column-Oriented Database Systems (3/6) - Cache Conscious Joins**
- 19.10.2022: Lecture 7: **Column-Oriented Database Systems (4/6) - “Vectorized Execution”**
- 26.10.2022: **No lecture!**
- 02.11.2022: Lecture 8: **DuckDB: An embedded database for data science (1/2) (guest lecture & hands-on)**  
(plus Assignment 3 [individual; 2 weeks]: Analysing NYC Cab dataset with DuckDB)
- 09.11.2022: Lecture 9: **DuckDB: An embedded database for data science (2/2) (guest lecture & hands-on)**
- 16.11.2022: Lecture 10: **Branch Misprediction & Predication**  
(plus Assignment 4 [individual; 2 weeks]: Predication)
- 23.11.2022: Lecture 11: **Column-Oriented Database Systems (5/6) - Adaptive Indexing**
- 30.11.2022: Lecture 12: **Column-Oriented Database Systems (6/6) - Progressive Indexing**

**Bring your  
own laptop!**

# ADM: Literature (5/6)



- **DuckDB: An embedded database for data science**
  - “DuckDB: an Embeddable Analytical Database”. Mark Raasveldt & Hannes Mühleisen. SIGMOD’19. Demo.
  - “Data Management for Data Science - Towards Embedded Analytics”. Mark Raasveldt & Hannes Mühleisen. CIDR’20.
  - “Integrating Analytics with Relational Databases”. Mark Raasveldt. PhD Thesis, Leiden University & CWI, 2020.
  - <https://duckdb.org>