

Similarity Search 2

Presentation partially based on slides from:

<http://infolab.stanford.edu/~ullman/mining/2009/index.html>

Anand Rajaraman & Jeff Ullman

Agenda

- Cosine Similarity and Random Projections (sketches)
- Euclidean distance and Random Projections
- An application of LSH to searching similar fingerprints
- *(Optional) General theory of LSH*
 - *LSH families*
 - *AND & OR constructs*

Cosine Similarity, Cosine Distance

- In many applications, objects are represented not by sets but as vectors in R^d (or in $\{0,1\}^d$ or in N^d , ...)
- **Examples:**
 - a document can be represented by a vector of frequencies (or **TF.IDF** values) of words that occur in the document
 - an image can be represented by a vector of pixel intensities (e.g., 256 counts, if 1 byte per pixel)
 - a sound (segment) can be represented by its frequency spectrum (e.g., FFT -> 128 floats)
 - ... a vector of ratings given by a user to movies he/she watched (Netflix Challenge: ~17.000 numbers)
- **Cosine Distance between two vectors, v_1 and v_2 , is defined as *the angle between v_1 and v_2***

Study
Chapter 1!


Thus, we don't care about the lengths of vectors, only about the angle between them.
The smaller the angle the more similar the vectors are.

Cosine Distance: the θ

- Think of a point as a vector from the origin $(0,0,\dots,0)$ to its location.
- Two points' vectors make an angle, whose cosine is the normalized dot-product of the vectors: **$\cos(\theta) = \mathbf{p}_1 \cdot \mathbf{p}_2 / |\mathbf{p}_1| |\mathbf{p}_2|$** .
 - **Example:** $\mathbf{p}_1 = 00111$; $\mathbf{p}_2 = 10011$.
 - $\mathbf{p}_1 \cdot \mathbf{p}_2 = 2$; $|\mathbf{p}_1| = |\mathbf{p}_2| = \sqrt{3}$.
 - $\cos(\theta) = 2/3$; θ is about **48 degrees**.



The $\cos(\theta) = p_1 \cdot p_2 / |p_2| |p_1|$ formula

- ❑ $p_1 \cdot p_2$ is a vector dot product: $p_1 \cdot p_2 = \sum_{n=1}^{dim} (p_{1,n} \cdot p_{2,n})$
- ❑ $|p_2| \cdot |p_1|$ is the product of lengths of p_1 and p_2
- ❑ There are many proofs of this equation, e.g.,:
 - <http://www.mit.edu/~hlb/StantonGrant/18.02/details/tex/lec1snip2-dotprod.pdf>
 - https://www.youtube.com/watch?v=PnJoKGynu_U
- ❑ Note that $p_1 \cdot p_2$ might be negative!
- ❑ $\cos(\theta)$ takes values between -1 and 1.
- ❑ If components of p_1 and p_2 are non-negative (as in A2) then $\cos(\theta)$ is also non-negative.



Cosine Distance and Cosine Similarity

cosine distance: $d(p_1, p_2) = \theta = \arccos(p_1 \cdot p_2 / |p_2| |p_1|)$



Wikipedia: $\text{cosine similarity}(p_1, p_2) = \cos(\theta) = p_1 \cdot p_2 / |p_2| |p_1|$; values in $[-1, 1]$

BUT: in the textbook*: $\text{cosine similarity}(p_1, p_2) = 1 - \theta/180$; values in $[0, 1]$

*Actually, the term "cosine similarity" is not used in the textbook at all!

Why C.D. Is a Distance Measure?


- ❑ $d(x,x) = 0$ because $\arccos(1) = 0$.
- ❑ $d(x,y) = d(y,x)$ by symmetry.
- ❑ $d(x,y) \geq 0$ because angles are chosen to be in the range 0 to 180 degrees.
- ❑ **Triangle inequality**: physical reasoning. If I rotate an angle from x to z and then from z to y , I can't rotate less than from x to y .
- ❑ *Cosine similarity ranges between 0 and 1 (Textbook!)*

Estimating θ in 2 dimensions

Consider two vectors A and B with angle $\theta < 180^\circ$.

Select at random a line (blue) $v_1x + v_2y = 0$.

What is the probability that A and B are on the same side of this line?



$$\text{Prob}[A \text{ and } B \text{ on different sides of the line}] = \theta/180$$


$$\text{Prob}[A \text{ and } B \text{ on the same side of the line}] = 1 - \theta/180$$

Let v a normal vector that defines the line. Then $v \cdot B > 0$ ("x is above") and $v \cdot A < 0$ ("y is below")

Textbook
Cosine
Similarity

Estimating θ in n -dimensions ...

Consider two vectors A and B in n -dimensions with angle $\theta < 180^\circ$.
These 2 vectors define (or span) a 2-dimensional plane!
Let's plot this plane with both vectors (see below) ...
Now take a random vector \mathbf{v} (in n -dims) and intersect
the hyperplane $\mathbf{v}\mathbf{x}=\mathbf{0}$ with our plane (the blue line):



Repeat the reasoning from the previous slide (+ some imagination: textbook page 118)

LSH for Cosine Similarity

- LSH is based on the idea of hash functions:
“the chance that two objects will hash to the same value is the same as similarity of these objects”
- Instead of “minhashing” we now use “random projections”: hash functions which return 1 or -1 depending on whether x is below or above some randomly chosen hyperplanes.

Random Projections

- Pick a random vector v , which determines a hash function h_v with two buckets.
- $h_v(x) = +1$ if $v \cdot x > 0$; -1 if $v \cdot x < 0$.
- **Claim:** $\text{Prob}[h(x)=h(y)] = \text{cosine_sim}(x,y)$
- Proof: look at slide 7!

Signatures for Cosine Distance


- ❑ Pick some number of random vectors, and hash your data for each vector.
- ❑ The result is a signature (*sketch*) of +1's and -1's that can be used for LSH like the minhash signatures (with the banding trick) for the Jaccard similarity!
- ❑ *If dimensionality is high, it suffices to consider only vectors v consisting of +1 and -1 components (eliminates multiplications - we sum up "positives" and subtract "negatives"!)*

LSH for Euclidean Distance


- **Key Idea:** if two points are close to each other then no matter from which direction we look at them, they always seem to be close; if they are far apart, then it's difficult to find such a direction
- **Hash functions are defined by random directions (lines) that are partitioned into buckets (intervals) of size a .**
- Hash each point to the bucket containing its projection onto the line:
 $h_{w,a}(x)$ = the id of the bucket that contains $w \cdot x$
- Nearby points are often close (in the same bucket); distant points are rarely in the same bucket => **probability of hashing x and y to the same bucket is proportional to the Euclidean similarity of x and y .**
- *More details in the textbook (3.7.4)*

Projection of Points

If $d \gg a$, θ must be close to 90° for there to be any chance points go to the same bucket.




If $d \ll a$, then the chance the points are in the same bucket is at least $1 - d/a$.



Application of LSH: Fingerprint Matching


- Fingerprints are represented by the set of *minutiae* features.



LSH for Fingerprints

- Place a grid on a fingerprint.
 - Normalize so identical prints will overlap.
- Set of grid points where minutiae are located represents the fingerprint.
 - Possibly, treat minutiae near a grid boundary as if also present in adjacent grid points.

Discretizing Minutiae



Assumptions:

- ❑ Suppose typical fingerprints have minutiae in 20% of the grid points.
- ❑ Suppose fingerprints from the same finger agree in at least 80% of grid points.
- ❑ Suppose we have 1.000.000.000 fingerprints
- ❑ **Challenge:**
speed up the process of searching this data base with help of LSH!

Applying LSH to Fingerprints

- ❑ Make a bit vector for each fingerprint's set of grid points with minutiae.
- ❑ We could minhash the bit vectors to obtain signatures.
- ❑ But since there probably aren't too many grid points, we can work from the bit-vectors directly.

LSH/Fingerprints (2)

- ❑ Pick 1024 (?) sets of 3 (?) grid points, randomly.
- ❑ For each set of points, prints with 1 for all three points are candidate pairs.
- ❑ Thus we have 1024 “hash functions” and every fingerprint is mapped into a string of 1024 bits - “signatures” of fingerprints (128 bytes long)
- ❑ Apply the “banding technique” with $b=1024$ (one row per band)...

LSH/Fingerprints (3)

- ❑ Typical fingerprints have minutiae in 20% of the grid points.
- ❑ Fingerprints from the same finger agree in at least 80% of grid points.
- ❑ Probability two random fingerprints each have 1 in all three points = $(0.2)^6 = .000064$.
- ❑ Why ^6? Both fingers should return 3x1 and they are coming from different persons

LSH/Fingerprints (4)

First image
has 1 in a point

Second image
of same finger
also has 1.

□ Probability two fingerprints from the same finger each have 1's in three given points = $((0.2)(0.8))^3 = .004096$.

□ Prob. for at least one of 1024 sets of three points = $1-(1-.004096)^{1024} = .985$.

1.5% false
negatives

□ But for random fingerprints:
 $1-(1-.000064)^{1024} = .063$

6.3% false
positives

Example: Reducing the false positives

- ❑ Build two systems with 1024 buckets each.
- ❑ For an input fingerprint return the intersections of sets returned by both systems
- ❑ Prob. of false negatives = $(1 - 0.985^2) = 3\%$
- ❑ Prob. of false positives:
 $0.063^2 = 0.4\%$ (16 times better!)