

Mining Data Streams

Partially based on slides from:

<http://infolab.stanford.edu/~ullman/mining/2009/index.html>

Anand Rajaraman & Jeff Ullman

Contents

☐ Data Stream Model

- Random Sampling of Data Streams
- Filtering Data Streams (Bloom Filter)
- Counting Distinct Elements
- ~~■ Estimating Moments~~

☐ ~~A sliding window model~~

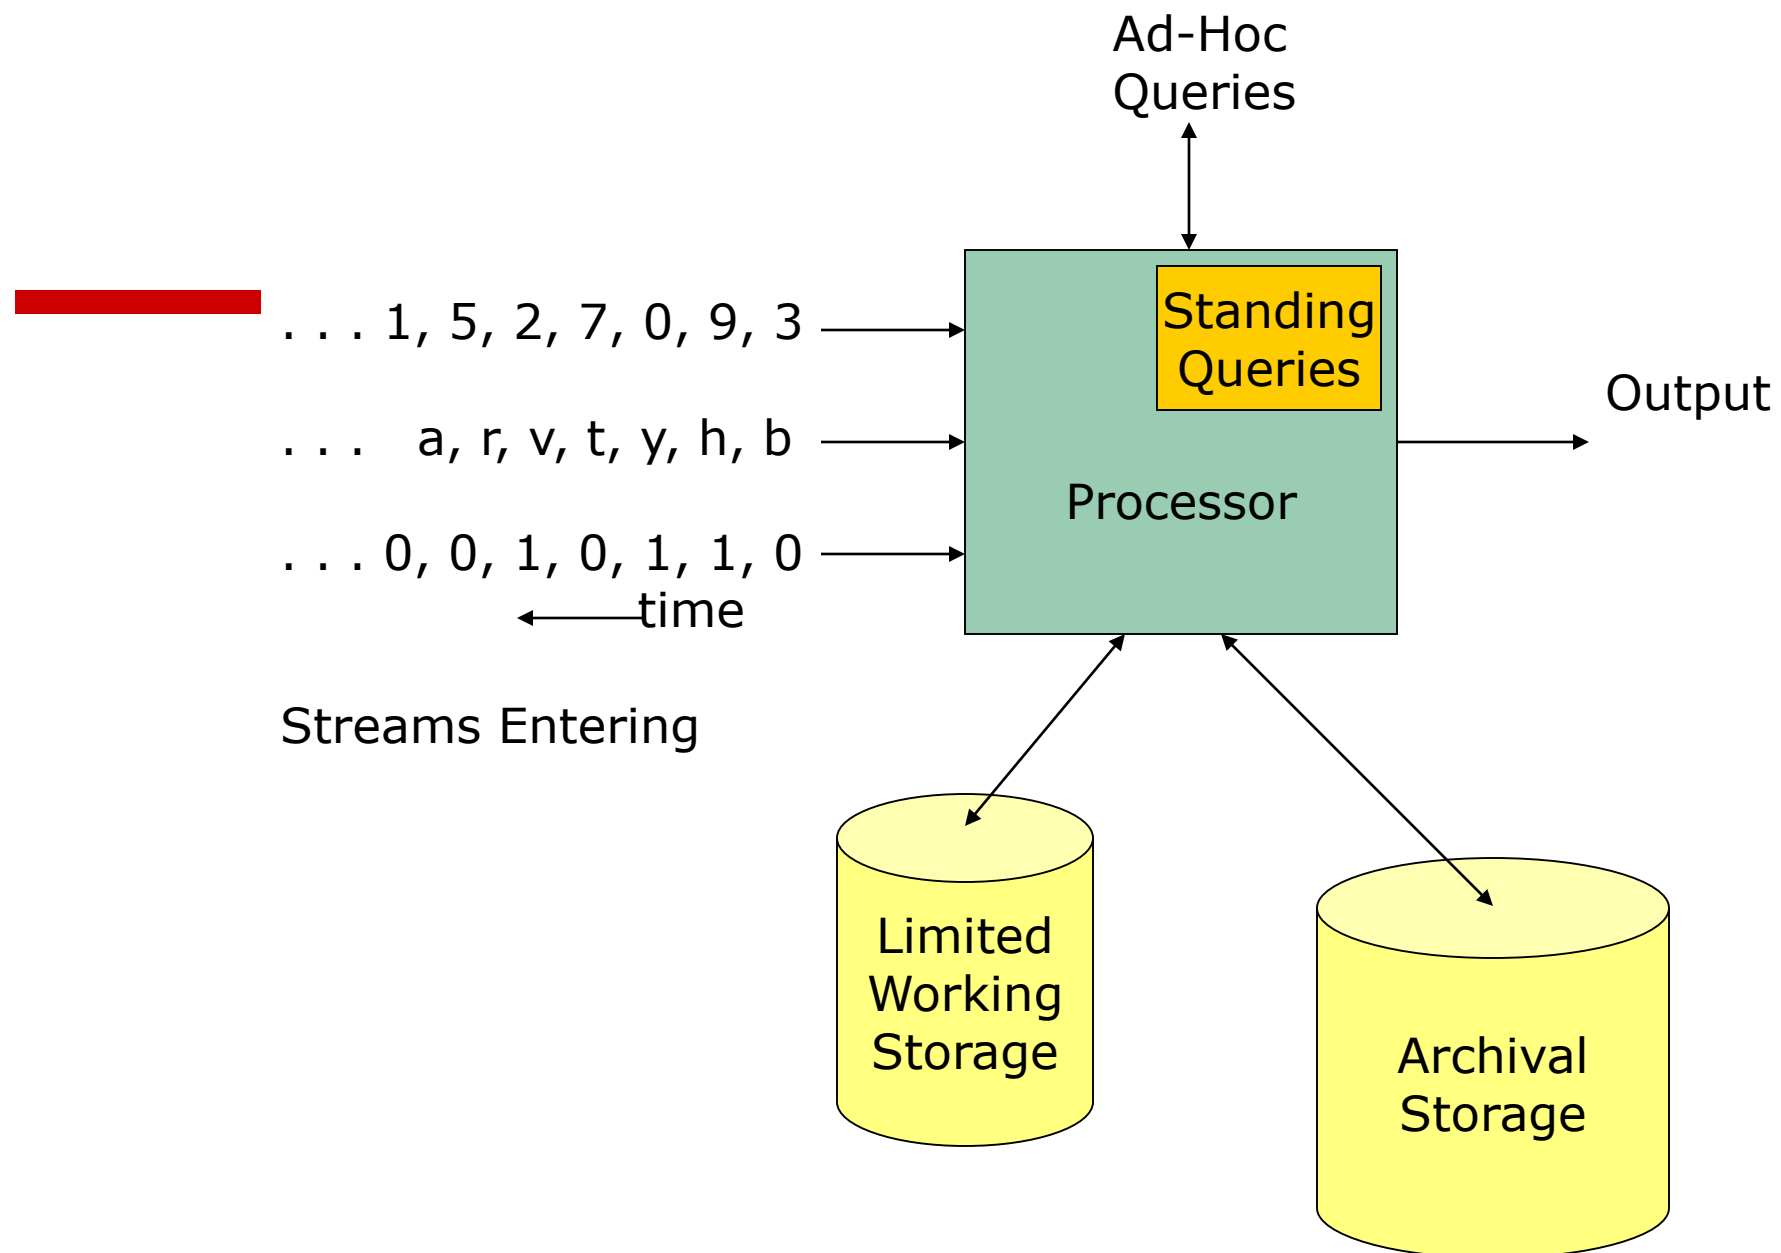
- ~~■ Counting 1's~~
- ~~■ Summing integers~~

Conventional Data Mining

- ❑ All data stored on slow hard disks/archives.
- ❑ Slow access, slow processing.
- ❑ Off-line analysis of data.
- ❑ Instant action on results impossible.
- ❑ However, trained MODELS (rules, decision trees, neural networks, etc.) can be deployed to work in real time.

The Stream Model

- ❑ Data enter the system at a rapid rate, at one or more input ports
- ❑ The system cannot store the entire stream accessibly
- ❑ How do you make critical calculations about the stream using a limited amount of fast (RAM) memory?
- ❑ We might be interested in queries over :
 - “the last N records” (a sliding window model), or
 - “everything seen so far (or the last N days)”



Applications – (1)

- Mining Twitter's tweets
 - Detecting breaking news, disasters, scandals, ...

- Mining query streams
 - Google wants to know what queries are more frequent today than yesterday

- Mining click streams
 - Yahoo wants to know which of its pages are getting an unusual number of hits in the past hour

Applications – (2)

- ❑ Sensors of all kinds need monitoring, especially when there are many sensors of the same type, feeding into a central controller.
 - Detecting/predicting traffic jams
 - Tsunami alerts

- ❑ Telephone call records are summarized into customer bills.

Applications – (3)

- IP packets can be monitored at a switch
 - Gather information for optimal routing
 - Detect denial-of-service attacks
 - Filter spam or phishing attacks

1. Sampling a data stream

- Suppose a bank would like to have (at any moment) a “representative” sample of 3% of all their data, so it would fit into RAM and could be queried interactively (quick response)
- **Example query:** *what is the ratio of clients with exactly 1 trx to clients which made exactly 2 trxs during the last weekend?*

Naïve Approaches

- ❑ Scan the stream of all transactions; the n-th transaction is put into the sample if
$$r = (n \bmod 100) < 3$$
(0, 1, 2, 100, 101, 102, ...)
- ❑ Use a random selection with $p=3\%$
- ❑ **Will not work:**
 - Many clients with 2, 3, ... transactions would be “viewed” as those which made a single transaction (see 4.2.1)

Better Approach

- ❑ Make, off-line, a list of all clients, take a 3% sample of them, put them into a hash table (or a binary search tree) and store only transactions from these selected clients.
- ❑ (If needed, use not only client_id as the key, but also other fields, like card type, terminal type, etc.)
- ❑ Problems:
 - Cost time and memory (to keep the hash table)
 - What about new clients (we will never see them!)?

An Even Better Approach

- - instead of a hash table, use a hash function to map each client to an integer 1, ..., 100;
 - transactions of clients with number 1, 2, 3 enter the sample; others are ignored

- Advantage:
 - No need to create and keep any hash table!
 - New clients are also included in the sampling process

- Problem: what if the sample size exceeds available RAM?

Sampling a stream in limited RAM

- ❑ Instead of a pre-specified percentage of clients that we want to cover (e.g., 3%), specify the amount of RAM that can be used for storing the sample.
- ❑ Use a hash function with, say, $L=10.000$ buckets (as before) and put all the records from buckets $1, 2, \dots, L$ to your sample. L is initially set to 10.000 .
- ❑ Whenever the size of your samples reaches the size of available RAM, **remove** all the records from **bucket L** and **set L to $L-1$** .

Reservoir Sampling

- Suppose your RAM can store **s records**; each record has the same size
- How could you sample an infinite stream of records, in such a way that at any moment your RAM keeps a **random, uniformly distributed, sample of s records**?
(uniformly distributed: every record from the stream has the same chance of being kept in RAM, i.e.,
 $s/(\text{\#records seen so far})$)
- **Initialization:** the first **s records are stored** in RAM.
- *How to proceed with further records, $s+1, s+2, \dots$ in such a way that at any moment, any element of the stream has the same chance of being in RAM?*

A Solution: Reservoir Sampling

□ Initialization:

Store the first s records of the stream in your RAM. At this moment $n=s$ and the probability of an element entering RAM the is s/n (accidentally, it's 1!)

□ Inductive Step:

- When the $(n+1)^{\text{th}}$ element arrives, decide with probability $s/(n+1)$ to keep the record in RAM (otherwise, ignore it)
- If you choose to keep it, throw one of the previously stored record out, selected with equal probability, and use the freed space for the new record.

□ Prove by induction that at any moment all records enter RAM with probability s/n (Chapter 4.5.5, page 181?)

Outline of the proof


- When the $(n+1)^{\text{th}}$ element arrives it is chosen to be stored with probability $s/(n+1)$ - that's what we wanted!
- However, all previous positions were chosen with probability s/n (by induction hypothesis) - that's bad - we need $s/(n+1)$!
- Fortunately, our procedure will delete one of the stored elements to create space for the newly selected element. This “random deletion” will modify the probabilities of “surviving” from s/n to $s/(n+1)$ - exactly what we need!

- $$\left(1 - \frac{s}{n+1}\right) \left(\frac{s}{n}\right) + \left(\frac{s}{n+1}\right) \left(\frac{s-1}{s}\right) \left(\frac{s}{n}\right) = \frac{s}{n+1}$$

Outline of the proof



n-th+1 element



s = buffer size

Every element has chance of ***s/n*** of being in buffer at moment ***n***; a new element arrives...

$$\underbrace{\left(1 - \frac{s}{n+1}\right) \left(\frac{s}{n}\right)}_{\text{Don't keep}} + \underbrace{\left(\frac{s}{n+1}\right) \left(\frac{s-1}{s}\right) \left(\frac{s}{n}\right)}_{\text{Keep}}$$

Survived

Simple implementation of Reservoir Sampling


(* S has items to sample, R will contain the result *)

ReservoirSample(S[1..n], R[1..s])

// fill the reservoir array

for i := 1 to s

 R[i] := S[i]



// replace elements with gradually decreasing probability

for i := s+1 to n

 (* randomInteger(a, b) generates a uniform integer from the inclusive range {a, ..., b} *)

 j := randomInteger(1, i)

 if j <= s


 R[j] := S[i]

2. Fast filtering of “bad” records

- Sometimes we know which records are “legal” and want to filter out all (or almost all) “illegal” records from a stream
- Examples:
 - A bank wants to ignore card numbers of clients which are not in their database
 - A mailing system wants to block all incoming e-mails from “non-registered” e-mail addresses (spam?)
 - Amazon.com wants to allow access only to their registered clients
 - Scenario: 1 billion “legal” card numbers, each 28 chars long

Bloom Filter

- ❑ Scenario: $N=1$ billion “legal” card numbers, each 28 chars long; how could we quickly filter out *almost all* “illegal” numbers?
- ❑ A hash table? How big should it be? **56GB RAM.**
- ❑ **Bloom Filter (one hash function):**
 - Consider a vector of **8 billion bits** (just **1GB of RAM!**) initialized to 0's and a hash function **$h(x)$** with values in $1, \dots, 8 \cdot 10^9$.
 - For each “legal” card number **x** , **set bit $h(x)$ to 1.**
 - An incoming **y** is considered to be legal if the **$h(y)$ bit is 1.**
- ❑ What is the percentage of *false positives*? (illegal cards that would pass through the filter)



Bloom Filter: accuracy

□ How many bits are set to 1?

■ $n=10^9$ legal cards

■ $N=8*10^9$ buckets

□ What is the chance that a randomly selected bit is **not set to 1**?

$$((N-1)/N)^n = [(1-1/N)^N]^{n/N} = e^{-1/8}$$

□ Thus the chance randomly hitting 1 is $1-\exp(-1/8)=\mathbf{0.1175}$, so the rate of false positive is about **11.75%**

□ Suppose that instead of one hash function we would use **k hash functions, h_1, \dots, h_k** and demand that y passes the filter when **all bits $h_1(y), \dots, h_k(y)$ are set to 1** ...

Bloom Filter with k hash functions

- How many bits are set to 1?
 - $n=10^9$ legal cards
 - $N=8*10^9$ buckets
 - $k=2$ the number of hash functions

- What is the chance that a randomly selected bit is **not set to 1**?


$$((N-1)/N)^{kn} = ((1-1/N)^N)^{kn/N} \approx e^{-k \frac{n}{N}}$$

- Thus the **chance of hitting 1 by k independent hash functions** is:


$$(1 - ((1-1/N)^N)^{kn/N})^k \approx (1 - e^{-k \frac{n}{N}})^k$$

- $k=2, n/N=1/8 \Rightarrow (1-\exp(-1/4))^2=0.0493$, so **the false positives rate is about 5%**

- What number of hash functions k is optimal? **$k=(\#bits/\#elements)*\ln 2$**



Accuracy of Bloom Filter for $n=10^9$; $N=8*10^9$



Final Remarks on Bloom Filters

- ❑ **Cascading**: using two or more filters one after another reduces errors “exponentially fast” (e.g., $5\% * 5\% = 0.25\%$).
- ❑ **Inserting** new elements to a filter is **easy**.
- ❑ **Removal of elements from a filter is “almost impossible”** (*Why?*)
- ❑ If two Bloom filters F_A and F_B represent sets A and B then the **bitwise AND** of F_A and F_B represents the **intersection** of A and B and the **bitwise OR** of F_A and F_B represents the **union** of A and B .
(*Why? What about false positive rates for such constructions?*)

More applications of Bloom Filters:

- ❑ Google [BigTable](#) and [Apache Cassandra](#) use Bloom filters to reduce the disk lookups for non-existent rows or columns. Avoiding costly disk lookups considerably increases the performance of a database query operation.
- ❑ The [Google Chrome](#) web browser used to use a Bloom filter to identify malicious URLs. Any URL was first checked against a local Bloom filter, and only if the Bloom filter returned a positive result was a full check of the URL performed (and the user warned, if that too returned a positive result).
- ❑ The [Squid Web Proxy Cache](#) uses Bloom filters for [cache digests](#).^[10]
- ❑ [Bitcoin](#) uses Bloom filters to speed up wallet synchronization.^{[11][12]}
- ❑ The [Venti](#) archival storage system uses Bloom filters to detect previously stored data.^[13]

3. Counting Distinct Elements

- ❑ **Problem**: a data stream consists of elements chosen from a set of size **n (n very big!)**.
How to maintain the count of the number of distinct elements seen so far?
- ❑ **Obvious approach**: maintain the set of elements seen (costs $O(n)$ memory!)
- ❑ *Use less memory (and accept loss of accuracy)*

Applications

- How many different URLs have we seen so far?
- How many different words are found among the Web pages being crawled at a site?
 - Unusually low or high numbers could indicate “artificial pages”
- How many different Web pages does each customer requests in a week?
- How many distinct elements in a column of a table?
(optimization of the join operation of two tables)
- How many distinct <source, destination> pairs through a router?
(detection of DoS attacks)

Using Small Storage

- **Real Problem:**
what if we do not have space to store the complete set?
- Estimate the count in an unbiased way.
- Accept that the count may be in error, but limit the probability that the error is large.

A simple idea: MinTopK estimate

- *Hash incoming objects into doubles from the interval $[0, 1]$ and count them shrinking the interval if needed.*
- Due to limited memory, maintain only the K biggest values (“TopK”), say, $K=1000$.
- Let s denote the minimum of our set (MinTopK)
- The number of distinct elements $\approx K/(1-s)$
- *What about the accuracy? The number of bits?*

