

# Probabilistic Counting

---

Partially based on slides from:

<http://infolab.stanford.edu/~ullman/mining/2009/index.html>

**Anand Rajaraman & Jeff Ullman**

# Counting Distinct Elements

---

- ❑ **Problem**: a data stream consists of elements chosen from a set of size  **$n$  ( $n$  very big!)**.  
*How to maintain the count of the number of distinct elements seen so far?*
- ❑ **Obvious approach**: maintain the set of elements seen (costs  $O(n)$  memory!)
- ❑ *Use less memory (and accept loss of accuracy)*

# Applications

---

- How many different URLs have we seen so far?
- How many different words are found among the Web pages being crawled at a site?
  - Unusually low or high numbers could indicate “artificial pages”
- How many different Web pages does each customer requests in a week?
- How many distinct elements in a column of a table?  
(optimization of the join operation of two tables)
- How many distinct <source, destination> pairs through a router?  
(detection of DoS attacks)

# Using Small Storage


---

- **Real Problem:**  
what if we do not have space to store the complete set?
- Estimate the count in an unbiased way.
- Accept that the count may be in error, but limit the probability that the error is large.

# A simple idea: MinTopK estimate

---

- *Hash incoming objects into doubles from the interval  $[0, 1]$  and count them shrinking the interval if needed.*
- Due to limited memory, maintain only the  $K$  biggest values (“TopK”), say,  $K=1000$ .
- Let  $s$  denote the minimum of our set (MinTopK)
- The number of distinct elements  $\approx K/(1-s)$
- *What about the accuracy? The number of bits?*



# Flajolet-Martin Approach

$X_1 \rightarrow 100011110 \rightarrow 1$   
 $X_2 \rightarrow 01010010 \rightarrow 1$   
 $X_3 \rightarrow 10011011 \rightarrow 0$   
 $X_4 \rightarrow 00101000 \rightarrow 3$   
 $X_5 \rightarrow 01011110 \rightarrow 1$

---

...  $\rightarrow$  .....  $\rightarrow$  ...  
#bits #bits  
 $N \rightarrow \log_2(N) \rightarrow \log_2(\log_2(N))$   
*eg.:  $\log_2(\log_2(1.000.000)) < 5$*

- Key idea:
  - hash passing elements into short bitstrings,
  - store only the length of the longest tail of 0's,
  - “the more distinct elements” the longer the longest tail of 0's.
- Pick a hash function  $h$  that maps each of the  $m$  elements to  $\log_2 m$  bits.
- For each stream element  $a$ , let  $r(a)$  be the number of trailing 0's in  $h(a)$ .
- Record  $R =$  the maximum  $r(a)$  seen.
- Estimate the number of distinct elements as  $2^R$ . **WHY?**

# Why does it Work?

---

- The probability that a given  $h(a)$  ends in at least  $r$  0's is  $2^{-r}$ .
- If there are  $m$  different elements, the probability that  $R \geq r$  is  $1 - (1 - 2^{-r})^m$

Prob. all  $h(a)$ 's  
end in fewer than  
 $r$  0's.


Prob. a given  $h(a)$   
ends in fewer than  
 $r$  0's.

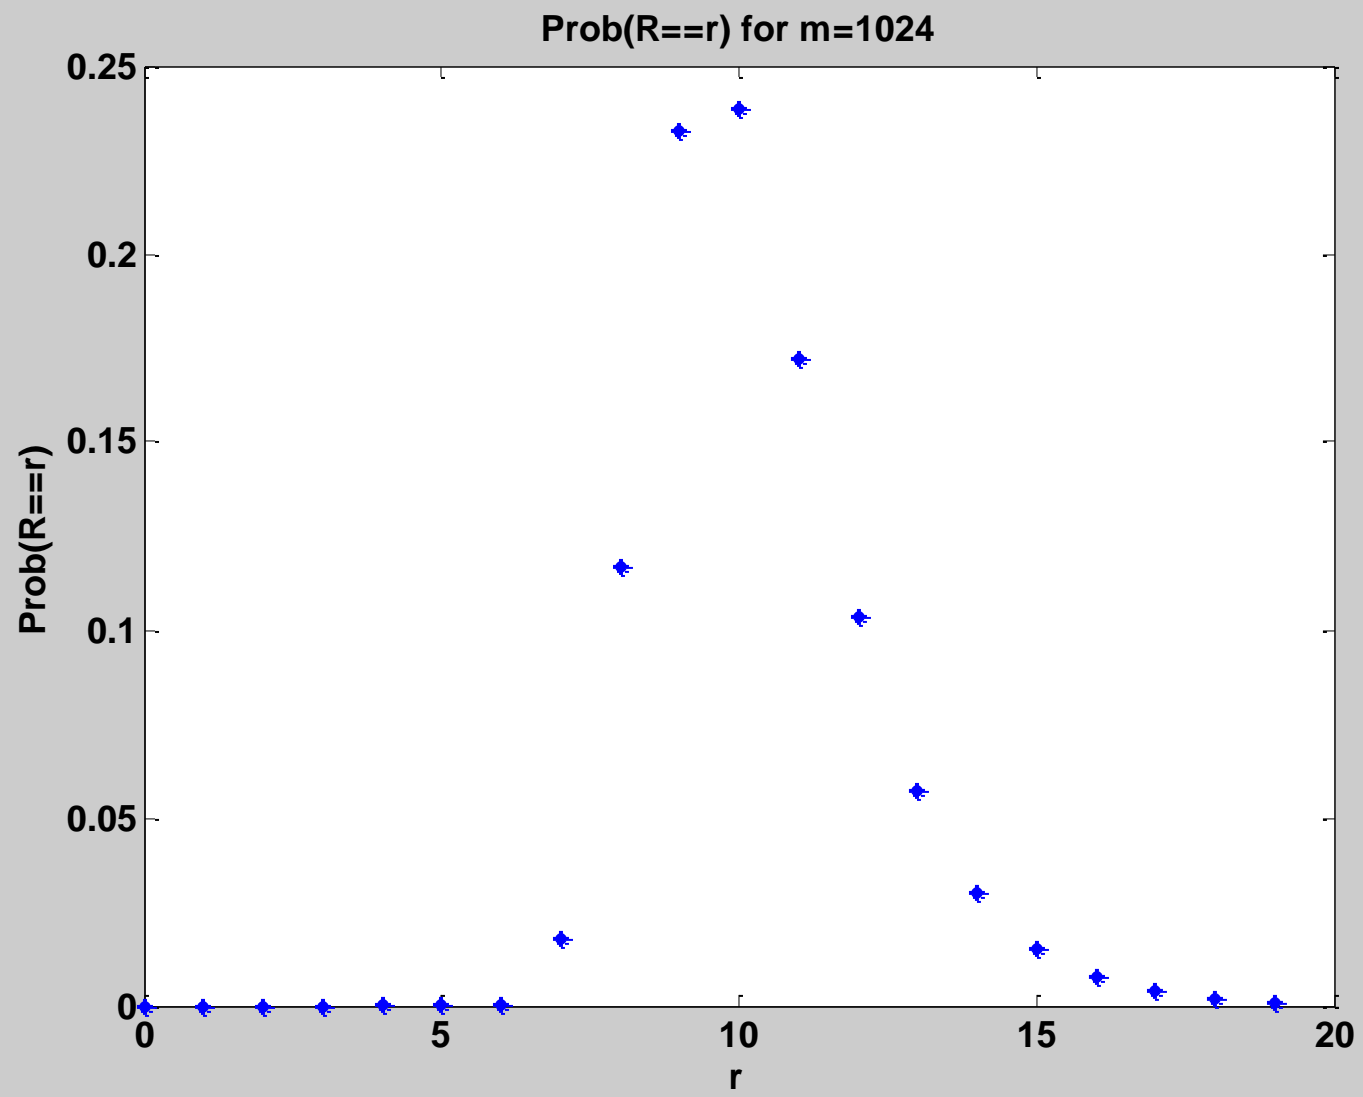
## Why does it Work – (2)

---

- Since  $2^{-r}$  is small,  $1 - (1 - 2^{-r})^m \approx 1 - e^{-m/2^r}$
- If  $2^r \gg m$ ,  $1 - (1 - 2^{-r})^m \approx 1 - 1 = 0$
- If  $2^r \ll m$ ,  $1 - (1 - 2^{-r})^m \approx 1 - 0 = 1$
- Thus,  $2^R$  will almost always be around  $m$







# Why It Doesn't Work?

---

- ❑  $2^R$  is always a power of 2 ....
- ❑ Bad luck can result in huge errors...
- ❑ Workaround: run in parallel several copies of this algorithm, using different hash functions and average the results
- ❑ How do we average results?
  - MEAN? What if one very large value?
  - MEDIAN? All values are a power of 2!

# Solution

---

- ❑ The Book:
  - Partition your hash functions into several groups.
  - Calculate the average of each group.
  - Then take the median of the averages.
  
- ❑ Does it really work?
- ❑ Check the original papers of Fajolet et al.
- ❑ Many variants (about 20), 1983-....

# Durand, Flajolet: the LogLog algorithm

---

- Key ideas: *stochastic averaging + calibration*
  - Partition your samples into  $n=2^l$  groups, using the first  $l$  bits of the hash function as a selector
  - Calculate  $R_1, \dots, R_n$ , ( $R_i$  for group  $i$ ) and return:  
 $a_n * n * 2^{\text{mean}(R_1, \dots, R_n)}$  where  $a_n$  = “a correction factor” (pre-computed)
  - $n=1024$  (*10 extra bits!*)  $\Rightarrow$  relative error 3% -4%

# Complexity

---

## ☐ Required memory

- We work with bit strings of **length  $\log_2 n$**
- We only have to maintain  
“**the length of the longest tail of 0’s**” -  **$\log_2(\log_2 n)$**   
- almost nothing:
  - ☐ **8 bits**  $\Rightarrow 2^{256} = 10^{77}$  objects
  - ☐ **5 bits**  $\Rightarrow 2^{32} = 10^9$  objects
- Multiply it by the number of hash functions:
  - ☐  **$1024 * 5 = 640$  bytes of memory**

## ☐ Time

- processing an element requires computing values of **ONE hash function** (linear in the length of input).

# Summary (Durand&Flajolet, 2003):

---

*The basic LogLog counting algorithm makes it possible to estimate cardinalities till  $10^8$  with a standard error of 4% using **1024 registers of 5 bits each**, that is, a table of **640 bytes in total**.*

```
ghfffghfghgghggggghghheehfhfhhgghghghhfgffffhhhiigfhhffgfiihfhhh  
igigighfgihfffghigihghigfhhgeegeghgghhhgghhhfidiigihighihehhhfgg  
hfgighigffghdieghhhggghhfhghhfiieffghghihifgggffihgihfggighgiiif  
fjgfgjhhjiifhjgehgghfhhfhjhiggghghihigghhihihgiighgfhlgjfgjjjml
```

The LOGLOG Algorithm with  $m = 256$  condenses the whole of Shakespeare's works to a table of 256 "small bytes" of 4 bits each. The estimate of the number of distinct words is here  $n^\circ = 30897$  (true answer:  $n = 28239$ ), i.e., a relative error of +9.4%.

## Recommended papers:

---

- ❑ P. Flajolet:  
**Counting by Coin Tossings**
- ❑ M. Durand and P. Flajolet:  
**Loglog Counting of Large Cardinalities**
- ❑ A. Metwally, D. Agrawal and A. El Abbadi:  
**Why Go Logarithmic if We Can Go Linear?**