

# Hadoop: HDFS and Map Reduce

---

Wojtek Kowalczyk  
(*Chapter 2 MMDS book*)


Based on:

<http://infolab.stanford.edu/~ullman/mining/2009/mapreduce.ppt>

[https://teaching.csse.uwa.edu.au/units/CITS3402/labs/project2-2018/amp\\_mapreduce.pdf](https://teaching.csse.uwa.edu.au/units/CITS3402/labs/project2-2018/amp_mapreduce.pdf)

# Single-node architecture

---



**Machine Learning, Statistics**

**"Classical" Data Mining**

# Commodity Clusters

---


- Web data sets can be very large
  - Tens to hundreds of terabytes or petabytes
- Cannot mine on a single server (why?)
- Standard architecture emerging:
  - Cluster of commodity Linux nodes
  - Gigabit ethernet interconnect
- How to organize computations?
- How to handle hardware failures?

# Google's dilemma: Computer cluster vs. Supercomputer ?

---



# Cluster Architecture



# Stable storage

---


- Typical usage pattern
  - Huge files (100s of GB to TB)
  - Data is rarely updated in place
  - Reads and appends are common
  - No overwrites (!)
  
- Main problem: nodes can fail -> how can we prevent data loss?
  
- Answer: Distributed File System
  - Provides global file namespace; a dedicated namenode(s)
  - Chunks, replicas, hashes, self-monitoring/healing mechanism
  - Google GFS; Kosmix KFS; **Hadoop HDFS**



# Distributed File System

---

- Files split into 128MB *blocks*
- Blocks replicated across several *datanodes* (usually 3)
- *Namenode* stores metadata (file names, locations, etc)
- Optimized for large files, sequential reads
- Files are append-only



# Warm up: Word Count (1)

---

- Input:
  - a large file of words, one word per line
  
- Task:
  - count the number of times each distinct word appears in the file
  
- Sample application:
  - analyze web server logs to find popular URLs (*what is a webserver log?*)

```
127.0.0.1 - frank [10/Oct/2000:13:55:36 -0700] "GET /apache_pb.gif HTTP/1.0"  
200 2326 "http://www.example.com/start.html" "Mozilla/4.08 [en] (Win98; I ;Nav)"
```



# Word Count (2)

---

- ❑ Case 1: Entire file fits in memory (RAM) -> trivial
- ❑ Case 2: File too large for RAM but all <word, count> pairs fit in RAM -> trivial
- ❑ Case 3: File on disk, too many distinct words to fit in memory
  - `sort datafile | uniq -c`  
*Sorting on HD is slow but "doable" -> check how?*

# Word Count (3)

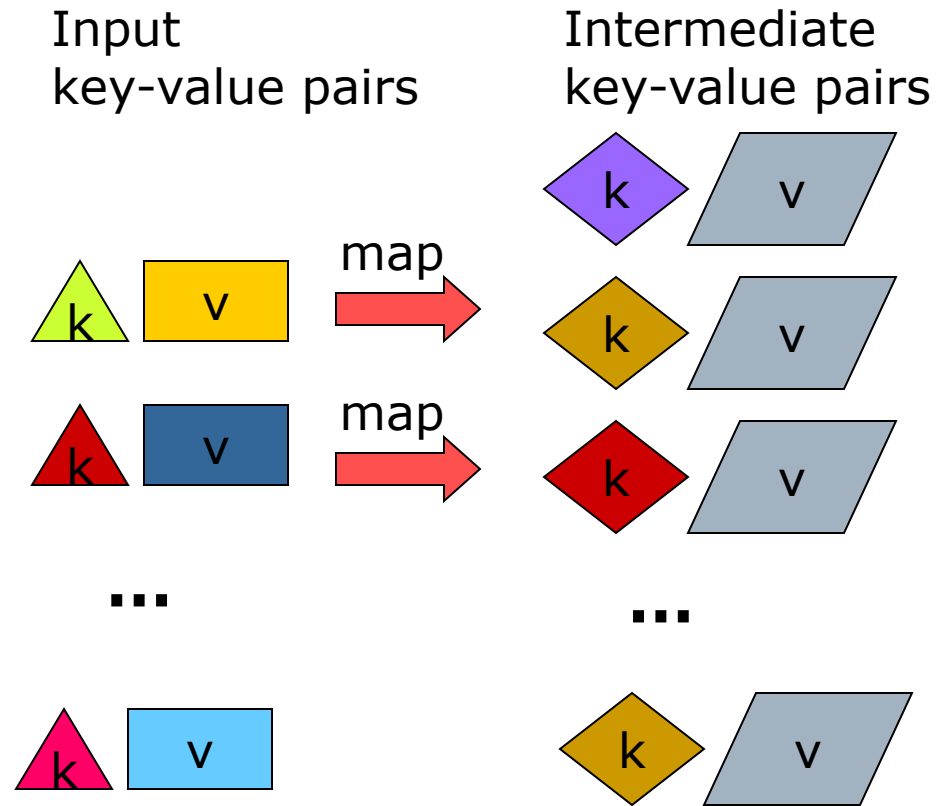
---

- ❑ To make it slightly harder, suppose we have a large corpus of documents
- ❑ Count the number of times each distinct word occurs in the corpus
  - `words(docs/*) | sort | uniq -c`
  - where `words` is a function that takes a file and outputs the words in it, one to a line
- ❑ The above captures the essence of MapReduce
  - Great thing is it is naturally parallelizable

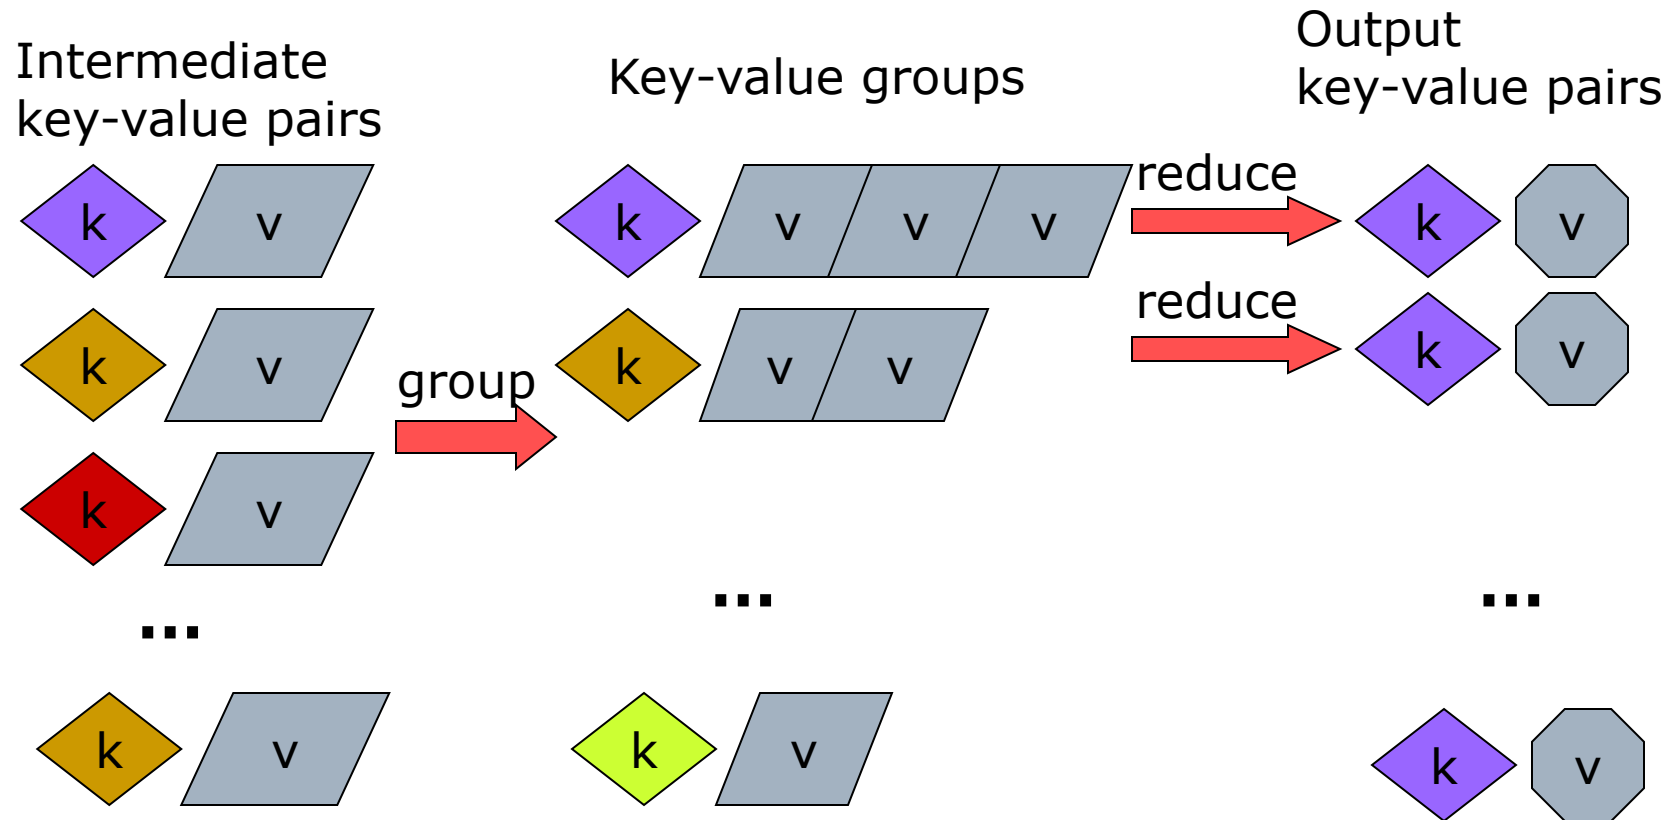
*Will not work when the harddisk is too small!*

# MapReduce: The Map Step

---



# MapReduce: The Reduce Step



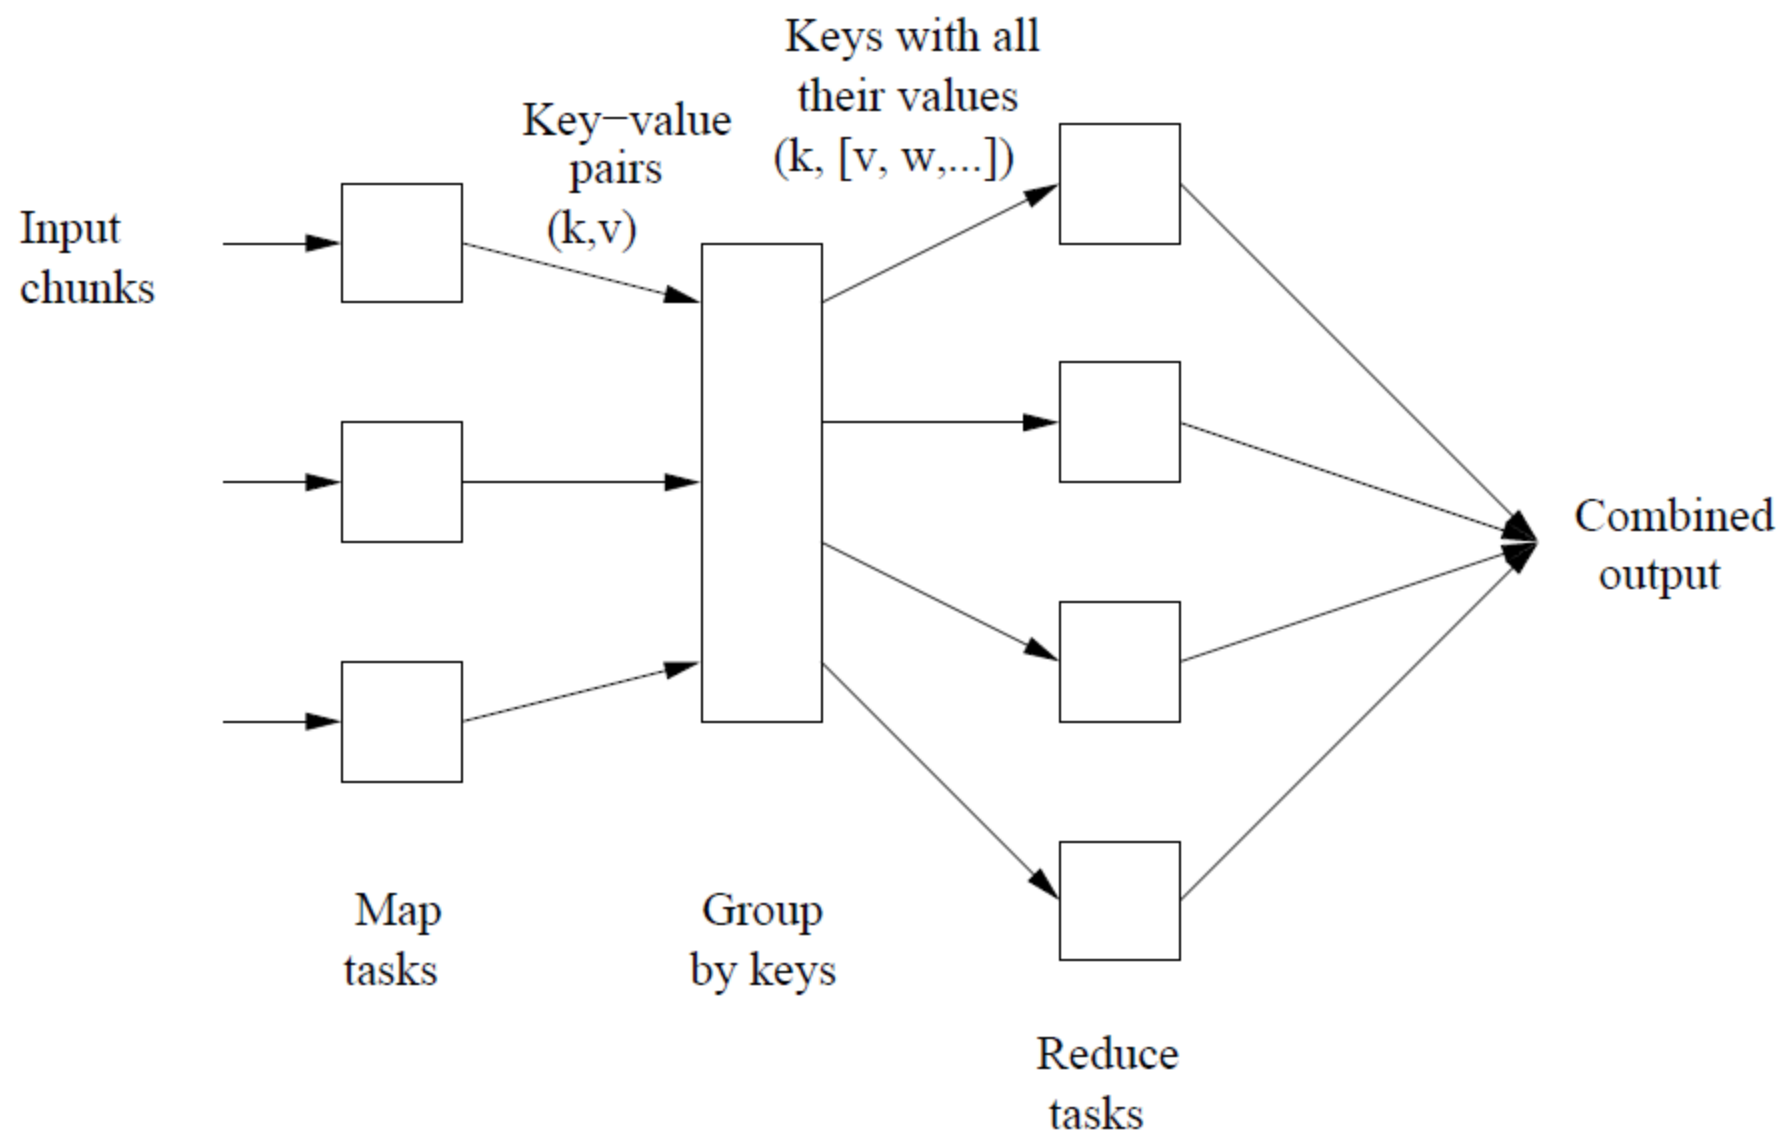


Figure 2.2: Schematic of a MapReduce computation

# MapReduce: an abstract model

---

- ❑ Input: a set of key/value pairs
- ❑ User supplies two functions:
  - $\text{map}(k,v) \rightarrow \text{list}(k_1,v_1)$
  - $\text{reduce}(k_1, \text{list}(v_1)) \rightarrow (k_1,v_2)$
- ❑  $(k_1,v_1)$  is an intermediate key/value pair
- ❑ Output is the set of  $(k_1,v_2)$  pairs

# Word Count using MapReduce

---

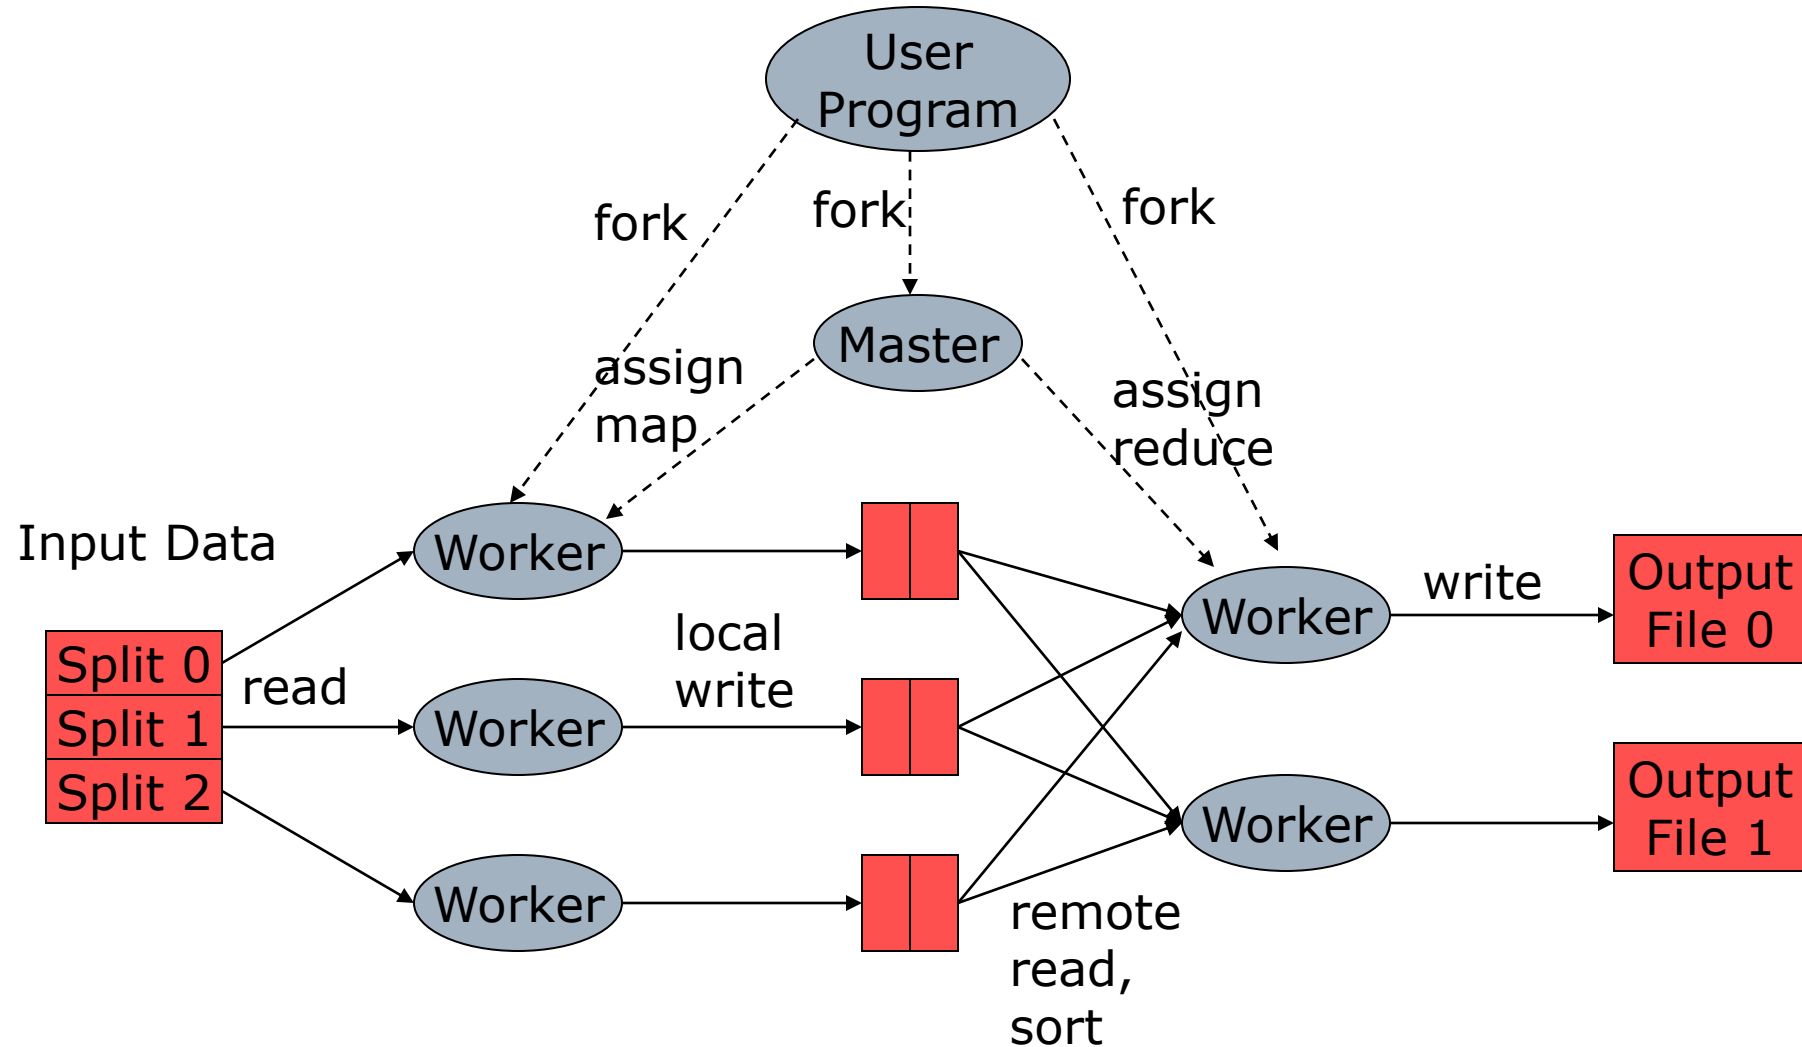
```
map(key, value):  
  // key: document name; value: text of document  
  for each word w in value:  
    emit(w, 1) //write to a harddisk (hdfs)
```

***"Magic": distribute & aggregate by keys [w, 1,1,1,1,1,...,1]***

```
reduce(key, values):  
  // key: a word; values: an iterator over counts  
  result = 0  
  for each count v in values:  
    result += v  
  emit(result) //write to a harddisk (hdfs)
```



# Distributed Execution Overview



# Data flow

---

- ❑ Input, final output are stored on DFS  
(a distributed file system)
  - Scheduler tries to schedule map tasks “close” to physical storage location of input data
- ❑ Intermediate results are stored on local FS of map and reduce workers
- ❑ Output is often input to another map reduce task

# Coordination

---

- Master data structures
  - Task status: (idle, in-progress, completed)
  - Idle tasks get scheduled as workers become available
  - When a map task completes, it sends the master the location and sizes of its R intermediate files, one for each reducer
  - Master pushes this info to reducers
- Master pings workers periodically to detect failures

# Failures

---

## ❑ Map worker failure

- Map tasks completed or in-progress at worker are reset to idle
- Reduce workers are notified when task is rescheduled on another worker

## ❑ Reduce worker failure

- Only in-progress tasks are reset to idle

## ❑ Master failure

- MapReduce task is aborted and client is notified

# How many Map and Reduce jobs?

---

- ❑ M map tasks, R reduce tasks
- ❑ Rule of thumb:
  - Make M and R much larger than the number of nodes in cluster
  - One DFS chunk per map is common
  - Improves dynamic load balancing and speeds recovery from worker failure
- ❑ Usually R is smaller than M, because output is spread across R files

# Combiners

---

- ❑ Often a map task will produce many pairs of the form  $(k, v_1), (k, v_2), \dots$  for the same key  $k$ 
  - E.g., popular words in Word Count
- ❑ Can save network time by pre-aggregating at mapper
  - $\text{combine}(k_1, \text{list}(v_1)) \rightarrow v_2$
  - Usually same as reduce function
- ❑ Works only if reduce function is commutative and associative, *like addition*:  $(a+b=b+a) \ \& \ (a+b)+c=a+(b+c)$

# Partition Function

---

- ❑ Inputs to map tasks are created by contiguous splits of input file
- ❑ For reduce, we need to ensure that records with the same intermediate key end up at the same worker
- ❑ System uses a default partition function e.g.,  $\text{hash}(\text{key}) \bmod R$
- ❑ Sometimes useful to override
  - E.g.,  $\text{hash}(\text{hostname}(\text{URL})) \bmod R$  ensures URLs from a host end up in the same output file



# Exercise 1: Host size

---

- Suppose we have a large web corpus
- Let's look at the metadata file:
  - Lines of the form [URL, host, size, date, ...]
- For each host, find the total number of bytes
  - i.e., the sum of the page sizes for all URLs from that host

# Host size: Solution

---

- $\text{map}(k,v) \rightarrow \text{list}(k1,v1)$
- $\text{reduce}(k1, \text{list}(v1)) \rightarrow (k1,v2)$
- $\text{map}(\text{line\_id}, \text{line\_content}) \rightarrow [\text{host}, \text{page\_size}]$
- $[\text{host}_1, s_{11}, s_{12}, \dots], [\text{host}_2, s_{21}, s_{22}, s_{32}, \dots], \dots$
- $\text{reduce}(\text{host}, [\text{sizes}]) \rightarrow [\text{host}, \text{sum}(\text{sizes})]$

# Exercise 2: Distributed Grep

---

- ❑ Find all occurrences of the given pattern in a very large set of files

## Unix “grep” function:

grep searches the named input FILES (or standard input) for lines containing a match to the given PATTERN (regular expression). By default, grep prints the matching lines.

- ❑ Example.:

```
[kowalczykwj@gold ~]$ history | grep das5
```

```
512 ssh wojtek@fs1.das5.liacs.nl
689 ssh das5
690 history |grep das5
837 ssh kowalczykwj@fs1.das5.liacs.nl
1005 history | grep das5
```

# Distributed Grep: solution

---

- ❑ `map(document_id, document)`  
    `->[document_id, 'grep output']`  
    (if non-empty)
- ❑ `reduce(document_id, ['grep outputs'])`  
    `->[document_id, ['grep outputs']]`

# Exercise 3: Graph reversal

---

- Given a directed graph as an adjacency list:  
    src1: dest11, dest12, ...  
    src2: dest21, dest22, ...
- Construct the graph in which all the links are reversed


# Example

---

□ Given a directed graph:

a: b,d

c: a,b,d




□ Construct the graph in which all the links are reversed:

a: c

b: a, c

d: a, d



# Graph reversal: solution

---

- Given a directed graph as an adjacency list:

src1: dest11, dest12, ...

src2: dest21, dest22, ...

.... : ....

- `map(src, [dest1, ..., dest_k])`  
-> `[[dest1, src], [dest2, src], ..., [dest_k, src]]`
- `reduce(dest, [src_1, ..., src_n])`  
-> `[dest, [src_1, ..., src_n]]`



# Exercise 4: Frequent Pairs

---

- Given a large set of market baskets, find all frequent pairs (i.e., frequency > threshold)
- Remember definitions from Association Rules lectures:  
  
given a set of “itemsets” (transactions)  $T$ , a set of items is called frequent if it occurs in  $>s\%$  of transactions ( $s$  is a parameter (support))

# Frequent Pairs: Solution

---

- `map(line_id, item_set) -> [(i1, i2), 1]`  
(all possible pairs from the item\_set)
- `reduce((i1, i2), [1, 1, 1,...1]) ->`  
if `sum([1, ...1]) > threshold`  
`emit([(i1,i2), sum(1's)])`;

## Exercise 5: Matrix-Vector Multiplication

---

- Given a large  $n \times n$  matrix  $M$ , represented by a list of  $(i, j, m_{ij})$  triplets and a vector  $v$  of length  $n$ , calculate  $w = Mv$

$$w_i = \sum_j m_{ij} v_j$$

- Scenario 1:  $M$  huge,  $v$  small enough to be kept in RAM
- Scenario 2: both  $M$  and  $v$  too big to be kept in RAM
- Motivation: the PageRank algorithm!!!

# Matrix-Vector Multiplication: Solution

---

- Scenario1:  $v$  fits in RAM
- $\text{map}(\text{line\_id}, (i, j, m_{ij}) \rightarrow [i, m_{ij}v_j]$
- $\text{reduce}(i, [m_{i1}v_1, m_{i2}v_2, \dots])$   
     $\rightarrow (i, \text{sum}(m_{ij}v_j))$   
    [the  $i$ -th element of  $Mv$ ]

$M$ =million by million  $\rightarrow$  8TB;

$v$ =million by 1  $\rightarrow$  8MB

# Matrix-Vector Multiplication: Solution

---

- Scenario2:  $v$  does not fit in RAM  
split  $M$  into stripes of  $k$  columns and  $v$  into chunks of length  $k$  and apply algorithm from the previous slide, summing up the results:




Figure 2.4: Division of a matrix and vector into five stripes

## Exercise 6: Matrix-Matrix Multiplication

---

- Given two large matrices  $M$ ,  $N$  calculate  $P=MN$ , given by:

$$p_{ik} = \sum_j m_{ij} n_{jk}$$

- *input:*  
(“ $M$ ”,  $i$ ,  $j$ ,  $m_{ij}$ ) or (“ $N$ ”,  $j$ ,  $k$ ,  $n_{jk}$ )

# Matrix-Matrix Multiplication: Solution

---

- *For each  $m_{ij}$  or  $n_{jk}$ :*
- *$map(("M", i, j, m_{ij})) \rightarrow [j, ("M", i, m_{ij})]$*
- *$map(("N", j, k, n_{jk})) \rightarrow [j, ("N", k, n_{jk})]$*
  
- *$reduce(j, [("M", i, m_{ij}), ("N", k, n_{jk})])$   
 $\rightarrow emit([(i, k), m_{ij}n_{jk}]), \text{ for all possible } (i, k)$*
  
- *To get  $p_{ik}$ , run another MapReduce to sum up all the terms with key  $(i, k)$ !*



# Limitations of MapReduce

---

- ❑ Usually a ***sequence of MapReduce operations*** is needed (complex SQL queries, construction of a decision tree, ... )
- ❑ Very ***limited/limiting "conceptual basis"*** for programmers: thinking in terms of "Map" and "Reduce"
- ❑ We need ***"higher level constructs"*** that can be incorporated into the Hadoop/MapReduce framework ...
- ❑ Performance: ***make better use of CPU's and RAM's!***
- ❑ ==> ***Pig, Hive, Hbase, Mahout, Storm, Spark, ...***
- ❑ [www.cloudera.com/products/open-source/apache-hadoop.html](http://www.cloudera.com/products/open-source/apache-hadoop.html)

# Building on Hadoop Map Reduce

---

Wojtek Kowalczyk

# Hadoop and MapReduce

---

- Hadoop Distributed File System (hdfs)
  - chunks, replicas, “read-only”, “write/append once”
  - unlimited scalability (million of nodes)
  - very robust, can run on a cluster/grid/WAN
- MapReduce:
  - **Map**: “process chunks of data”
  - *shuffle and sort (implicit, pre-programmed)*
  - **Reduce**: “aggregate partial results”

# Example algorithms in MapReduce

---

- word count
- distributed grep
- inverted index (documents -> “word index”)
- matrix multiplication
- PageRank
- atomic “database” operations: join, merge, group by, ...
- Locality Sensitive Hashing (LSH)
- ...
- *check the MMDS book!*

# Limitations of MapReduce

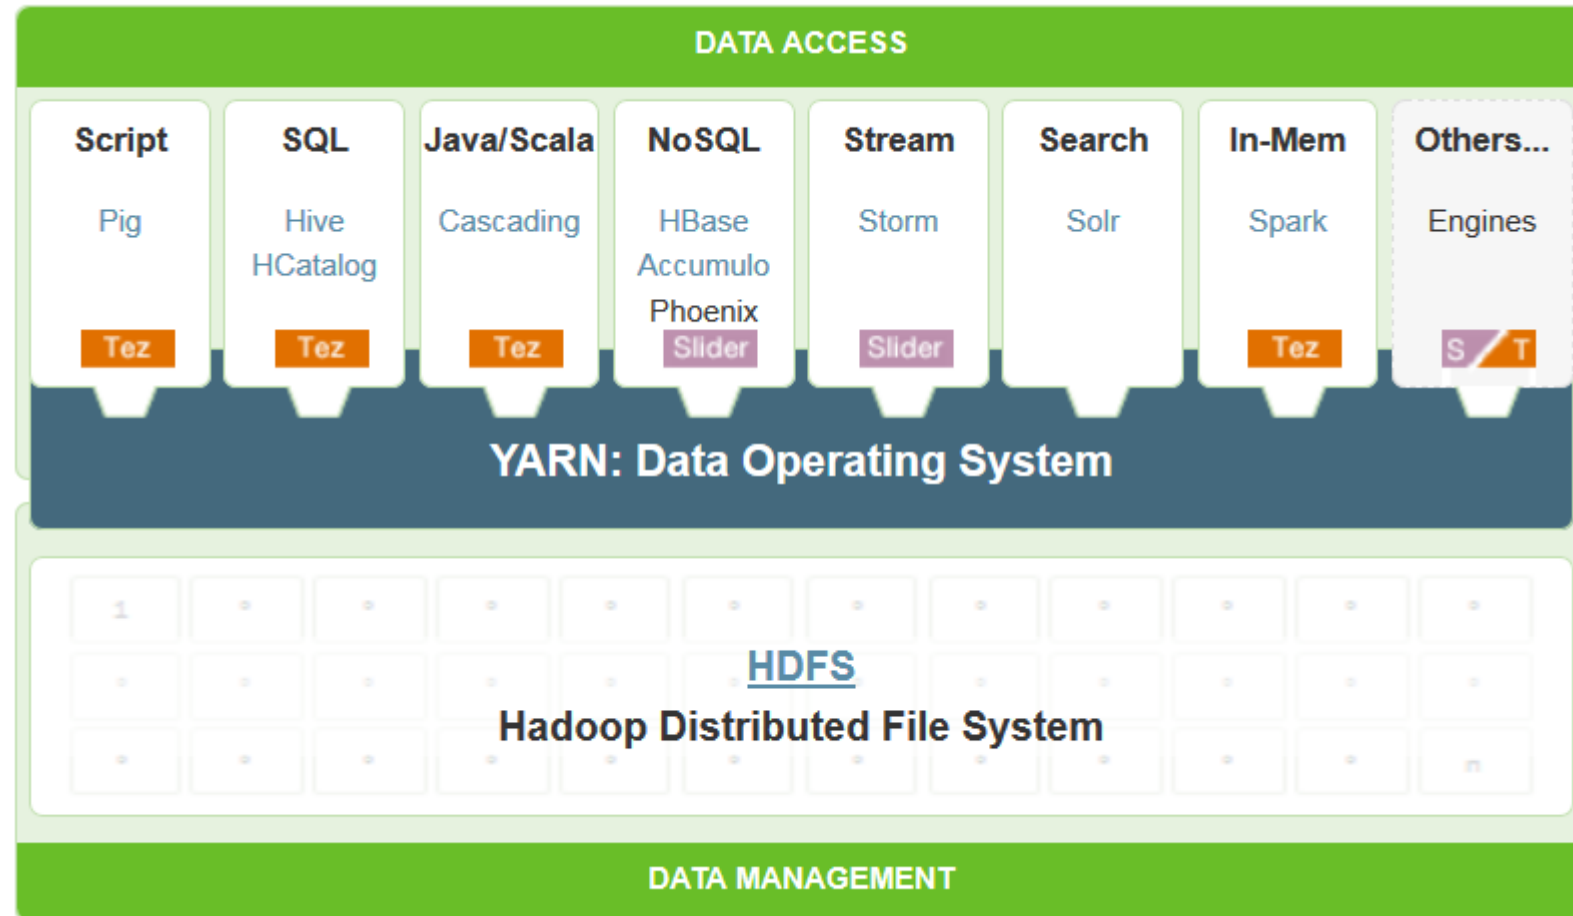
---

- ❑ Usually a ***sequence of MapReduce operations*** is needed (complex SQL queries, construction of a decision tree, ... )
- ❑ Very ***limited/limiting "conceptual basis"*** for programmers: thinking in terms of "Map" and "Reduce"
- ❑ We need ***"higher level constructs"*** that can be incorporated into the Hadoop/MapReduce framework ...
- ❑ Performance: ***make better use of CPU's and RAM's!***
- ❑ ==> ***Pig, Hive, Hbase, Mahout, Storm, Spark, ...***


# Hadoop Data Platform (→ Cloudera.com)

<http://hortonworks.com/hdp/>

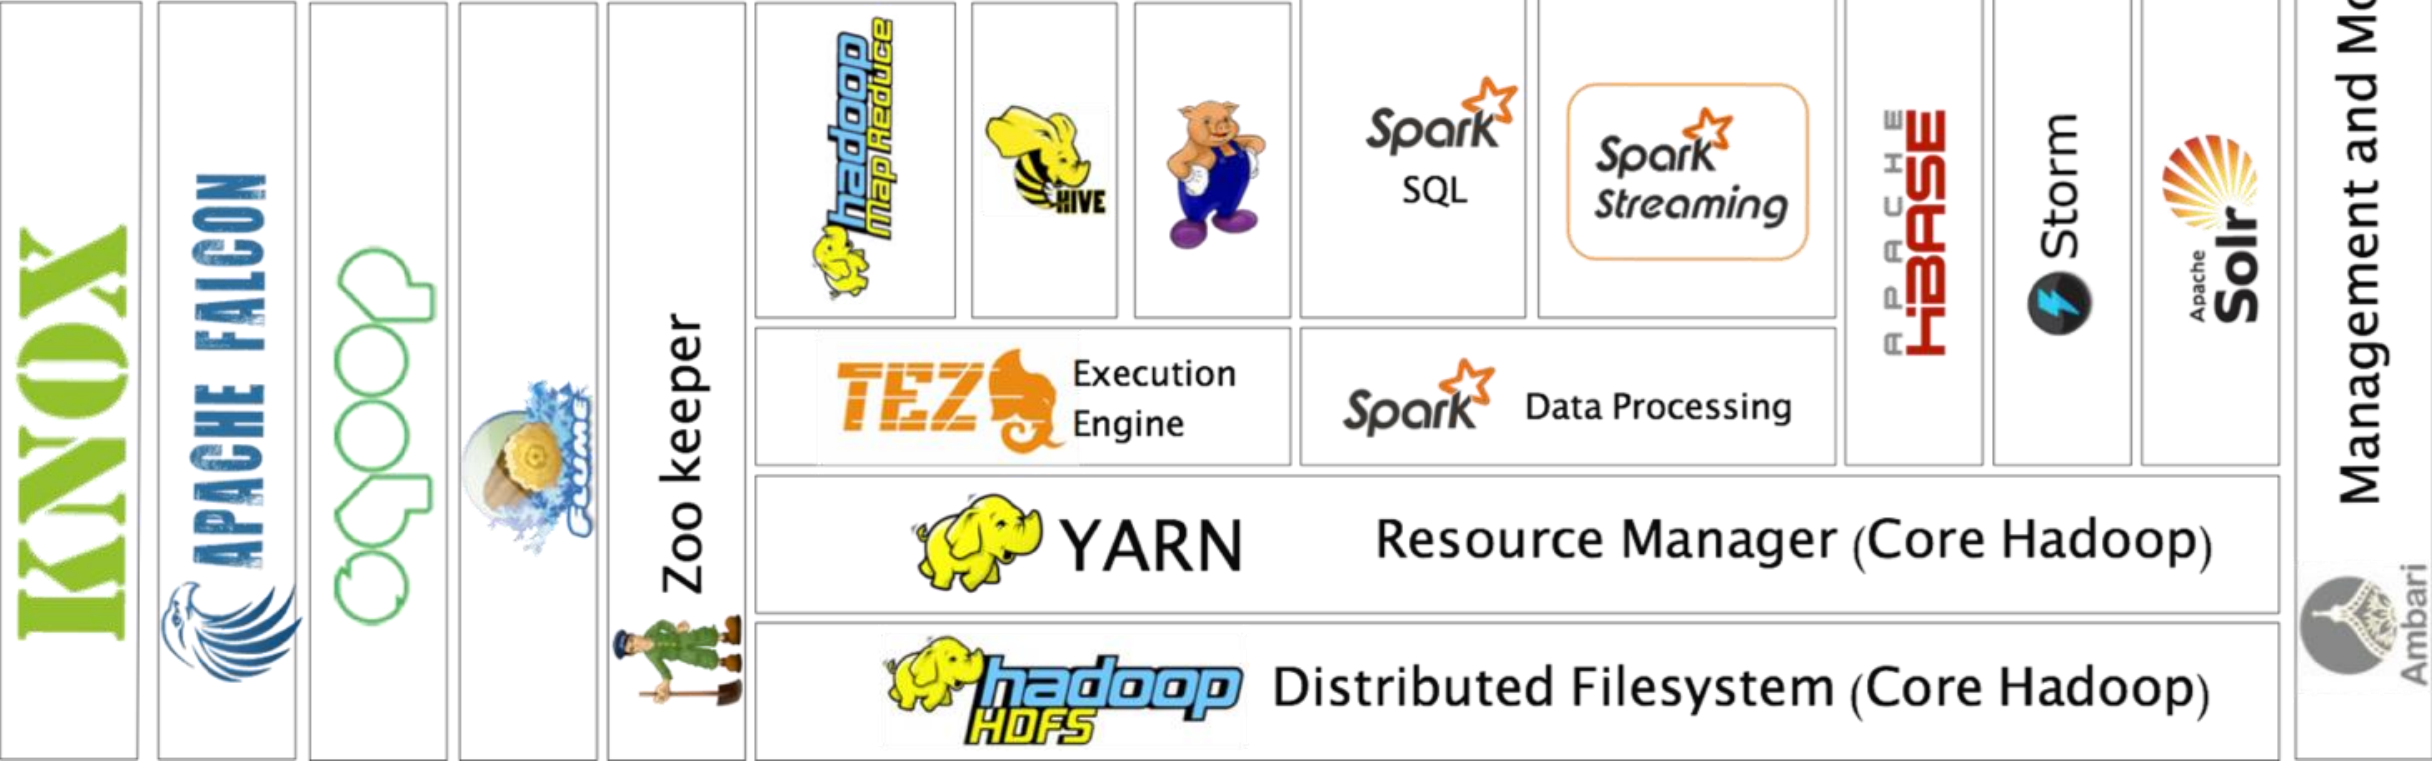
---



# Hadoop Overview



Hadoop Ecosystem



# Components

---

- **Core Hadoop Ecosystem**

(definitions and use cases)

- ☐ HDFS
- ☐ YARN
- ☐ Mesos
- ☐ Zookeeper
- ☐ MapReduce
- ☐ Spark
- ☐ Storm
- ☐ Pig
- ☐ Hive
- ☐ Sqoop
- ☐ Oozie
- ☐ Kafka\*\*
- ☐ Flume
- ☐ Flink
- ☐ Ambari

- **Query Engines**

- ☐ Drill
- ☐ Phoenix
- ☐ Presto
- ☐ Hue
- ☐ Zeppelin
- ☐ \* Impala

- **External Data Storage**

- ☐ MySQL
  - ☐ HBase
  - ☐ Cassandra
  - ☐ MongoDB
-



<https://spark.apache.org/>

---

