

AiDM 9: From Ensembles to XGboost

- Bagging
- Boosting
- Stacking
- Random Forests + Boosting + Bagging

Chapters 15,16 (ESLI book)

<http://statweb.stanford.edu/~tibs/ElemStatLearn>

<https://www.youtube.com/watch?v=wPqtzj5VZus>

The XGBoost algorithm

<https://xgboost.readthedocs.io/en/stable/tutorials/model.html>

<https://www.youtube.com/watch?v=Vly8xGnNiWs>

Many ML algorithms for Classification/Regression

- Decision/regression Trees
- Distance based: Nearest-Neighbor, kernel methods
- Statistical (Linear/Logistic Regression, NaïveBayes, BayesianNetworks, MixtureModels, ...)
- SVMs, SVMs with kernels => <http://mmds.org/mmds/v2.1/ch12-ml2.pdf>
- ...
- The Weka Book/Software/Site:
www.cs.waikato.ac.nz/ml/weka/mooc/dataminingwithweka/

Ensembles of models

Main idea:

instead of building a single model build
several models and combine their outputs

Fact:

It ***usually works better*** than a single model!

How to do it?

Breiman, 96: Bagging

Given: a training set \mathbf{T} and a learning method (e.g. C4.5, Backprop, Naive Bayes,).

- 1) Create, say, 100 versions of \mathbf{T} by ***resampling it with replacement***: $\mathbf{T}_1, \dots, \mathbf{T}_{100}$
- 2) For each \mathbf{T}_i build a classifier
- 3) Return an “ensemble” of classifiers:

Classification: majority voting

Regression: average

Bagging: Advantages

- + No thinking, no tuning
- + Better accuracy (almost for free) / “regularization”
- More computations => easily to run on parallel
- Loss of interpretability

Bagging=bootstrap aggregating

$$\text{error} = \text{bias} + \text{variance}$$

bias - limitation of the learning method

variance- limitation of the training data

bagging reduces variance

makes sense when:

- the “basic learner” is sensitive to small changes in data
- the “basic learner” overfits the data (e.g. a full tree)
- the data is scarce

Boosting=learning on errors of others

Main idea:

build a sequence of classifiers C_1, C_2, \dots, C_k , as follows:

C_1 is trained on the original data

C_2 “pays more attention” to cases misclassified by C_1 ,

C_3 “pays more attention” to cases misclassified by C_1 and C_2 ,
etc.

“pay more attention” \Rightarrow “try to correct errors” \Rightarrow

\Rightarrow “attach bigger weights to misclassified cases”

In other words, boosting develops a number of ‘experts’ that specialize in different regions of the data (the more difficult case the more attention it gets).

Freund, Schapire, Breiman: Boosting

Construct an ensemble of classifiers C_1, \dots, C_k , as follows:

Let C_1 - a “normal” classifier trained on dataset T
(C_1 becomes initial ensemble E_1)

For $i=2:k$

- assign to every case in T a **weight** that is “**proportional to the error**” made on this case by the current ensemble;
- train C_i tries to correct errors made by E_{i-1} ; $E_i = E_{i-1} + C_i$

“consecutive classifiers are trained to correct errors made by the previous ensemble”


Many strategies of weighting errors and weighting C_i 's !

Boosting: properties

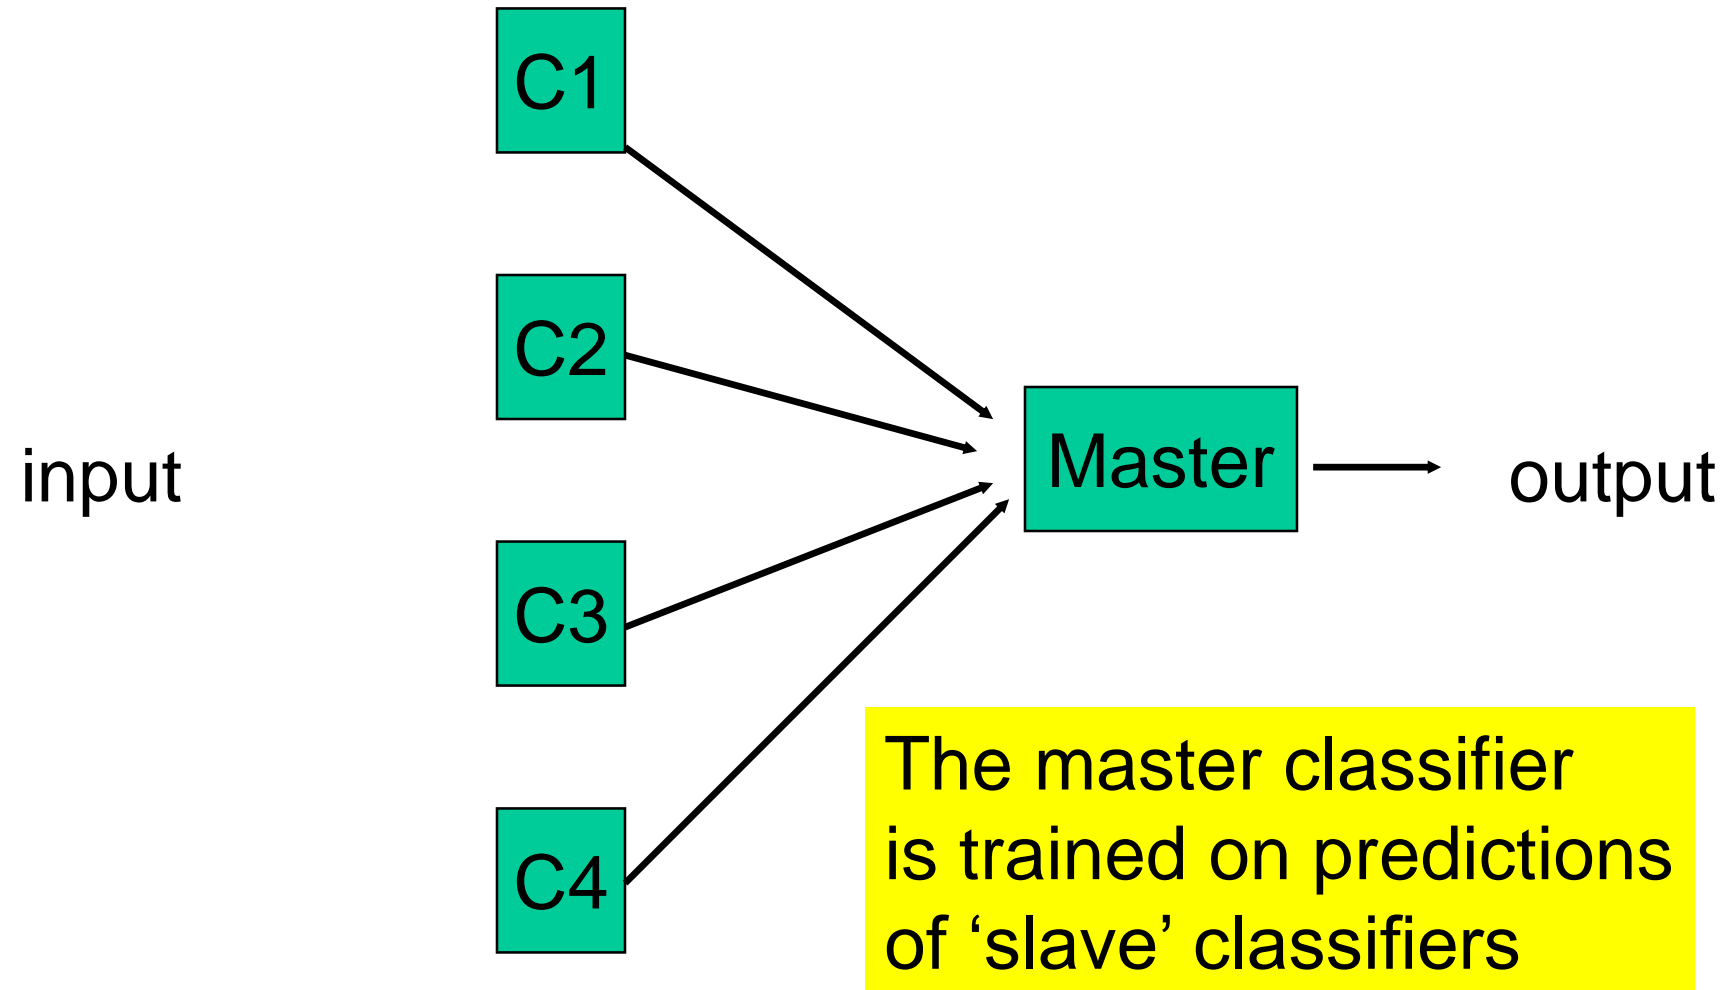
- Reduces exponentially fast the error on the training set
- Does not overfit the training set ! (sometimes...)
- Most successful with primitive “base classifiers”, e.g., decision stamps, linear regression
- Models expensive to build & difficult to interpret

Example:

separate points within the circle from points outside
with help of “decision stamps”: $x > \text{“a value”}$ or $y > \text{“a value”}$



Stacking: learning how to combine experts



Benchmark Comparisons (many years ago)




- Bagging > boosting > single classifiers
- Boosting > bagging > single classifiers
- SVM>single classifiers
- SVM=Boosting=Bagging

Boosting, Bagging and SVMs are usually better than single classifiers (about 10-20% relative error reduction)

Random Forests (L. Breiman, 2001)

Algorithm 15.1 *Random Forest for Regression or Classification.*



1. For $b = 1$ to B :
 - (a) Draw a bootstrap sample \mathbf{Z}^* of size N from the training data. (looks like bagging!)
 - (b) Grow a random-forest tree T_b to the bootstrapped data, by recursively repeating the following steps for each terminal node of the tree, until the minimum node size n_{min} is reached.
 - i. Select m variables at random from the p variables. → (add more randomness: best out of m instead of best of all!)
 - ii. Pick the best variable/split-point among the m .
 - iii. Split the node into two daughter nodes.
2. Output the ensemble of trees $\{T_b\}_1^B$.

To make a prediction at a new point x :

Regression: $\hat{f}_{rf}^B(x) = \frac{1}{B} \sum_{b=1}^B T_b(x)$.

Classification: Let $\hat{C}_b(x)$ be the class prediction of the b th random-forest tree. Then $\hat{C}_{rf}^B(x) = \text{majority vote } \{\hat{C}_b(x)\}_1^B$.

Advantages:



- Superior accuracy (comparable with SVM's, GBDT, NN)
- No cross-validation needed (Out-Of-Bag error estimate)
- Few parameters to tune; highly robust (not very sensitive)
- Trivial to parallelize (?)
- Provides a heuristic measure of importance of variables

Conclusions



- Ensembles usually outperform single models
- Ensembles of “weak” classifiers often outperform single sophisticated classifiers
- Boosting may overfit the data
- Gentle AdaBoost:
one of many ways of combining/weighting classifiers

Case Study: Benelearn99 competition

A classification problem:

- 5823 cases for training; 4000 for validation
- “Yes”/“No” rate: 6%
- 85 variables

Objective:

select 800 cases from the validation set
with the highest chance of being “Yes”

STEP 1: Framework

Different data sets require different methods ...
How can we decide which one is most appropriate?

Cross-validation:

- data split into train and test set; 6:4
- stratified sampling (the same response rate)
- performance = percentage of correctly predicted “1’s” in top 20%
- each experiment repeated 10 times



average accuracy

stability

Conventional Approach

I tried:

Decision trees (C4.5, CHAID, CART, ...)

Naive Bayes Classifier

Logistic Regression

Results:

Logistic Regression > Naive Bayes > Decision Trees

Logistic Regression: best accuracy & smallest variance

Accuracy varying from **11%** to **13%** (expected **88-104** hits)

How could I make it better ???

Boosting Accuracy: Gentle AdaBoost

Gentle AdaBoost (Friedman, Hastie, Tibshirani, 98):

fits an ***additive logistic regression*** model with
a Newton optimization algorithm:

$$score = \sum_i \sum_j a_{ij}(x_i = v_j)$$

Results:

15% on raw data ! (solution 1)

15.5% on pre-processed data (solution 2)

pre-processing: visual inspection of cross-tables 9 variables selected;
one variable manually discretized recoded

Boosting Gentle AdaBoost

- interactions of second order: $(x_i=v) \& (x_j=w)$
=> overfitting
- “clever” weighting strategies
=> no improvement
- “better” pre-processing (recoding)
=> no improvement
- other error measures
=> no improvement

Impressions ...



Gentle AdaBoost is a magic algorithm that:

- outperforms classical methods
- is very fast (7 minutes on pre-processed data)
- no parameters involved (number of iterations?)
- no overfitting (!)
- no preprocessing necessary
- very easy to implement (**27** lines of MATLAB code !!!)

Result: a triple winner!

Benelearn99:

- 1) Boosting with pre-processed data (129)
 - 2) Boosting with original data (124)
 - 3) Rough Data Models (118)
- Others: 62-112 hits

Baselines:

- random selection: about 48 hits
- “straightforward approach”: around 96 hits

Update 2019

**My implementation of boosting contained
“critical conceptual errors”!**
(a wrong weight-update mechanism)

Main conclusion:

Always start with a good “model evaluation scheme”!

State-of-the-art: XGBoost


- Many variants of boosting (10++)
- Friedman (2001): Gradient Boosted Machines

<https://www.jstor.org/stable/pdf/2699986.pdf>

- **The best implementation: XGBoost**

<https://xgboost.readthedocs.io/en/stable/tutorials/model.html>

Homework:



Study the XGBoost site: you will need it for A3!