

Autoencoders and GANs (Ch.17)

- Autoencoders: the key idea
- Variants:
 - AE and PCA
 - Denoising
 - Sparse
 - Variational (VAE)
 - Convolutional, recurrent, ...

Study (and play) with the tutorial:


<https://blog.keras.io/building-autoencoders-in-keras.html>

Generative Adversarial Networks:

http://slazebni.cs.illinois.edu/spring17/lec11_gan.pdf

The simplest autoencoder

- Consider a 4:2:4 network that computes identity function



- Inputs: $(0\ 0\ 0\ 1)$, $(0\ 0\ 1\ 0)$, $(0\ 1\ 0\ 0)$, $(1\ 0\ 0\ 0)$
- Outputs: $(0\ 0\ 0\ 1)$, $(0\ 0\ 1\ 0)$, $(0\ 1\ 0\ 0)$, $(1\ 0\ 0\ 0)$

Autoencoders




- The network "learns" the concept of binary representation of integers!
- General case: $2^n:n:2^n$ networks
- Somehow the network is able to "extract" the most concise representation of the input!
- The same trick works for "real numbers": PCA networks
- The same for "bigger networks" with "coder" and "decoder" parts being a convolutional or recurrent networks!
- It can be used for data compression, dimensionality reduction, visualization, anomaly detection, data generator(!)

PCA and Autoencoders

- m inputs, n ($n \ll m$) hidden nodes, m outputs
- Activation function: $f(x)=x$ (identity function)
- the network reduces the dimensionality of data (from m to n) while preserving as much information as possible (MSE minimized)
- Equivalent to Principal Component Analysis
- What happens when activations are non-linear? More hidden layers?

Principal Components Analysis



Principal Components

Given: n variables x_1, x_2, \dots, x_n

PCs are linear combinations of x_1, \dots, x_n such that:

- 1) they are orthogonal to each other
- 2) they maximize the “variance of projections”
- 3) the first PC explains most of the variance, second PC less, ...

In practice the first few PCs (2-3) explain most of the variance

Remark:

PCs (may) have nothing to do with classification !!!

Finding Principal Components

Given: n variables x_1, x_2, \dots, x_n

1) calculate the matrix of covariances (or correlations)

$$C = (\text{cov}(x_i, x_j)) , i, j = 1, \dots, n$$

2) Calculate the eigen structure of the covariance matrix
 $[V, D] = \text{eig}(C)$, such that D is a diagonal matrix of eigen values
and V is a matrix of eigen vectors.

3) Choose k eigen vectors with the biggest eigen values to form k
principal components:

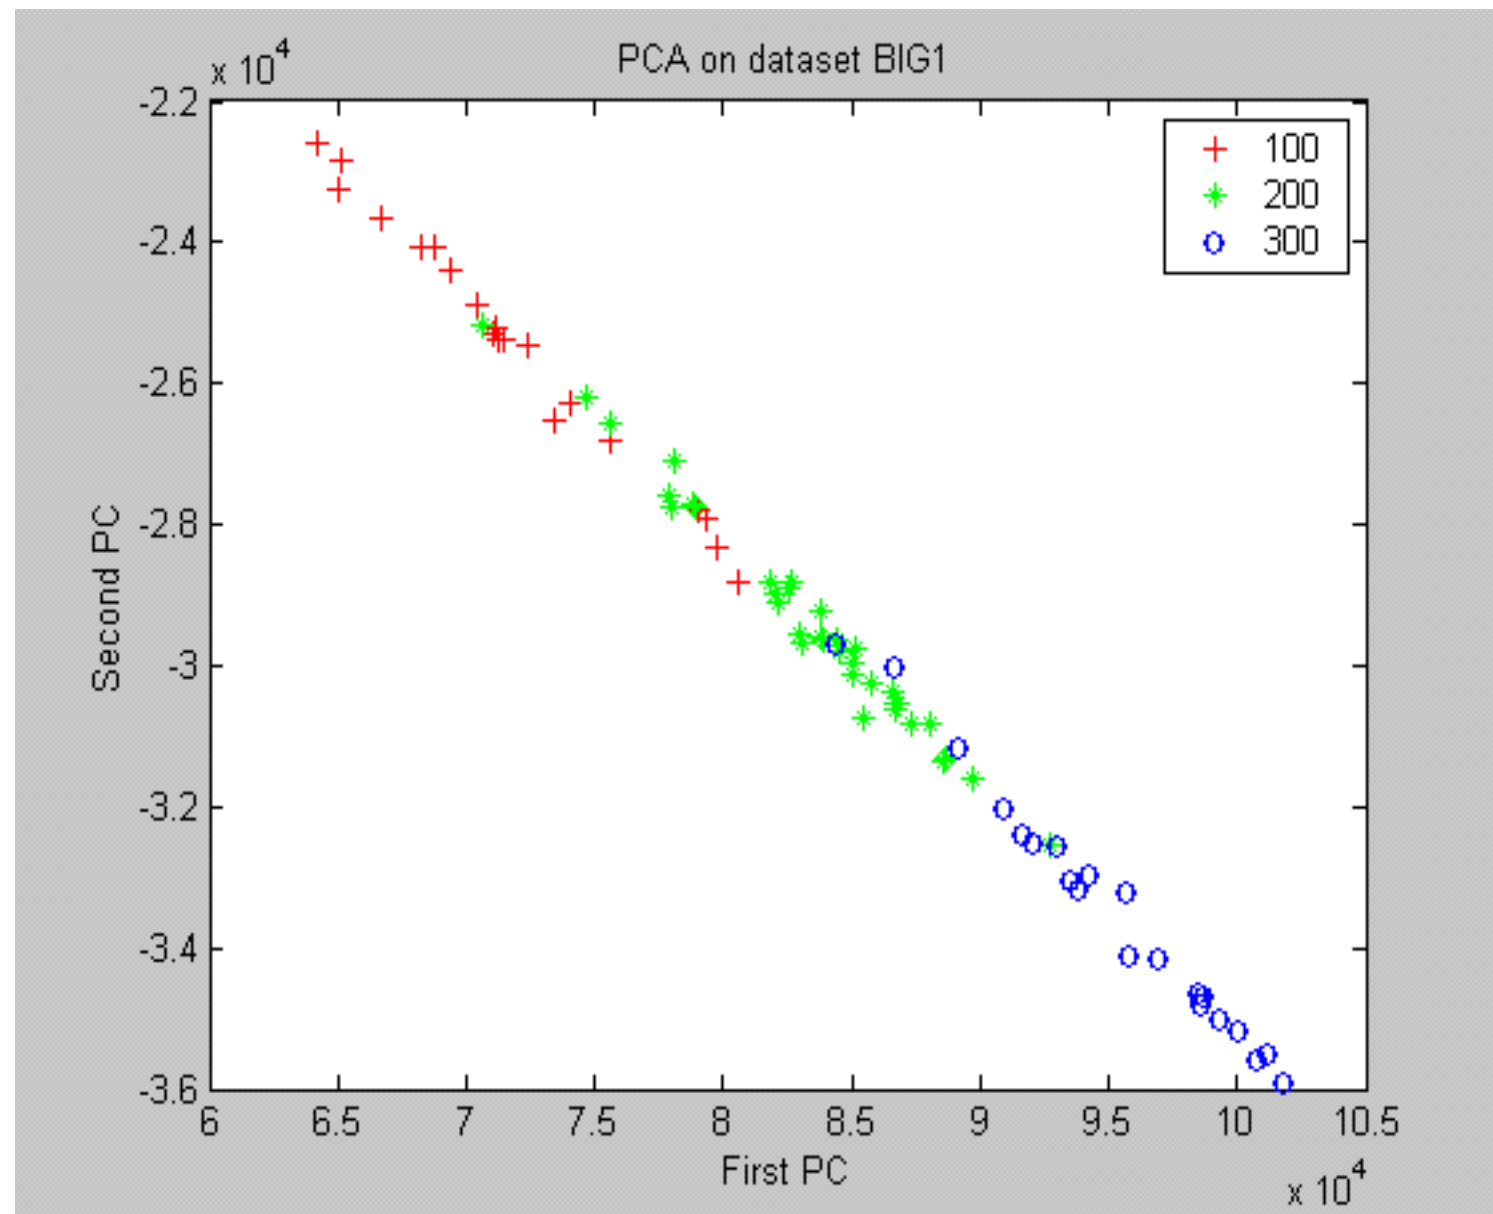
$$x \rightarrow x^* [v_1, v_2, \dots, v_k]' \quad (n \text{ dimensions} \rightarrow k \text{ dimensions})$$

Example application of PCA's: flower quality estimation

- a digital camera takes several photos of a flower
- some features are extracted from images (sizes, proportions, etc): numbers
- some flowers are labeled as “nice” others as “ugly”

Problem:

develop a system that would automatically estimate flower quality from 30 features



Denoising AutoEncoders

- Key idea:
- Given a set of “clean” images add noise to them and train an autoencoder on:
 - Input: a noisy image
 - Output: a clear image
- In this way the network “learns” what is essential in an image and what is not!

Sparse AutoEncoders

The key idea:

- The bottleneck layer computes some “essential features” that are needed for input reconstruction.
- Usually, only a few features are really important; the rest represent noise.
- => impose some constraints on the bottleneck layer:
e.g., “on average, activations are close to 0”.
- => **Kullback-Leibler divergence**: measures the distance between two probability distributions; used as a part of the loss function.

Variational AutoEncoders



The key idea:

- The bottleneck layer represents a *latent space*: vectors of features that are needed for reconstruction
- “to reconstruct an output we can slightly vary the values of the these essential features”
- Assume that the latent features are “normally distributed”

=> **Extend the bottleneck layer to include “means” and “std’s”.**

Mastering Autoencoders



Study the tutorial:

blog.keras.io/building-autoencoders-in-keras.html

- - get the code working (TF1.0 -> TF2.x)
- - run various experiments, playing with different parameters
- - play with alternative data sets (eg. FashionMNIST)

References on GANs

- Short explanation:

<https://www.slideshare.net/xavigiro/deep-learning-for-computer-vision-generative-models-and-adversarial-training-upc-2016>

- A longer one:

http://slazebni.cs.illinois.edu/spring17/lec11_gan.pdf

- Several lectures on GANs:

<https://sites.google.com/view/cvpr2018tutorialongans/>

- <https://thispersondoesnotexist.com/>