

# Modern Game AI Algorithms

## Monte Carlo Tree Search

Mike Preuss | LIACS



Universiteit  
Leiden  
The Netherlands



# roadmap

lecture plan (Thursdays 14:15), <https://smart.newrow.com/#/room/rzm-514>

Feb	18	introduction	Apr	8	procedural content generation II
	25	learning and optimization		15	player experience modeling
Mar	4	procedural content generation I	Apr	22	Georgios: spatial data analysis
	11	Monte Carlo tree search		29	believable behavior
	18	experimentation / ANN / GAN	May	6	team AI and mass behavior
	25	Vanessa: 1 year game industry...		13	realtime strategy AI
Apr	1	learning from pixels	May	20	free project presentations 1

labs plan (Thursdays 16:15)

Feb	18	no lab	Apr	8	free projects (assignment 3)
	25	single assignment: map generation		15	free projects (assignment 3)
Mar	4	single assignment (week 2)	Apr	22	free projects (assignment 3)
	11	single ass. deadline / start ass 2		29	free projects (assignment 3)
	18	assignment 2	May	6	free projects (assignment 3)
	25	assignment 2		13	deadline free projects
Apr	1	deadline ass. 2 / start free proj.	May	20	free project presentations 2

# assessment

what is graded:

- single assignment 25%, smaller group assignment 30%,  
free project assignment 45% (including presentation), no written exam

the single assignment is issued today (shall be on Brightspace already):

- simple programming task requiring a bit creativity, Python

smaller group project (largely March):

- group work, most likely in a competition-based fashion on the GVGAI (new version)
- groups of 2 intended

free group project (free means you choose the topic):

- working in small groups (4-5 people, exceptions possible)
- during the second half of the tutorials, and in your preparation time

# topics of today

- game trees, MIN-MAX and alpha-beta
- Monte Carlo Tree Search
- using it for „real“ problems (example)
- GGP and GVGAI
- Rolling Horizon EA



picture from janeb13 on Pixabay

# MCTS in Total War: Rome II (Creative Assembly)

- since the ROME II edition, MCTS is used for planning
- what are the advantages?  
randomness (surprise), avoid mistakes (alternatives tried), anytime algorithm
- what is it used for?  
resource allocation to tasks selection of actions (according to resource availability)
- domain knowledge leads to aggressive pruning and biased search (balancing/design)
- industry reluctant to take over, but many will follow

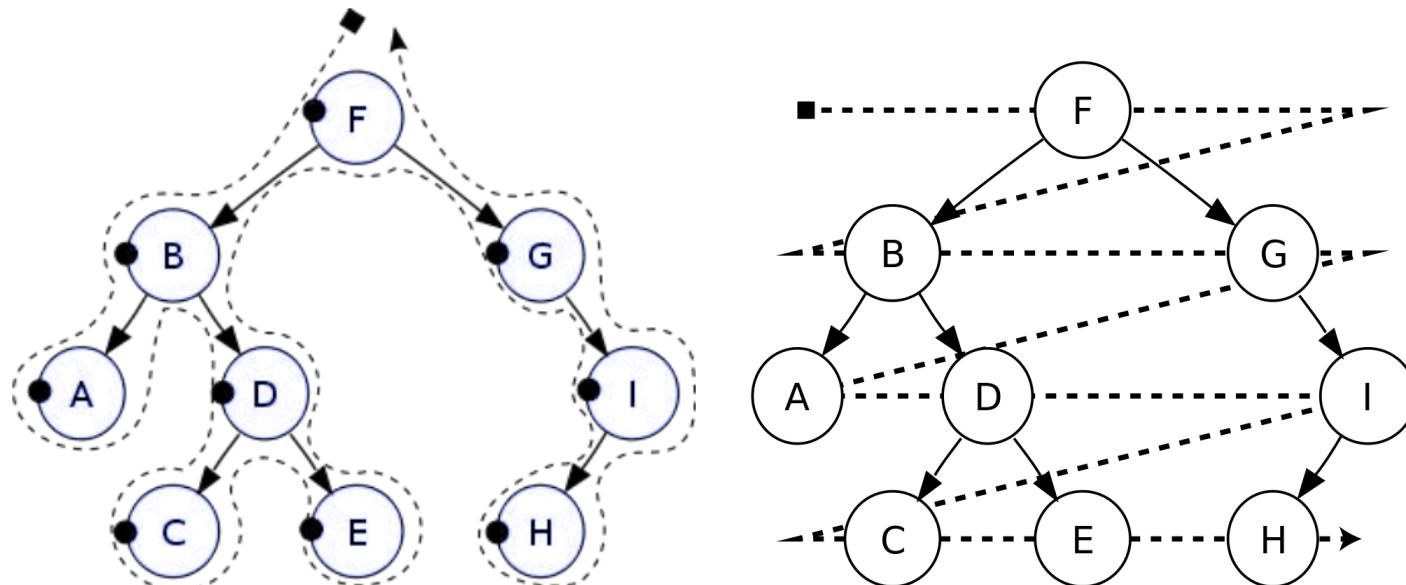
see Tommy Thompson's report:

[https://www.gamasutra.com/blogs/TommyThompson/20180212/314399/Revolutionary\\_Warfare\\_\\_The\\_AI\\_of\\_Total\\_War\\_Part\\_3.php](https://www.gamasutra.com/blogs/TommyThompson/20180212/314399/Revolutionary_Warfare__The_AI_of_Total_War_Part_3.php)



# (game) tree search

- any method that traverses (examines or updates), expands (creates) and approximates a (game) tree
- result of search: a strategy (optimal hopefully) that wins the game
- it requires end state outcomes and/or utilities/heuristics
- best-first, A\*, Dijkstra, breadth-first are examples...



# defeating Kasparov

- the Chess world champion was defeated by a min-max tree search algorithm with alpha-beta pruning in 1997
- seemingly partly due to a bug :-)
- can't we just transfer this to other types of games? not so easy
- Kasparov has interesting position toward AI nowadays:

<https://www.wired.com/story/defeated-chess-champ-garry-kasparov-made-peace-ai/>

"Technology is the main reason why so many of us are still alive to complain about technology."

Garry Kasparov



# what we can solve with min-max tree search

perfect information

Chess, Checkers  
Go, Othello  
Atari, GGP

partial information

Battleship

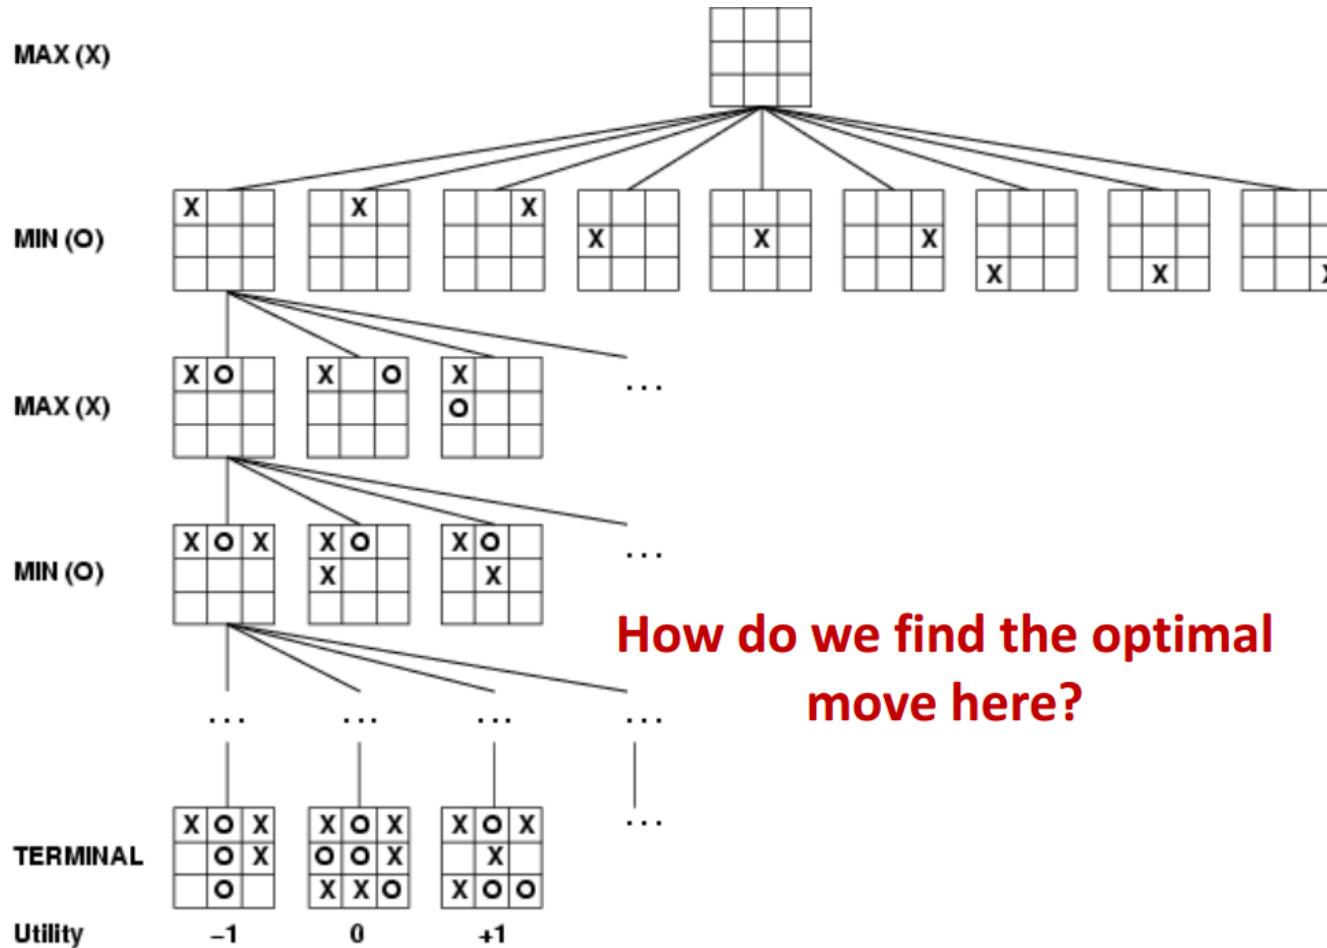
Backgammon  
GVGAI, sports games

StarCraft, Poker  
Hanabi, Bridge,  
MOBA, RTS, FPS

deterministic

non-deterministic

# tic-tac-toe: two-player, deterministic, turn-based



# MIN and MAX are the players

## MAX

- wants to **maximize** the result of the utility function
- winning strategy if, on MIN's turn, a win is obtainable for MAX for all moves that MIN can make

## MIN

- wants to **minimize** the result of the evaluation function
- winning strategy if, on MAX's turn, a win is obtainable for MIN for all moves that MAX can make

# utility functions

a utility function:

- estimates how good the current board configuration is for a player
- $\text{utility(state)}$

examples:

- typical values from -infinity (loss) to +infinity (win) or  
-1 (MIN wins), 0 (draw) , +1 (MAX wins)
- Othello: number of white pieces - number of black pieces
- Chess: value of all white pieces - value of all black pieces
- weighted sum of factors (e.g. Chess)

$$\text{utility}(S) = w_1 f_1(S) + w_2 f_2(S) + \dots + w_n f_n(S)$$

–  $f_1(S)$  = (number of white queens) – (number of black queens),  $w_1 = 9$

–  $f_2(S)$  = (number of white rooks) – (number of black rooks),  $w_2 = 5$

– ...

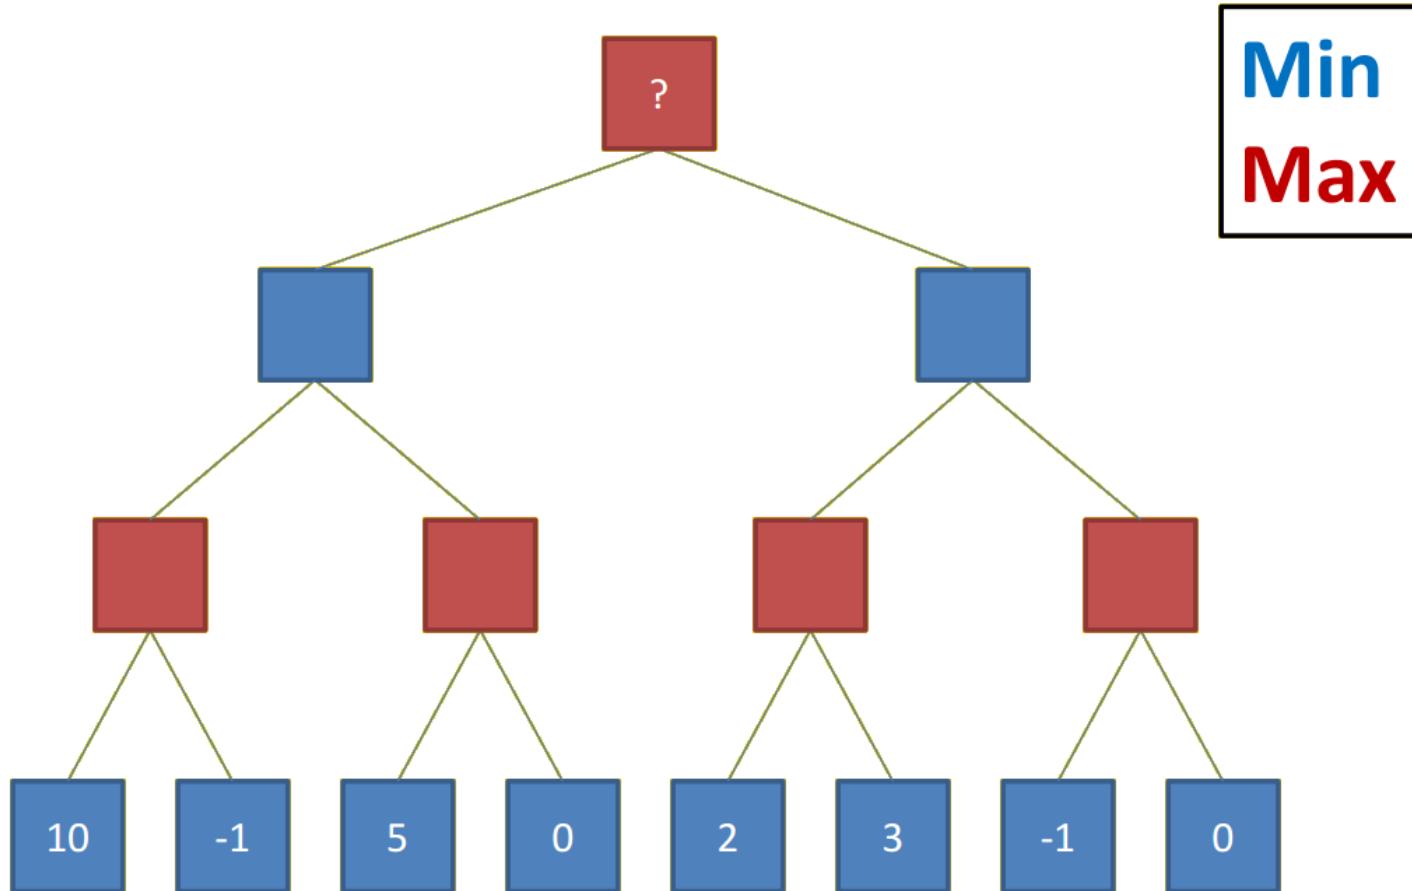
- if the utility is X for a player, it is -X for the opponent  
– “zero-sum game”

# MIN-MAX tree search

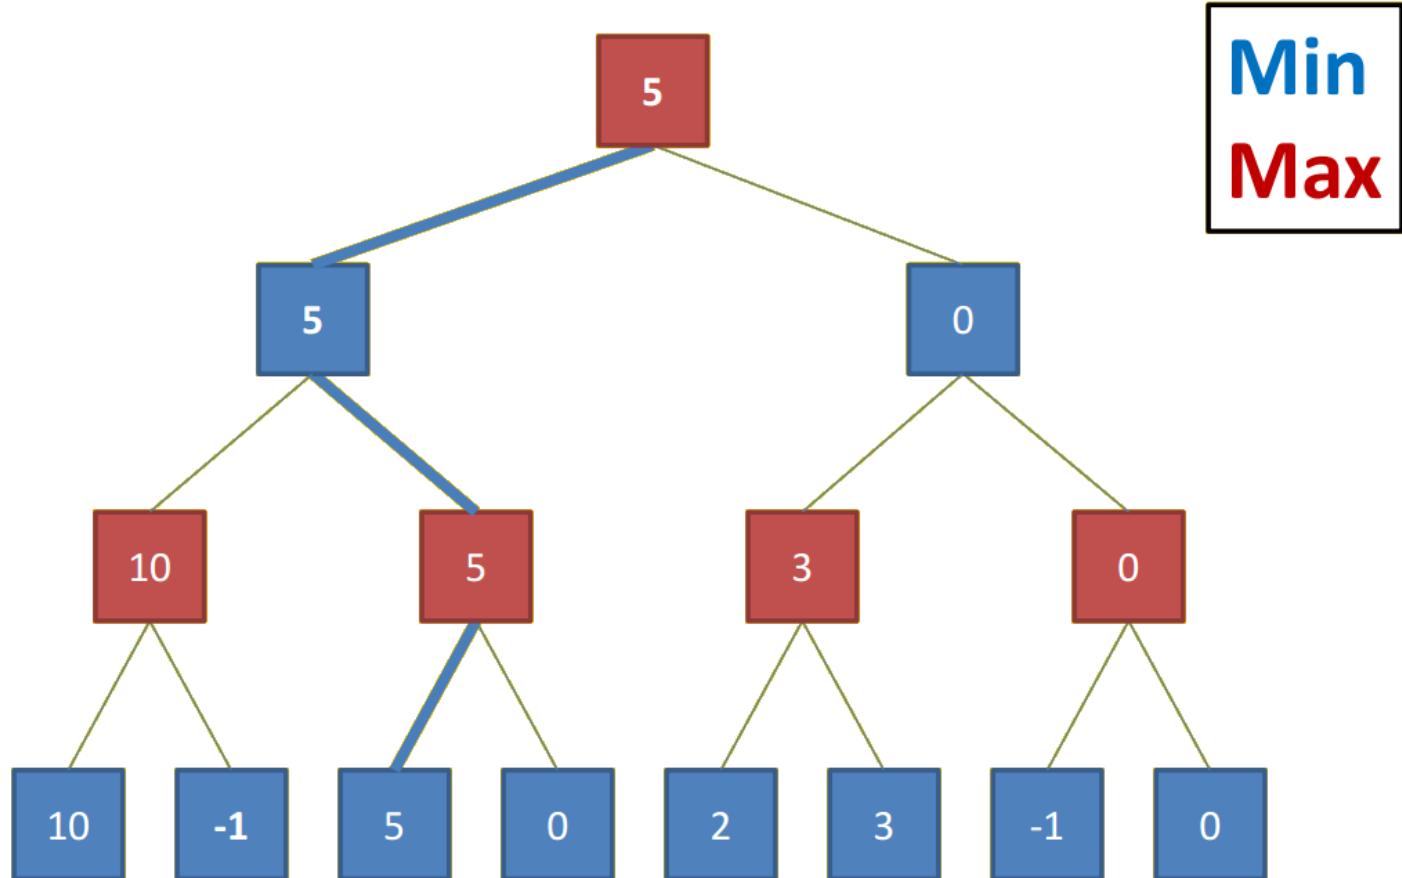
designed to find the **optimal strategy for MAX** and find **best move (policy, strategy)**:

1. generate the whole game tree, down to the leaves
2. apply utility function to each leaf
3. back-up values from leaves through branch nodes:
  - a MAX node computes the MAX of its child values
  - a MIN node computes the MIN of its child values
4. at root: choose the move leading to the child of highest value

# MIN-MAX tree search: an example



# MIN-MAX tree search: an example



# MIN-MAX tree search algorithm

```
function MINIMAX-DECISION(state) returns action
```

```
    inputs: state // current state in game
```

```
    return the a in actions(state) maximizing MIN-VALUE(forward(state,a))
```

```
function MAX-VALUE(state) returns utility
```

```
    if TERMINAL-TEST(state) then return UTILITY(state)
```

```
    v = -infinity
```

```
    for a, s in SUCCESSORS(state) do v = MAX(v, MIN-VALUE(s))
```

```
    return v
```

```
function MIN-VALUE(state) returns utility
```

```
    if TERMINAL-TEST(state) then return UTILITY(state)
```

```
    v = infinity
```

```
    for a, s in SUCCESSORS(state) do v = MIN(v, MAX-VALUE(s))
```

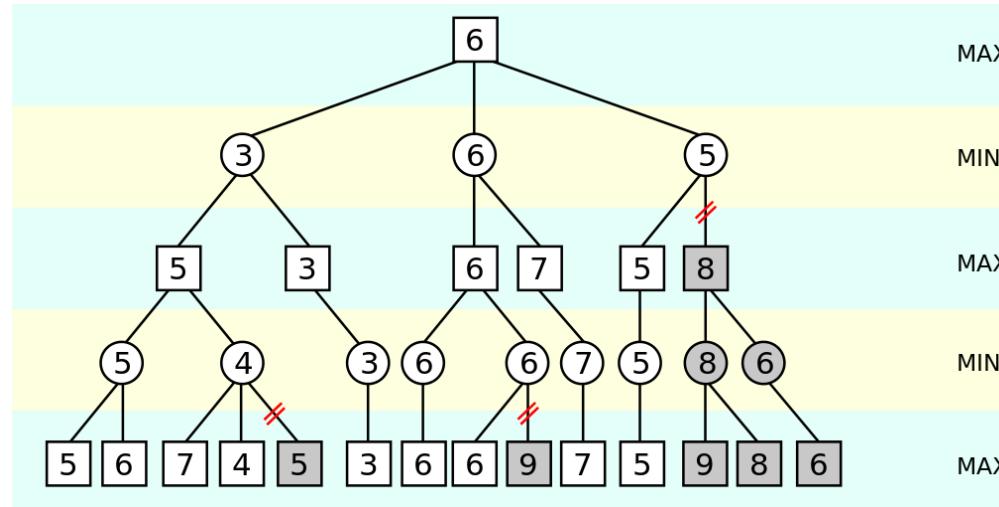
```
    return v
```

# game tree sizes

- Tic-Tac-Toe
  - $9! = 362,880$
  - exact solution quite reasonable**
- Chess
  - $b \approx 35$  (approximate average branching factor)
  - $d \approx 100$  (depth of game tree for “typical” game)
  - $bd \approx 35^{100} \approx 10^{154}$  nodes!!
  - exact solution infeasible**
- StarCraft
  - conservative estimate of state number from 400 units (2-player)  
on 128x128 map  $16384^{400} \approx 10^{1685}$
- it is **usually** impossible to develop the whole search tree
  - need of good utility functions and/or pruning techniques

# alpha-beta pruning

- we can disregard a part of the tree (prune)
- if a move is worse than an already evaluated one, stop evaluating it (we save building part of the tree and even computing some leaves)
- alpha = minimum score MAX can get, beta = maximum score MIN can get
- tree evaluated from left, if  $\alpha \geq \beta$  we can discard that subtree (will never be played)



(remember we are starting with an empty tree and compute values on demand from left to right)  
demo: <http://homepage.ufp.pt/jtorres/ensino/ia/alfabeta.html>

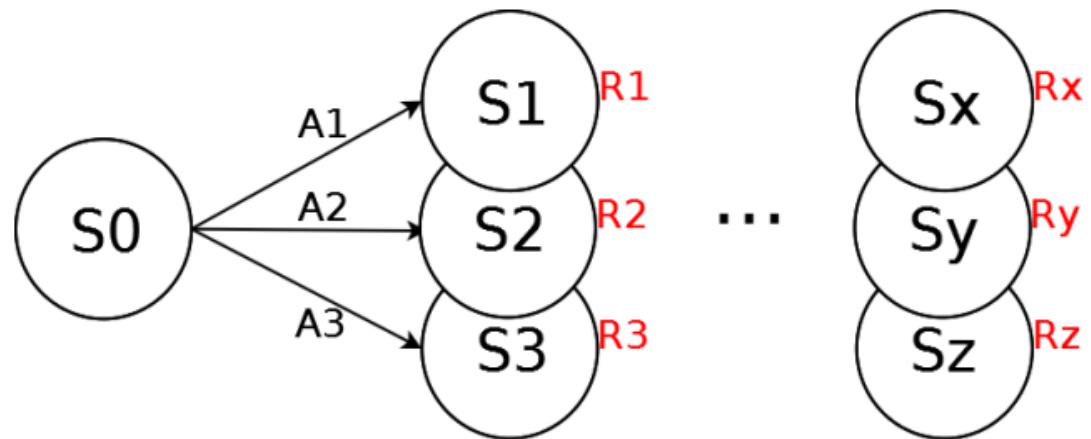
# what we can solve with Monte Carlo tree search

perfect information	partial information
deterministic	Battleship
non-deterministic	StarCraft, Poker Hanabi, Bridge, MOBA, RTS, FPS

# states and actions

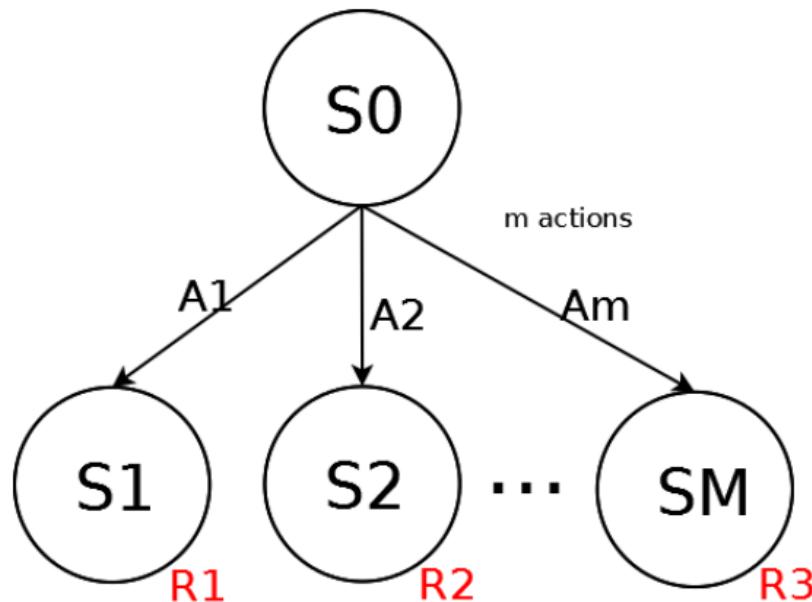
- a Markov decision process (MDP) is a tuple  $< S, A, P, R, \gamma >$   
 $S$ : finite set of all possible states of the game  
 $A$ : finite set of all possible actions in the game  
 $P$ : a state transition probability matrix  $P^a_{ss'} = P[s' | s, a]$   
 $R$ : a reward Function,  $R(s, a) = E[r|s, a]$   
 $\gamma$ : a discount factor  $\gamma \in [0, 1]$

- we need to learn a table
- “finite set” is problematic



# one step look ahead

- try all actions, pick the one with the highest reward
- problem: we do not reach a goal state, need heuristics



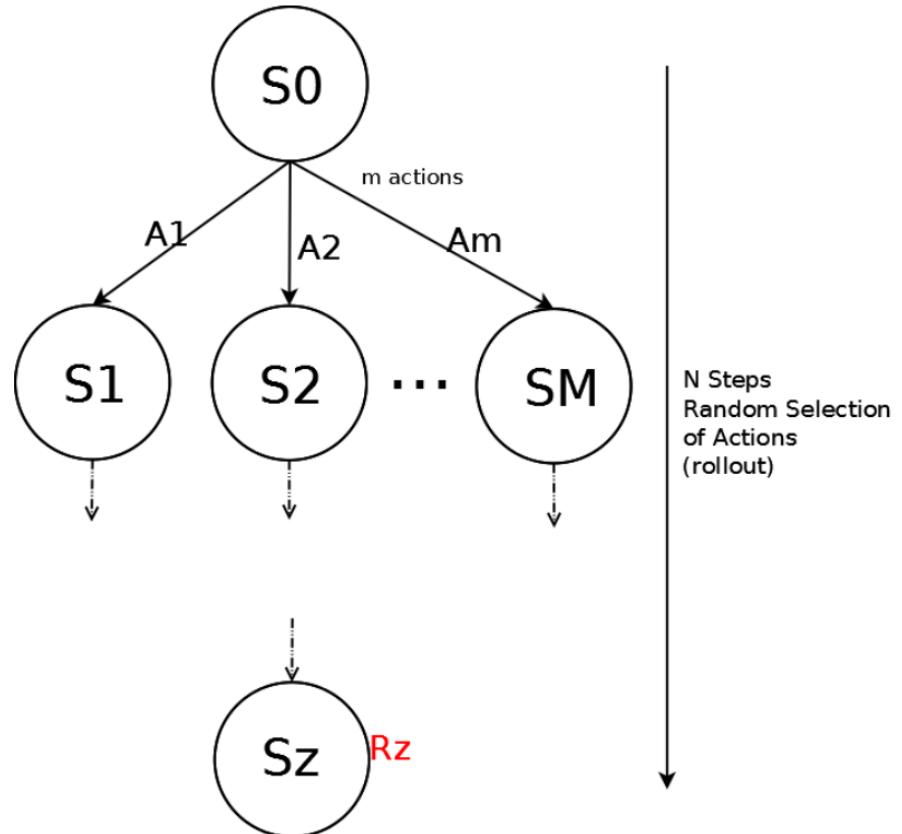
# flat Monte Carlo

- iteratively, apply  $N$  consecutive actions picked randomly
- pick the action that leads to:
  - the highest reward after  $N$  steps
  - the highest **average** reward after  $N$  steps, etc.

problems:

- this is similar to randomized brute force
- how large is  $N$ ? when do we get a non-0 reward?

can we somehow “intelligently” select interesting paths?

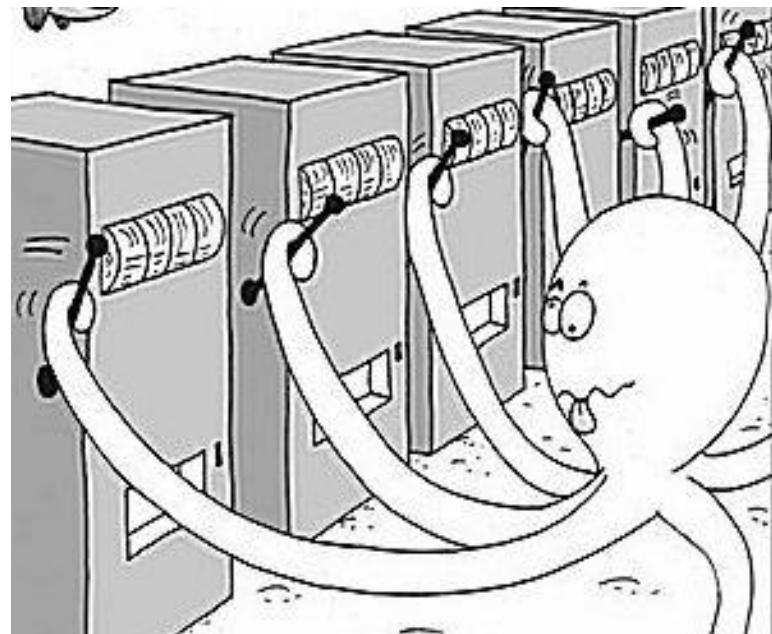


# multi-armed bandit problem

- given that there are several choices, which one do we take?
- this relates to a well-known decision making problem:  
the Multi-Armed Bandit Problem (MAB)
- at each step pull one arm
- noisy/random reward signal

pick the arm so as to:

- minimize regret (expected loss  
due to not picking the best arm)
- maximize expected return



# Monte Carlo (uniform)

most simple approach: **Monte Carlo**

- share trials uniformly between arms
- don't keep any outcome information

almost as simple:  **$\epsilon$ -greedy**

- $P(1-\epsilon)$  – pick the best arm so far
- $P(\epsilon)$  – pick an arm randomly
- exploitation vs. exploration



# upper confidence bound (UCB)

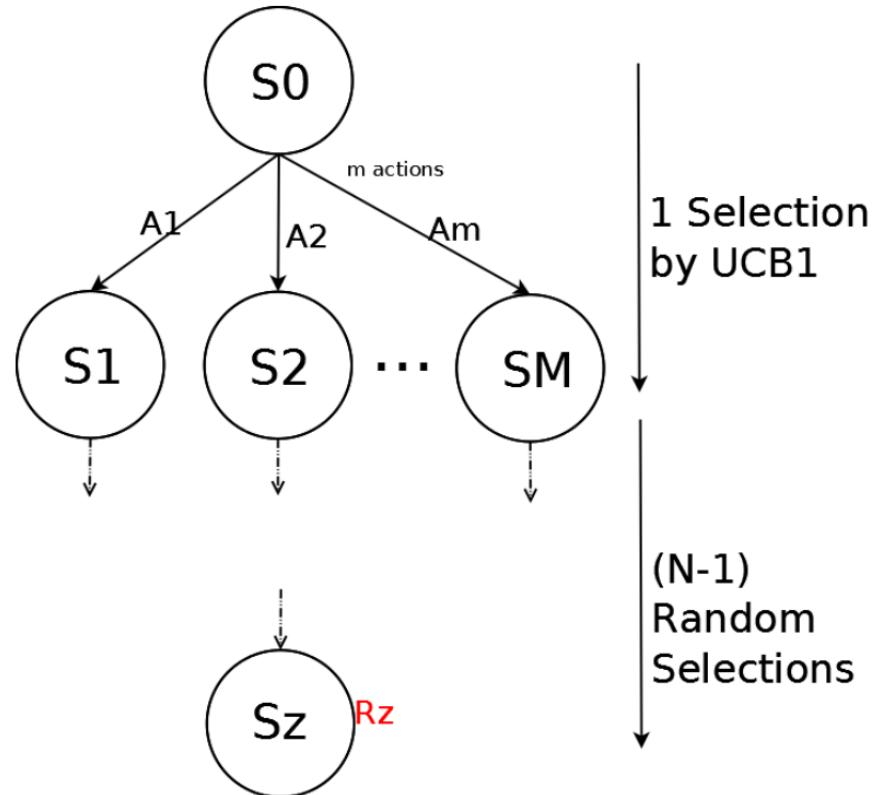
- balance exploitation and exploration
- an example: UCB1

$$a^* = \operatorname{argmax}_{a \in A(s)} \left\{ Q(s, a) + C \sqrt{\frac{\ln N(s))}{N(s, a)}} \right\}$$

- $Q(s, a)$ : average of rewards after taking action  $a$  from state  $s$
- $N(s)$ : times the state  $s$  has been visited
- $N(s, a)$ : times the action  $a$  has been picked from state  $s$
- $C$ : constant that balances exploitation ( $Q$ ) and exploration terms
  - application dependent
  - typical value for single player games with rewards in  $[0,1]$ :  $\sqrt{2}$

# flat UCB

- use UCB1 to select action from the root node (or current state)
- pick  $(N - 1)$  actions at random
- repeat iteratively, return action with (one of):
  - the highest reward after  $N$  steps
  - the highest **average** reward after 1 step ( $Q(s, a)$ )
  - the most visited node after 1 step (highest  $a$  for  $N(s, a)$ ).
  - the highest UCB1 value after 1 step, etc.



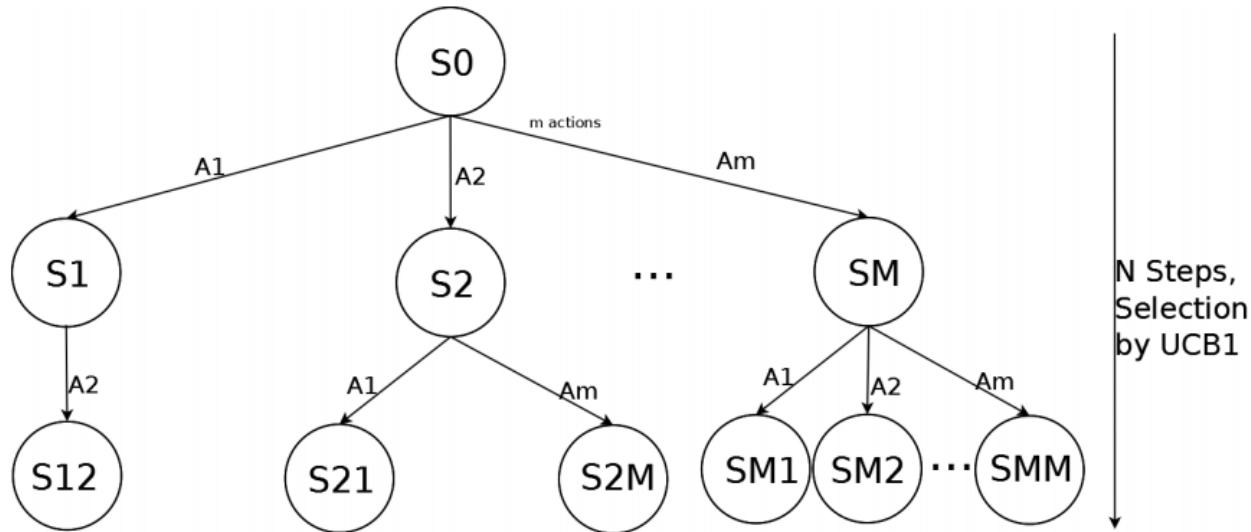
# UCB1 - discussion

- UCB1 formula minimises regret
  - bound grows logarithmically with the total number of actions taken:  $\ln(N(s))$
  - bound shrinks with the number of times we have taken this particular action:  $N(s,a)$
  - ensures: actions are tried infinitely often but still a good balance of exploration/exploitation
- needs limited information:
  - game rules (states, actions)
  - terminal state evaluation (reward)
- super simple to implement
- surprisingly effective, but...

... doesn't look ahead !!!

# building a tree with UCB1

- build a tree: search  $N$  steps in the future
- the search is **not** exhaustive: tree grows asymmetrically
- repeat iteratively, return action with:
  - the highest reward after  $N$  steps
  - the highest **average** reward after 1 step ( $Q(s, a)$ )
  - the most visited node after 1 step (highest  $a$  for  $N(s, a)$ ).
  - the highest UCB1 value after 1 step, etc.



# Monte Carlo tree search (MCTS)

depends on two concepts:

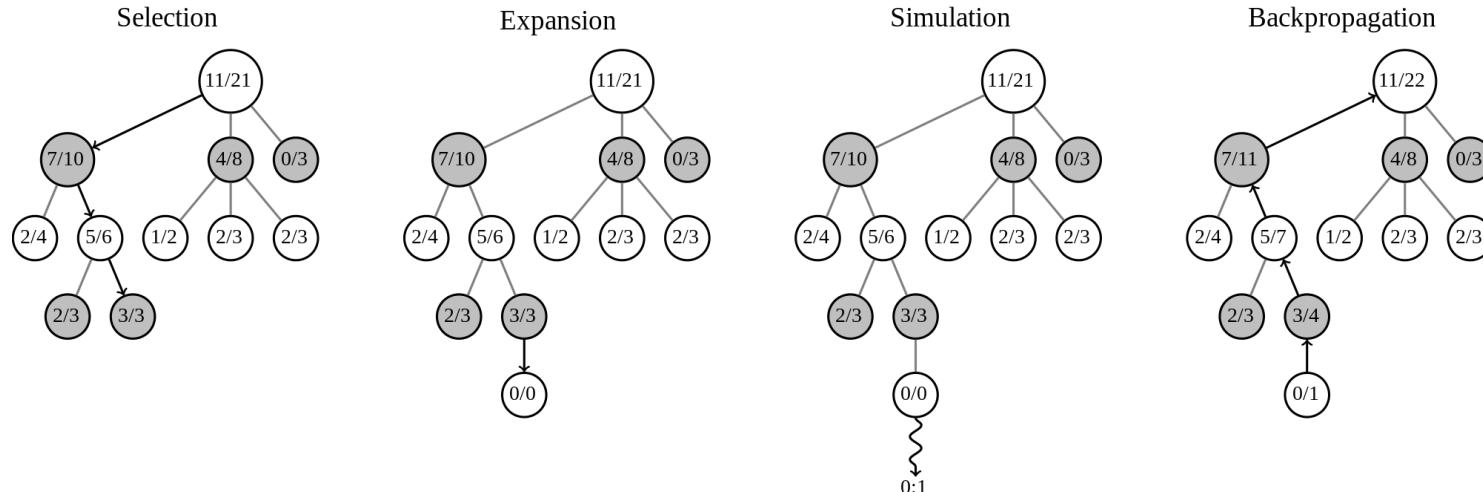
1. true value of action can be approximated via random **simulations**
2. the values may be used efficiently to adjust the policy towards a **best-first** strategy (brute force is too expensive)

advantages:

- anytime algorithm – stop whenever you like
- needs only game rules:
  - moves
  - terminal state evaluation (win, loss, draw, score)
- no need for a heuristic function, ...
  - but can be enhanced with domain knowledge in various ways
  - key advantage over MIN MAX

# MCTS – the big picture

- selection: select promising node within the tree (by means of UCB1)
  - expansion: add new leaf
  - simulation: play out the game until we reach a terminal state (and obtain a reward)
  - backpropagation: inform the “higher” nodes about move potential
- 
- selection is also called “tree policy”, simulation the “default policy”
  - for each node, we store: incoming action a, associated state s, total simulation reward Q, visit count N



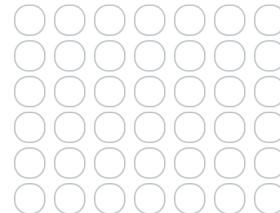
# MCTS plays "connect-4" (against you)

- in connect-4 you can only "throw in" a coin in one of the columns and it falls down
- the aim is to get 4 coins in a row (any form, diagonal included)
- start to play against 100 MCTS iters (last parameter)
- if you win, proceed with 1000, then with 10000
- do you realize a difference in AI playing strength?
- btw, program does not always recognize the human winning (you have to)

<https://sethpipho.github.io/monte-carlo-tree-search-js/demo/connect-4/>

## Connect-N

A generalized connect-4 game vs. an AI powered by Monte-Carlo-Tree Search. [Github](#)



Human's Turn



New Game

<input type="text" value="7"/>	- Board Width
<input type="text" value="6"/>	- Board Height
<input type="text" value="4"/>	- Number to connect
<input type="text" value="100"/>	- MCTS Iters

# MCTS main algorithm

- BestChild simply picks best child node of root according to some criteria: e.g. best mean value
- in our pseudo-code BestChild is called from TreePolicy and from MctsSearch, but different versions can be used
  - e.g., final selection can be the max value child or the most frequently visited one

```
function MCTSSEARCH( $s_0$ )
    create root node  $v_0$  with state  $s_0$ 
    while within computational budget do
         $v_l \leftarrow \text{TREEPOLICY}(v_0)$ 
         $\Delta \leftarrow \text{DEFAULTPOLICY}(s(v_l))$ 
        BACKUP( $v_l, \Delta$ )
    return  $a(\text{BESTCHILD}(v_0, 0))$ 
```

# tree policy

- note that node selected for expansion does not need to be a leaf of the tree
  - the *nonterminal* test refers to the game state

```
function TREEPOLICY( $v$ )
    while  $v$  is nonterminal do
        if  $v$  not fully expanded then
            return EXPAND( $v$ )
        else
             $v \leftarrow \text{BESTCHILD}(v, Cp)$ 
    return  $v$ 
```

# expand

- note: need to choose an appropriate data structure for the children (array, ArrayList, Map)

**function** EXPAND( $v$ )

choose  $a \in$  untried actions from  $A(s(v))$

add a new child  $v'$  to  $v$

with  $s(v') = f(s(v), a)$

and  $a(v') = a$

**return**  $v'$

# best child (UCT)

- this is the standard UCB equation  
(we rewrite the UCB1 C slightly here)
  - used in the tree
- higher values of  $c$  lead to more exploration
- other terms can be added, and usually are

**function** BESTCHILD( $v, c$ )

**return**  $\arg \max_{v' \in \text{children of } v} \frac{Q(v')}{N(v')} + c \sqrt{\frac{2 \ln N(v)}{N(v')}}$

# default policy

- each time a new node is added to the tree, the default policy randomly rolls out from the current state until a terminal state of the game is reached
- the standard is to do this uniformly randomly
  - but better performance may be obtained by biasing with knowledge

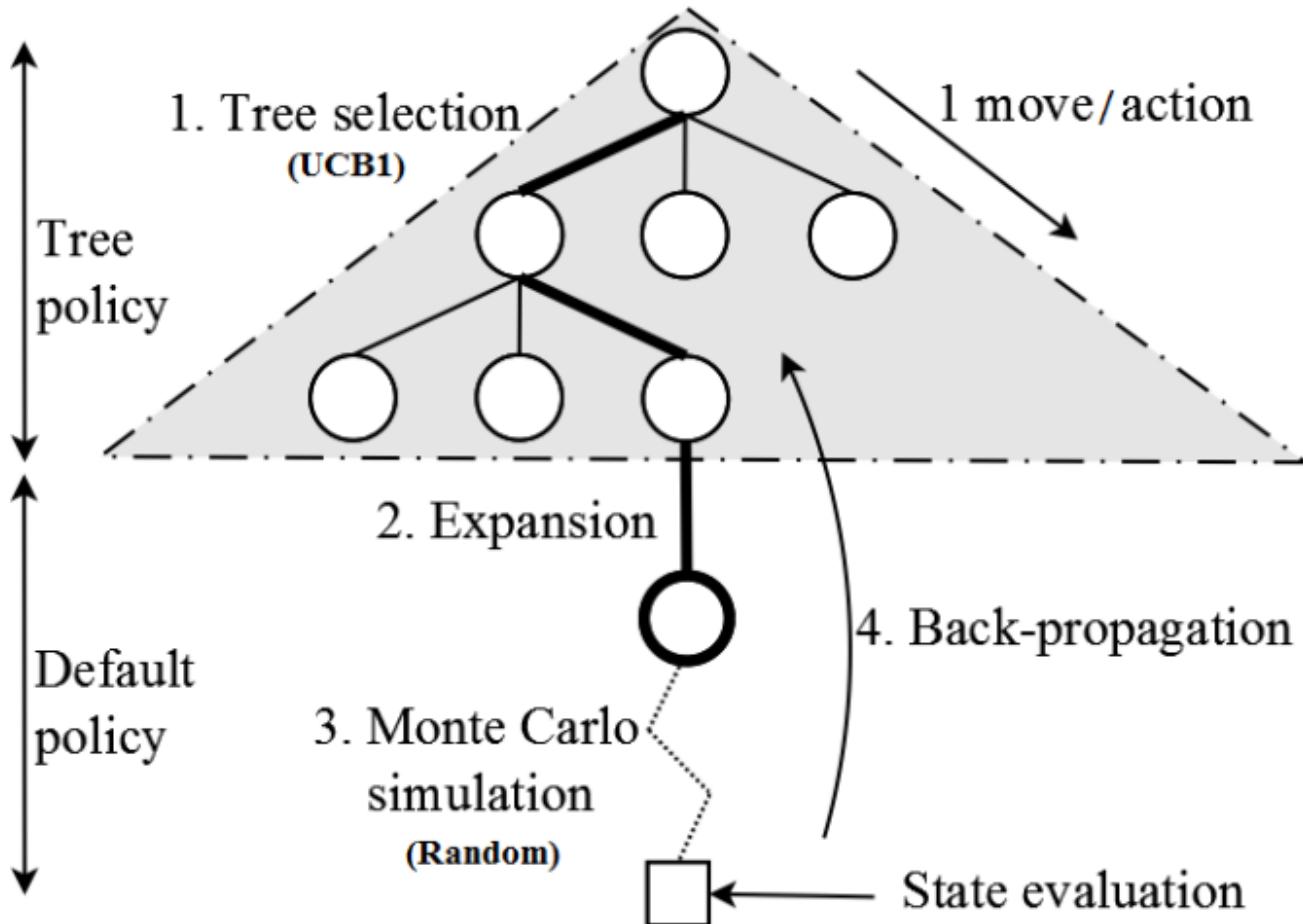
```
function DEFAULTPOLICY( $s$ )
  while  $s$  is non-terminal do
    choose  $a \in A(s)$  uniformly at random
     $s \leftarrow f(s, a)$ 
  return reward for state  $s$ 
```

# backup

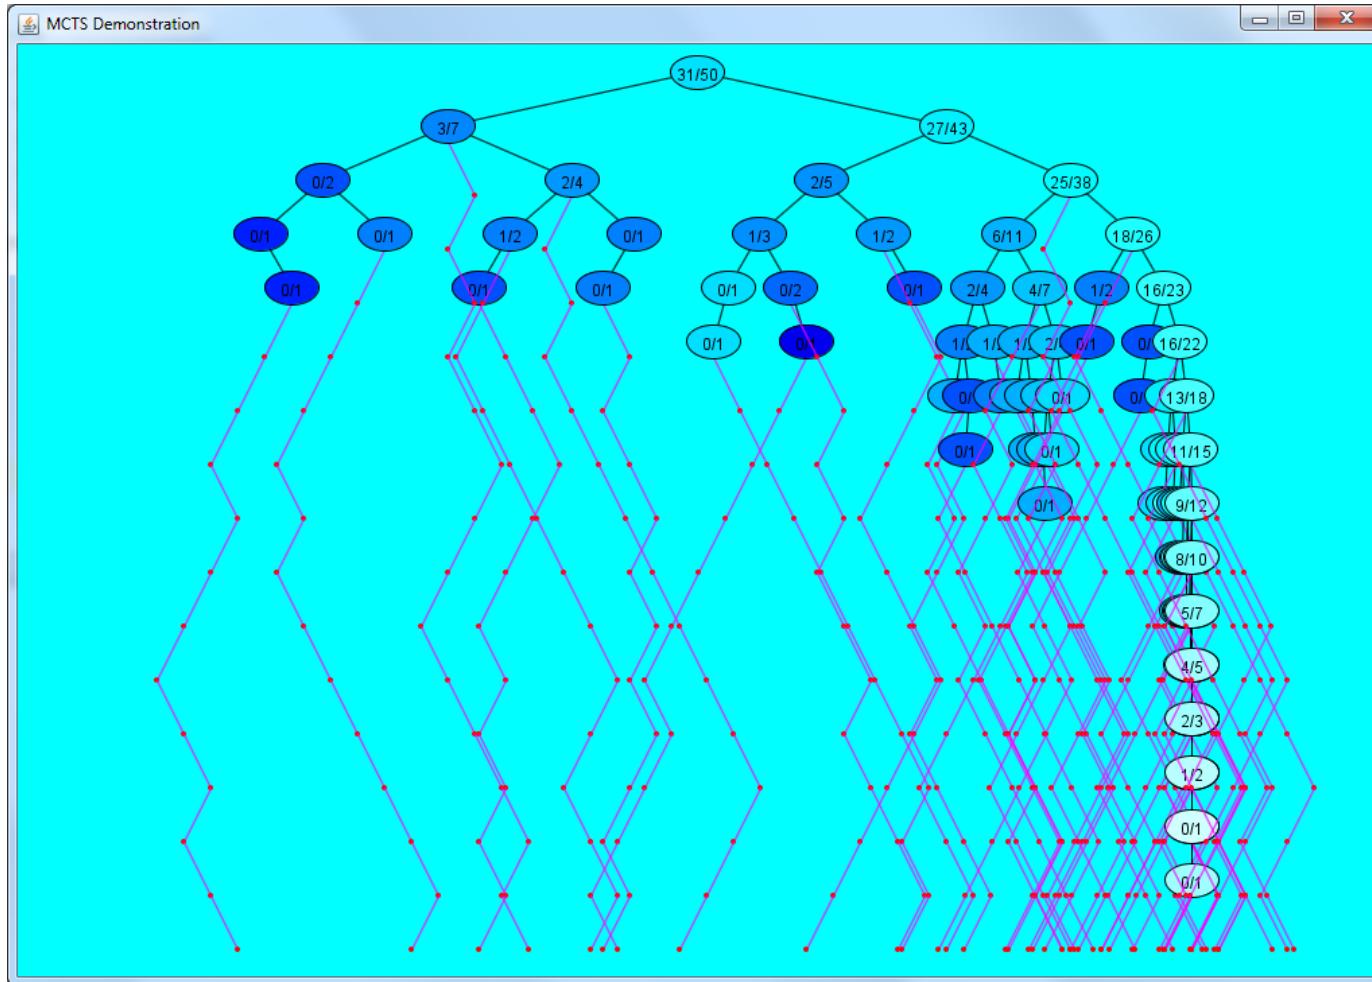
- note that  $v$  is the new node added to the tree by the tree policy
- back up the values from the added node up the tree to the root

```
function BACKUP( $v, \Delta$ )
  while  $v$  is not null do
     $N(v) \leftarrow N(v) + 1$ 
     $Q(v) \leftarrow Q(v) + \Delta(v, p)$ 
     $v \leftarrow$  parent of  $v$ 
```

# MCTS: the even bigger picture



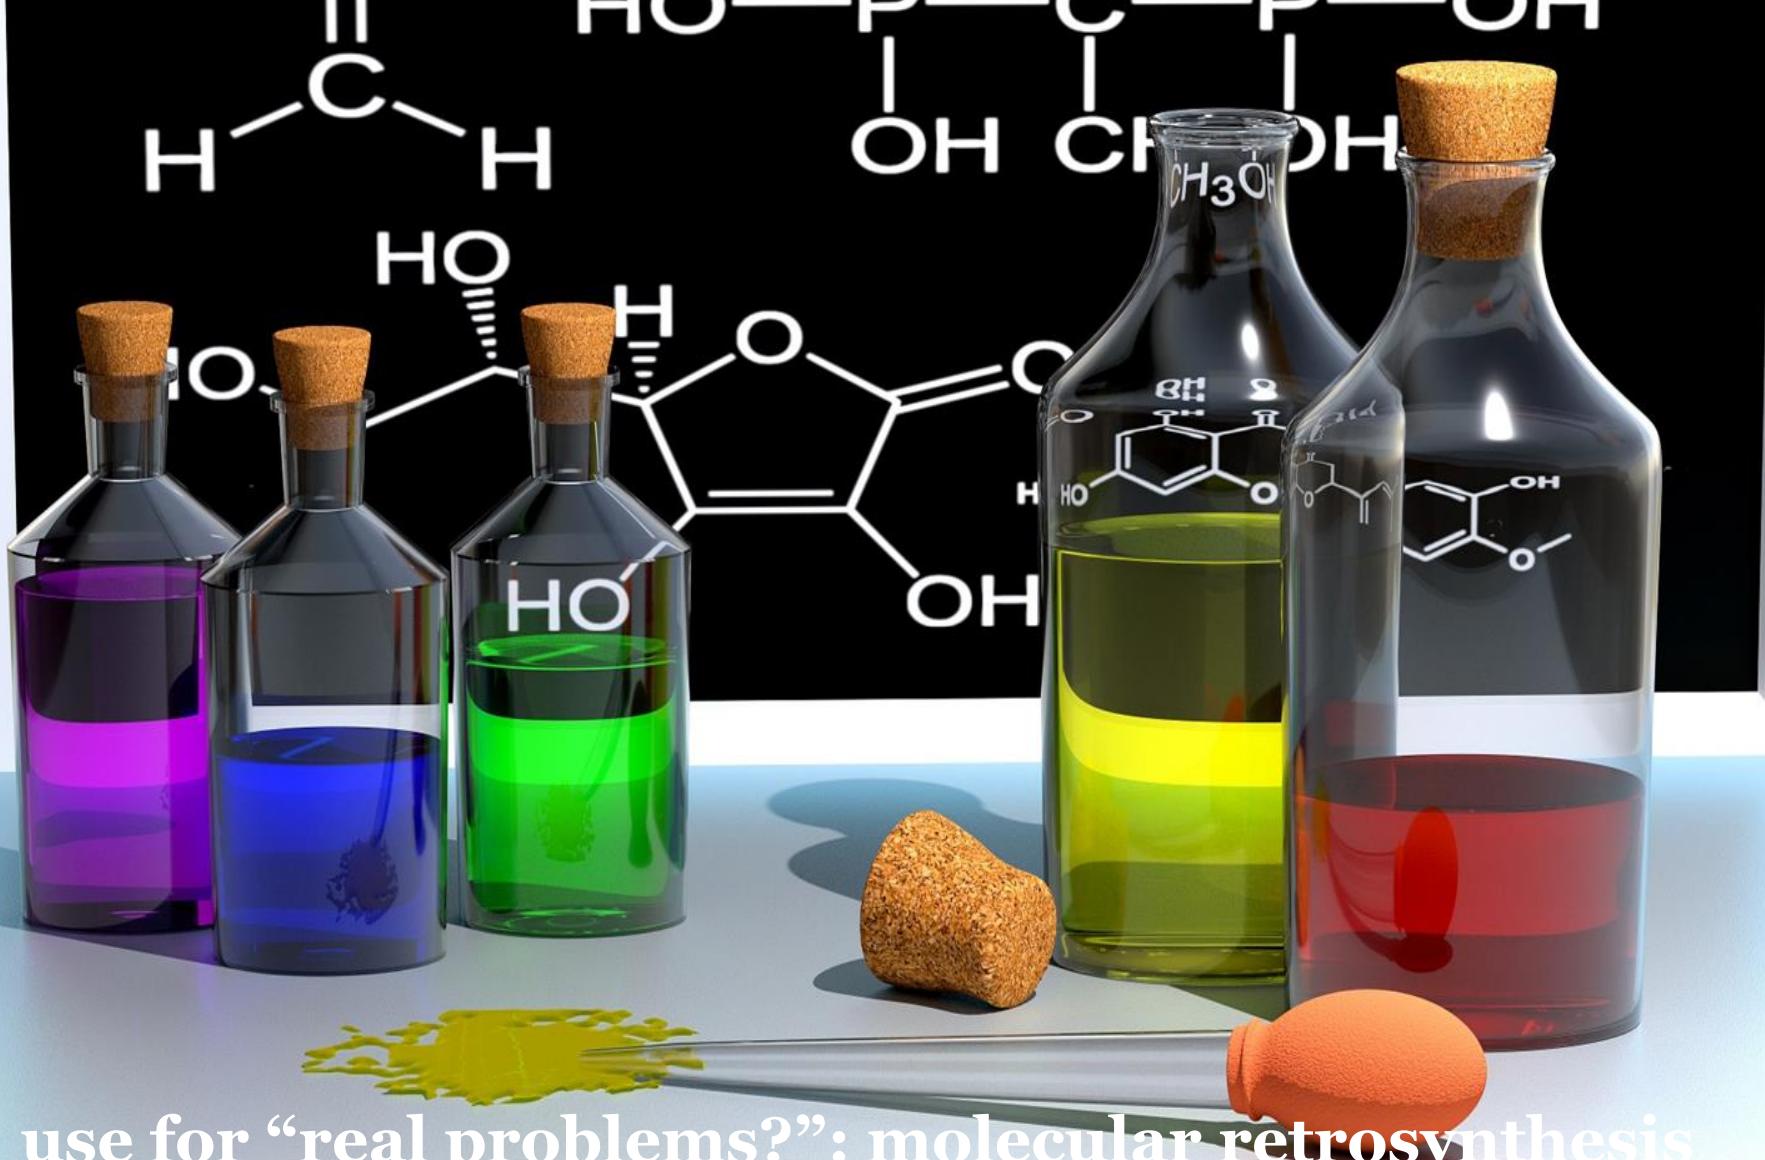
# MCTS builds asymmetric trees



# MCTS goes real-time

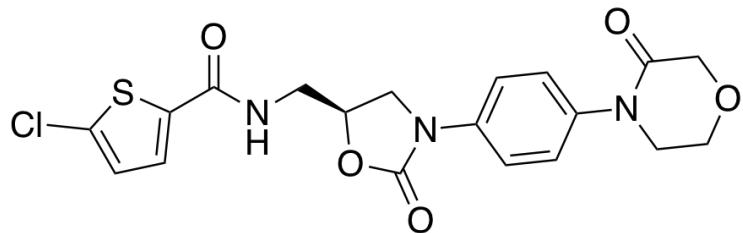
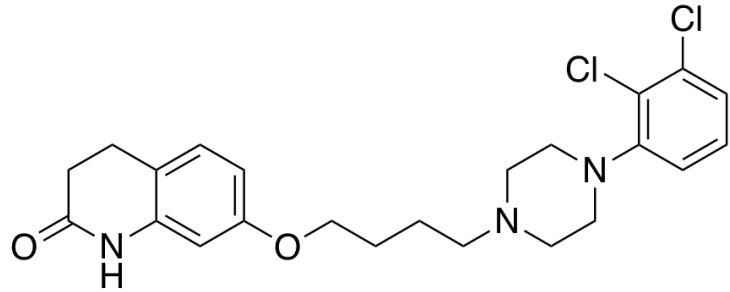
- limited roll-out budget
  - heuristic knowledge becomes important
  - we will only see a small part of the tree
- action space is fine-grained
  - take *macro actions* otherwise planning will be very short-term
  - currently unclear how to "generate" macro actions
- maybe no terminal node in sight
  - use a heuristic
  - tune simulation depth
- or use "heavy" playouts (use heuristic in playouts)?





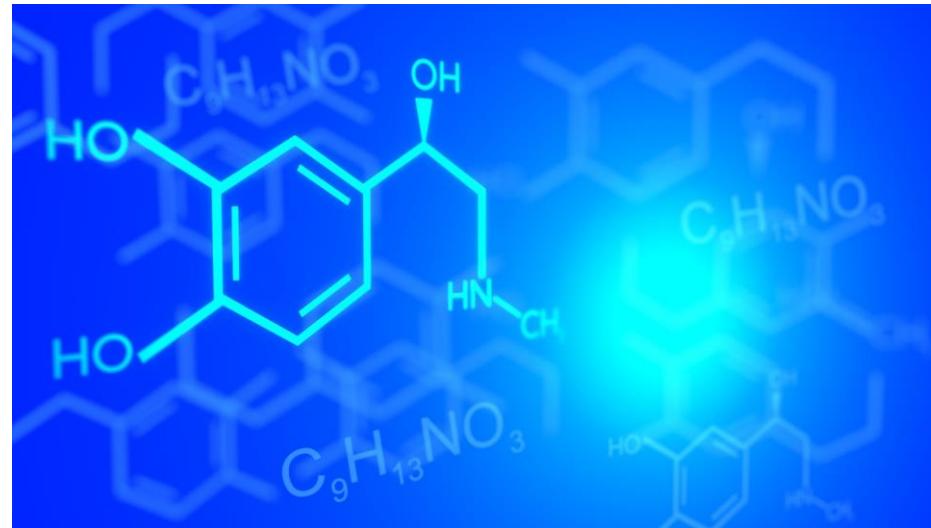
# chemical synthesis

- many important drugs are small organic molecules
- synthesis is a significant bottleneck
- 60 years history in computer aided synthesis planning, not much success



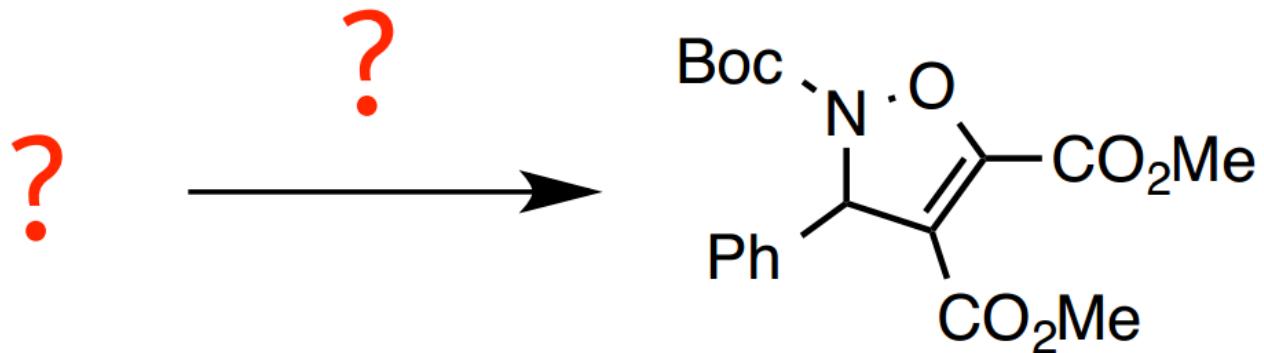
# retrosynthesis is a key problem

- much of Chemistry is “theoretical” (in a computer) now
- there are tools for predicting properties of unknown substances
- but being able to model a substance does not mean that we know how actually make it
- retrosynthesis can find such ways
- this is a multi-modal problem  
we want several different solutions
- this is a multi-objective problem,  
there are multiple objectives
  - cost
  - yield
  - conflicting required properties  
(bind to x but not to y)



# can we frame this in the MCTS/DNN context?

- the problem is not to take the right move
- but to take the right decisions for synthesizing a molecule
- both share a lot:
  - high branching factor
  - sequence length not exactly known
  - full tree cannot be enumerated
  - no strong heuristics



# solving different problems

	Chess	Go	Synthesis
System	DeepBlue	AlphaGo	?
Scoring	Heuristics	MCTS+DNN, RL	?
Rules	Trivial & Fixed	Trivial & Fixed	Complex, unknown, expanding
Branching Factor	very small	small	large
Experiments	cheap & fast	cheap & fast	takes days (lab/simulation)

# games vs the real world



**Julian Togelius**  
@togelius

Following



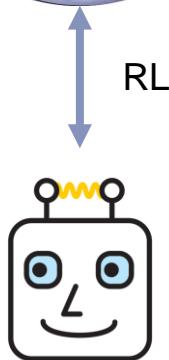
Another thing to remember is that we know the rules of the games and can simulate them very fast. In other words, we have a perfect and lightweight forward model. This is one way in which board games differ from most real-world problems.

9:32 PM - 5 Dec 2017

# similar but different

## AlphaGo (Zero):

take perfect world simulator (forward model),  
then RL/planning with lots of compute



Use agent in  
real world

A horizontal arrow points from the robot head towards the Earth image.

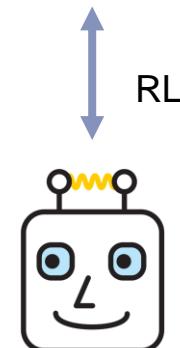
## This work:

Learn generalizing(!) world simulator (forward model), then RL/planning



Learn from  
past  
experiments

A horizontal arrow points from the test tubes towards the Earth image.



Use agent in  
real world

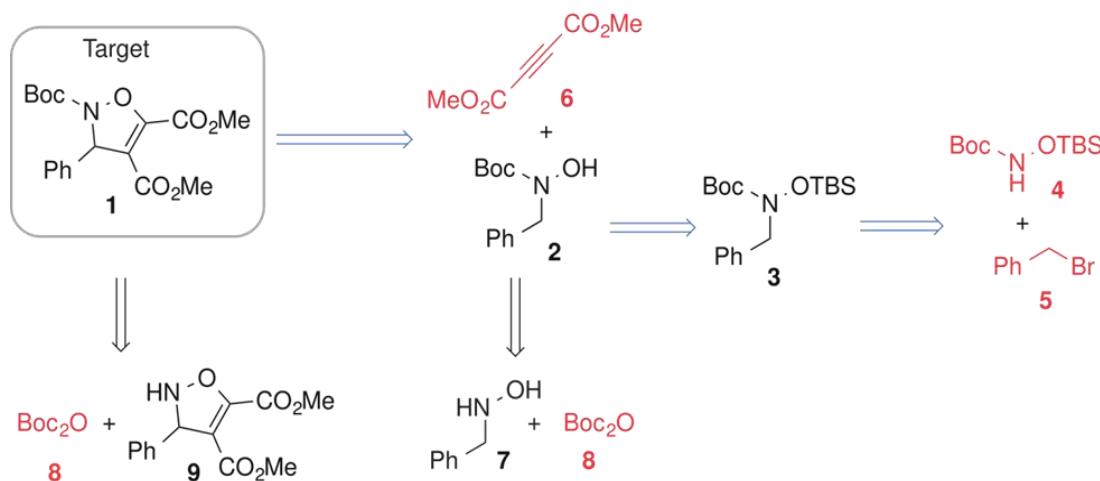
A horizontal arrow points from the robot head towards the Earth image.

# what problems need to be solved?

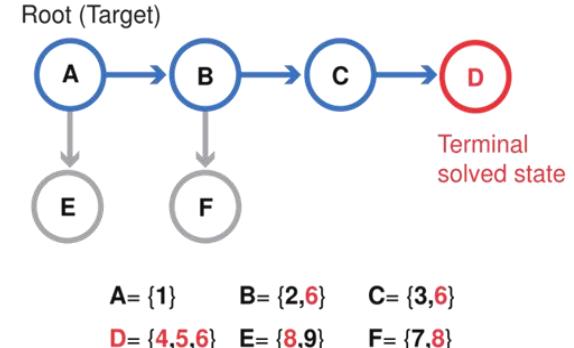
1. we need a way to encode the available molecules
2. we need a way to encode reaction rules
3. we need to **obtain** existing reaction rules
4. we need to **find out** which reaction rules are most suitable for a specific molecule (part)

finally, it should look like this:

a) Chemical Representation of the Synthesis Plan



b) Search Tree Representation



# obtaining reaction rules

- we have representations for 1) and 2)
- for 3), there have been manual attempts, but this does not scale:
- chemical knowledge doubles every decade, currently around 11 million reactions known
- solution: we extract reaction rules from an existing database (Reaxys),  
this results in 301,671 rules

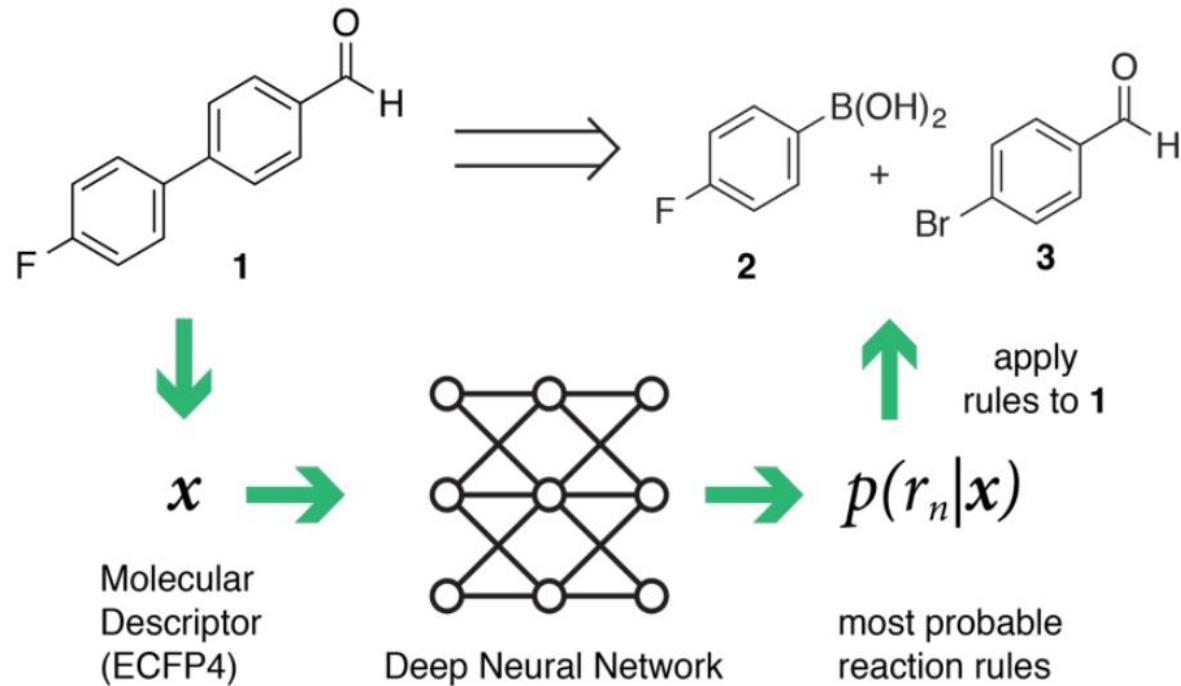
thus we can construct search trees with

- depth from around 5 up to around 20
- branching factors of approx. 46,000 (avg. applicable rules)
- for 5 steps, the node count is already a number with 23 digits
- **this is far too much, even for MCTS!**

**help! how do we get the branching factor down?**

# deep neural network for reaction preferences

- from all reaction rules we learn a preference for suggesting most likely step
- the top50 accuracy is already 0.73 for the whole ruleset
- this reduces the branching factor to 50

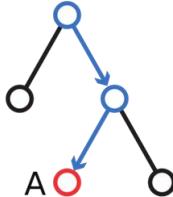


Segler, Waller, Chem. Eur. J. 2017, DOI: 10.1002/chem.201605499

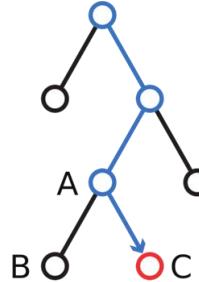
# putting everything together

## a) Synthesis Planning with Monte Carlo Tree Search

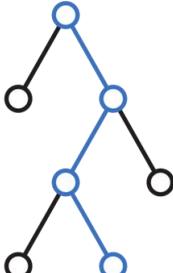
1) Selection  
pick most promising position



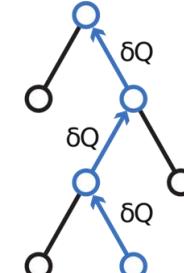
2) Expansion  
retroanalyze, add new nodes to tree by *expansion procedure* (panel b)



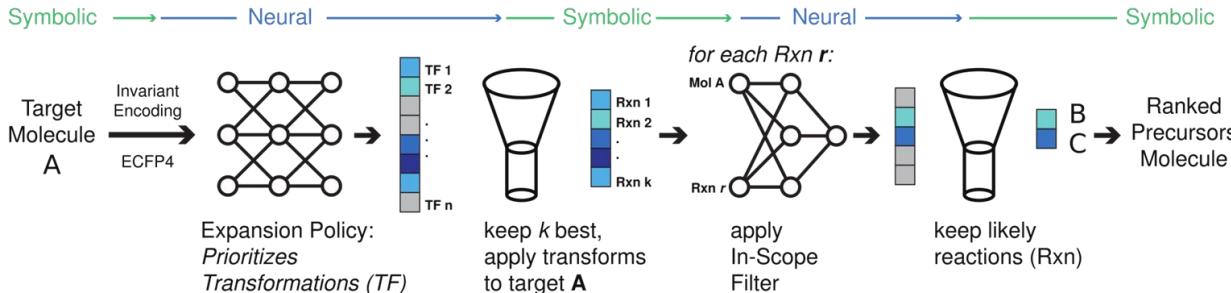
3) Rollout  
pick and evaluate new position



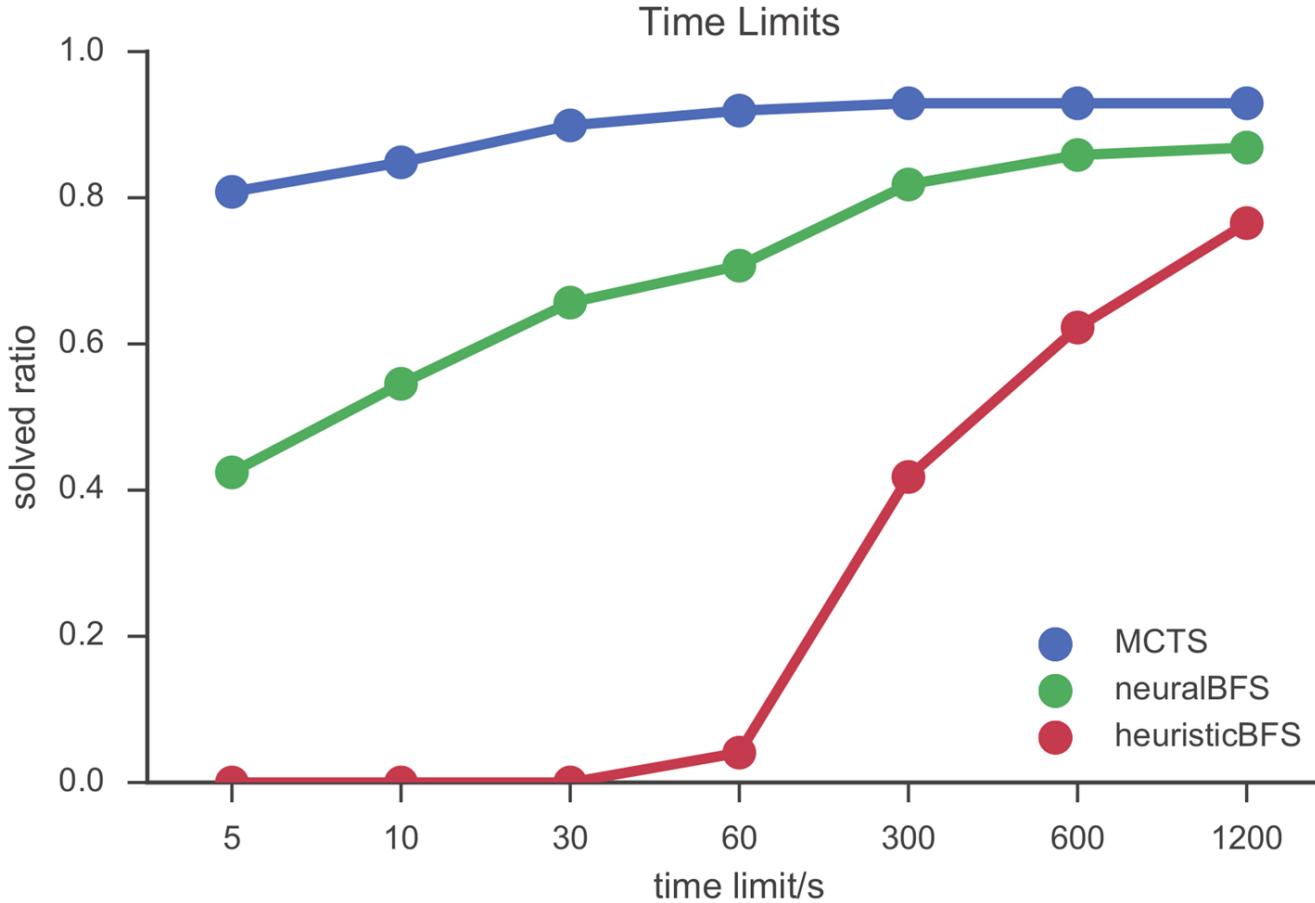
4) Update  
incorporate evaluation in the search tree



## b) Expansion procedure



# how does it perform?



nature.com > nature > articles > article

# nature

International journal of science

Access provided by Uni Muenster

Altmetric: 459 [More detail >>](#)

Article | Published: 28 March 2018

## Planning chemical syntheses with deep neural networks and symbolic AI

Marwin H. S. Segler , Mike Preuss & Mark P. Waller 

Nature 555, 604–610 (29 March 2018) | Download Citation 

### Abstract

To plan the syntheses of small organic molecules, chemists use retrosynthesis, a problem-solving technique in which target molecules are recursively transformed into increasingly simpler precursors. Computer-aided retrosynthesis would be a valuable tool but at present it is slow and provides results of unsatisfactory quality. Here we use Monte Carlo tree search and symbolic artificial intelligence (AI) to discover retrosynthetic routes. We combined Monte Carlo tree search

a natureresearch journal

Search E-alert Submit Login

 PDF

Editorial Summary

Computers teach themselves to make molecules

Chemical reaction databases that are automatically filled from the literature have made the planning of chemical syntheses, whereby... [show more](#)

Sections Figures References

Abstract

Main

Training the expansion and rollout policies

Prediction with the in-scope filter network

Integrating neural networks and MCTS

Evaluating the performance characteristics of 3N-MCTS

Quantitative evaluation

# General Game Playing (GGP)



General Game Playing

General  
Artificial  
Intelligence

- one player can play several games
- without actually knowing much about the game
- considered important step toward AGI (or GAI? :-))
- however, games very simple
- important publications:

Pell, Barney. "METAGAME: A new challenge for games and learning." (1992).

Genesereth, Michael, Nathaniel Love, and Barney Pell. "General game playing: Overview of the AAAI competition." AI magazine 26.2 (2005): 62-62.

<http://gpp.stanford.edu/ggp>



Gamemaster  
Leaderboard

Sign In

Player	Matches	Score
<a href="#">mozart_the_cat</a>	10492	40
<a href="#">indy</a>	14436	35
<a href="#">lara</a>	14337	35
<a href="#">alloy</a>	2491	21
<a href="#">sanche</a>	628	18
<a href="#">drop_table_teams</a>	315	16
<a href="#">ocminimax</a>	591	15
<a href="#">ocmc5</a>	1037	4
<a href="#">aaa</a>	0	0
<a href="#">afflach</a>	0	0
<a href="#">alague_rhyme</a>	0	0
<a href="#">ares</a>	0	0
<a href="#">blinky</a>	0	0
<a href="#">dfadsfadfadfc</a>	0	0
<a href="#">egghead</a>	0	0
<a href="#">perceval_ens</a>	0	0
<a href="#">ravat</a>	0	0

NB: Scores and counts do not include matches with more than 3 recorded errors.

Comments and complaints to [genesereth@stanford.edu](mailto:genesereth@stanford.edu).



Gamemaster  
Games

Sign In

Id	Name	Roles	Creator
<a href="#">alquerque</a>	Alquerque	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">asylum</a>	Asylum	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">badconnectfour</a>	Bad Connect Four	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">badskirmish</a>	Bad Skirmish	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">badtictactoe</a>	Bad Tic Tac Toe	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">battleofnumbers</a>	Battle of Numbers	2	<a href="mailto:dustinsfink@gmail.com">dustinsfink@gmail.com</a>
<a href="#">bestbuttonsandalights</a>	Best Buttons and Lights	1	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">bestknightsTour</a>	Best Knights Tour	1	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">bestsukoshi</a>	Best Sukoshi	1	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">breakthrough</a>	Breakthrough	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">breakthroughsmall</a>	Breakthrough Small	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">buttonsandalights</a>	Buttons and Lights	1	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">checkersonabarrelno kings</a>	Checkers on a Barrel No Kings	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">chinesecheckers</a>	Chinese Checkers	6	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">chinesecheckers2</a>	Chinese Checkers 2	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">chinesecheckers3</a>	Chinese Checkers 3	3	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">chinesecheckers4</a>	Chinese Checkers 4	4	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">chinook</a>	Chinook	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">connectfour</a>	Connect Four	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">connectfour3p</a>	Connect Four 3 Player	3	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">connectfourteam</a>	Connect Four Team	4	<a href="mailto:chuchro3@stanford.edu">chuchro3@stanford.edu</a>
<a href="#">donttouch</a>	Don't Touch	2	<a href="mailto:bhz@stanford.edu">bhz@stanford.edu</a>
<a href="#">dualhunter</a>	Dual Hunter	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">dualrainbow</a>	Dual Rainbow	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">dudoku</a>	Dudoku	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">duikoshi</a>	Duijoshi	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">eighthpuzzle</a>	Eight Puzzle	1	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">farmingquandries</a>	Farming Quandries	2	<a href="mailto:ptr.pham@gmail.com">ptr.pham@gmail.com</a>
<a href="#">firesheep</a>	Fire Sheep	2	<a href="mailto:ptr.pham@gmail.com">ptr.pham@gmail.com</a>
<a href="#">freeforall</a>	Freeforall	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">hamilton</a>	Hamilton	1	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">hex</a>	Hex	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">horseshoe</a>	Horseshoe	2	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">hunter</a>	Hunter	1	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>
<a href="#">hunterbig</a>	Hunter Big	1	<a href="mailto:bhz@stanford.edu">bhz@stanford.edu</a>
<a href="#">jointbuttonsandalights</a>	Joint Buttons and Lights	1	<a href="mailto:genesereth@stanford.edu">genesereth@stanford.edu</a>

# GVGAI environment

- general video-game playing (GVGP) succeeds general game-playing
- targeted at more complex, non-deterministic games
- language first shaped during Dagstuhl seminar CI and AI in games, 2012
- important publications:

John Levine, Clare B. Congdon, Michal Bída, Marc Ebner, Graham Kendall, Simon Lucas, Risto Miikkulainen, Tom Schaul, and Tommy Thompson. General Video Game Playing. Dagstuhl Follow-up, 6:1–7, 2013.

Tom Schaul. A Video Game Description Language for Model-based or Interactive Learning. In Proceedings of the IEEE Conference on Computational Intelligence in Games, Niagara Falls, 2013. IEEE Press

<http://www.gvgai.net> still exists, but is not actively used any more (last competition 2019)

the new implementation is here: [https://github.com/SUSTechGameAI/GVGAI\\_GYM](https://github.com/SUSTechGameAI/GVGAI_GYM)

# „old“ GVGAI website

- competitions ran for 5 years at major games conferences
- currently discontinued but we are using the “old” Java code with the new sprites and run only locally.
- learning (not planning) competitions are still running

<http://www.aingames.cn/>

<https://github.com/SUSTechGameAI>

- new sprite set:



## The General Video Game AI Competition - 2018

Check out our **new GVGAI Book** here: <https://gaigresearch.github.io/gvgabook/>

**IMPORTANT:** We have **migrated the code** to a **NEW** repository, please check [here](#). This code is used for the Planning and Generation tracks.  
The code for the GVGAI Single-Player Learning track is here: [https://github.com/rubenrtorrado/GVGAI\\_GYM](https://github.com/rubenrtorrado/GVGAI_GYM).

Looking for papers about GVGAI? Try [here](#) and our [survey](#) (under review at IEEE ToG).

2019 Roadmap

CoG 2019

Submission: 1st August

Results: 21st August

Available Tracks

[Single Player Planning](#)

[Level Generation](#)

[Rule Generation](#)

[Two-Player Planning](#)

[Single Player Learning](#)

Follow @gvgai

And join our Google Group



Sponsored by:

DeepMind

Welcome to the General Video Game AI Competition webpage. The GVG-AI Competition explores the problem of creating controllers for general video game playing. How would you create a single agent that is able to play any game it is given? Could you program an agent that is able to play a wide variety of games, without knowing which games are to be played? Can you create an automatic level generation that designs levels for any game is given?

# learning or planning?

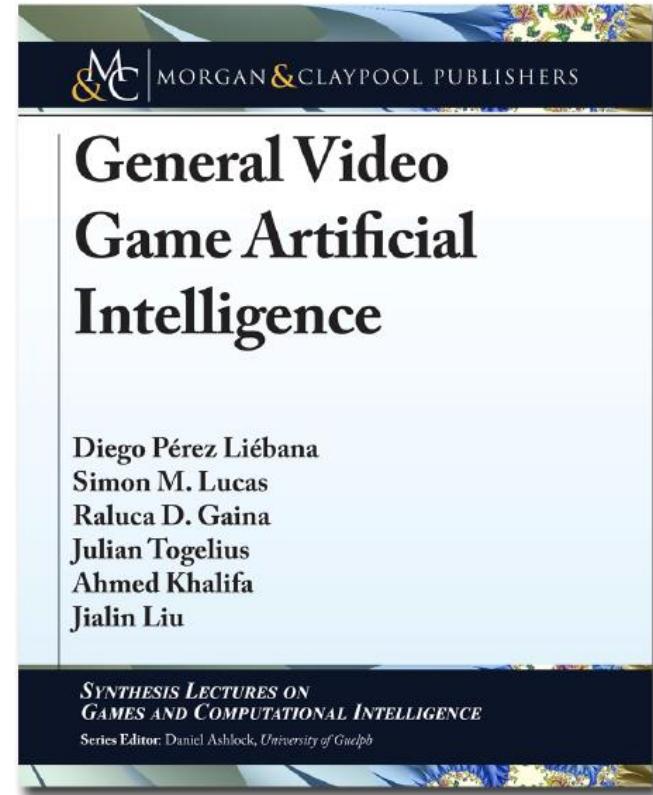
- boundaries between these two sometimes a bit fuzzy
- in game playing, we usually mean learning when we have no means to look into the future
- more specifically, when we have no "forward model" that can be used as simulation
- if we have a forward model we usually speak of planning
- here, we can compare the consequences of different actions and then roll back
- however, the forward model could be "learned" from past experiences
- or it is just a very sophisticated heuristic? ;-)



picture from Chuck Underwood on Pixabay

# GVGAI book

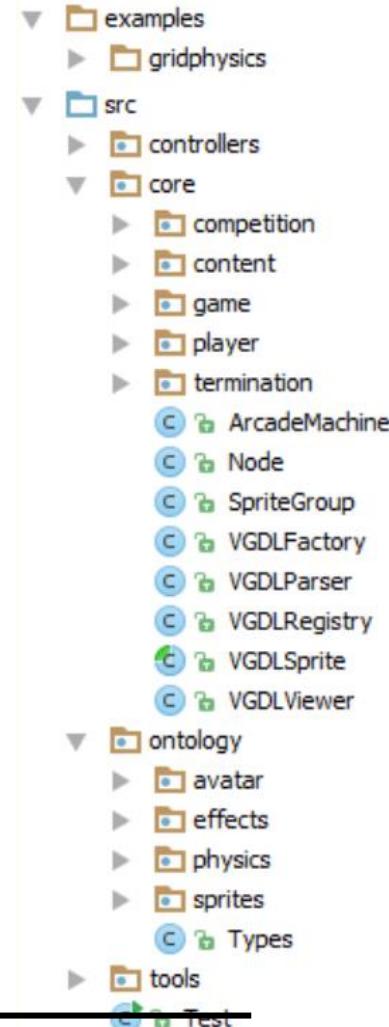
- summarizes what happened in GVGAI from around 2013 to 2019
- different types of competitions:
  - planning
  - learning
  - procedural content generation
- <https://gaigresearch.github.io/gvgaibook/>  
(chapters available for download)



# code overview

on the github server <https://github.com/SUSTechGameAI>  
go into GVGAI\_GYM/gym\_gvgai/envs/gvgai/

- examples → gridphysics: game and levels files
- controllers: sample working controllers
- core: core codebase of the framework
- core.competition: competition parameters
- core.content: game and sprite creation
- core.game: game engine, forward model and observations
- core.player: abstract class for players
- core.termination: game termination conditions
- ontology: definitions for sprites, avatars, physics and events
- tools: several useful classes
- **Test.java**: entry point to run the framework, now at  
src/tracks/singlePlayer/Test.java



# VGDL – game definitions

```
1  BasicGame
2      SpriteSet
3          base > Immovable    color=WHITE img=base
4          avatar > FlakAvatar  stype=sam
5          missile > Missile
6              sam > orientation=UP   color=BLUE singleton=True img=spaceship
7              bomb > orientation=DOWN  color=RED speed=0.5 img=bomb
8          alien > Bomber       stype=bomb   prob=0.01 cooldown=3 speed=0.8 img=alien
9          portal >
10             portalSlow > SpawnPoint  stype=alien cooldown=16 total=20 img=portal
11             portalFast > SpawnPoint stype=alien cooldown=12 total=20 img=portal
12
13     LevelMapping
14         0 > base
15         1 > portalSlow
16         2 > portalFast
17
18     TerminationSet
19         SpriteCounter      stype=avatar           limit=0 win=False
20         MultiSpriteCounter stype1=portal stype2=alien limit=0 win=True
21
22     InteractionSet
23         avatar EOS > stepBack
24         alien  EOS > turnAround
25         missile EOS > killSprite
26         missile base > killSprite
27         base bomb > killSprite
28         base sam > killSprite scoreChange=1
29         base alien > killSprite
30         avatar alien > killSprite scoreChange=-1
31         avatar bomb > killSprite scoreChange=-1
32         alien sam > killSprite scoreChange=2
```

# VGDL – level definitions

1 wwwwwwwwwwwwwwwwwwwwwwwwwwwwwww  
2 w w  
3 w1 w  
4 w000 w  
5 w000 w  
6 w w  
7 w w  
8 w w  
9 w w  
10 w 000 000000 000 w  
11 w 00000 00000000 00000 w  
12 w 0 0 00 00 00000 w  
13 w A w  
14 wwwwwwwwwwwwwwwwwwwwwwwwwwwwwww

# a sample (random) controller

- realtime conditions important:
  - constructor max 1 sec
  - every action max 40 msec
- we do not need to have an idea of what actions actually mean
- but we get the action list and have to decide which to take

```
package random; //The package name is the same as the username in the web.

public class Agent extends AbstractPlayer {

    protected Random randomGenerator;

    //Constructor. It must return in 1 second maximum.
    public Agent(StateObservation so, ElapsedCpuTimer elapsedTimer)
    {
        randomGenerator = new Random();
    }

    //Act function. Called every game step, it must return an action in 40 ms maximum.
    public Types.ACTIONS act(StateObservation stateObs, ElapsedCpuTimer elapsedTimer) {

        //Get the available actions in this game.
        ArrayList<Types.ACTIONS> actions = stateObs.getAvailableActions();

        //Determine an index randomly and get the action to return.
        int index = randomGenerator.nextInt(actions.size());
        Types.ACTIONS action = actions.get(index);

        //Return the action.
        return action;
    }
}
```

# StateObservation (I)

- allows the agent to query the state of the game:
  - `StateObservation.getGameScore();`
  - `StateObservation.getGameTick();`
  - `StateObservation.getGameWinner();`
  - `StateObservation.isGameOver();`
  - `StateObservation.getWorldDimension()`
- ... the state of the avatar:
  - `StateObservation.getAvatarPosition();`
  - `StateObservation.getAvatarSpeed();`
  - `StateObservation.getAvatarOrientation();`
  - `StateObservation.getAvatarResources();`
- ... the available actions in the game:
  - `StateObservation.getAvailableActions();`
- ... the history of events (collisions) in the game:
  - `StateObservation.getEventsHistory();`

# StateObservation (II)

- ... Observations in the game:
  - `StateObservation.getObservationGrid();`
  - `StateObservation.getNPCPositions(position?);`
  - `StateObservation.getImmovablePositions(position?);`
  - `StateObservation.getMovablePositions(position?);`
  - `StateObservation.getResourcesPositions(position?);`
  - `StateObservation.getPortalsPositions(position?);`
  - `StateObservation.getFromAvatarSpritesPositions(position?);`

what is an *Observation*? it is an object that contains:

- int itype: type of sprite of this observation
- int category: category of this observation (static, resource, npc, etc.)
- int obsID: unique ID for this observation
- Vector2 position: position of the observation
- Vector2 reference: reference (pivot) position
- double sqDist: distance from this observation to the reference

# StateObservation (III)

in StateObservation:

- 1 ArrayList <Observation>[][] getObservationGrid();
- bi-dimensional array, matching the level grid
- each ArrayList contains list of observations in that position



- 1 ArrayList <Observation>[] getNPCPositions();
- 2 ArrayList <Observation>[] getNPCPositions(Vector2d reference);
- returns a list of observations of NPC in the game
- as there can be NPCs of different type, each entry in the array corresponds to a sprite type
- every ArrayList contains a list of objects of type Observation
- if “reference” is given, observations are sorted by distance to the reference’s position

# the forward model

also, in StateObservation:

```
StateObservation copy();
```

- create a copy of the StateObservation object

```
void advance(Types.ACTIONS action);
```

- advances the StateObservation object, applying the action supplied
- allows to simulate the effects of applying actions
- the StateObservation updates itself to reflect the next state
- important: the games are stochastic in nature!
  - the next state must be considered as a possible future state when applying a certain action.
  - the agent has the responsibility to deal with these inaccuracies.

# **evolutionary controller for the GVGAI**

following a “classical” approach for General Video Game Playing:

- each individual is a player: determines how action decisions are made

this is offline training:

- EA determines the best agent first
- the submitted agent does not use online evolution

this is the GA-Eater approach:

train a rule set, apply the learned rules

# closed loop – open loop

- until now we have always represented a state with a node
- this is called **closed loop**
- we can also represent the whole path that leads to a specific state in a node
- this is called **open loop**
- especially useful if our game is non-deterministic, and much smaller trees
- however, usually performing worse than closed loop if enough time
  
- another way to cope with chance events is to use AMAF value estimates / RAVE / GRAVE
- the key idea here is to generalize samples over more moves to get reasonable values quicker



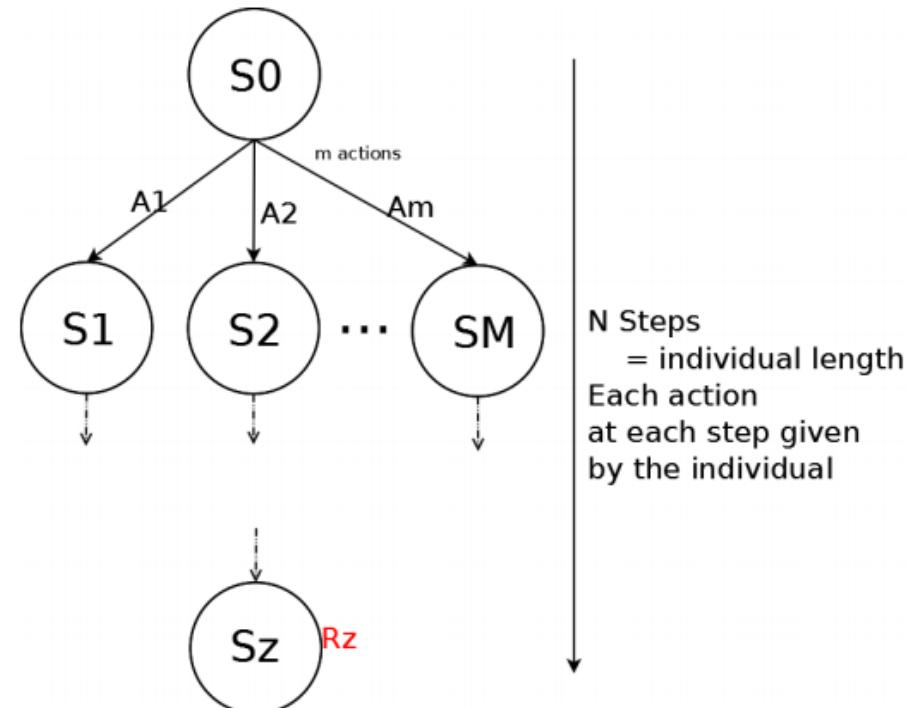
picture from Free Photos on Pixabay

# rolling horizon EA

basic idea:

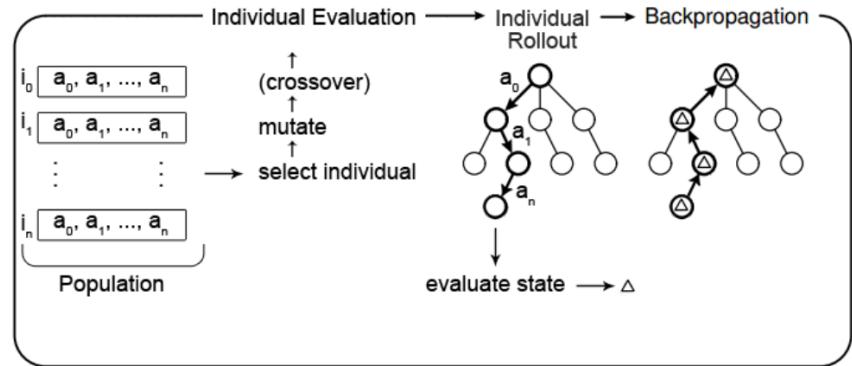
no offline training, evolution is used online, while playing the game

- each individual is a sequence of actions to apply from the current state (open loop control)
- fitness: evaluation of the state reached after applying the sequence of actions
- within the real-time constraints, the RH-EA evolves the best sequence of actions
- evolve the population normally: mutation, crossover, elitism, etc.
- when the process is over, apply first action of the best individual
- SampleGA sample controller



# combine RHEA and MCTS (or other enhancements)

- can't we use a "statistical tree" to memorize our experiences?
- or just add rollouts as in MCTS?
- or shift the buffer (move with applied action) and keep the rest of the precomputed path?



- yes, shift buffer and added rollouts work well, on par with MCTS on 20 GVGAI test games, but cannot say in advance which will work for which game (not very well understood)

#	Algorithm	Points	Avg-Wins	G-0	G-4	G-13	G-15	G-18	G-22	G-25	G-29	G-36	G-46	G-49	G-50	G-58	G-60	G-61	G-67	G-75	G-77	G-84	G-91
1	EA-shift-roll	430	42.05 (2.48)	18	<b>25</b>	18	<b>25</b>	18	18	18	<b>25</b>	<b>25</b>	<b>25</b>	18	<b>25</b>	18	<b>25</b>	<b>25</b>	<b>25</b>	18	18	18	<b>25</b>
2	MCTS	430	41.30 (1.76)	<b>25</b>	18	<b>25</b>	18	<b>25</b>	<b>25</b>	<b>25</b>	18	18	18	<b>25</b>	18	<b>25</b>	18	18	18	<b>25</b>	<b>25</b>	<b>25</b>	18

Table IV: Configuration 10-14,  $R = 5$ . Best algorithm found compared with MCTS. In this order, the table shows the rank of the algorithms, their name, total F1 points, average of victories and F1 points achieved on each game.

Raluca D. Gaina, Simon M. Lucas, Diego Perez Liebana: Rolling horizon evolution enhancements in general video game playing. CIG 2017: 88-95

# readings

very good intro paper that already explains a lot of variants:

- Cameron Browne, Edward Powley, Daniel Whitehouse, Simon Lucas, Peter Cowling, Philipp Rohlfshagen, Stephen Tavener, Diego Perez, Spyridon Samothrakis and Simon Colton, **A Survey of Monte Carlo Tree Search Methods**, *IEEE Transactions on Computational Intelligence and AI in Games* (2012), pp. Vol. 4:1 143-191.

multi-objective MCTS:

- Diego Perez Liebana, Sanaz Mostaghim, Spyridon Samothrakis, Simon M. Lucas: Multiobjective Monte Carlo Tree Search for Real-Time Games. *IEEE Trans. Comput. Intell. AI Games* 7(4): 347-360 (2015)

self-adapting RHEA:

- Raluca D. Gaina, Diego Perez Liebana, Simon M. Lucas, Chiara F. Sironi, Mark H. M. Winands: Self-Adaptive Rolling Horizon Evolutionary Algorithms for General Video Game Playing. *CoG* 2020: 367-374
- start here for more: <https://dblp.org/pid/164/9035.html>  
(Diego is probably the guy who did most GVGAI papers)

# take home

- game trees as primary tools for board game planning/learning
- Monte Carlo tree search necessary for larger games where min-max is too slow
- MCTS can be combined with heuristics
- for very large state spaces, MCTS is too slow: merge states (macro states), focus on subareas
- for planning, a forward model is necessary; we have to learn one if we do not have one
- Rolling Horizon EA competitive to plain MCTS in GVGAI (surprisingly)



picture from OpenClipart-Vectors on Pixabay