

# TEXT MINING

## L02. PREPROCESSING

SUZAN VERBERNE 2021

# TODAY'S LECTURE

- Quiz about week 1
- Go from raw text to clean text
- Character encoding
- Edit distance (+ exercise)
- Regular expressions
- Tokenization and sentence splitting
- Lemmatization and stemming

# QUIZ ABOUT WEEK 1

- What is Optical Character Recognition (OCR)?
  - a. A technique for recognizing the correct character encoding
  - b. A technique for converting handwritten text to digital text
  - c. A technique for converting the image of a printed text to a digital text
  - d. A technique for converting a formatted text to a plain text

# QUIZ ABOUT WEEK 1

- If we use words as features in a text classification task, the resulting vectors are
  - a. Low-dimensional and dense
  - b. Low-dimensional and sparse
  - c. High-dimensional and dense
  - d. High-dimensional and sparse

# QUIZ ABOUT WEEK 1

- What does the long-tail distribution for text data refer to?
  - a. When we add a document to a collection, the number of unique terms will increase
  - b. In a given document, most of the terms will have a frequency of zero
  - c. In a given collection, there are many terms with a low frequency and few terms with a high frequency

# QUIZ ABOUT WEEK 1

- For which type of text processing task are capitalization and punctuation more useful?
  - a. For sequence labelling
  - b. For classification

# QUIZ ABOUT WEEK 1

- Which evaluation metric would you prioritize for the task of identifying terrorist threats on Twitter, and why?
  - a. Precision, because we want to be sure that we found all threats
  - b. Recall, because we want to be sure that we found all threats
  - c. Precision, because we don't want to accuse someone wrongly
  - d. Recall, because we don't want to accuse someone wrongly

# WHO FOUND A TEAM MATE?



# GO FROM RAW TEXT TO CLEAN TEXT

# GATHERING RAW TEXT

- Written text
  - Digitized (scanned) documents (We need **OCR = optical character recognition**)
  - Born-digital documents
    - text, html, pdf, MS Word documents
- All text needs clean-up of some kind

# DIGITAL INPUT NOT CLEAN

- Scanned text and born-digital PDFs might contain:
  - photos, tables, graphics
  - layout or design information
  - disclaimers, copyright statements
  - headers, footers
  - column and page breaks
  - OCR errors
  - character encoding errors
- Semi-structured text: text with markup (HTML, XML, docx)

# OCR EXAMPLE (HISTORICAL)

TO THE TRVLVY HO-  
NORABLE AND RIGHT WOR-  
THY KNIGHT SIR THOMAS SMITH,

*TREASVRER* for the Colonies and Com-  
panies of VIRGINIA: and *Gouernour* of Mus-  
couia, East-India, North-west Passage,  
and SOMMER Ilands  
*Companies.*



HONORABLE SIR, the wisest of Men, or rather the wisdom of God tells vs, that *there is a time for all things*: and that the great God, who at his owne will beganne Time it selfe, doth at his owne time beginne all things else: the foolishnesse of men may aske and muse why was this so soone, and that so late? but the wisdom of God knowes what is fit for euery time: And surely amongst the sensible signes, and euident demonstrations of Gods all-gouerning prouidence, this is not the least, that he brings not forth his mightie works altogether, but *makes every thing beautifull in his time.* *Eccles. 3. 1.* And as in his creation he made not al at once, but produced them in their feuerall daies: so in his gubernation, he reuileth not the knowledge of all things in one Age, but discouers them in the feuerall ages of the World. And if man aske why God doth thus, holy *David* giues the answer; *The Lord hath so done his mar- uailous works, that they should be had in remembrance*; for were they all in one age (such is our corruption) they would bee lesse obserued and sooner forgotten, but being declared in their feuerall times, euery Age finds matter to magnifie God; And therefore He *whose glorious name is to be praised for euer*, reuails some meruailous thing in euery generation, that so his name may be praised from Generation to Generation. *Psal. 111. 4.* *Psal. 72. 19.*

TO THE TRVLVY HO.

NORABLE AND RIGHT Wok  
THY KNIGHT SIR Tiiom^s SMITHY  
YR E A S F R E R for the Colonies and Conj  
panics of V i & G 1 N 1 A: aod Gouernsur of Mtif.  
couia, Eaft-lildiaNorth-weltPaffagc,  
and S O M M E 9 Ilands  
COM

N O R A B L 11 S 1 It, the W121 Of Men, Ot  
rather the wifedome of God tells vs, thatibere  
*~r a time for a# tkings*: and that the great Gor4,  
who at his owne will beganne Time it fcl&  
dothathisowne timebeginne all things else:  
the-foolithneffe ofinen may askc and tnufo why  
was this fo roone, and that so late? but the wifedome of God  
knowes what is fit for euery time: And furely among(l the fen  
fible signes, and cuidcnt demonfirations of Gods all-gouerning  
prouidence, this is not the least, that he brings not fbrth his migh  
tie works altogether, but makgrvffY *thing boc#tifs.11 in Ur time. rs.*  
And as in his creation he made not al at once, but produced them  
in their feuerall daies: fo in his gubernation, hcreuedethnotthc  
knowledge of all things in one Age, but difcouers them in the  
fcuerall ages ofthe World. And if man aske why God doth thus,  
holy *David* giues the answerc; *The Lord both fo done Ur mar-*

# HTML SOURCE

```

73 site_slug: 'nunlv2';
74 };
```

```

75 </script>
76 <!-- DPG Ad Framework -->
77 <script async src="https://advertising-cdn.dpgmedia.cloud/web-advertising/11/8/1/advert-xandr.js" type="application/javascript"></script><script async data-advert-src="https://acdn.adnxs.com/ast/ast.js" data-advert-vendor=
78 <script id="gtm_dataLayer">
79 window.dataLayer = window.dataLayer || [];
80 dataLayer.push({"site_name": "NU.nl", "site_location": "prod", "site_country": "NL", "page_zone": "muziek-artikel", "page_type": "article", "page_category": "entertainment", "page_subcategory": "muziek", "page_catego
81 const pageViewId = window.crypto.getRandomValues(new Uint32Array(4)).join('');
82 window.pageViewId = pageViewId;
83 dataLayer.push({ pageViewId });
84 </script>
85 <!-- Google Tag Manager -->
86 <script>(function(w,d,s,l,i){w[l]=w[l]||[];w[l].push({'gtm.start':
87 new Date().getTime(),event:'gtm.js'});var f=d.getElementsByTagName(s)[0],
88 j=d.createElement(s),dl=l!='dataLayer'?'&l='+l:'';j.async=true;j.src=
89 'https://www.googletagmanager.com/gtm.js?id='+i+dl;f.parentNode.insertBefore(j,f);
90 })(window,document,'script','dataLayer','GTM-N2NWCN7');

```

# CLEAN TEXT STORAGE

- **Markup**: meta-information in a text file that is clearly distinguishable from the textual content
  - In the case of **XML and json**, markup often provides useful information in text processing
  - We typically convert PDF to XML, using pdf-to-xml convertors
  - Also, benchmark data is often stored as XML
- **Character Encoding**: the way that a computer displays text in a way that humans can understand.

# CHARACTER ENCODING

# ASCII

- Character encoding: translates a string of 0s and 1s to a character
- ASCII is a 7-bit encoding based on the English alphabet
  - 1100001 a
  - 1100010 b
  - 1100011 c
- ASCII: **American** Standard Code for Information Interchange



# HOW ABOUT THE REST OF THE WORLD?

- corpus linguistics (English)
- corpuslinguïstiek (Dutch)
- कोष भाषा विज्ञान (Hindi)
- اللسانيات الإحصاء (Arabic)
- 語料庫語言學 (Chinese)
- בלשנות קורפוס (Hebrew)
- การ ศึกษา ภาษาศาสตร์ (Thai)
- Цорпус Лингвистицс (Serbian)



# UNICODE

- Universal standard for all writing systems (>100,000 characters)
  - Independent of platform, software, vendor
- Interpretation of the character is done by the implementation (e.g. **UTF-8**) in the software (e.g. editor, printer or web browser) that determines the actual rendering (size, shape, font, style)
- For maximum compatibility (forward and backward) we encode text in UTF-8 when we read and write them

# READ AND WRITE UTF-8 IN PYTHON 3

```
with open(filename, 'r', encoding='utf-8') as raw:  
    text = raw.read()
```

```
with open(filename, 'w', encoding='utf-8') as clean:  
    clean.write(text)
```

<https://docs.python.org/3/library/functions.html#open>



# DATA CLEANING IN PRACTICE

- Digitalization, data conversion & cleaning are the first steps in the text mining process. These steps are:
  - necessary
  - time consuming
  - error prone
  - often complicated
- When building a text mining pipeline for new (raw) data, data collection and cleaning is long and tedious step that is not to be underestimated!

# EDIT DISTANCE

(SECTION 2.5 IN J&M)



# WHY EDIT DISTANCE

For measuring string similarity:

- **Spelling correction/normalization** (think about normalizing the content of doctor's notes, user-generated content, or search queries)
  - E.g. 'graffe' – what word was meant?
  - 'giraffe' differs by only one letter  $\Rightarrow$  most likely
  - 'grail' or 'graf' differ in more letters
- In bio-informatics: **Align two sequences** of nucleotides
  - AGGCTATCACCTGACCTCCAGGCCGATGCCC
  - TAGCTATCACGACCGCGGTTCGATTTGCCCGAC

# MINIMAL EDIT DISTANCE

- the **minimum edit distance** between two strings is defined as the minimum number of editing operations (insertion, deletion, substitution) needed to transform one string into another
- **Levenshtein distance**: insertion, deletion and substitution all have a cost of 1.
  - dog-do: 1
  - cat-cart: 1
  - cat-cut: 1
  - cat-act: ?
  - cat-act: 2

# COMPUTING MINIMAL EDIT DISTANCE

- “The space of all possible edits is enormous, so we can’t search naively.
- However, lots of distinct edit paths will end up in the same state (string), so rather than recomputing all those paths, we could just remember the shortest path to a state each time we saw it.
- We can do this by using dynamic programming.
- **Dynamic programming** is the name for a class of algorithms that apply a table-driven method to solve problems by combining solutions to sub-problems.”

(J&M, section 2.5.1)



# LEVENSHTEIN DISTANCE: ALGORITHM

		f	a	s	t
c					
a					
t					
s					

- Operations: insert (cost 1), delete (cost 1), substitute (cost 1), copy (cost 0)

# LEVENSHTTEIN DISTANCE: ALGORITHM

➤ Initialization

$$D(i, 0) = i$$

$$D(0, j) = j$$

➤ Recurrence Relation:

For each  $i = 1 \dots M$

For each  $j = 1 \dots N$

$$D(i, j) = \min \begin{cases} D(i-1, j) + 1 \\ D(i, j-1) + 1 \\ D(i-1, j-1) + \end{cases} \begin{cases} 1; & \text{if } X(i) \neq Y(j) \\ 0; & \text{if } X(i) = Y(j) \end{cases}$$

➤ Termination:

$D(N, M)$  is distance

$X$  = string 1

$Y$  = string 2

$i$  = index in  $X$

$j$  = index in  $Y$

$D(i, j)$  = value of cell  $(i, j)$

# LEVENSHTEIN DISTANCE: ALGORITHM

		f	a	s	t
c					
a					
t					
s					

- Operations: insert (cost 1), delete (cost 1), substitute (cost 1), copy (cost 0)

# LEVENSHTEIN DISTANCE: ALGORITHM

		f	a	s	t
	0	1	2	3	4
c	1	1	2	3	4
a	2	2	1	2	3
t	3	3	2	2	2
s	4	4	3	2	3

- Operations: insert (cost 1), delete (cost 1), substitute (cost 1), copy (cost 0)

# LEVENSHTTEIN DISTANCE: ALGORITHM

➤ Initialization

$$D(i, 0) = i$$

$$D(0, j) = j$$

➤ Recurrence Relation:

For each  $i = 1 \dots M$

For each  $j = 1 \dots N$

$$D(i, j) = \min \begin{cases} D(i-1, j) + 1 \\ D(i, j-1) + 1 \\ D(i-1, j-1) + \begin{cases} 1; & \text{if } X(i) \neq Y(j) \\ 0; & \text{if } X(i) = Y(j) \end{cases} \end{cases}$$

➤ Termination:

$D(N, M)$  is distance

$X$  = string 1

$Y$  = string 2

$i$  = index in  $X$

$j$  = index in  $Y$

$D(i, j)$  = value of cell  $(i, j)$

# LEVENSHTEIN DISTANCE: ALGORITHM

cost	operation	input	output
1	substitute	c	f
0	(copy)	a	a
1	substitute	t	s
1	substitute	s	t
Total: 3			

cost	operation	input	output
1	substitute	c	f
0	(copy)	a	a
1	deletion	t	*
0	(copy)	s	s
1	insertion	*	t
Total: 3			

# LEVENSHTEIN DISTANCE: EXERCISE

		a	n		a	c	t
	0	1	2	3	4	5	6
a	1						
	2						
c	3						
a	4						
t	5						

# LEVENSHTEIN DISTANCE: EXERCISE

		a	n		a	c	t
	0	1	2	3	4	5	6
a	1	0	1	2	3	4	5
	2	1	1	1	2	3	4
c	3	2	2	2	2	2	3
a	4	3	3	3	2	3	3
t	5	4	4	4	3	3	3



# LEVENSHTEIN DISTANCE: EXERCISE

cost	operation	input	output
0	(copy)	a	a
1	insert	*	n
0	(copy)		
1	substitute	c	a
1	substitute	a	c
0	(copy)	t	t

# REGULAR EXPRESSIONS

SECTION 2.1 IN J&M



# WHY REGULAR EXPRESSIONS

- For finding patterns, e.g.
  - “[0-9][0-9][0-9][0-9] [A-Z][A-Z]”
  - “https?:\/\/\S+”
  - “\S+@\S+”

# WHY REGULAR EXPRESSIONS

- [illegible]

# EXAMPLE (FROM THE BOOK)

- Suppose we want to write a regular expression to find all occurrences of the English article *the*
- A simple (but incorrect) pattern might be: `/the/`
- Why is this incorrect?

# EXAMPLE (FROM THE BOOK)

- Suppose we want to write a regular expression to find all occurrences of the English article *the*
- A simple (but incorrect) pattern might be: `/the/`
- Why is this incorrect?
  - We want to find both the and The
  - But not: their, apothecary, etc.
- Instead, if we would really want to cover all occurrences, and no non-relevant occurrences of the string, we need:

`/ (^ | [^a-zA-Z]) [tT]he ([^a-zA-Z] | $) /`

# EXAMPLE (ALTERNATIVES)

```
from nltk.tokenize import word_tokenize  
tokens = word_tokenize(raw_text)  
print(tokens.count('the'))
```

```
import spacy  
spacy_nlp = spacy.load("en_core_web_sm")  
spacy_doc = spacy_nlp(raw_text.lower())  
print(sum(token.text == 'the' for token in spacy_doc))
```

# TOKENIZATION AND SENTENCE SPLITTING

SECTION 2.2-2.4 & 2.4.5 IN J&M



# DEFINITIONS

- **Token** – An instance of a word or term occurring in a document
- **Term** – A token when used as feature (or in an index), generally in normalized form (e.g. lowercased)
- **Token count** is the number of words (running) in a collection/document; this includes duplicates
- **Vocabulary size** is the number of unique terms; the feature size when we use words as features

# TOKENIZATION

- Tokenization: Split text in tokens
- (1) remove punctuation; (2) split on whitespaces characters
- Question: *how would you want to tokenize these strings?*
  1. Hewlett-Packard
  2. State-of-the-art
  3. aren't
  4. C++
  5. cheap San Francisco-Los Angeles fares
  6. 20/03/97
  7. 071 527 7043

# TOKENIZATION WITH NLTK

```
from nltk.tokenize import word_tokenize  
tokens = word_tokenize(raw_text, language= 'english')
```

<https://www.nltk.org/api/nltk.tokenize.html>

# TOKENIZATION WITH SPACY

➤ <https://spacy.io/usage/linguistic-features#how-tokenizer-works>

➤ Customization for your task possible

Editable Code spaCy v3.0 · Python 3 · via Binder

```
import spacy

nlp = spacy.load("en_core_web_sm")
doc = nlp("Apple is looking at buying U.K. startup for $1 billion")
for token in doc:
    print(token.text)
```

RUN

```
Apple
is
looking
at
buying
U.K.
startup
for
$
1
billion
```

# STOP WORDS

- Stop words: extremely common words that don't carry any content
- Examples: *a, an, and, are, as, at, be, by, for, from, has, he, in, is, it, its, of, on, that, the, to, was, were, will, with*
- Stop word elimination commonly used in retrieval and classification
- But... For which cases is this problematic?

# STOP WORDS

- Stop words: extremely common words that don't carry any content
- Examples: *a, an, and, are, as, at, be, by, for, from, has, he, in, is, it, its, of, on, that, the, to, was, were, will, with*
- Stop word elimination commonly used in retrieval and classification
- But... For which cases is this problematic?
  - “to be or not to be”
  - You might need stop words for multi-word terms, e.g. “King of Denmark”, “Marks and Spencer”

# STOP WORDS

- **Advantages** of stop word removal in classification and retrieval:
  - reduces dimensionality (number of features)
  - removes the noise of highly frequent words (better model)
- **Disadvantages:**
  - in some cases stop words do carry important information
  - stop words are needed for phrases
- Rule of thumb: remove stop words for large data sets and long documents; keep them in small data sets and short documents

# SENTENCE SPLITTING

- Many cases can be solved with a relatively simple approach:
  - Sentences end with . ? ! followed by whitespace
- Challenges:
  - Abbreviations ('Mrs. Doe')
  - Names/initials that include punctuation marks ('S. Verberne')
  - Sentences without punctuation markers (headers/titles)
  - Line endings inside sentences (PDF conversion)



# SENTENCE SPLITTING WITH NLTK

```
from nltk.tokenize import sent_tokenize  
sentences = sent_tokenize(document, language='english')
```

<https://www.nltk.org/api/nltk.tokenize.html>

# SENTENCE SPLITTING WITH SPACY

➤ <https://spacy.io/usage/linguistic-features#sbd>

Editable Code

spaCy v3.0 · Python 3 · via Binder

```
import spacy

nlp = spacy.load("en_core_web_sm")
doc = nlp("This is a sentence. This is another sentence.")
assert doc.has_annotation("SENT_START")
for sent in doc.sents:
    print(sent.text)
```

RUN

```
This is a sentence.
This is another sentence.
```



# LEMMATIZATION AND STEMMING

(SECTION 2.4.4 IN J&M)

# BASIC WORD FORMS

- We might want to normalize specific word forms to the same term:
  - For example, it can be useful to have the term *bicycle* for each occurrence of either *bicycle* or *bicycles*

# BASIC WORD FORMS

- We might want to normalize specific word forms to the same term:
  - For example, it can be useful to have the term *bicycle* for each occurrence of either *bicycle* or *bicycles*
- Advantages:
  - reduces the number of features
  - generalizes better, especially for small datasets

# BASIC WORD FORMS

- We might want to normalize specific word forms to the same term:
  - For example, it can be useful to have the term *bicycle* for each occurrence of either *bicycle* or *bicycles*
- Advantages:
  - reduces the number of features
  - generalizes better, especially for small datasets
- Two types of basic word forms:
  - Lemma
  - Stem

# LEMMA

- Lemma: dictionary form of a word
- For example:
  - Verbs: infinitive
    - ‘think’ for ‘thinks’, ‘thinking’, ‘thought’
  - Nouns: singular form
    - ‘mouse’ for ‘mice’
    - ‘computer’ for ‘computers’

# STEM

- Stem: the portion of a word that:
  - is common to a set of (inflected) forms when all affixes are removed
  - is not further analyzable into meaningful elements, being morphologically simple
- The stem is the part of the word that 'never' changes even when morphologically inflected. It is not necessarily an existing word:
  - 'comput' for forms 'computer', 'computing', 'computers', 'compute'



# STEM VS LEMMA

- ‘produced’
  - what is the lemma?
  - what is the stem?

# STEM VS LEMMA

- ‘produced’
  - what is the lemma?
  - what is the stem?
- Answers:
  - the lemma is ‘produce’
  - the stem is ‘produc’

(think of ‘producing’ as one of the morphological forms of the verb)
- We almost always prefer lemmas over stems. Stemming can be effective for very small collections.

# EXAMPLE

- **Sample text:** Such an analysis can reveal features that are not easily visible from the variations in the individual genes and can lead to a picture of expression that is more biologically transparent and accessible to interpretation
- **Lemmatizer:** Such an analysis can reveal feature that be not easily visible from the variation in the individual gene and can lead to a picture of expression that be more biologically transparent and accessible to interpretation
- **Stemmer:** such an analysi can reveal featur that ar not easili visibl from the variat in the individu gene and can lead to a pictur of express that is more biolog transpar and access to interpret

# CONCLUSIONS

SUZAN VERBERNE 2021

# PYTHON PACKAGES AND TOOLS

- Sklearn (<http://scikit-learn.org>) has built-in functionality for:
  - Tokenization
    - [https://scikit-learn.org/stable/modules/generated/sklearn.feature\\_extraction.text.CountVectorizer.html](https://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.CountVectorizer.html)
  - Stop word removal (with option to supply your own list)
- NLTK (<http://www.nltk.org/>) and Spacy (<https://spacy.io/>) have functionality for:
  - Sentence splitting
  - Lemmatization and stemming
  - and additional pre-processing steps

# HOMework

- Read Jurafsky & Martin chapter 2: “Regular Expressions, Text Normalization, Edit Distance” (Brightspace)
- Complete this week’s exercise: preprocessing with spacy.  
<https://course.spacy.io/en/chapter1> **(note: only chapter 1)**
- Contact address for the teaching assistants:  
[tmcourse@liacs.leidenuniv.nl](mailto:tmcourse@liacs.leidenuniv.nl)

# AFTER THIS LECTURE...

- you know what issues to take into account when converting raw text to clean plain text
- you can explain the the difference between ASCII and [Unicode](#); and why Unicode exists
- you can use [regular expressions](#) in Python and you know what the possibilities and challenges of regular expressions are
- you can describe the challenges of [tokenization and sentence splitting](#)
- you can explain the considerations for [stop word removal](#)
- you can define the difference between [stemming and lemmatization](#)
- you can compute the [minimal edit distance](#) between two strings using the matrix algorithm

